

Theory and Methods in Causal Inference

Lecture Notes

Jae Kwang Kim

Department of Statistics

Iowa State University

STAT 5900B

Last updated: April 23, 2026

Contents

Preface	1
Organization	1
Appendices	1
Notation	2
 I Foundations	 3
1 Introduction	5
1.1 Motivation: Why Causal Inference Is Hard	5
1.1.1 The Central Question	5
1.2 The Causal Trinity: Three Languages for One Idea	6
1.2.1 Language 1: Structural Equation Models	7
1.2.2 Language 2: Directed Acyclic Graphs	8
1.2.3 Language 3: Potential Outcomes	8
1.2.4 Roadmap: How the Trinity Organizes These Notes	8
1.2.5 Historical Roots of the Causal Trinity	8
1.3 The Core Distinction	9
1.3.1 Two Scenarios: Observed vs. Unobserved Confounder	10
1.4 The Gaussian Linear Confounded Model	10
1.4.1 Two Explicit Densities	11
1.4.2 The Endogeneity Gap	11
1.4.3 Lab: Simulating the Two Densities	12
1.5 The Two-Step Paradigm	12
1.5.1 Worked Example: Flu Vaccination and Infection (Simpson’s Paradox)	12
1.6 Summary	13
1.7 Problems	14
 2 DAGs and d-Separation	 15
2.1 Directed Acyclic Graphs	15
2.1.1 Basic Definition	15
2.1.2 Structural Relationships	16
2.2 Paths, Blocking, and d-Separation	16
2.2.1 Path Structures	17
2.2.2 Blocking Rules	17
2.2.3 The d-Separation Criterion	17
2.2.4 Practical Ways to Check d-Separation	18
2.3 Collider Bias and the IV DAG	19
2.3.1 Berkson’s Bias	19
2.3.2 Full d-Separation Analysis of the IV DAG	19
2.4 The Markov Property and Factorization	20
2.5 Worked Example: The Education–Earnings DAG	21
2.6 The Big Picture	22
2.7 Summary	22
2.8 Problems	22

3	The Do-Calculus and Identification Criteria	25
3.1	From d-Separation to Intervention: Intervention Graphs	25
3.1.1	What Does Intervention Mean?	25
3.1.2	The Two Graph Operations	26
3.1.3	Why the Intervention Graph Works	27
3.2	The Back-Door Criterion	27
3.2.1	Motivation: When the Confounder Is Unobserved	27
3.2.2	Definition and Theorem	27
3.2.3	The Adjustment Formula: Standardization	29
3.2.4	Application to the Education–Earnings DAG	29
3.3	The Front-Door Criterion	30
3.3.1	The Configuration and Criterion	30
3.3.2	The Front-Door Formula	31
3.3.3	Interpretation: Two-Stage Causal Accounting	32
3.4	The Three Rules of Do-Calculus	33
3.4.1	Statement of the Rules	34
3.4.2	Intuition for Each Rule	34
3.4.3	Proofs of the Identification Formulas	36
3.5	The Do-Calculus is Complete	36
3.6	The Big Picture	38
3.7	Summary	38
3.8	Problems	39
II	Identification	41
4	Potential Outcomes and Adjustment	43
4.1	Motivation: A Third Language for Causality	43
4.2	The Neyman–Rubin Potential Outcomes Framework	44
4.2.1	The Potential Outcome	44
4.2.2	SUTVA and Consistency	44
4.2.3	Connection to the Do-Operator	44
4.3	Causal Estimands	45
4.3.1	The Average Treatment Effect and Its Relatives	45
4.3.2	Causal Estimands as Functionals of the Structural Equation	45
4.4	Ignorability, Positivity, and Adjustment	46
4.4.1	Strong Ignorability	46
4.4.2	Adjustment Formula under Ignorability	46
4.4.3	Back-Door Interpretation	46
4.4.4	Overlap and Positivity	47
4.5	Where the Frameworks Agree and Diverge	47
4.6	Summary	48
4.7	Problems	49
5	Randomization and Back-Door Adjustment	51
5.1	Randomized Experiments	51
5.1.1	The Do-Calculus of Randomization	51
5.1.2	Estimation of the ATE in a Randomized Experiment	52
5.2	Ignorability and the Back-Door Formula	53
5.2.1	Randomization vs. Observational Studies	53
5.2.2	Definition	53
5.2.3	Graphical Interpretation	53
5.2.4	From Ignorability to Identification	54
5.3	Regression Adjustment	54
5.3.1	The Outcome Regression (OR) Estimator	54
5.3.2	Standardization as a Special Case	54
5.3.3	Regression Adjustment in Randomized Experiments: Lin (2013)	55
5.4	Stratification	55
5.4.1	The Idea	55

5.4.2	The Stratified Estimator	56
5.4.3	Limitations	56
5.5	Simpson's Paradox	56
5.5.1	When Should You <i>Not</i> Condition on X ?	56
5.6	Lab: Simulation Study of the Outcome Regression Estimator	57
5.7	Where the Propensity Score Fits	58
5.8	Summary	58
5.9	Problems	58
6	Propensity Score Methods	61
6.1	Motivation: The Curse of Dimensionality	61
6.2	The Propensity Score	61
6.2.1	Definition	61
6.2.2	The Balancing Property	62
6.2.3	The Propensity Score Theorem	62
6.2.4	The Propensity Score is the Coarsest Balancing Score	63
6.3	Estimation of the Propensity Score	63
6.4	Matching on the Propensity Score	64
6.5	Inverse Probability Weighting	64
6.5.1	The IPW Identification Formula	64
6.5.2	The Horvitz–Thompson and Hájek Estimators	65
6.6	Overlap and Positivity	65
6.6.1	The Strict Overlap Condition	65
6.6.2	Practical Consequences of Near-Violations	65
6.6.3	Trimming Strategies	65
6.7	Lab: Simulation Study of IPW and Matching Estimators	66
6.7.1	Part 1: Correctly Specified Propensity Score	66
6.7.2	Part 2: Robustness to PS Model Misspecification	67
6.8	The Limits of Propensity Score Methods	68
6.9	Summary	68
6.10	Problems	68
7	Instrumental Variables	71
7.1	Why Instrumental Variables?	71
7.1.1	The Endogeneity Problem	71
7.1.2	The IV Idea	72
7.2	Graphical Setup and Core Assumptions	72
7.2.1	The IV DAG	72
7.2.2	The Three Assumptions in Three Languages	72
7.2.3	Relevance	73
7.2.4	Exogeneity	74
7.2.5	Exclusion	74
7.3	Identification in the Linear Homogeneous-Effect Model	74
7.3.1	The Linear Structural Model	75
7.3.2	The OLS Bias	75
7.3.3	Derivation of the Wald Estimand	75
7.4	Why the IV Assumptions Matter	75
7.5	Lab: OLS vs. IV Across Instrument Strengths	76
7.6	Multiple Instruments and Overidentification	77
7.7	Heterogeneous Treatment Effects and the LATE Framework	77
7.7.1	Compliance Types	77
7.7.2	The Monotonicity Assumption	78
7.7.3	The LATE Theorem	78
7.8	Interpreting IV Estimands	78
7.8.1	When Does LATE Equal ATE?	78
7.8.2	Different Instruments, Different Estimands	79
7.8.3	The Policy Relevance of LATE	79
7.9	IV versus Back-Door Adjustment	79
7.10	Applied Example: Charter School Lotteries and the KIPP Lynn Study	79

7.11	Summary	80
7.12	Problems	81
8	Mediation and Front-Door Identification	83
8.1	Motivation: Mechanisms	83
8.2	The Mediation DAG	84
8.2.1	The Prototype Graph	84
8.2.2	What Makes Mediation Harder Than Total Effect Estimation	84
8.2.3	Working Assumption for §§ 3–6	85
8.3	Total Causal Effect	85
8.4	Controlled Direct Effect	85
8.4.1	Definition	85
8.4.2	Identification of the CDE	86
8.4.3	The CDE Does Not Have a Complementary Indirect Effect	86
8.5	Natural Direct and Indirect Effects	86
8.5.1	Motivation: Cross-World Counterfactuals	86
8.5.2	Definitions	86
8.6	Identification of Natural Effects	87
8.6.1	Sequential Ignorability	87
8.6.2	The Mediation Formula	88
8.7	The Linear Mediation Model: Baron–Kenny	89
8.7.1	The Three-Equation System	89
8.7.2	The Product and Difference Formulas	89
8.7.3	Inference: The Sobel Test and Bootstrap	90
8.7.4	The Baron–Kenny Assumptions: Two Distinct Categories	90
8.8	Front-Door Identification	90
8.8.1	The Front-Door DAG	90
8.8.2	The Three Front-Door Conditions	91
8.8.3	Derivation of the Front-Door Formula	91
8.8.4	Intuition: Routing Around Confounding	92
8.9	Mediation vs. Instrumental Variables	92
8.10	Summary	93
8.11	Problems	94
9	Estimating Equations and Influence Functions	97
9.1	Why Estimation Needs Its Own Theory	97
9.2	A Running Example: The Average Treatment Effect	98
9.3	Estimating Equations	98
9.4	From Estimating Equations to Asymptotic Linearity	98
9.5	Influence Functions: Intuition	99
9.5.1	Statistical Functionals	99
9.5.2	Influence Functions	99
9.6	Influence Functions for Simple Causal Estimators	100
9.6.1	Known Outcome Regression	100
9.6.2	Known Propensity Score	100
9.7	Z-Estimation with Nuisance Parameters	101
9.7.1	The Stacked Estimating Equation Framework	101
9.7.2	Working Example: IPW with Estimated Propensity Score	102
9.8	Variance Estimation and Confidence Intervals	102
9.9	Toward Efficiency: Semiparametric Models and the EIF	103
9.9.1	Semiparametric Models	103
9.9.2	Regular Estimators	103
9.9.3	The Semiparametric Efficiency Bound and the EIF	103
9.9.4	The Efficient Influence Function for the ATE	104
9.10	Chapter Summary	104
9.11	Problems	105

III	Estimation	107
10	Doubly Robust Estimation and Semiparametric Efficiency	109
10.1	Why Combine Outcome Regression and Weighting?	109
10.2	The Prediction Estimator and Its Bias	110
10.3	The Augmented IPW Estimator	110
10.4	A Class of Augmented Estimators	111
10.5	The Projection Interpretation	111
10.6	Projection in the Causal Inference Setting	112
10.7	The Efficient Influence Function and Semiparametric Efficiency	112
10.8	Doubly Robust Regression: Weighted and Augmented Approaches	113
10.9	Lab: Simulation Study	114
10.10	Asymptotic Inference with Estimated Nuisance Functions	115
10.11	Comparison of Regression, IPW, and AIPW	116
10.12	Chapter Summary	116
10.13	Exercises	117
11	Flexible Nuisance Estimation, Orthogonal Scores, and Cross-Fitting	119
11.1	Why Flexible Nuisance Estimation Is Both Attractive and Dangerous	119
11.2	Why Naive Plug-In Estimation Can Fail	120
11.3	Orthogonal Scores	120
11.4	The Orthogonal Score for the Average Treatment Effect	121
11.5	Why Reusing the Same Data Can Be Problematic	122
11.6	Sample Splitting	122
11.7	Cross-Fitting	123
11.8	Double Machine Learning for the Average Treatment Effect	123
11.9	Rate Conditions and Asymptotic Normality	124
11.10	Lab: Simulation Study of the DML Estimator	125
11.11	Variance Estimation and Confidence Intervals	126
11.12	A Practical Workflow	126
11.13	What Machine Learning Does Not Solve	127
11.14	Chapter Summary	127
11.15	Problems	128
	References	131
	Appendices	135
A	Graphical Intuition for Conditional Independence and d-Separation	135
A.1	Conditional Independence: The Probabilistic Language Behind Graphs	135
A.2	Three Basic Motifs: Chain, Fork, and Collider	136
A.2.1	Chain	136
A.2.2	Fork	137
A.2.3	Collider	137
A.2.4	Summary of the Three Motifs	137
A.3	d-Separation: When Does Conditioning Block a Path?	137
A.3.1	A Confounding Example	138
A.3.2	A Collider Warning	138
A.3.3	A Practical Checklist	138
A.4	DAG Factorization and the Markov Property	139
A.5	Optional: Moralization as an Alternative Criterion	140
A.6	Summary	140
B	Single World Intervention Graphs	141
B.1	The SWIG Construction	141
B.2	The Fixed Half as a Source Node	141
B.3	When Ignorability Fails: Hidden Confounding	142
B.4	What SWIGs Achieve	143

B.5	Biconditional Equivalence under Faithfulness	144
B.6	Problems	144
C	Pathwise Differentiability and Efficient Influence Functions	147
C.1	Regular Parametric Submodels and Scores	147
C.2	Pathwise Differentiability	148
C.3	Non-Uniqueness of Influence Functions	148
C.4	The Tangent Space	149
C.5	The Canonical Gradient and the Efficiency Bound	150
	Bibliographic Notes	151

Preface

These lecture notes accompany the graduate course **STAT 5900B: Theory and Methods in Causal Inference** at the Department of Statistics, Iowa State University (Version 2.0).

The course targets first-year PhD students in statistics. It assumes familiarity with probability theory, mathematical statistics, and linear models at the level of a graduate sequence, but no prior exposure to causal inference.

[Download full PDF](#)

Organization

The notes are organized in three parts, with a fourth part on policy learning deferred to a later release.

Part I — Foundations (Chapters 1–3) builds the graphical and conceptual foundation of causal inference. Chapter 1 introduces the three languages of causality — structural equation models (SEMs), directed acyclic graphs (DAGs), and potential outcomes — and explains how they relate. Chapter 2 develops DAGs and d-separation as tools for reading off conditional independence from causal structure. Chapter 3 introduces the do-calculus and the main identification criteria (back-door, front-door, do-calculus rules) that translate causal queries into statistical ones.

Part II — Identification (Chapters 4–9) bridges the graphical framework and observed data. Chapter 4 connects the potential outcomes framework to the do-calculus and establishes the consistency, exchangeability, and positivity assumptions. Chapter 5 covers randomization and the back-door adjustment formula. Chapter 6 develops propensity score methods — weighting, matching, and the balancing property. Chapter 7 develops instrumental variables: the Wald estimand, the local average treatment effect (LATE) for compliers, and the three IV assumptions. Chapter 8 covers mediation analysis and the front-door criterion, distinguishing the controlled direct effect from the natural direct and indirect effects. Chapter 9 introduces the estimation framework — estimating equations, asymptotic linearity, influence functions, and the semiparametric efficiency bound — that underlies the methods of Part III.

Part III — Estimation (Chapters 10–13) covers semiparametric efficiency theory and its modern applications. Chapter 10 develops the augmented inverse probability weighted (AIPW) estimator from three complementary perspectives — bias correction, optimal augmentation, and the efficient influence function — and establishes double robustness and semiparametric efficiency. Chapter 11 addresses flexible nuisance estimation with orthogonal scores and cross-fitting, culminating in the double machine learning (DML) estimator. Chapters 12 and 13 cover estimation under instrumental variables and further extensions (to be added).

Part IV — Policy Learning and Sequential Decision Making (Chapters 14–16, deferred) will cover policy evaluation, optimal policy learning, and dynamic treatment regimes.

Appendices

Appendix A provides graphical intuition for conditional independence and d-separation. **Appendix B** introduces Single World Intervention Graphs (SWIGs). **Appendix C** develops the formal foundations of pathwise differentiability and efficient influence functions for readers who want the semiparametric geometry underlying Chapter 9.

Notation

Throughout these notes:

- $\text{do}(T = t)$ denotes the do-operator (intervention).
- $Y(t)$ denotes the potential outcome under $T = t$.
- $\mathbb{E}[\cdot]$ denotes expectation; $\perp\!\!\!\perp$ denotes conditional independence.
- $\mathcal{G}$ denotes a DAG; $\mathcal{G}_{\overline{T}}$ denotes the mutilated graph after intervening on T .
- $\pi(x) = P(T = 1 \mid X = x)$ denotes the propensity score.
- $\mu_t(x) = \mathbb{E}(Y \mid T = t, X = x)$ denotes the outcome regression.

Part I

Foundations

Chapter 1

Introduction

Learning Objectives

By the end of this chapter, students should be able to:

1. Articulate the distinction between an interventional distribution $P(y \mid \text{do}(T=t))$ and an observational conditional distribution $P(y \mid T=t)$, and explain why they differ whenever T is endogenous.
2. Describe the same causal question in all three languages — SEM, DAG graph surgery, and potential outcomes — and begin to translate between them.
3. Derive the two densities explicitly in the Gaussian linear confounded model, and identify the endogeneity bias $\alpha\gamma\sigma_U^2/\sigma_T^2$ as the gap between them.
4. State the two-step paradigm (identification then estimation) and locate each step in the causal inference pipeline.

This chapter is an orientation, not a complete treatment. It introduces four ideas — the interventional/observational distinction, the causal trinity, the Gaussian confounded model as a running example, and the identification-versus-estimation paradigm — that will be developed rigorously in the chapters that follow.

1.1 Motivation: Why Causal Inference Is Hard

1.1.1 The Central Question

Causal inference is concerned with a deceptively simple question: what would happen to Y if we were to *set* $T = t$? This is fundamentally different from the question that standard statistical procedures answer: what is the distribution of Y among units *observed* to have $T = t$?

Definition: Interventional vs. Observational Distribution

- The *interventional distribution* $P(y \mid \text{do}(T=t))$ is the distribution of Y in a hypothetical world where T has been *externally set* to t , regardless of the usual data-generating mechanism for T .
- The *observational conditional distribution* $P(y \mid T=t)$ is the distribution of Y among the subpopulation of units for whom $T = t$ was *observed* to hold.

These two distributions coincide when T is *exogenous* — that is, when T shares no common causes with Y . Whenever T is endogenous, the two distributions differ, and the difference is *endogeneity bias*.

Example: Drug Trial with Unmeasured Severity

Does a new drug ($T = 1$) reduce blood pressure (Y)? Suppose sicker patients are more likely to take the drug, so unmeasured severity U is a common cause of T and Y . Then:

- *Observational study*: $\mathbb{E}[Y \mid T=1] > \mathbb{E}[Y \mid T=0]$. Drug users have higher blood pressure on

average. A naïve analyst concludes the drug is harmful.

- *Randomized trial*: $\mathbb{E}[Y \mid \text{do}(T=1)] < \mathbb{E}[Y \mid \text{do}(T=0)]$. Forcing everyone to take the drug reduces blood pressure. The drug is beneficial.

Same data. Different questions. Opposite answers. The entire discrepancy is caused by the back-door path $T \leftarrow U \rightarrow Y$.

The Fundamental Problem of Causal Inference

We observe $P(y \mid T=t)$ directly from data. We want $P(y \mid \text{do}(T=t))$. The goal of *identification theory* is to find conditions under which the former can be used to recover the latter.

1.2 The Causal Trinity: Three Languages for One Idea

The same causal question can be expressed in three closely related languages: the *structural equation model* (SEM), the *directed acyclic graph* (DAG) with graph surgery, and the *potential outcomes* framework. Collectively, we call these the **causal trinity**. Each language illuminates a different facet of the same underlying idea, and fluency in all three is essential for modern causal reasoning.

```
\usetikzlibrary{arrows.meta, positioning, calc, shapes.geometric}
\newcommand{\Gcal}{\mathcal{G}}
\definecolor{isubblue}{RGB}{30,56,100}
\definecolor{accent}{RGB}{46,117,182}
\definecolor{defbg}{RGB}{238,244,251}
\definecolor{darkgrey}{RGB}{80,80,80}
\definecolor{semcolor}{RGB}{30,56,100}
\definecolor{dagcolor}{RGB}{26,122,58}
\definecolor{pocolor}{RGB}{150,50,150}
\begin{tikzpicture}[
  lang/.style={rectangle,rounded corners=6pt,draw,very thick,text width=3.0cm,align=center,minimum height=1.2cm},
  sem/.style={lang,draw=semcolor,fill=semcolor!10,text=semcolor},
  dag/.style={lang,draw=dagcolor,fill=dagcolor!10,text=dagcolor},
  po/.style={lang,draw=pocolor,fill=pocolor!10,text=pocolor},
  center/.style={ellipse,draw=accent,fill=defbg,thick,text width=2.6cm,align=center,font=\small\itshape},
  trans/.style={<->,thick,color=darkgrey,shorten >=4pt,shorten <=4pt},
  tlabel/.style={font=\footnotesize\itshape,color=darkgrey,fill=white,inner sep=2pt}
]
\node[po] (PO) at (90:3.8cm) {Potential\Outcomes\Y(t)};
\node[sem] (SEM) at (210:3.8cm) {SEM\Y=f_Y(T,U,\varepsilon)};
\node[dag] (DAG) at (330:3.8cm) {DAG \&\Graph Surgery\Gcal,\Gcal_{\overline{T}}};
\node[center] (C) at (0,0) {Set T=t?};
\draw[-{Stealth[length=5pt]},thick,color=pocolor!70,shorten >=4pt,shorten <=4pt] (PO)--(C);
\draw[-{Stealth[length=5pt]},thick,color=semcolor!70,shorten >=4pt,shorten <=4pt] (SEM)--(C);
\draw[-{Stealth[length=5pt]},thick,color=dagcolor!70,shorten >=4pt,shorten <=4pt] (DAG)--(C);
\draw[trans] (PO)--(SEM) node[tlabel,midway,left=2pt]{\parbox{2.2cm}{centering consistency\+ no in}};
\draw[trans] (SEM)--(DAG) node[tlabel,midway,below=4pt]{\parbox{2.2cm}{centering equation\leftarrow}};
\draw[trans] (DAG)--(PO) node[tlabel,midway,right=2pt]{\parbox{2.2cm}{centering SUTVA\+ graph structure}};
\end{tikzpicture}
```




Figure 1.1: The **causal trinity**: three languages that express the same causal question from different vantage points. Potential outcomes define *what* we want to estimate; the DAG with graph surgery specifies *how* to identify it from observed data; the SEM provides the *generative mechanism* for likelihood construction. Double-headed arrows show the translation assumptions connecting adjacent languages.

1.2.1 Language 1: Structural Equation Models

A SEM represents the causal data-generating mechanism as a system of equations, one per variable, with a joint distribution over exogenous disturbances. A variable is *exogenous* if it is determined outside the system; a variable is *endogenous* if its value is produced by a structural equation within the model.

The running setup for this chapter is a three-variable confounded SEM over (U, T, Y) :

$$T = f_T(U, \delta), \quad Y = f_Y(T, U, \varepsilon), \quad (1.1)$$

where δ and ε are mutually independent exogenous disturbances and U is an unobserved exogenous confounder.

Remark: Error Independence in the NPSEM-IE

The assumption that the disturbances δ and ε are *mutually independent* is a substantive restriction. Because U is unobserved and enters both f_T and f_Y , error independence does not remove the T – $Y(t)$ dependence induced by U ; it restricts only the additional channel through the equation-specific errors (δ, ε) . Richardson and Robins (2014) call SEMs with this structure *Nonparametric Structural Equation Models with Independent Errors* (NPSEM-IEs). These notes adopt the NPSEM-IE framework throughout: it ensures the Markov condition holds and integrates cleanly with the graphical identification theory of later chapters.

Observational regime. Because U enters both equations, conditioning on T changes the distribution of U , so $P(y \mid T=t) \neq P(y \mid \text{do}(T=t))$. An analyst who regresses Y on T without observing U faces a residual correlated with T — even though the structural error ε is independent of T by NPSEM-IE.

Interventional regime. The do-operator replaces the equation for T with a constant t , yielding the *mutilated system*: $U \sim p(u)$, $T = t$ (fixed), $Y = f_Y(t, U, \varepsilon)$. We define the *potential outcome* as $Y(t) := f_Y(t, U, \varepsilon)$; by construction, $Y(t) \sim P(y \mid \text{do}(T=t))$.

1.2.2 Language 2: Directed Acyclic Graphs

Definition: Graph Surgery

In the observational DAG, arrows into T encode the mechanisms that determine T . Under $\text{do}(T=t)$, all arrows into T are *deleted*, producing the mutilated graph $\mathcal{G}_{\overline{T}}$. This severs spurious back-door paths while preserving directed causal pathways from T to its descendants.

1.2.3 Language 3: Potential Outcomes

Definition: Potential Outcome [rubin1974estimating]

In the SEM framework, $Y(t) = f_Y(t, U, \varepsilon)$ is the outcome that *would have been observed* had T been set to t , possibly contrary to fact.

Definition: SUTVA

The *stable unit treatment value assumption* (SUTVA) has two components: (i) *no interference* — unit i 's potential outcome depends only on unit i 's own treatment; and (ii) *no hidden versions of treatment* — each treatment level corresponds to a single, well-defined intervention. Under SUTVA, the *consistency* relation $Y_i = Y_i(T_i)$ holds. SUTVA is developed further in Chapter 4.

1.2.4 Roadmap: How the Trinity Organizes These Notes

Part I — Foundations: the graphical language (Chapters 1–3). The DAG and the do-operator are the primary language here because they make the interventional/observational distinction *syntactically explicit*: confounding is visible as an open back-door path, and identifying assumptions are visible as graph criteria.

Part II — Identification: designs and research strategies (Chapters 4–9). Part II introduces the potential outcomes framework (Chapter 4) and then studies the specific research designs through which observational and experimental data support identification: randomization and back-door adjustment (Chapter 5), propensity score methods (Chapter 6), instrumental variables (Chapter 7), mediation and front-door identification (Chapter 8), and sensitivity analysis (Chapter 9). Ignorability ($Y(t) \perp\!\!\!\perp T \mid X$) is the potential-outcomes counterpart of the back-door criterion.

Part III — Estimation: semiparametric and machine-learning methods (Chapters 10–13). Once identification has been established, Part III turns to estimation. The semiparametric efficiency framework — estimating equations, influence functions, doubly robust estimators, and double machine learning — is largely language-neutral. Chapter 13 applies these tools to estimation under instrumental variables.

1.2.5 Historical Roots of the Causal Trinity

The three languages did not emerge from a single research program but grew independently in different disciplines over more than a century.

Structural Equation Models: genetics and econometrics. Sewall Wright (1921) introduced *path analysis* to decompose correlations among hereditary traits. The framework was transplanted into economics by Haavelmo (1943), who argued that economic relationships should be modeled as autonomous structural equations robust to policy interventions. The Cowles Commission formalized simultaneous equation systems throughout the 1940s–50s. Lucas's (1976) critique is essentially a statement of the interventional/observational distinction, decades before that phrase existed.

Potential Outcomes: experimental statistics and epidemiology. Neyman (1923) introduced the notation $Y_i(t)$ in the context of randomized agricultural experiments. The crucial extension to *observational* studies was made by Rubin (1974), who formalized the assignment mechanism and stated SUTVA explicitly. Holland's (1986) *JASA* paper brought the framework to mainstream statistics. Robins

(1986) extended potential outcomes to longitudinal treatments via g -computation. Imbens and Angrist (1994) connected the framework to instrumental variables and defined the LATE.

DAGs and do-calculus: computer science and philosophy. Pearl developed Bayesian networks through the 1980s. The pivotal step came in Pearl (1995), where the *do-operator* was introduced. The 2000 monograph *Causality* synthesized DAGs, the do-calculus, identification theory, and connections to potential outcomes. Spirtes, Glymour, and Scheines (1993) — working from philosophy at Carnegie Mellon — developed constraint-based algorithms for learning causal structure from data.

Convergence. The three frameworks are closely related and often intertranslatable, but exact equivalence requires additional assumptions. These notes work mainly in an NPSEM-IE setting, where translation between the languages is especially clean.

1.3 The Core Distinction

The single most important inequality of this course is:

$$\underbrace{P(y \mid x, \text{do}(T=t))}_{\text{interventional}} \neq \underbrace{P(y \mid x, T=t)}_{\text{observational}} \quad \text{whenever } T \text{ is endogenous.} \quad (1.2)$$

Definition: Conditional Exchangeability

Given observed covariates X and confounders U , the treatment assignment is *conditionally exchangeable with the intervention* if:

$$Y(t) \perp\!\!\!\perp T \mid X, U, \quad \text{or equivalently,} \quad P(y \mid x, T=t, U=u) = P(y \mid x, \text{do}(T=t), U=u).$$

Once we condition on all common causes (X, U) , observing $T=t$ and intervening to set $T=t$ produce the same distribution of Y .

Definition: Positivity

The treatment assignment satisfies *positivity* if, for every treatment value t of interest, $p_{T \mid X, U}(t \mid x, u) > 0$ for almost all (x, u) . This ensures that $P(y \mid x, T=t, U=u)$ is well-defined for every treatment value of interest.

Proposition: Back-Door Adjustment Formula

Under conditional exchangeability and positivity:

$$P(y \mid x, \text{do}(T=t)) = \int P(y \mid x, T=t, U=u) p(u \mid x) du. \quad (1.3)$$

Proof

By positivity, $P(y \mid x, T=t, U=u)$ is well-defined. Applying the law of total probability to the interventional density:

$$P(y \mid x, \text{do}(T=t)) = \int P(y \mid x, U=u, \text{do}(T=t)) p(u \mid x) du.$$

The weight is $p(u \mid x)$, not $p(u \mid x, T=t)$, because under $\text{do}(T=t)$ the treatment is set externally and carries no information about U . By conditional exchangeability, $P(y \mid x, U=u, \text{do}(T=t)) = P(y \mid x, T=t, U=u)$. Substituting completes the proof. $\square$

Remark: Confounding Discrepancy

What OLS estimates is the observational density $P(y \mid x, T=t) = \int P(y \mid x, T=t, U=u) p(u \mid x, T=t) du$. The kernel $P(y \mid x, T=t, U=u)$ is the same as in Equation 1.3, but averaged over the *selection-distorted* weight $p(u \mid x, T=t)$ rather than the marginal $p(u \mid x)$. The confounding discrepancy:

$$P(y \mid x, T=t) - P(y \mid x, \text{do}(T=t)) = \int P(y \mid x, T=t, U=u) [p(u \mid x, T=t) - p(u \mid x)] du$$

is nonzero when U and T are dependent *and* U affects the conditional distribution of Y .

1.3.1 Two Scenarios: Observed vs. Unobserved Confounder

Scenario 1: U is observed (no unmeasured confounding). When every component of U is recorded, both $P(y \mid x, T=t, U=u)$ and $p(u \mid x)$ can be estimated, and the back-door formula Equation 1.3 is a functional of the observable distribution. A special case is *unconfoundedness*: if $U \perp\!\!\!\perp T \mid X$ already holds, then $p(u \mid x, T=t) = p(u \mid x)$, and the formula reduces to $\mathbb{E}[Y \mid X=x, T=t]$. This is the key assumption underlying regression adjustment and propensity score methods (Chapters 5 and 6).

Scenario 2: U is unobserved (unmeasured confounding). When U is latent, the back-door formula is not usable from data. The main identification strategies are **instrumental variables** (Chapter 7), the **front-door criterion** (Chapter 3, applied in Chapter 8), and **sensitivity analysis** (Chapter 9).

	Scenario 1: U observed	Scenario 2: U unobserved
Confounding type	No unmeasured confounding	Unmeasured confounding
Back-door formula usable?	Yes: both terms estimable	No: U latent
Identification route	Back-door adjustment	IV, front-door, ...
Key assumption	$Y(t) \perp\!\!\!\perp T \mid X, U$	IV, front-door, or partial-identification assumptions
Chapters in this course	Chs. 3–6 (back-door); Chs. 10–11 (IPW, DR)	Ch. 7 (IV); Ch. 8 (front-door); Ch. 9 (sensitivity)

1.4 The Gaussian Linear Confounded Model

Example: Gaussian Linear Confounded Model

The structural equations are:

$$Y = \beta T + \gamma U + \varepsilon, \quad T = \alpha U + \delta, \quad (1.4)$$

where $U \sim \mathcal{N}(0, \sigma_U^2)$, $\varepsilon \sim \mathcal{N}(0, \sigma_\varepsilon^2)$, $\delta \sim \mathcal{N}(0, \sigma_\delta^2)$ are mutually independent. Here β is the causal effect of T on Y , γ is the direct effect of U on Y , and α is the direct effect of U on T .

The endogeneity of T is immediate: $\text{Cov}(T, \gamma U + \varepsilon) = \gamma \alpha \sigma_U^2 \neq 0$ whenever $\alpha \neq 0$ and $\gamma \neq 0$. Consequently, OLS regressing Y on T does not estimate β .

```
\usetikzlibrary{arrows.meta, positioning}
\newcommand{\doop}{\mathrm{do}}
\definecolor{isubblue}{RGB}{30,56,100}
\definecolor{defbg}{RGB}{238,244,251}
\definecolor{darkgrey}{RGB}{80,80,80}
\tikzset{
  node/.style={circle,draw=isubblue,fill=defbg,thick,minimum size=7mm,font=\small\bfseries},
  unode/.style={circle,draw=darkgrey,fill=gray!10,dashed,thick,minimum size=7mm,font=\small\bfseries},
  mnode/.style={rectangle,rounded corners=3pt,draw=darkgrey,fill=gray!15,thick,minimum size=7mm,font=\small\bfseries},
  edge/.style={-{Stealth[length=5pt]}},thick,color=isubblue},
```

```

    dedge/.style={--{\Stealth[length=5pt]},thick,color=darkgrey,dashed}
}
\begin{tikzpicture}[node distance=1.8cm]
  \node[node] (T1) at (0.5,0)   {\$T\$};
  \node[node] (Y1) at (2.5,0)   {\$Y\$};
  \node[unode] (U1) at (1.5,-1.3) {\$U\$};
  \draw[edge] (T1) -- (Y1);
  \draw[dedge] (U1) -- (T1);
  \draw[dedge] (U1) -- (Y1);
  \node[font=\small\itshape,color=darkgrey] at (1.5,-2.1) {Observational DAG};
  \node[mnode] (T2) at (5.5,0)   {\$t\$};
  \node[node] (Y2) at (7.5,0)   {\$Y\$};
  \node[unode] (U2) at (6.5,-1.3) {\$U\$};
  \draw[edge] (T2) -- (Y2);
  \draw[dedge, opacity=0.25] (U2) -- (T2);
  \draw[dedge] (U2) -- (Y2);
  \node[font=\small\itshape,color=darkgrey] at (6.5,-2.1) {Mutilated DAG: $\backslash$doop(T{=}=t)$};
\end{tikzpicture}

```




Figure 1.2: Graph surgery on the Gaussian confounded model. The faded arrow in the mutilated DAG has been severed by the intervention $\text{do}(T=t)$. The back-door path $T \leftarrow U \rightarrow Y$ is eliminated; only the causal path $T \rightarrow Y$ remains active.

1.4.1 Two Explicit Densities

Write $\sigma_T^2 = \alpha^2 \sigma_U^2 + \sigma_\delta^2$ for the marginal variance of T .

Interventional density. Set $T = t$ externally; the back-door path through U is severed, so $U \sim \mathcal{N}(0, \sigma_U^2)$ independently:

$$P(y \mid \text{do}(T=t)) = \mathcal{N}(\beta t, \gamma^2 \sigma_U^2 + \sigma_\varepsilon^2). \quad (1.5)$$

Observational density. Since (Y, T) is jointly normal, $\text{Cov}(Y, T) = \beta \sigma_T^2 + \gamma \alpha \sigma_U^2$, giving:

$$P(y \mid T=t) = \mathcal{N}\left(\left(\beta + \frac{\gamma \alpha \sigma_U^2}{\sigma_T^2}\right) t, \frac{\gamma^2 \sigma_U^2 \sigma_\delta^2}{\sigma_T^2} + \sigma_\varepsilon^2\right). \quad (1.6)$$

1.4.2 The Endogeneity Gap

	Interventional	Observational
Mean	βt	$\left(\beta + \frac{\gamma \alpha \sigma_U^2}{\sigma_T^2}\right) t$
Variance	$\gamma^2 \sigma_U^2 + \sigma_\varepsilon^2$	$\frac{\gamma^2 \sigma_U^2 \sigma_\delta^2}{\sigma_T^2} + \sigma_\varepsilon^2$
Residual	$(\gamma U + \varepsilon) \perp\!\!\!\perp T$	$(\gamma U + \varepsilon)$ correlated with T
Identified by	RCT (or IV, Chapter 7)	OLS
Equal when	$\alpha = 0$ or $\gamma = 0$ (no confounding)	

The *endogeneity bias* of OLS is $\gamma \alpha \sigma_U^2 / \sigma_T^2$. With $\alpha = \gamma = 1$ and $\sigma_U^2 = \sigma_\delta^2 = 1$: $\text{bias} = 1/(1+1) = 0.5$.

1.4.3 Lab: Simulating the Two Densities

The analytical results were verified by simulation using $n = 10,000$ draws with parameters $\beta = 2$, $\alpha = 1$, $\gamma = 1$, $\sigma_U = \sigma_\varepsilon = \sigma_\delta = 1$. (R code: `chapter1_lab.R`.)

Experiment 1: OLS bias. By the omitted-variable bias formula, the OLS probability limit is $\beta + \gamma\alpha\sigma_U^2/\sigma_T^2 = 2 + 0.5 = 2.5$.

	Simulated	Theory	Formula
OLS slope	2.5147	2.5000	$\beta + \gamma\alpha\sigma_U^2/\sigma_T^2$
True β	—	2.0000	β
Bias	0.5147	0.5000	$\gamma\alpha\sigma_U^2/\sigma_T^2$

OLS overestimates the causal effect by 25%.

Experiment 2: Two-sample comparison at $t_0 = 1$. The exact observational conditional distribution is $\mathcal{N}(2.5, 1.5)$; the interventional distribution is $\mathcal{N}(2, 2)$.

	Mean (Sim / Theory)	SD (Sim / Theory)
$\mathbb{E}[Y \mid T=1]$	2.497 / 2.500	1.229 / $\sqrt{1.5} = 1.225$
$\mathbb{E}[Y \mid \text{do}(T=1)]$	2.001 / 2.000	1.428 / $\sqrt{2} = 1.414$

The variance reduction in the observational distribution reflects $\sigma_{\text{obs}}^2/\sigma_{\text{int}}^2 = 1.5/2 = 0.75$: conditioning on $T=1$ pins down $\alpha U + \delta$, restricting the spread of Y through the correlation with γU .

Experiment 3: Density overlay for varying α . As α increases from 0 to 1.0, the observational density shifts rightward (growing bias) and narrows (reduced variance). The interventional density $\mathcal{N}(2, 2)$ is invariant to α .

1.5 The Two-Step Paradigm

Causal inference is a two-step discipline. The two steps are logically distinct and require different tools.

Definition: Identification

A causal quantity $P(y \mid \text{do}(T=t))$ is *identified* from the observational distribution P if there exists a functional Ψ such that $P(y \mid \text{do}(T=t)) = \Psi(P)$, and $\Psi(P)$ depends only on the observable joint distribution.

Step 1: Identification is a purely mathematical question about the causal model: can the interventional density be expressed as a function of the observed-data distribution? It does not depend on sample size. The main tools are: the back-door criterion (Chapter 3, applied in Chapters 4–6), the front-door criterion (Chapter 3), the three rules of the do-calculus (Chapter 3), and IV assumptions (Chapter 7).

Step 2: Estimation. Once $\Psi(P)$ has been identified, the statistical problem is to estimate the functional $\Psi(P)$ from n observations as efficiently as possible. The main tools are efficient influence functions (Part III), IPW (Part III), doubly robust estimators (Part III), and double machine learning (Part III).

1.5.1 Worked Example: Flu Vaccination and Infection (Simpson’s Paradox)

A public health agency observes 1,000 individuals. Each person either received a flu vaccine ($T = 1$) or not ($T = 0$), and either became infected ($Y = 1$) or not. Age group X (elderly = E , young = Y) is recorded. Elderly people are both more likely to be vaccinated and more susceptible to infection, so X confounds the T – Y relationship. The causal DAG has paths: $T \rightarrow Y$ (causal) and $T \leftarrow X \rightarrow Y$ (back-door).

Observed data:

	Vaccinated ($T = 1$)	Unvaccinated ($T = 0$)	Total
Elderly ($X = E$)	360, 30% infected	40, 50% infected	400
Young ($X = Y$)	60, 5% infected	540, 10% infected	600
Total	420, 26.4% infected	580, 12.8% infected	1,000

Simpson's paradox:

Stratum	Vaccinated	Unvaccinated	Difference	Conclusion
Elderly	30%	50%	−20 pp	vaccine helps
Young	5%	10%	−5 pp	vaccine helps
Aggregate	26.4%	12.8%	+13.6 pp	vaccine harms??

Within every stratum the vaccine reduces infection. Yet in the aggregate the vaccinated group has a *higher* infection rate, because 86% of vaccinated are elderly vs. 7% of unvaccinated.

Step 1 (Identification). The set $\{X\}$ satisfies the back-door criterion:

$$P(Y=1 \mid \text{do}(T=t)) = \sum_{x \in \{E, Y\}} P(Y=1 \mid T=t, X=x) P(X=x).$$

Step 2 (Estimation). With $\hat{P}(X=E) = 0.4$ and $\hat{P}(X=Y) = 0.6$:

$$\hat{P}(Y=1 \mid \text{do}(T=1)) = 0.30 \times 0.4 + 0.05 \times 0.6 = 0.15,$$

$$\hat{P}(Y=1 \mid \text{do}(T=0)) = 0.50 \times 0.4 + 0.10 \times 0.6 = 0.26.$$

The estimated causal risk difference is $0.15 - 0.26 = -0.11$: the vaccine causally reduces infection probability by 11 percentage points.

The Two Steps, Concretely

Step 1 (identification) established — from the DAG alone — that the back-door formula expresses the causal quantity as a functional of observable data. **Step 2 (estimation)** replaced the population probabilities in that formula with sample proportions.

These two steps are logically separate: Step 1 is a mathematical argument about the causal model that does not depend on sample size; Step 2 is a statistical problem conditional on the identification formula. A researcher who skips Step 1 and reports the raw difference conflates the two and draws the wrong conclusion.

1.6 Summary

- Interventional vs. observational distribution.** Causal inference answers questions about interventions $P(y \mid \text{do}(T=t))$, not about observations $P(y \mid T=t)$. In the Gaussian confounded model, the two densities coincide only when $\alpha = 0$ or $\gamma = 0$.
- The causal trinity.** The same causal question can be expressed in three languages — SEM (equation surgery), DAG (graph surgery), and potential outcomes ($Y(t)$). In the SEM framework, $Y(t)$ is defined as the solution for Y in the mutilated system.
- Endogeneity bias in the Gaussian confounded model.** The coefficient of T in $\mathbb{E}[Y \mid T=t]$ differs from the causal coefficient β by $\gamma\alpha\sigma_U^2/\sigma_T^2$. The observational variance $\gamma^2\sigma_U^2\sigma_\delta^2/\sigma_T^2 + \sigma_\varepsilon^2$ is smaller than the interventional variance $\gamma^2\sigma_U^2 + \sigma_\varepsilon^2$.
- Two-step paradigm.** Causal inference is: *identification* (expressing $\Psi(P)$ as a functional of observable data — a mathematical question) followed by *estimation* (constructing $\hat{\Psi}_n$ efficiently — a statistical question). The two steps are analytically separate.

Causal inference is the study of when the observational distribution contains enough information to recover the interventional distribution.

1.7 Problems

1. Do-notation fundamentals. For each statement below, decide whether it refers to an interventional or an observational distribution, and rewrite it unambiguously using do-notation where appropriate.

- (a) “The probability that a patient recovers given that they took the drug.”
- (b) “The probability that a patient would recover if we prescribed the drug to everyone.”
- (c) “Among students who attended tutoring, the average exam score was 85.”
- (d) “If we enrolled all students in tutoring, the average exam score would be 85.”

2. Deriving the observational density. In the Gaussian confounded model (Equation 1.4), verify Equation 1.6 by carrying out the following steps. Let $\sigma_T^2 = \alpha^2 \sigma_U^2 + \sigma_\delta^2$.

- (a) Show that (Y, T) is jointly normal by writing both as linear combinations of the independent normals (U, ε, δ) .
- (b) Compute $\text{Cov}(Y, T)$ and $\text{Var}(T)$.
- (c) Apply the conditional normal formula to derive $\mathbb{E}[Y \mid T=t]$ and $\text{Var}[Y \mid T=t]$, and hence verify Equation 1.6.
- (d) Show that the OLS probability limit is $\beta + \gamma \alpha \sigma_U^2 / \sigma_T^2$, and interpret this as an omitted-variable bias formula.

3. The causal trinity. Consider the following verbal causal claim: “Aspirin (T) reduces the risk of heart attack (Y) because it inhibits platelet aggregation (M), but patients with pre-existing cardiovascular disease (U , unobserved) are both more likely to take aspirin and more likely to have a heart attack.”

- (a) Draw the causal DAG implied by this description.
- (b) Write down the recursive SEM with appropriate structural functions f_T, f_M, f_Y .
- (c) State the SUTVA assumption and discuss whether it is plausible in this setting.
- (d) Explain why $\mathbb{E}[Y \mid T=1] - \mathbb{E}[Y \mid T=0]$ does not identify the causal effect $\mathbb{E}[Y \mid \text{do}(T=1)] - \mathbb{E}[Y \mid \text{do}(T=0)]$ in this DAG.

Chapter 2

DAGs and d-Separation

Learning Objectives

By the end of this chapter, students should be able to:

1. Construct a DAG from a verbal description of a causal model and identify parents, descendants, ancestors, and non-descendants of any node.
2. Classify every intermediate node on a path as a chain, fork, or collider, and apply the blocking rule for each.
3. Apply the d-separation criterion to determine all conditional independence relationships implied by a DAG.
4. Recognize and avoid collider bias, including the case where a descendant of a collider is conditioned upon.
5. State the Markov factorization and use it to write the joint density of a DAG in terms of local conditional densities.
6. Interpret d-separation results as conditional independence assumptions on observed variables, and recognize their role as the graphical expression of causal assumptions.

Readers who want a gentler introduction to conditional independence, the three basic path motifs, and the Markov-property interpretation may consult Appendix A before or alongside this chapter.

2.1 Directed Acyclic Graphs

DAGs allow us to translate qualitative causal assumptions into quantitative statistical restrictions. Once we draw a graph encoding our causal assumptions, d-separation tells us which conditional independences are implied by the graph and helps identify candidate adjustment sets for removing confounding.

Remark

In this chapter, the word *graph* does not mean a plot of data or a graph of a function. It means a collection of nodes and edges used to represent relationships among variables. Our objective is to understand how such graphs encode statistical structure — especially conditional independence, confounding, and path blocking.

2.1.1 Basic Definition

Definition: Directed Acyclic Graph

A *directed acyclic graph* (DAG) $\mathcal{G} = (\mathcal{V}, E)$ consists of a finite set of nodes $\mathcal{V}$ and a set of directed edges $E \subseteq \mathcal{V} \times \mathcal{V}$ such that there is no directed cycle. Each node represents a variable and each directed edge $A \rightarrow B$ encodes that A is a direct cause of B relative to the variables included in the graph.

The acyclicity condition guarantees a topological ordering so that the joint density admits the recursive factorization $p(v_1, \dots, v_k) = \prod_{i=1}^k p(v_i \mid \text{Pa}(v_i))$.

2.1.2 Structural Relationships**Definition: Structural Relationships in a DAG**

For a node $v \in \mathcal{V}$ in $\mathcal{G}$:

- $\text{Pa}(v) = \{w \in \mathcal{V} : w \rightarrow v \in E\}$ are the *parents* of v .
- $\text{De}(v)$ = all nodes reachable from v by directed paths are the *descendants* of v .
- $\text{An}(v)$ = all nodes with a directed path to v are the *ancestors* of v .
- $\text{Nd}(v) = \mathcal{V} \setminus (\text{De}(v) \cup \{v\})$ are the *non-descendants* of v .

Note that $\text{Nd}(v)$ includes the parents of v .

Quick Terminology Summary

If $A \rightarrow B$, then A is a *parent* of B and B is a *child* of A . Two nodes are *adjacent* if connected by an edge. A *path* is any sequence of connected nodes regardless of arrow direction; a *directed path* has arrows all pointing in the same direction. A *cycle* is a directed path that returns to its starting node.

Example: Education, Earnings, and Family Background

Education (E) affects earnings (Y); family background (B) is a common cause of both; neighborhood (N) affects education but has no direct effect on earnings. The Markov factorization is $p(n, b, e, y) = p(n)p(b)p(e \mid n, b)p(y \mid e, b)$. Reading off: $\text{Pa}(E) = \{N, B\}$, $\text{Pa}(Y) = \{E, B\}$, $\text{Pa}(N) = \text{Pa}(B) = \emptyset$.

Example: The IV DAG

The canonical IV model involves Z (instrument), T (treatment), Y (outcome), U (unobserved confounder), with $\text{Pa}(T) = \{Z, U\}$, $\text{Pa}(Y) = \{T, U\}$, $\text{Pa}(Z) = \text{Pa}(U) = \emptyset$. The descendants of Z are $\{T, Y\}$; U and Z are non-descendants of each other.

Example: The Mediation DAG

Treatment T affects outcome Y both directly and through mediator M . An unobserved confounder U creates a back-door path $T \leftarrow U \rightarrow Y$. $\text{Pa}(T) = \{U\}$, $\text{Pa}(M) = \{T\}$, $\text{Pa}(Y) = \{M, T, U\}$. The total effect of T on Y travels along two paths: $T \rightarrow Y$ (direct) and $T \rightarrow M \rightarrow Y$ (indirect). Note: this is a standard mediation DAG, not a front-door DAG. For the front-door criterion, every directed path from T to Y must pass through M .

2.2 Paths, Blocking, and d-Separation

A DAG encodes direct causal relationships through its edges, but variables can also be statistically related through longer routes. The concept of a *path* formalizes these routes, and d-separation answers whether a path transmits or blocks statistical dependence.

2.2.1 Path Structures

Definition: Path

A *path* between nodes X and Y in $\mathcal{G}$ is any sequence of distinct nodes $X = V_0, V_1, \dots, V_k = Y$ such that each consecutive pair is connected by an edge in either direction.

Every intermediate node V_m on a path plays one of three structural roles. In a **chain** ($V_{m-1} \rightarrow V_m \rightarrow V_{m+1}$), V_m is a causal intermediary. In a **fork** ($V_{m-1} \leftarrow V_m \rightarrow V_{m+1}$), V_m is a common cause. In a **collider** ($V_{m-1} \rightarrow V_m \leftarrow V_{m+1}$), both arrows point *into* V_m ; this asymmetric case has the opposite behavior under conditioning from the other two.

2.2.2 Blocking Rules

Structure	Pattern	Effect of conditioning	Key intuition
Chain	$A \rightarrow M \rightarrow B$	Blocks the path	Causal transmission
Fork	$A \leftarrow M \rightarrow B$	Blocks the path	Common cause
Collider	$A \rightarrow M \leftarrow B$	Opens the path	Common effect

The critical asymmetry: colliders are *closed by default* and *opened* by conditioning, while chains and forks are *open by default* and *closed* by conditioning.

Three-Node Motifs at a Glance

In a *chain* $A \rightarrow M \rightarrow B$: conditioning on M blocks the path. In a *fork* $A \leftarrow M \rightarrow B$: conditioning on M removes the spurious association. In a *collider* $A \rightarrow M \leftarrow B$: the path is blocked by default, but conditioning on M — or any descendant of M — opens it. These three motifs are the local building blocks of d-separation.

Chain: Smoking $\rightarrow$ Tar $\rightarrow$ Cancer. Marginally, smokers have elevated cancer rates (path is open). Conditioning on tar level blocks the path: once tar is fixed, the causal channel is fully accounted for.

Fork: Poverty $\rightarrow$ Poor Diet; Poverty $\rightarrow$ Lack of Exercise. Unconditionally, diet quality and exercise are correlated through poverty. Conditioning on income level removes the spurious association.

Collider: Accident $\rightarrow$ Hospitalization $\leftarrow$ Cancer. In the general population, $A \perp\!\!\!\perp B$. Among hospitalized patients — conditioning on M — the two become negatively associated. This is Berkson's bias (Berkson 1946).

2.2.3 The d-Separation Criterion

Definition: d-Blocking and d-Separation

A path π is *d-blocked* by a set $\mathbf{S}$ if:

1. π contains a chain $A \rightarrow M \rightarrow B$ or fork $A \leftarrow M \rightarrow B$ with $M \in \mathbf{S}$; or
2. π contains a collider $A \rightarrow C \leftarrow B$ with $C \notin \mathbf{S}$ and no descendant of C is in $\mathbf{S}$.

Nodes X and Y are *d-separated* by $\mathbf{S}$, written $(X \perp\!\!\!\perp Y \mid \mathbf{S})_{\mathcal{G}}$, if every path between X and Y in $\mathcal{G}$ is d-blocked by $\mathbf{S}$.

Example: Applying the Criterion to the Three Toy Settings

- (i) **Chain.** $A \rightarrow M \rightarrow B$: $\mathbf{S} = \emptyset$ gives $A \not\perp\!\!\!\perp B$; $\mathbf{S} = \{M\}$ gives $A \perp\!\!\!\perp B \mid M$.
- (ii) **Fork.** $A \leftarrow M \rightarrow B$: $\mathbf{S} = \emptyset$ gives $A \not\perp\!\!\!\perp B$; $\mathbf{S} = \{M\}$ gives $A \perp\!\!\!\perp B \mid M$.
- (iii) **Collider.** $A \rightarrow M \leftarrow B$: $\mathbf{S} = \emptyset$ gives $A \perp\!\!\!\perp B$ (collider blocks the path); $\mathbf{S} = \{M\}$ gives $A \not\perp\!\!\!\perp B \mid M$ (conditioning opens the path — Berkson's bias).

Example: A Direct Probability Proof for the Chain

Consider the chain $X_1 \rightarrow X_2 \rightarrow X_3$ with factorization $p(x_1, x_2, x_3) = p(x_1)p(x_2 | x_1)p(x_3 | x_2)$. Conditioning on X_2 :

$$p(x_1, x_3 | x_2) = \frac{p(x_1)p(x_2 | x_1)p(x_3 | x_2)}{p(x_2)} = p(x_1 | x_2)p(x_3 | x_2).$$

Hence $X_1 \perp\!\!\!\perp X_3 | X_2$. The analogous fork case $X_1 \leftarrow X_2 \rightarrow X_3$ is left as an exercise.

Theorem: Soundness of d-Separation [pearl2009causality, Theorem 1.2.4]

If $(X \perp\!\!\!\perp Y | \mathbf{S})_{\mathcal{G}}$, then $X \perp\!\!\!\perp Y | \mathbf{S}$ in every distribution that is Markov with respect to $\mathcal{G}$.

Remark: Soundness, Not Converse

The theorem states that d-separation *implies* conditional independence for every distribution Markov with respect to $\mathcal{G}$. The converse need not hold without additional assumptions such as *faithfulness*: a conditional independence may occur because of special parameter values even when the corresponding nodes are not d-separated in the graph.

2.2.4 Practical Ways to Check d-Separation

1. Bayes-Ball Intuition. Imagine releasing a ball from X asking whether it can reach Y , with nodes in $\mathbf{S}$ shaded. The ball: passes through chain/fork nodes not in $\mathbf{S}$; stops at chain/fork nodes in $\mathbf{S}$; stops at colliders not in $\mathbf{S}$ (with no descendant in $\mathbf{S}$); passes through colliders in $\mathbf{S}$ (or with a descendant in $\mathbf{S}$). The algorithm is formalized in Shachter (1998).

2. Moral Graph Transformation. (1) Restrict to the ancestral set $\text{An}(X \cup Y \cup \mathbf{S})$. (2) Moralize: for every collider $A \rightarrow C \leftarrow B$, add an undirected edge $A - B$. (3) Remove all edge directions. (4) Delete all nodes in $\mathbf{S}$. If X and Y are disconnected, then $(X \perp\!\!\!\perp Y | \mathbf{S})_{\mathcal{G}}$. See Lauritzen (1996, Ch. 3) for a full treatment.

Example: Checking d-Separation in a Four-Node DAG

Consider the DAG with edges $Z \rightarrow T, U \rightarrow T, T \rightarrow Y, U \rightarrow Y$.

Query 1. Is $(Z \perp\!\!\!\perp U)_{\mathcal{G}}$?

Bayes-ball: The only path is $Z \rightarrow T \leftarrow U$. Node T is a collider with $T \notin \emptyset$, so the ball stops. Hence $(Z \perp\!\!\!\perp U)_{\mathcal{G}}$.

Moral graph: Ancestral set of $\{Z, U\}$ is $\{Z, U\}$ (neither has parents); T and Y are discarded. No edges, no colliders. Z and U are disconnected $\Rightarrow (Z \perp\!\!\!\perp U)_{\mathcal{G}}$.

Query 2. Is $(Z \perp\!\!\!\perp U | T)_{\mathcal{G}}$?

Bayes-ball: Same path $Z \rightarrow T \leftarrow U$. Now $T \in \{T\}$: the collider is observed, so the ball passes through. Hence $Z \not\perp\!\!\!\perp U | T$.

Moral graph: Ancestral set of $\{Z, U, T\}$ includes T . Moralize: add edge $Z - U$ for the collider $Z \rightarrow T \leftarrow U$. After removing T , the remaining graph has edge $Z - U$. Connected $\Rightarrow Z \not\perp\!\!\!\perp U | T$. Comparing: Z and U are marginally independent (instrument is exogenous) but become dependent once we condition on T — Berkson's bias.

Example: Practice DAG — Eight-Node Graph

Consider the DAG with nodes $W, X_1, T, M_1, M_2, Y, C, D$ and edges $W \rightarrow X_1, W \rightarrow T, W \rightarrow Y, T \rightarrow M_1, M_1 \rightarrow Y, M_1 \rightarrow M_2, X_1 \rightarrow C, M_2 \rightarrow C, C \rightarrow D$.

(1) Is $X_1 \perp\!\!\!\perp Y$? No. Two open paths: $X_1 \leftarrow W \rightarrow Y$ (fork at W , unblocked) and $X_1 \leftarrow W \rightarrow T \rightarrow M_1 \rightarrow Y$ (also open). The path $X_1 \rightarrow C \leftarrow M_2$ ends at a collider C with $C \notin \emptyset$, so it is blocked.

(2) Is $X_1 \perp\!\!\!\perp Y | W$? Yes. Both open paths above pass through the fork W ; conditioning on W

blocks them. The remaining path $X_1 \rightarrow C \leftarrow M_2 \leftarrow \dots$ has collider $C \notin \{W\}$ — blocked. All paths blocked.

(3) Is $T \perp\!\!\!\perp M_2 \mid M_1$? Yes. The direct path $T \rightarrow M_1 \rightarrow M_2$ is a chain with $M_1 \in \{M_1\}$ — blocked. Other paths through W or Y also contain blocked colliders. All paths blocked.

(4) Does conditioning on C create collider bias? Yes. C is a collider on $X_1 \rightarrow C \leftarrow M_2$. Marginally $X_1 \perp\!\!\!\perp M_2$ (blocked by the collider). Conditioning on C opens the path, inducing spurious association.

(5) Does conditioning on D (a descendant of C) open the path $X_1 \rightarrow C \leftarrow M_2$? Yes. By clause (2) of the d-blocking definition, a collider path is unblocked whenever the collider *or any of its descendants* is in the conditioning set.

2.3 Collider Bias and the IV DAG

2.3.1 Berkson's Bias

Definition: Collider Bias

Collider bias is the spurious association induced between two variables when one conditions on their common effect (a collider), or on a descendant of that collider. In a path $A \rightarrow C \leftarrow B$, conditioning on C or any descendant of C can make A and B statistically dependent even if they are marginally independent.

Collider Bias vs. Confounding

Unlike confounding, which is *removed* by conditioning on a common cause, collider bias is *created* by conditioning on a common effect. Adding a collider or a descendant of a collider to the adjustment set introduces bias rather than reducing it.

Example: Collider Bias through Selection — Talent and Wealth

Suppose both talent (A) and wealth (B) increase the probability of admission to an elite school, so admission (S) is a collider: $A \rightarrow S \leftarrow B$. In the general population, $A \perp\!\!\!\perp B$ (the collider blocks the path). Among admitted students — conditioning on S — lower talent makes higher wealth more likely, and vice versa. This is collider bias: restricting to admitted students is precisely conditioning on the common effect.

2.3.2 Full d-Separation Analysis of the IV DAG

We work through the IV DAG with edges $Z \rightarrow T$, $T \rightarrow Y$, $U \rightarrow T$, $U \rightarrow Y$, with U *unobserved*. Because U cannot be conditioned on, the back-door path $T \leftarrow U \rightarrow Y$ cannot be blocked by adjustment, making the instrument Z the only route to identification.

Classifying paths by type:

Endpoints	Path	Type	Why
$Z \leftrightarrow Y$	$Z \rightarrow T \rightarrow Y$	<i>Causal</i>	Directed; carries the IV signal
$Z \leftrightarrow Y$	$Z \rightarrow T \leftarrow U \rightarrow Y$	<i>Associational</i>	Collider at T ; activates back-door via U
$Z \leftrightarrow U$	$Z \rightarrow T \leftarrow U$	<i>Associational</i>	Collider at T ; dormant unless T conditioned on

Question 1. Is $Z \perp\!\!\!\perp Y$? No. The directed path $Z \rightarrow T \rightarrow Y$ is open (chain, not conditioned on T). The path $Z \rightarrow T \leftarrow U \rightarrow Y$ has a collider at T with $T \notin \emptyset$ — blocked.

Question 2. Is $Z \perp\!\!\!\perp Y \mid T$? No. The causal path $Z \rightarrow T \rightarrow Y$ is blocked (conditioning on T in the chain). But the collider at T in $Z \rightarrow T \leftarrow U \rightarrow Y$ is now opened. Hence $Z \not\perp\!\!\!\perp Y \mid T$.

Question 3. Is $Z \perp\!\!\!\perp Y \mid \{T, U\}$? Yes. The causal path is blocked by conditioning on T . The path through U has the collider at T opened (by conditioning on T), but then blocked at U (by conditioning on U). Both paths blocked.

Question 4. Is $Z \perp\!\!\!\perp U$? Yes. The only path is $Z \rightarrow T \leftarrow U$ with a collider at T . Since $T \notin \emptyset$ and no descendant of T is in $\emptyset$, the path is blocked.

Interpretation. Question 1 establishes *relevance*: Z has an open causal path to Y through T . Question 4 establishes *exogeneity*: the instrument is independent of unmeasured confounding. Question 3 establishes the *exclusion restriction*: once T and U are held fixed, Z carries no residual information about Y . Note that this requires observing U , which is unavailable by assumption; the IV strategy exploits the exclusion restriction indirectly through the moment condition $\mathbb{E}[\varepsilon \cdot Z] = 0$ (Chapter 7). Question 2 is the *warning*: conditioning on T alone simultaneously closes the causal channel and opens the confounded channel — worse than no adjustment at all.

The IV DAG encodes the substantive assumptions behind instrument validity graphically, but does not make them testable in any strong sense.

2.4 The Markov Property and Factorization

Up to this point, we have used DAGs qualitatively. We now connect the graph to probability algebra: the same parent structure that governs d-separation also determines how the joint distribution factorizes.

Without structural assumptions, any joint distribution can always be written by repeated conditioning, but that generic representation is too high-dimensional to reveal much structure. A DAG becomes statistically meaningful because, together with the Markov property, it replaces the generic factorization by a sparse one involving only the parents of each node.

Proposition: Conditional Independence in the Three-Node Chain

Let $X_1 \rightarrow X_2 \rightarrow X_3$ be a chain with Markov factorization $p(x_1, x_2, x_3) = p(x_1)p(x_2 \mid x_1)p(x_3 \mid x_2)$. Then $X_1 \perp\!\!\!\perp X_3 \mid X_2$. (Proved in the chain example above; the fork case is analogous.)

Definition: Local Markov Property

A distribution P satisfies the *local Markov property* with respect to $\mathcal{G}$ if, for every node $V_i \in \mathcal{V}$:

$$V_i \perp\!\!\!\perp [\text{Nd}(V_i) \setminus \text{Pa}(V_i)] \mid \text{Pa}(V_i).$$

Once we condition on the direct causes of a node, that node is independent of all variables that are neither its descendants nor its parents.

Remark

The *global Markov property* is the statement that every d-separation in $\mathcal{G}$ implies a conditional independence in P — precisely the content of the Soundness theorem. For DAGs, the local and global Markov properties are equivalent (Lauritzen 1996, Proposition 3.27): the local property implies the global one via the Markov factorization.

An immediate consequence is the **Markov factorization**:

$$p(v_1, \dots, v_k) = \prod_{i=1}^k p(v_i \mid \text{Pa}(v_i)). \quad (2.1)$$

Example: Markov Factorization in the Education DAG

For the DAG $N \rightarrow E, B \rightarrow E, B \rightarrow Y, E \rightarrow Y$: $p(n, b, e, y) = p(n)p(b)p(e | n, b)p(y | e, b)$. A regression of Y on E alone is confounded by B . Conditioning on B blocks the back-door path $E \leftarrow B \rightarrow Y$ and identifies $P(y | \text{do}(E=e))$ via the back-door formula (Chapter 3).

Example: Markov Factorization in the IV DAG

Including the latent U : $p(z, t, y, u) = p(z)p(u)p(t | z, u)p(y | t, u)$. The observable factorization is obtained by marginalizing over U :

$$p(z, t, y) = \int p(z)p(u)p(t | z, u)p(y | t, u) du.$$

This cannot be simplified to $p(z)p(t | z)p(y | t)$ when U is a common cause of T and Y — this is precisely the endogeneity problem.

2.5 Worked Example: The Education–Earnings DAG

This section is the template for how we will use DAGs throughout the course: specify the DAG, read off the factorization, use d-separation to identify the implied conditional independences, and interpret in causal terms.

Practice. Before working through the example below, return to the eight-node DAG example and rework each of the five queries from scratch, following three steps: (1) list every path; (2) classify every intermediate node as chain, fork, or collider; (3) determine which paths are blocked or open.

The causal story. Education (E) affects earnings (Y); family background (B) is a common cause of both; neighborhood (N) affects education but has no direct effect on earnings. All four variables are observed (no latent confounders).

Step 1 — Markov factorization. $\text{Pa}(N) = \text{Pa}(B) = \emptyset$, $\text{Pa}(E) = \{N, B\}$, $\text{Pa}(Y) = \{E, B\}$:

$$p(n, b, e, y) = p(n)p(b)p(e | n, b)p(y | e, b).$$

Step 2 — d-Separation. There are two paths between N and Y :

- Path 1: $N \rightarrow E \rightarrow Y$ (a chain through E).
- Path 2: $N \rightarrow E \leftarrow B \rightarrow Y$ (a collider at E , followed by the fork leg $B \rightarrow Y$).

(i) Is $(N \perp\!\!\!\perp Y)_{\mathcal{G}}$? Path 1 is a chain with nothing conditioned on, so it is open. Path 2 has a collider at E (not conditioned on), so it is blocked. One open path: $N \not\perp\!\!\!\perp Y$.

(ii) Is $(N \perp\!\!\!\perp Y | E)_{\mathcal{G}}$? Path 1 is blocked (conditioning on E in the chain). Path 2 has a collider at E opened by conditioning, and the activated path continues through B (not conditioned on), so it is open. Hence $N \not\perp\!\!\!\perp Y | E$ — a collider trap: conditioning on E introduces bias, not removes it.

(iii) Is $(N \perp\!\!\!\perp Y | \{E, B\})_{\mathcal{G}}$? Path 1 is blocked by E . Path 2 is opened at collider E but then blocked at B . All paths blocked: $(N \perp\!\!\!\perp Y | E, B)_{\mathcal{G}}$.

(iv) Is $(N \perp\!\!\!\perp B)_{\mathcal{G}}$? The only path $N \rightarrow E \leftarrow B$ has a collider at E (not conditioned on). Path is blocked: $N \perp\!\!\!\perp B$.

Step 3 — Conditional Independence. By the Soundness theorem:

$$N \perp\!\!\!\perp B, \quad N \perp\!\!\!\perp Y | E, B, \quad N \not\perp\!\!\!\perp Y, \quad N \not\perp\!\!\!\perp Y | E.$$

The statement $N \not\perp\!\!\!\perp Y | E$ is an important warning: conditioning only on education when studying the neighborhood–earnings association *introduces* bias through the activated collider path.

Step 4 — Identification. We wish to identify $P(y | \text{do}(E=e))$. The only back-door path is $E \leftarrow B \rightarrow Y$ (a fork at B).

- $\mathbf{S} = \{B\}$: blocks the fork $E \leftarrow B \rightarrow Y$, and B is not a descendant of E . Valid.
- $\mathbf{S} = \{N\}$: does not block $E \leftarrow B \rightarrow Y$. Invalid.

With $\mathbf{S} = \{B\}$, the back-door formula (Chapter 3) gives:

$$P(y \mid \text{do}(E=e)) = \sum_b P(y \mid e, b) P(b),$$

identifying the causal effect entirely from observational data.

2.6 The Big Picture

The central logic runs from a qualitative causal graph to an identification formula:

structural assumptions $\implies$ DAG $\implies$ d-separation relations $\implies$ conditional independences (Markov) $\implies$ identification for

Chapter 2 establishes the first four links in this chain. Chapters 3 and beyond use the same graphical machinery to derive specific identification results: back-door adjustment, front-door adjustment, and do-calculus formulas.

2.7 Summary

1. **DAGs as causal structure.** A DAG encodes which variables are causally connected, which are potential confounders, and which variables may block or open paths.
2. **Three-node motifs.** Every path is built from chains, forks, and colliders. Chains and forks are open by default and are blocked by conditioning on the middle node. Colliders are blocked by default and are opened by conditioning on the collider or on one of its descendants.
3. **d-Separation and conditional independence.** X and Y are d-separated by $\mathbf{S}$ exactly when every path between them is blocked by $\mathbf{S}$. Under the Markov property, d-separation implies the corresponding conditional independence. Some implications may be assessed empirically, but assumptions involving unobserved variables are generally not testable.
4. **Markov factorization.** The local Markov property yields $p(v_1, \dots, v_k) = \prod_{i=1}^k p(v_i \mid \text{Pa}(v_i))$, the bridge from graphical structure to probability calculus.
5. **Collider bias.** Conditioning on a collider — or on a descendant of a collider — can induce a spurious association between otherwise independent variables (Berkson 1946). This is the opposite of confounding adjustment.
6. **A practical workflow.** To analyze a DAG: identify the graph structure, determine which paths are open or blocked, translate d-separation statements into conditional independences under the Markov property, and interpret in light of the causal question.

2.8 Problems

1. **Warm-up: a single collider.** Consider the DAG $X \rightarrow Y \leftarrow Z$.
 - (a) Identify the structural role of Y on the path $X \rightarrow Y \leftarrow Z$.
 - (b) Is $X \perp\!\!\!\perp Z$? Apply the d-separation criterion and state which blocking rule applies.
 - (c) Is $X \perp\!\!\!\perp Z \mid Y$? Explain what happens to the path when Y is conditioned on, and describe in one sentence the real-world phenomenon this illustrates.
2. **d-Separation practice.** Consider the DAG: $A \rightarrow B \rightarrow D$, $A \rightarrow C \rightarrow D$, $B \rightarrow E$, $C \rightarrow E$.
 - (a) List all paths between A and E . (*Hint*: there are four paths; two pass through D .)
 - (b) For each path, identify the role (chain, fork, collider) of each intermediate node.
 - (c) Does $\{B, C\}$ d-separate A and E ?
 - (d) Does $\{D\}$ d-separate B and C ? What type of node is D on the path $B \rightarrow D \leftarrow C$?

3. Berkson's bias. Suppose X and Y are independent standard normal variables, and let $S = \mathbf{1}[X + Y > 0]$.

- (a) Verify analytically that $\text{Cov}(X, Y \mid S=1) < 0$.
- (b) Draw the DAG for (X, Y, S) and identify S as a collider.
- (c) Explain in one sentence why restricting the analysis to the subsample with $S = 1$ biases estimates of any association between X and Y .

4. Markov factorization and collider activation. Consider the DAG with edges $A \rightarrow E$, $A \rightarrow W$, $F \rightarrow E$, $E \rightarrow W$, where A = ability, F = family income, E = education, W = wages.

- (a) Write down the Markov factorization $p(a, f, e, w)$.
- (b) Is $(F \perp\!\!\!\perp W)_{\mathcal{G}}$? List all paths between F and W and determine which are open.
- (c) Is $(F \perp\!\!\!\perp W \mid E)_{\mathcal{G}}$? Identify the role of E on each path.
- (d) A researcher regresses W on E and F , omitting A . Is the coefficient on E a causal effect? Explain using the graph.

5. Soundness theorem for the collider. Consider the collider $A \rightarrow M \leftarrow B$ with factorization $p(a, m, b) = p(a)p(b)p(m \mid a, b)$.

- (a) By marginalizing over M , show that $A \perp\!\!\!\perp B$ in the joint distribution. This verifies the Soundness theorem for $\mathbf{S} = \emptyset$.
- (b) Show that conditioning on M breaks this independence: write out $p(a, b \mid m)$ and explain why it does *not* factorize into $p(a \mid m)p(b \mid m)$ in general.
- (c) Explain in one sentence why parts (a) and (b) together are consistent with the Soundness theorem. (*Hint*: the theorem is a one-directional statement.)

6. Terminology check. Consider the DAG with edges $U \rightarrow X$, $U \rightarrow Z$, $X \rightarrow W$, $Z \rightarrow W$, $W \rightarrow Y$.

- (a) Identify the parents, children, ancestors, descendants, and non-descendants of node W .
- (b) List all pairs of adjacent nodes.
- (c) Which pairs of nodes are connected by a directed path? List every such pair and the corresponding path.
- (d) Write the Markov factorization $p(u, x, z, w, y)$.

7. Markov factorization and local Markov property. Consider the DAG with edges $X_1 \rightarrow X_2$, $X_1 \rightarrow X_3$, $X_2 \rightarrow X_4$, $X_3 \rightarrow X_4$.

- (a) Write the joint density $p(x_1, x_2, x_3, x_4)$ implied by the Markov factorization.
- (b) State the local Markov property for each of the four nodes.
- (c) Is $(X_2 \perp\!\!\!\perp X_3)_{\mathcal{G}}$? Is $(X_2 \perp\!\!\!\perp X_3 \mid X_1)_{\mathcal{G}}$? Justify each answer by listing all paths.

8. Toy proof: conditional independence in the fork. Consider the fork $X_1 \leftarrow X_2 \rightarrow X_3$ with factorization $p(x_1, x_2, x_3) = p(x_2)p(x_1 \mid x_2)p(x_3 \mid x_2)$.

- (a) Show directly, by conditioning on $X_2 = x_2$, that $X_1 \perp\!\!\!\perp X_3 \mid X_2$.
- (b) Is $X_1 \perp\!\!\!\perp X_3$ marginally? Justify both graphically and algebraically.
- (c) Explain in one sentence what the fork represents substantively and why conditioning on the common cause removes the association.

9. d-Separation in the practice DAG. Refer to the eight-node DAG in the Practice DAG example (nodes $W, X_1, T, M_1, M_2, Y, C, D$; edges $W \rightarrow X_1, W \rightarrow T, W \rightarrow Y, T \rightarrow M_1, M_1 \rightarrow Y, M_1 \rightarrow M_2, X_1 \rightarrow C, M_2 \rightarrow C, C \rightarrow D$). For each query, state whether it is true or false and justify by listing all relevant paths.

- (a) $X_1 \perp\!\!\!\perp Y$
- (b) $X_1 \perp\!\!\!\perp Y \mid W$
- (c) $T \perp\!\!\!\perp M_2 \mid M_1$
- (d) Does conditioning on C induce collider bias between X_1 and M_2 ? Identify the specific path opened.
- (e) Does conditioning on D (a descendant of C) open the path $X_1 \rightarrow C \leftarrow M_2$? State which clause of the d-blocking definition applies.

Chapter 3

The Do-Calculus and Identification Criteria

Learning Objectives

By the end of this chapter, students should be able to:

1. Define the two intervention graph operations $\mathcal{G}_{\overline{X}}$ and $\mathcal{G}_{\underline{X}}$, explain the causal meaning of $\mathcal{G}_{\overline{X}}$, and describe the technical role of $\mathcal{G}_{\underline{X}}$.
2. Apply the back-door criterion to determine whether a set $\mathbf{S}$ identifies a causal effect, and write the back-door adjustment formula.
3. Apply the front-door criterion in graphs where the confounder is unobserved but a suitable mediator exists.
4. State all three rules of do-calculus, identify the graph condition that licenses each rule, and use them to prove the back-door and front-door formulas.
5. State the completeness theorem of Shpitser and Pearl (2006) and explain its practical implication: if the do-calculus cannot identify a quantity from the observational distribution relative to the assumed graph, then no purely observational method can do so without additional assumptions or new data.

Remark: How to Read This Chapter

This chapter has two pedagogical layers. Sections 1–3 develop the two most important concrete identification criteria — back-door and front-door adjustment — using intervention graphs and direct graphical reasoning. These sections are the core material for a first reading. Sections 4–5 then place these criteria inside the more general theory of the do-calculus and identifiability. These later sections are conceptually important, but more abstract and may be read as a second pass after the concrete criteria are understood.

3.1 From d-Separation to Intervention: Intervention Graphs

In Chapter 2 we learned to read conditional independence from a DAG using d-separation. This chapter takes the next step: translating the graphical language into *identification formulas* — expressions that write the interventional distribution $P(y \mid \text{do}(T=t))$ entirely in terms of quantities observable from data.

3.1.1 What Does Intervention Mean?

Conditioning versus intervening. The conditional distribution $P(Y \mid T=t)$ describes the subpopulation of units for whom T was *observed* to equal t . Because treatment assignment may be influenced by confounders, this subpopulation is not representative of the full population. The interventional distribution $P(Y \mid \text{do}(T=t))$, by contrast, describes the population that would result if T were *externally set* to t for

everyone — removing the natural process that determines T and replacing it with the forced value. The two distributions coincide only when T has no unmeasured common causes with Y .

A simple illustration: suppose temperature Z causes both ice-cream sales T and crime Y . The quantity $P(Y | T=t)$ is the crime rate on days when ice-cream sales happen to equal t — confounded by Z . The quantity $P(Y | \text{do}(T=t))$ is the crime rate we would observe if we *fixed* sales at level t by decree. Forcing sales does not change the temperature, so crime remains driven only by Z , not by the sales level we mandated.

The structural basis of the do-operator. In the SEM framework of Chapter 1, an intervention $\text{do}(T=t)$ **replaces the entire structural equation** for T with the constant $T = t$. The replacement severs T 's dependence on its parents $\text{Pa}(T)$, while all other structural equations remain unchanged.

Graph surgery as the graphical realization. Because $\text{Pa}(T)$ no longer affects T after the intervention, every arrow pointing *into* T in the DAG is removed. The resulting graph — the *intervention graph* $\mathcal{G}_{\overline{T}}$ — represents the post-intervention world. D-separation in $\mathcal{G}_{\overline{T}}$ encodes conditional independence in $P(\cdot | \text{do}(T=t))$, not in the original observational distribution P .

3.1.2 The Two Graph Operations

Definition: Intervention Graphs $\mathcal{G}_{\overline{X}}$ and $\mathcal{G}_{\underline{X}}$

Let $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ be a DAG and $X \subseteq \mathcal{V}$ a set of nodes.

- $\mathcal{G}_{\overline{X}}$ (*intervention graph*, or *mutilated graph*) is obtained by *deleting all arrows into* X . This represents the intervention $\text{do}(X=x)$, which replaces the structural equation for X by the constant x , severing X 's dependence on all its former parents.
- $\mathcal{G}_{\underline{X}}$ (*auxiliary observation graph*) is obtained by *deleting all arrows out of* X . This is a *technical device* used in graphical conditions for certain do-calculus steps. It does *not* represent a probabilistic model of conditioning on X ; it is simply the graph in which the relevant d-separation condition must be checked.

Two Graph Operations at a Glance

- **Delete arrows into X** ($\mathcal{G}_{\overline{X}}$): models the intervention $\text{do}(X=x)$. D-separation in $\mathcal{G}_{\overline{X}}$ encodes independence in the post-intervention distribution $P(\cdot | \text{do}(X=x))$.
- **Delete arrows out of X** ($\mathcal{G}_{\underline{X}}$): a technical device used in graphical conditions for certain do-calculus steps. It is *not* a model of conditioning on X .

Figure 3.1 illustrates both graph operations on the IV DAG $\{Z \rightarrow T, T \rightarrow Y, U \rightarrow T, U \rightarrow Y\}$.




Figure 3.1: The two graph operations on the IV DAG. Faded arrows have been removed. $\mathcal{G}_{\overline{T}}$ models external intervention $\text{do}(T=t)$: the back-door path $T \leftarrow U \rightarrow Y$ is severed. $\mathcal{G}_{\underline{T}}$ is a technical device used in graphical conditions for certain do-calculus steps; it does not represent conditioning on T .

3.1.3 Why the Intervention Graph Works

The Intervention Principle

In the original DAG $\mathcal{G}$, the node T is determined by its parents $\text{Pa}(T)$ via the structural equation $T = f_T(\text{Pa}(T), \varepsilon_T)$. Intervening on T removes the structural mechanism that normally determines T ; therefore the parents of T no longer influence it. The intervention $\text{do}(T=t)$ replaces this equation with the constant $T = t$, making T independent of all its former parents. Graphically, this severs every arrow into T . The remaining structure of the graph is unchanged: all other structural equations still hold.

A d-separation statement in $\mathcal{G}_{\overline{T}}$ is therefore a conditional independence statement in the *post-intervention* distribution $P(\cdot \mid \text{do}(T=t))$, not in the observational distribution P .


3.2 The Back-Door Criterion

The back-door criterion is the most widely used identification strategy. It gives a simple graphical condition under which conditioning on an observed set $\mathbf{S}$ is sufficient to identify the causal effect of T on Y .

3.2.1 Motivation: When the Confounder Is Unobserved

If all common causes of T and Y were observed, identification would be trivial: adjust for all of them. The interesting case is when some confounders are *unobserved*, yet an observed set $\mathbf{S}$ can still block every spurious path.

Consider the following DAG: a drug T affects recovery Y . Socioeconomic status U is an unobserved confounder ($U \rightarrow T, U \rightarrow Y$), but U also affects an observed proxy W ($U \rightarrow W$), which in turn affects Y ($W \rightarrow Y$).



The only back-door path is $T \leftarrow U \rightarrow W \rightarrow Y$. Although U is unobserved, W sits on this path and *is* observed. Conditioning on W blocks the path at the chain node W , and W is not a descendant of T , so $\mathbf{S} = \{W\}$ satisfies the back-door criterion. The causal effect is identified even though U is never measured. The criterion does not require adjusting for the confounders themselves, only for an observed set that *intercepts* every back-door path.

3.2.2 Definition and Theorem

Definition: Back-Door Criterion

A set of observed variables $\mathbf{S}$ satisfies the *back-door criterion* for the causal effect of T on Y in $\mathcal{G}$ if:

1. No node in $\mathbf{S}$ is a descendant of T .
2. $\mathbf{S}$ blocks every *back-door path* from T to Y , i.e., every path that begins with an arrow pointing into T .

Remark

The terminology reflects the geometry of the graph relative to T . A *back-door path* enters T from behind — it represents a confounding route through which T and Y become spuriously associated without any causal mechanism from T to Y . The front-door criterion of Section 8.8 is named for the complementary idea: it identifies the causal effect by intercepting the *front-door* (causal) paths at a mediator M , rather than blocking the back-door (confounding) paths directly. Standing at T , confounders sneak in through the back door; the causal effect walks out through the front.

Remark

Condition 2 is equivalent to requiring $(Y \perp\!\!\!\perp T \mid \mathbf{S})_{\mathcal{G}_T}$: $\mathbf{S}$ d-separates T and Y in the graph $\mathcal{G}_T$ obtained by *deleting all arrows out of T* . Removing T 's outgoing edges eliminates the causal paths $T \rightarrow \dots \rightarrow Y$, leaving only back-door paths between T and Y .

The back-door criterion adds two constraints: $\mathbf{S}$ must consist of *observed* variables only, and it must contain no descendant of T . Every valid adjustment set is therefore a valid d-separator in $\mathcal{G}_T$, but the converse does not hold.

Example: A Single Confounder

Consider the DAG $T \leftarrow C \rightarrow Y$ with $T \rightarrow Y$, where C is an observed confounder. There is one back-door path: $T \leftarrow C \rightarrow Y$.

Does $\mathbf{S} = \{C\}$ satisfy the back-door criterion?

1. C is not a descendant of T .
2. $\{C\}$ blocks $T \leftarrow C \rightarrow Y$ (fork at C).

The back-door formula gives: $P(y \mid \text{do}(T=t)) = \sum_c P(y \mid t, c) P(c)$.

Example: Two Back-Door Paths — Both Must Be Blocked

Consider treatment T , outcome Y , observed variables C_1 and C_2 , and an unobserved confounder U : $C_1 \rightarrow T$, $C_1 \rightarrow Y$, $U \rightarrow T$, $U \rightarrow C_2$, $C_2 \rightarrow Y$, $T \rightarrow Y$.

There are exactly two back-door paths from T to Y :

1. $T \leftarrow C_1 \rightarrow Y$ (direct confounding by the observed C_1),
2. $T \leftarrow U \rightarrow C_2 \rightarrow Y$ (indirect confounding via the unobserved U , routing through the observed proxy C_2).

$\mathbf{S} = \{C_1\}$ fails: it blocks path 1 but leaves path 2 open. $\mathbf{S} = \{C_2\}$ fails: it blocks path 2 but leaves path 1 open. $\mathbf{S} = \{C_1, C_2\}$ succeeds: it blocks both paths, and neither C_1 nor C_2 is a descendant of T .

M-Bias: More Variables Can Be Worse

Consider a graph where T and Y each have their own separate unobserved confounders, U_1 and U_2 , that happen to share a common observed descendant C : $U_1 \rightarrow T$, $U_2 \rightarrow Y$, $U_1 \rightarrow C$, $U_2 \rightarrow C$, $T \rightarrow Y$ (with U_1, U_2 unobserved).

Marginally, the only back-door path $T \leftarrow U_1 \rightarrow C \leftarrow U_2 \rightarrow Y$ is *blocked* by the collider C . The causal effect is already identified without any adjustment!

Conditioning on C *activates* the collider, opening the path. The set $\{C\}$ therefore fails Condition 2: it does not block the back-door path; it creates one.

This phenomenon is called *M-bias* because the graph has an M-shaped skeleton. The practical lesson is that Condition 1 alone (“include all pretreatment variables”) is not sufficient; every candidate adjustment variable must also be verified not to open a previously closed back-door path.

Theorem: Back-Door Adjustment Formula

(Pearl 1993) If $\mathbf{S}$ satisfies the back-door criterion for the effect of T on Y in $\mathcal{G}$, and $P(T=t \mid \mathbf{S}=\mathbf{s}) > 0$ for all $\mathbf{s}$ with $P(\mathbf{S}=\mathbf{s}) > 0$, then:

$$P(y \mid \text{do}(T=t)) = \int P(y \mid T=t, \mathbf{S}=\mathbf{s}) dP(\mathbf{s}). \quad (3.1)$$

Equivalently, for any measurable function h :

$$\mathbb{E}[h(Y) \mid \text{do}(T=t)] = \mathbb{E}_{\mathbf{S}}[\mathbb{E}[h(Y) \mid T=t, \mathbf{S}]]. \quad (3.2)$$

Remark

The positivity condition requires every value of the treatment $T=t$ to be observable within each stratum of $\mathbf{S}$: if some stratum $\mathbf{s}$ never receives treatment t in the data, the conditional distribution $P(y | t, \mathbf{s})$ is undefined and Equation 5.3 cannot be computed. Positivity is a *data-support* condition, not a graphical one. When positivity fails on a set of measure zero, the formula still identifies $P(y | \text{do}(t))$ because the offending strata contribute zero mass to the outer integral.

3.2.3 The Adjustment Formula: Standardization

Equation 5.3 is also known as the *standardization formula* or, in the epidemiological literature, the *g-formula* (Robins 1986). The interpretation is important:

- $P(y | t, \mathbf{s})$ is the stratum-specific conditional distribution of Y given $T=t$ and $\mathbf{S}=\mathbf{s}$. Within each stratum, T is not confounded by any back-door path (all are blocked by $\mathbf{S}$), so this conditional distribution *is* causal.
- The outer integration $\int \dots dP(\mathbf{s})$ weights the strata by the *population marginal* distribution of $\mathbf{S}$, rather than the treatment-conditional distribution $dP(\mathbf{s} | t)$ that a naïve analysis would use.
- In a randomized experiment, $\mathbf{S} = \emptyset$ satisfies the back-door criterion (no confounding), and Equation 5.3 reduces to $P(y | \text{do}(t)) = P(y | t)$: the observational and interventional distributions coincide.

Example: Back-Door Adjustment — A Numerical Walkthrough

Use the single-confounder graph ($T \leftarrow C \rightarrow Y, T \rightarrow Y$) with all variables binary: C = disease severity, T = treatment, Y = recovery. Severity is equally prevalent, $P(C=1) = 0.5$, and sicker patients are preferentially treated.

C	$P(C)$	$P(T=1 C)$	$P(Y=1 T=0, C)$	$P(Y=1 T=1, C)$
0	0.50	0.20	0.10	0.40
1	0.50	0.80	0.50	0.80

The true ATE within each stratum is 0.30.

Naïve estimand (ignoring confounding): $P(Y=1 | T=1) = 0.72$, $P(Y=1 | T=0) = 0.18$. Naïve difference: **0.54**.

Back-door adjusted estimand:

$$P(Y=1 | \text{do}(T=1)) = 0.40 \times 0.5 + 0.80 \times 0.5 = 0.60,$$

$$P(Y=1 | \text{do}(T=0)) = 0.10 \times 0.5 + 0.50 \times 0.5 = 0.30.$$

Adjusted difference: **0.30** — matching the true ATE.

The naïve estimate (0.54) overstates the causal effect by 0.24 because treated patients are predominantly severe ($C=1$), who recover at higher rates *even without treatment*. The back-door formula removes this confounding by imposing the *same* covariate distribution $P(C)$ on both potential worlds.

3.2.4 Application to the Education–Earnings DAG

Example: Back-Door Identification in the Education Example

This is the Education–Earnings DAG analyzed in full in Chapter 2 (Section 5). The DAG has edges $N \rightarrow E$, $B \rightarrow E$, $B \rightarrow Y$, $E \rightarrow Y$, where N = neighborhood, B = family background, E = education, Y = earnings. We wish to identify $P(y | \text{do}(E=e))$.

The only back-door path is $E \leftarrow B \rightarrow Y$.

- Does $\{B\}$ satisfy the criterion? B is not a descendant of E ; $\{B\}$ blocks $E \leftarrow B \rightarrow Y$. By

Equation 5.3: $P(y \mid \text{do}(E=e)) = \sum_b P(y \mid e, b) P(b)$.

- Does $\{N\}$ satisfy the criterion? $\{N\}$ does *not* block $E \leftarrow B \rightarrow Y$ (the fork at B remains open). Fails condition 2.
- Does $\{N, B\}$ satisfy the criterion? Yes: neither N nor B is a descendant of E , and $\{N, B\}$ blocks the only back-door path.

Remark

The Education example illustrates three important points.

(a) Multiple valid sets may exist, and they identify the same quantity. Both $\{B\}$ and $\{N, B\}$ satisfy the back-door criterion here. Any such set yields the same interventional distribution; the identification target is unique even though the adjustment sets are not.

(b) Supersets of a valid adjustment set are not automatically valid. Adding a variable to a valid set can violate Condition 2 if that variable is a collider on some back-door path. In this graph $\{N, B\}$ happens to remain valid, but this must be verified from the graph, not assumed.

(c) Smaller valid sets are often preferable in practice. Adjusting for fewer variables — the minimal sufficient set — is generally more efficient and more robust, because every variable added to $\mathbf{S}$ must also satisfy positivity.

3.3 The Front-Door Criterion

The back-door criterion requires that all confounders be observed. The front-door criterion provides identification in a different configuration: when there exists a measured *mediator* M on the causal path from T to Y with specific independence properties.

3.3.1 The Configuration and Criterion

Definition: Front-Door Criterion

A set of variables $\mathbf{M}$ satisfies the *front-door criterion* for the effect of T on Y in $\mathcal{G}$ if:

1. $\mathbf{M}$ intercepts all directed paths from T to Y (all causal paths pass through $\mathbf{M}$).
2. There are no unblocked back-door paths from T to $\mathbf{M}$ (the effect of T on $\mathbf{M}$ is not confounded).
3. All back-door paths from $\mathbf{M}$ to Y are blocked by T (conditioning on T closes all confounding of the $\mathbf{M} \rightarrow Y$ effect).

Example: A Single Mediator with Unobserved Confounding

Consider the DAG $T \rightarrow M \rightarrow Y$ with $U \rightarrow T$ and $U \rightarrow Y$, where M is observed and U is unobserved. Does M satisfy the front-door criterion?

1. M intercepts all directed paths from T to Y . The only directed path is $T \rightarrow M \rightarrow Y$, which passes through M .
2. No unblocked back-door paths from T to M . There is no edge $U \rightarrow M$, so no back-door path exists.
3. All back-door paths from M to Y are blocked by T . The only back-door path is $M \leftarrow T \leftarrow U \rightarrow Y$. Conditioning on T blocks this path (chain at T).

All three conditions hold, so the front-door formula applies:

$$P(y \mid \text{do}(T=t)) = \sum_m P(m \mid t) \sum_{t'} P(y \mid t', m) P(t').$$

Example: When the Front-Door Criterion Fails

Each condition of the definition is load-bearing.

Case (a): Condition 2 fails — the $T \rightarrow M$ link is confounded. Add the edge $U \rightarrow M$ so the unobserved confounder also directly affects the mediator. Condition 1 still holds, but condition 2 now fails: the path $T \leftarrow U \rightarrow M$ is an unblocked back-door path from T to M . Stage 1 of the front-door formula uses $P(M=m \mid T=t)$ as if the $T \rightarrow M$ link were unconfounded. But with $U \rightarrow M$, the association between T and M is partly due to their shared cause U , not purely the causal effect $T \rightarrow M$. Stage 1 no longer estimates $P(m \mid \text{do}(T=t))$.

Case (b): Condition 3 fails — the $M \rightarrow Y$ link has an extra unobserved confounder. Add a second unobserved variable V with $V \rightarrow M$ and $V \rightarrow Y$. Conditions 1 and 2 still hold. But condition 3 now fails: the path $M \leftarrow V \rightarrow Y$ is a back-door path from M to Y that is *not* blocked by conditioning on T . Stage 2 uses $\sum_{t'} P(y \mid t', m) P(t')$ as the back-door-adjusted effect of M on Y , but T is no longer a valid adjustment set for Stage 2.

Modification	Condition violated	Stage broken
Add $U \rightarrow M$	Cond. 2: $T \rightarrow M$ confounded	Stage 1
Add $V \rightarrow M$ and $V \rightarrow Y$	Cond. 3: $M \rightarrow Y$ confounded by V	Stage 2

The prototypical graph where the front-door criterion applies is: $T \rightarrow M \rightarrow Y$, $U \rightarrow T$, $U \rightarrow Y$, no direct $T \rightarrow Y$ edge. Figure 8.2 shows this graph.




Figure 3.2: The prototypical front-door graph. The confounder U is unobserved, so back-door adjustment on U is infeasible. However, the mediator M is fully observed and satisfies the three front-door conditions, enabling identification via the front-door formula.

3.3.2 The Front-Door Formula

The front-door strategy splits the total effect of T on Y into two identifiable pieces: the effect of T on the mediator M , and the effect of M on Y . Front-door identification does not remove confounding directly; it *routes around it* through an observed causal pathway.

Theorem: Front-Door Adjustment Formula

If M satisfies the front-door criterion for the effect of T on Y (Pearl 1995), and $P(t) > 0$ for all t , then:

$$P(y \mid \text{do}(T=t)) = \sum_m \underbrace{P(m \mid T=t)}_{\text{Stage 1}} \underbrace{\sum_{t'} P(y \mid T=t', M=m) P(t')}_{\text{Stage 2}}. \quad (3.3)$$

Equivalently, writing $\mu(t', m) = \mathbb{E}[Y \mid T=t', M=m]$:

$$\mathbb{E}[Y \mid \text{do}(T=t)] = \mathbb{E}_{M|T=t} \left[\mathbb{E}_T [\mu(T, M)] \right]. \quad (3.4)$$

The formula chains two stages: **Stage 1** $P(m \mid T=t)$ is the distribution of the mediator given $T=t$. Condition 2 guarantees no confounding between T and M , so this is directly identifiable. **Stage 2** $\sum_{t'} P(y \mid T=t', M=m) P(t')$ is the back-door-adjusted causal effect of M on Y , with T as the adjustment variable. Condition 3 guarantees that conditioning on T blocks all back-door paths from M to Y .

Derivation Sketch

The formula is the composition of two back-door applications.

1. **Expand over the mediator.** By condition 1, every directed path from T to Y passes through M :

$$P(y \mid \text{do}(T=t)) = \sum_m P(y \mid \text{do}(M=m)) P(m \mid \text{do}(T=t)).$$

2. **Identify the distribution of M under intervention.** By condition 2 there are no back-door paths from T to M , so T satisfies the (trivial) back-door criterion for its effect on M with empty adjustment set. Hence $P(m \mid \text{do}(T=t)) = P(m \mid T=t)$.
3. **Identify the effect of M on Y .** By condition 3, T is a valid back-door adjustment set for the $M \rightarrow Y$ effect:

$$P(y \mid \text{do}(M=m)) = \sum_{t'} P(y \mid T=t', M=m) P(t').$$

4. **Substitute back** to obtain Equation 8.13. $\square$

3.3.3 Interpretation: Two-Stage Causal Accounting




Figure 3.3: Why the front-door criterion works. The total effect of T on Y is not directly identifiable because U confounds T and Y . However, the effect of T on the mediator M is unconfounded (Stage 1), and the effect of M on Y becomes identifiable after conditioning on T , which blocks the back-door path through U (Stage 2).

Example: Smoking, Tar, and Cancer

The classic front-door example (Pearl 1995): T = smoking, M = tar deposits in lungs, Y = lung cancer, U = genetic predisposition (unobserved). The graph is $T \rightarrow M \rightarrow Y$ with $U \rightarrow T$ and $U \rightarrow Y$. There is no direct $T \rightarrow Y$ path.

Conditions 1–3 are satisfied: all causal paths from T to Y pass through M ; no back-door path exists from T to M ; and the only back-door path from M to Y is $M \leftarrow T \leftarrow U \rightarrow Y$, which is blocked by conditioning on T .

The front-door formula gives a double sum over (m, t') that is entirely computable from the observed joint distribution of (T, M, Y) , without any data on U .

Example: Front-Door Formula — Binary Numerical Computation

Let $T, M, Y \in \{0, 1\}$ with U unobserved and the graph $T \rightarrow M \rightarrow Y$, $U \rightarrow T$, $U \rightarrow Y$.

T	M	Y	$P(T, M, Y)$
0	0	0	0.360
0	0	1	0.040
0	1	0	0.050

0	1	1	0.050
1	0	0	0.070
1	0	1	0.030
1	1	0	0.120
1	1	1	0.280

The inner sum (Stage 2):

$$S_0 = 0.1 \times 0.5 + 0.3 \times 0.5 = 0.20, \quad S_1 = 0.5 \times 0.5 + 0.7 \times 0.5 = 0.60.$$

The outer sum (Stage 1):

$$P(Y=1 \mid \text{do}(T=1)) = 0.2 \times 0.20 + 0.8 \times 0.60 = 0.52,$$

$$P(Y=1 \mid \text{do}(T=0)) = 0.8 \times 0.20 + 0.2 \times 0.60 = 0.28.$$

The average causal effect is $0.52 - 0.28 = 0.24$. The naïve estimate $P(Y=1 \mid T=1) - P(Y=1 \mid T=0) = 0.62 - 0.18 = 0.44$ is almost twice the true effect, because U simultaneously increases T and Y .

Comparing Back-Door and Front-Door

The two criteria are *not* nested. Neither implies the other.

- The back-door criterion requires all confounders to be observed (or blocked by observed variables), but places no restriction on mediators.
- The front-door criterion requires an observed mediator intercepting all causal paths, but the confounder U can be entirely unobserved.

There exist graphs where neither criterion alone applies, but where the full do-calculus (or the ID algorithm) can still identify the causal effect.

3.4 The Three Rules of Do-Calculus

Having seen two concrete identification strategies, we now develop the general algebraic machinery that underlies both. Pearl (1995) introduced the three rules of do-calculus as a system for manipulating expressions involving the do-operator: each rule licenses a transformation that removes or exchanges terms, bringing the expression one step closer to a purely observational formula.

How to Use a Do-Calculus Rule

When applying a do-calculus rule, proceed in three steps:

1. **Choose the relevant modified graph.** Form the graph obtained by deleting the appropriate incoming or outgoing arrows.
2. **Check the d-separation condition in that graph.** Work purely graphically: list the relevant paths and determine whether they are blocked.
3. **Apply the corresponding algebraic transformation.** Once the graphical condition holds, use the rule to remove an observation, insert an intervention, or exchange an observation for an intervention.

The do-calculus is therefore a graph modification $\rightarrow$ d-separation check $\rightarrow$ probability transformation procedure.

Remark: First Reading / Second Reading

On a first reading, focus on the qualitative role of the three rules: Rule 1 adds or removes observations, Rule 2 exchanges an action for an observation, and Rule 3 adds or removes actions. The detailed graph subscripts and formal statements can be absorbed gradually.

3.4.1 Statement of the Rules

Each rule's graphical condition uses a specific modified graph. The table below gives the recipe for each subscript that appears.

Subscript	Step 1: delete	Step 2: delete
$\mathcal{G}_{\overline{X}Z}$	all arrows <i>into</i> X	all arrows <i>out of</i> Z
$\mathcal{G}_{\overline{X}\overline{Z(W)}}$	all arrows <i>into</i> X	arrows <i>into</i> Z except those from W
$\mathcal{G}_{\overline{X}\overline{Z(W)}}$	all arrows <i>into</i> X	arrows <i>into</i> Z that are not from ancestors of W in $\mathcal{G}_{\overline{X}}$

Theorem: The Three Rules of Do-Calculus

(Pearl 1995) Let $\mathcal{G}$ be a DAG over variables $\mathcal{V}$, and let $X, Y, Z, W \subseteq \mathcal{V}$ be disjoint sets. The rules hold for any distribution P Markov with respect to $\mathcal{G}$.

Rule 1 (Insertion/Deletion of Observations). *Remove or insert an observation when it becomes irrelevant in the modified graph.*

$$P(y \mid \text{do}(x), z, w) = P(y \mid \text{do}(x), w) \quad \text{if} \quad (Y \perp\!\!\!\perp Z \mid X, W)_{\mathcal{G}_{\overline{X}Z}}.$$

Rule 2 (Action/Observation Exchange). *Replace an intervention by an observation, or vice versa, when the modified graph makes them equivalent.*

$$P(y \mid \text{do}(x), \text{do}(z), w) = P(y \mid \text{do}(x), z, w) \quad \text{if} \quad (Y \perp\!\!\!\perp Z \mid X, W)_{\mathcal{G}_{\overline{X}\overline{Z(W)}}}.$$

Rule 3 (Insertion/Deletion of Actions). *Remove or insert an intervention when it becomes irrelevant in the modified graph.*

$$P(y \mid \text{do}(x), \text{do}(z), w) = P(y \mid \text{do}(x), w) \quad \text{if} \quad (Y \perp\!\!\!\perp Z \mid X, W)_{\mathcal{G}_{\overline{X}\overline{Z(W)}}}.$$

Remark

Each rule is a *conditional independence statement in a specific intervention graph*. The three rules cover all ways of adding or removing terms from a do-expression: Rule 1 handles observations; Rule 2 exchanges actions for observations; Rule 3 deletes actions. Any sequence of such operations that converts a do-expression into a purely observational expression is an *identification proof*.

Remark: Soundness of the Rules

The three rules are *sound*: each holds for every distribution P Markov with respect to $\mathcal{G}$. For Rule 2: the post-intervention distribution $P(\cdot \mid \text{do}(x))$ is Markov with respect to $\mathcal{G}_{\overline{X}}$. Further intervening on Z gives a distribution Markov with respect to $\mathcal{G}_{\overline{X}\overline{Z}}$. The d-separation condition $(Y \perp\!\!\!\perp Z \mid X, W)_{\mathcal{G}_{\overline{X}\overline{Z(W)}}}$ implies $Y \perp\!\!\!\perp Z \mid W$ in the post- $\text{do}(x, z)$ distribution. Since the distribution of Y is the same whether we *intervene* on Z or *observe* Z , the two expressions are equal. Full proofs appear in Pearl (1995).

3.4.2 Intuition for Each Rule

Rule 1 — Adding/Removing Observations. The condition says: in the post- $\text{do}(x)$ world, if we further deactivate Z 's causal influence (delete its outgoing arrows), Y and Z become independent given W . This means Z carries no information about Y that is not already in W , so it can be freely inserted or removed.

Example: Rule 1 — Deleting a Redundant Observation

Graph: $Z \rightarrow T \rightarrow Y$, no other edges. Suppose we intervene on T .

In $\mathcal{G}_{\overline{T}Z}$: delete arrows into T (removes $Z \rightarrow T$), then delete arrows out of Z (nothing remains).

Node Z is isolated, so $(Y \perp\!\!\!\perp Z \mid T)$ holds trivially. Rule 1 gives:

$$P(y \mid \text{do}(t), z) = P(y \mid \text{do}(t)).$$

The intervention $\text{do}(t)$ severs Z 's only route to Y . Knowing the former cause Z of T is irrelevant once T has been set externally.

Rule 2 — Action/Observation Exchange. When the condition holds, intervening on Z is equivalent to conditioning on Z (within the scope of $\text{do}(x)$ and W).

Example: Rule 2 — Exchanging Action for Observation

Graph: $X \rightarrow Z \rightarrow Y$, no other edges. Suppose we intervene on X and wish to show $P(y \mid \text{do}(x), \text{do}(z)) = P(y \mid \text{do}(x), z)$.

In $\mathcal{G}_{\overline{X}\overline{Z}}$: delete arrows into X (removes the incoming arrows of X) and arrows into Z (removes $X \rightarrow Z$). Node Z has no incoming arrows, so no back-door paths from Z to Y can arise. The only path is the causal $Z \rightarrow Y$, and $(Y \perp\!\!\!\perp Z)_{\mathcal{G}_{\overline{X}\overline{Z}}}$ holds in the sense that no spurious association arises. Rule 2 gives:

$$P(y \mid \text{do}(x), \text{do}(z)) = P(y \mid \text{do}(x), z).$$

Y depends on Z only through Z 's value, not through the mechanism that generated Z . Externally fixing Z and merely observing $Z=z$ are therefore equivalent for Y .

Rule 3 — Deleting an Intervention. The condition says that Z has no causal path to Y that is not blocked by X or W in the appropriate intervention graph. Setting Z to any value has no effect on Y , so $\text{do}(z)$ can be dropped.

Example: Rule 3 — Deleting an Irrelevant Intervention

Graph: $Z \rightarrow X \rightarrow Y$. Suppose we intervene on both X and Z .

In $\mathcal{G}_{\overline{X}\overline{Z}}$: delete arrows into X (removes $Z \rightarrow X$), then delete arrows into Z (none exist). Node Z is now isolated. Since there is no path from Z to Y , $(Y \perp\!\!\!\perp Z)_{\mathcal{G}_{\overline{X}\overline{Z}}}$ holds trivially. Rule 3 gives:

$$P(y \mid \text{do}(x), \text{do}(z)) = P(y \mid \text{do}(x)).$$

Once X is fixed externally, Z 's only route to Y (through X) is severed. Setting Z to any value cannot change Y , so $\text{do}(z)$ is superfluous.

Example: A Two-Rule Derivation — the Proxy Variable Graph

T has a direct effect on Y . An unobserved variable U is a common cause of T and an observed proxy S ; S in turn causes Y . U does *not* directly cause Y . Edges: $T \rightarrow Y$, $U \rightarrow T$, $U \rightarrow S$, $S \rightarrow Y$.

Goal: simplify $P(y \mid \text{do}(t))$ to a purely observational expression.

Step 1. Introduce S by the law of total probability:

$$P(y \mid \text{do}(t)) = \sum_s P(y \mid \text{do}(t), s) P(s \mid \text{do}(t)).$$

Step 2. Rule 3 applied to $P(s \mid \text{do}(t))$. Form $\mathcal{G}_{\overline{T}}$ by deleting the arrow $U \rightarrow T$. In this graph, T has no parents and no path connects T to S or to the U - S component. Hence $(S \perp\!\!\!\perp T)_{\mathcal{G}_{\overline{T}}}$, and Rule 3 gives $P(s \mid \text{do}(t)) = P(s)$.

Step 3. Back-door d-separation argument applied to $P(y \mid \text{do}(t), s)$. Form $\mathcal{G}_{\overline{T}}$ by deleting the arrow $T \rightarrow Y$. The only back-door path $T \leftarrow U \rightarrow S \rightarrow Y$ is blocked by conditioning on S . Hence $\{S\}$ satisfies the back-door criterion, so $P(y \mid \text{do}(t), s) = P(y \mid t, s)$.

Combining:

$$P(y \mid \text{do}(t)) = \sum_s P(y \mid t, s) P(s).$$

The causal effect is identified by standardizing over the population marginal $P(s)$, even though the confounder U is entirely unobserved.

3.4.3 Proofs of the Identification Formulas

Proof: Back-Door Formula via Rule 3 and the Back-Door d-Separation Argument

Graph structure: $T \rightarrow Y$, unobserved $U \rightarrow T$, $U \rightarrow Y$, and observed $\mathbf{S}$ blocking the back-door path $T \leftarrow U \rightarrow Y$.

Step 1: Introduce $\mathbf{S}$ by the law of total probability.

$$P(y \mid \text{do}(t)) = \int P(y \mid \text{do}(t), \mathbf{s}) dP(\mathbf{s} \mid \text{do}(t)).$$

Step 2: Simplify $P(\mathbf{s} \mid \text{do}(t))$. Because $\mathbf{S}$ contains no descendants of T (condition 1), the graphical condition for Rule 3 is $(\mathbf{S} \perp\!\!\!\perp T)_{\mathcal{G}_{\overline{T}}}$: in $\mathcal{G}_{\overline{T}}$, all arrows into T are deleted so no directed path from T reaches any node in $\mathbf{S}$. Therefore $(\mathbf{S} \perp\!\!\!\perp T)_{\mathcal{G}_{\overline{T}}}$ holds, and Rule 3 gives: $P(\mathbf{s} \mid \text{do}(t)) = P(\mathbf{s})$.

Step 3: Convert $P(y \mid \text{do}(t), \mathbf{s})$ to an observation. Condition 2 ensures that, within each stratum $\mathbf{S} = \mathbf{s}$, all back-door paths from T to Y are blocked. By the truncated factorization formula for $P(\cdot \mid \text{do}(t))$, the post-intervention conditional equals the observational conditional: $P(y \mid \text{do}(t), \mathbf{s}) = P(y \mid t, \mathbf{s})$.

Combining:

$$P(y \mid \text{do}(t)) = \int P(y \mid t, \mathbf{s}) dP(\mathbf{s}). \quad \square$$

Proof: Front-Door Formula via Three Applications of the Rules

Graph structure: $T \rightarrow M \rightarrow Y$, unobserved $U \rightarrow T$, $U \rightarrow Y$, no direct $T \rightarrow Y$ edge.

Step 1: Introduce M by the law of total probability.

$$P(y \mid \text{do}(t)) = \int P(y \mid \text{do}(t), m) dP(m \mid \text{do}(t)).$$

Step 2: Simplify $P(m \mid \text{do}(t))$. In $\mathcal{G}_{\overline{T}}$, there are no back-door paths from T to M by front-door condition 2. Rule 2 applies: $P(m \mid \text{do}(t)) = P(m \mid t)$.

Step 3a: $P(y \mid \text{do}(t), m) = P(y \mid \text{do}(m))$. In $\mathcal{G}_{\overline{T}}$, the arrow $U \rightarrow T$ is deleted, so U and T become independent. Since there is no arrow from U into M (condition 2), the conditional of Y given $M=m$ integrates over U freely: $P(y \mid \text{do}(t), m) = \int P(y \mid m, u) dP(u) = P(y \mid \text{do}(m))$.

Step 3b: By front-door condition 3, $\{T\}$ satisfies the back-door criterion for the effect of M on Y , so $P(y \mid \text{do}(m)) = \int P(y \mid m, t') dP(t')$.

Assembling:

$$P(y \mid \text{do}(t)) = \int \left[\int P(y \mid t', m) dP(t') \right] dP(m \mid t). \quad \square$$

3.5 The Do-Calculus is Complete

A natural question after seeing the three rules is: are they enough? Could there be a graph where the causal effect is identifiable in principle, but none of the rules — applied in any sequence — can derive a purely observational expression? Shpitser and Pearl (2006) answered this definitively.

Throughout Sections 1–4 we have represented latent common causes as explicit unobserved nodes. The completeness theorem uses a more general graph class that encodes latent common causes compactly as *bidirected edges*: a bidirected edge $X \leftrightarrow Y$ stands for an unobserved variable U with $U \rightarrow X$ and $U \rightarrow Y$. A DAG augmented with such bidirected edges is called a *semi-Markovian DAG*.

Theorem: Completeness of the Do-Calculus

(Shpitser and Pearl 2006) Let $\mathcal{G}$ be a semi-Markovian DAG. A causal quantity $P(y \mid \text{do}(t))$ is identifiable from P relative to $\mathcal{G}$ if and only if the do-calculus can derive a purely observational expression for it.

Example: The Bow Graph — The Simplest Non-Identifiable Structure

The *bow graph* is the semi-Markovian DAG with one directed edge $T \rightarrow Y$ and one bidirected edge $T \leftrightarrow Y$ (representing an unobserved U with $U \rightarrow T$ and $U \rightarrow Y$). The single back-door path is $T \leftarrow U \rightarrow Y$; no observed set can block it, so the back-door criterion cannot be satisfied. No observed mediator exists, so the front-door criterion also fails.

We show directly that $P(y \mid \text{do}(t))$ is not determined by $P(T, Y)$ alone, by constructing two models that agree on $P(T, Y)$ but disagree on $P(y \mid \text{do}(t))$.

Let $T, Y, U \in \{0, 1\}$ with $U \sim \text{Bernoulli}(1/2)$.

Model $\mathcal{M}_1$ (pure confounding; no causal effect of T): $T = U, Y = U$. Both T and Y are driven entirely by U . The observed joint distribution is $P(T=0, Y=0) = P(T=1, Y=1) = 1/2$ and $P(T \neq Y) = 0$. Under $\text{do}(T=1)$, $Y = U$ is unchanged, so $P_1(Y=1 \mid \text{do}(T=1)) = P(U=1) = 1/2$.

Model $\mathcal{M}_2$ (full causal effect; U acts on Y only through T): $T = U, Y = T$. Since $T = U$, we again have $Y = T = U$, so the observed joint distribution is *identical* to $\mathcal{M}_1$. Under $\text{do}(T=1)$, $Y = T = 1$ deterministically, so $P_2(Y=1 \mid \text{do}(T=1)) = 1$.

Conclusion. The two models produce identical observed distributions $P(T, Y)$, yet $P(Y=1 \mid \text{do}(T=1))$ equals $1/2$ in $\mathcal{M}_1$ versus 1 in $\mathcal{M}_2$. Therefore $P(y \mid \text{do}(t))$ is *not identified* from $P(T, Y)$ in the bow graph.

Proof Sketch and the ID Algorithm

The *ID algorithm* (Shpitser and Pearl 2006) takes as input a semi-Markovian DAG $\mathcal{G}$, a target intervention distribution $P(y \mid \text{do}(t))$, and the observed joint distribution P , and either returns an explicit observational formula or reports non-identifiability.

Sufficiency (do-calculus identifies whenever identification is possible): proved constructively via the ID algorithm. The key structural concept is the *c-component* (confounded component): a maximal set of nodes connected by bidirected edges. The ID algorithm proceeds recursively: (1) simplification — remove from $\mathcal{G}$ any node that is not an ancestor of Y in $\mathcal{G}_{\overline{T}}$; (2) decomposition by c-components — factorise $P(\mathbf{v} \setminus \mathbf{t} \mid \text{do}(t))$ over the c-components; (3) recursive identification — attempt to identify each c-component factor; (4) failure — if any subproblem cannot be reduced, the algorithm returns FAIL.

Necessity (do-calculus cannot identify non-identifiable quantities): proved by showing that whenever the ID algorithm fails, the causal effect is genuinely non-identifiable. The key graphical obstruction is the *hedge*: a pair of subgraphs (F, F') where F contains a back-door path from T to Y , and F' is a subgraph of $\mathcal{G}_{\overline{T}}$ in which Y remains connected to a latent common cause. The existence of a hedge allows explicit construction of two observationally equivalent models with different causal effects.

Combining: the do-calculus succeeds if and only if no hedge exists, which is if and only if $P(y \mid \text{do}(t))$ is identifiable. $\square$

Remark

The completeness theorem has a crucial practical implication: *if the ID algorithm reports non-identification, then no estimation method — however clever — can recover the causal effect from observational data alone*, given the assumed graph structure. Non-identification is not a limitation of a particular technique; it is a fundamental property of the causal model. Three remedies exist: impose additional structural assumptions (e.g., linearity or monotonicity) that go beyond the graph; collect additional data that break the obstruction (e.g., an instrument or a randomized experiment); or target a different, identified causal quantity instead.

3.6 The Big Picture

This chapter completes the *theoretical* identification machinery of the course. The chain of reasoning begun in Chapter 1 (the interventional/observational distinction) and formalized in Chapter 2 (d-separation) now terminates in a concrete identification formula. Chapters 4–7 then apply this machinery to specific research designs.




Figure 3.4: The six-step pipeline from causal assumptions to estimation.

The completeness theorem closes the loop: every causal effect that is identifiable from the observational distribution relative to the assumed graph can be derived by the do-calculus, and every effect that is not identifiable in that setting cannot be recovered by any observational method without additional assumptions or new data.

3.7 Summary

1. Two graph operations underpin the do-calculus: $\mathcal{G}_{\overline{X}}$ deletes arrows *into* X and models the intervention $\text{do}(x)$; $\mathcal{G}_{\underline{X}}$ deletes arrows *out of* X and is used as a technical device in graphical conditions for certain do-calculus steps — it does not model conditioning on X .
2. The back-door criterion identifies $P(y \mid \text{do}(t))$ when an observed set $\mathbf{S}$ blocks all back-door paths and contains no descendant of T . The general formula is $\int P(y \mid t, \mathbf{s}) dP(\mathbf{s})$, the standardization formula (the g-formula for $K = 1$).
3. The front-door criterion identifies $P(y \mid \text{do}(t))$ via an observed mediator M intercepting all causal paths from T to Y , even when the confounder U is entirely unobserved, provided the three front-door conditions hold.

4. The do-calculus has three rules, each licensed by a d-separation condition in an appropriate intervention graph. Rule 1 adds/removes observations; Rule 2 exchanges actions for observations; Rule 3 deletes actions. The back-door and front-door formulas are each derivable as three-step sequences of these rules.
5. **Completeness** (Shpitser and Pearl 2006): the do-calculus is complete for identification from observational data relative to the assumed graph. A causal effect is identifiable if and only if the do-calculus can derive a purely observational expression for it. If the do-calculus cannot identify the quantity in that setting, then no purely observational method can do so without additional assumptions or new data.

3.8 Problems

1. Warm-up: graph operations. Let $\mathcal{G}$ be the IV DAG with edges $Z \rightarrow T, T \rightarrow Y, U \rightarrow T, U \rightarrow Y$ (U unobserved).

- (a) Draw $\mathcal{G}_{\overline{T}}$ and identify all paths from Z to Y that remain open.
- (b) Draw $\mathcal{G}_{\underline{T}}$ and identify all paths from Z to Y that remain. For each path, classify it as a causal chain, fork, or collider path, and state whether it is open or closed by default.
- (c) In $\mathcal{G}_{\overline{T}}$, is $Z \perp\!\!\!\perp Y \mid \emptyset$? How does the absence of a direct edge $Z \rightarrow Y$ in the DAG relate to the exclusion restriction?

2. Back-door practice. Consider the DAG: $X \rightarrow T, X \rightarrow Y, T \rightarrow M, M \rightarrow Y, T \rightarrow Y$, where X is observed.

- (a) List all back-door paths from T to Y .
- (b) Does $\{X\}$ satisfy the back-door criterion for the effect of T on Y ? Write the resulting adjustment formula.
- (c) Does $\{M\}$ satisfy the back-door criterion? Explain why or why not.
- (d) Does $\{X, M\}$ satisfy the back-door criterion? Identify which condition of the criterion M violates.
- (e) Even if the formal criterion issue in (d) were somehow set aside, explain the substantive bias that arises from conditioning on a mediator M that lies on a causal path $T \rightarrow M \rightarrow Y$.

3. Front-door identification. Consider the DAG: $T \rightarrow M \rightarrow Y$, with $U \rightarrow T$ and $U \rightarrow Y$ (unobserved U , no direct $T \rightarrow Y$ edge).

- (a) Verify that M satisfies all three front-door conditions.
- (b) Derive the front-door formula Equation 8.13 step by step, citing the do-calculus rule used at each step.
- (c) Now add a direct edge $T \rightarrow Y$ to the graph. Does M still satisfy the front-door criterion? Explain which condition fails.

4. Identification or non-identification? For each of the following DAGs, determine whether $P(y \mid \text{do}(t))$ is identified non-parametrically from observational data. If identified, state the formula; if not, explain the structural obstruction.

- (a) $T \rightarrow Y, U \rightarrow T, U \rightarrow Y, U$ unobserved.
- (b) Same as (a), with instrument $Z \rightarrow T$ and $Z \perp\!\!\!\perp U$ (no direct $Z \rightarrow Y$ edge). Show that $P(y \mid \text{do}(t))$ is **not** identified non-parametrically from $P(Z, T, Y)$ by constructing two binary SEMs $\mathcal{M}_1$ and $\mathcal{M}_2$, each compatible with the IV graph, such that (i) $P_{\mathcal{M}_1}(Z, T, Y) = P_{\mathcal{M}_2}(Z, T, Y)$ and (ii) $P_{\mathcal{M}_1}(Y=1 \mid \text{do}(T=1)) \neq P_{\mathcal{M}_2}(Y=1 \mid \text{do}(T=1))$. (*Hint*: let $Z, U \sim \text{Bern}(1/2)$ independently and set $T = U$ in both models. Then set $Y = U$ in $\mathcal{M}_1$ and $Y = T$ in $\mathcal{M}_2$.)
- (c) $T \rightarrow M \rightarrow Y, U \rightarrow T, U \rightarrow Y, U \rightarrow M, U$ unobserved.
- (d) $T \rightarrow Y, X \rightarrow T, X \rightarrow Y, X$ observed.

5. Back-door formula: proof and uniqueness.

- (a) Let $\mathcal{M}$ be a structural causal model Markov with respect to $\mathcal{G}$, and let $\mathbf{S}$ satisfy the back-door criterion for (T, Y) in $\mathcal{G}$ with positivity. Prove that $P(y \mid \text{do}(t)) = \int P(y \mid t, \mathbf{s}) dP(\mathbf{s})$. (*Hint*: begin with $P(y \mid \text{do}(t)) = \int P(y \mid \text{do}(t), \mathbf{s}) P(\mathbf{s} \mid \text{do}(t)) d\mathbf{s}$. Apply the Rule 3 graphical condition $(\mathbf{S} \perp\!\!\!\perp T)_{\mathcal{G}_{\overline{T}}}$, using back-door condition 1, to simplify $P(\mathbf{s} \mid \text{do}(t))$. Then invoke the d-separation $(Y \perp\!\!\!\perp T \mid \mathbf{S})_{\mathcal{G}_{\underline{T}}}$, which back-door condition 2 implies, together with the truncated factorization, to argue $P(y \mid \text{do}(t), \mathbf{s}) = P(y \mid t, \mathbf{s})$.)

- (b) Suppose $\mathbf{S}_1$ and $\mathbf{S}_2$ both satisfy the back-door criterion for (T, Y) in $\mathcal{G}$ (with positivity). Deduce from part (a) that $\int P(y | t, \mathbf{s}_1) dP(\mathbf{s}_1) = \int P(y | t, \mathbf{s}_2) dP(\mathbf{s}_2)$. Interpret this result: the identification target $P(y | \text{do}(t))$ is unique even when the adjustment set is not.

6. Rule sequencing on an unseen graph. Consider the graph: $C_1 \rightarrow T, C_1 \rightarrow Y, U \rightarrow T, U \rightarrow C_2, C_2 \rightarrow Y, T \rightarrow Y$, with C_1 and C_2 observed and U unobserved.

Derive $P(y | \text{do}(t)) = \sum_{c_1, c_2} P(y | t, c_1, c_2) P(c_1, c_2)$ from first principles using the three rules. For each step: (i) name the rule applied, (ii) form the relevant intervention graph by stating which arrows are deleted and what edges remain, and (iii) verify the required d-separation condition by tracing all paths and confirming each is blocked.

7. Non-identifiability by construction. Consider the graph from the front-door failure example: $T \rightarrow M \rightarrow Y, U \rightarrow T, U \rightarrow M, U \rightarrow Y$ (front-door condition 2 violated). Let $T, M, Y, U \in \{0, 1\}$ with $U \sim \text{Bern}(1/2)$.

- (a) Construct two SEMs $\mathcal{M}_1$ and $\mathcal{M}_2$, each compatible with the graph, such that $P_{\mathcal{M}_1}(T, M, Y) = P_{\mathcal{M}_2}(T, M, Y)$ but $P_{\mathcal{M}_1}(Y=1 | \text{do}(T=1)) \neq P_{\mathcal{M}_2}(Y=1 | \text{do}(T=1))$. (*Hint*: set $T = U$ in both models. In $\mathcal{M}_1$ set $M = U$ and $Y = U$; in $\mathcal{M}_2$ set $M = T$ and $Y = M$.)
- (b) Explain precisely which identification strategy fails in this graph and why, referencing (i) the back-door criterion, (ii) the front-door criterion (which condition fails and which stage collapses), and (iii) the completeness theorem (what does the existence of your two models imply?).

Part II

Identification

Chapter 4

Potential Outcomes and Adjustment

Learning Objectives

By the end of this chapter, students should be able to:


1. Define the potential outcome $Y(t)$, state SUTVA precisely, and derive the consistency equation $Y = Y(T)$.
2. Define the ATE and ATT, explain how they relate to each other, and re-express them as functionals of the structural equation for Y .
3. Show that the linear SEM structural coefficient β equals the ATE under homogeneous effects, and explain why ATE and ATT diverge under treatment effect heterogeneity.
4. State the ignorability assumption, interpret it graphically as a back-door condition, and explain why overlap is a necessary companion assumption.
5. Derive the adjustment formula for the ATE under strong ignorability and positivity, and explain how the back-door criterion provides the graphical justification for conditional exchangeability.

4.1 Motivation: A Third Language for Causality

The first three chapters developed causal inference primarily in the languages of structural equations and directed acyclic graphs. In this chapter we introduce a third language: the potential outcomes framework of Neyman (1923) and Rubin (1974).

Potential outcomes provide a direct language for defining causal estimands. Quantities such as the average treatment effect, the average treatment effect on the treated, and related causal contrasts are most naturally written in terms of counterfactual outcomes $Y(1)$, $Y(0)$, and more generally $Y(t)$. The two frameworks should be viewed as complementary rather than competing: potential outcomes define the causal quantities of interest, while DAGs and the do-calculus clarify identification.

The purpose of this chapter is therefore not to replace the graphical viewpoint developed earlier, but to connect it to the potential-outcomes notation that dominates much of the applied literature. We first define potential outcomes and the standard causal estimands built from them. We then show how ignorability and positivity together identify the average treatment effect, and how the back-door criterion provides the graphical justification for that identification. For readers who want a formal graphical representation of counterfactual variables, Appendix B introduces single world intervention graphs (SWIGs).



4.2 The Neyman–Rubin Potential Outcomes Framework

4.2.1 The Potential Outcome

Definition: Potential Outcome

(Neyman 1923; Rubin 1974) For each unit i and each possible treatment value $t \in \mathcal{T}$, the *potential outcome* $Y_i(t)$ is the value of the outcome that unit i *would have exhibited* had its treatment been set to t , possibly contrary to fact.

The collection $\{Y_i(t) : t \in \mathcal{T}\}$ is called the *schedule of potential outcomes* for unit i . In the binary case $\mathcal{T} = \{0, 1\}$, the schedule is $(Y_i(0), Y_i(1))$.

$Y_i(1)$ is the outcome unit i would achieve *under treatment*; $Y_i(0)$ is the outcome under *control*. Only one of these is ever observed for any given unit. The unobserved potential outcome is called the *counterfactual*.

The Fundamental Problem of Causal Inference

(Holland 1986) For any unit i , at most one potential outcome $Y_i(t)$ is observed. The individual-level causal effect $Y_i(1) - Y_i(0)$ is therefore *never directly observable*. Causal inference is an inference problem precisely because this quantity must be recovered from a population of units rather than from a single unit.

4.2.2 SUTVA and Consistency

Definition: SUTVA

(Rubin 1980) The *stable unit treatment value assumption* has two components:

1. **No interference.** The potential outcome $Y_i(t)$ depends only on unit i 's own treatment, not on the treatments of other units: $Y_i(t_1, \dots, t_n) = Y_i(t_i)$.
2. **No hidden versions.** For each treatment level t , there is a single well-defined version.

The consistency equation. Under SUTVA, the observed outcome Y_i equals the potential outcome at the treatment actually received:

$$Y_i = Y_i(T_i). \quad (4.1)$$

This equation is the bridge between the potential outcome world and the observed data world. It fails whenever SUTVA fails: if treatment spills over between units (e.g., vaccination provides herd immunity), or if the treatment label $T = 1$ covers multiple distinct interventions.

4.2.3 Connection to the Do-Operator

The notation $Y(t)$ and the expression $P(Y \mid \text{do}(T=t))$ refer to the same causal idea viewed from two complementary angles. Under the structural causal framework adopted in these notes, together with consistency and no interference, these two descriptions coincide.

Proposition: Potential Outcomes and the Do-Operator

Under the structural causal model of Chapters 1–3, together with consistency and no interference (SUTVA),

$$Y(t) \stackrel{d}{=} Y \mid \text{do}(T=t), \quad (4.2)$$

so that $P(Y(t) \leq y) = P(Y \leq y \mid \text{do}(T=t))$ for every y . Equivalently, $\mathbb{E}[Y(t)] = \mathbb{E}[Y \mid \text{do}(T=t)]$.

Proof Sketch

In the mutilated SEM $\mathcal{G}_{\bar{T}}$, the structural equation for T is replaced by $T := t$ for every unit. The structural model then determines the value of Y for unit i from t and unit i 's own background variables — which is precisely what the potential outcomes framework defines as $Y_i(t)$. SUTVA's no-interference condition ensures $Y_i(t)$ does not depend on other units' treatment values, so the marginal distribution of Y in $\mathcal{G}_{\bar{T}}$ equals the marginal distribution of $Y(t)$. By definition of the do-operator, the former is $f(y \mid \text{do}(T=t))$. $\square$

This equivalence lets us move freely between two notational traditions. When defining estimands such as $\mathbb{E}[Y(1) - Y(0)]$, potential-outcome notation is often most natural. When proving identification results from a graph, the do-operator is often more transparent: $f(y \mid \text{do}(t)) \neq f(y \mid T=t)$ whenever T is endogenous, whereas $Y(t)$ carries no such syntactic marker.

4.3 Causal Estimands

4.3.1 The Average Treatment Effect and Its Relatives

Definition: ATE and ATT

For a binary treatment $T \in \{0, 1\}$:

- The *average treatment effect* (ATE) is $\tau_{\text{ATE}} = \mathbb{E}[Y(1) - Y(0)]$.
- The *average treatment effect on the treated* (ATT) is $\tau_{\text{ATT}} = \mathbb{E}[Y(1) - Y(0) \mid T=1]$.

The ATE and ATT coincide only when the treatment effect does not depend on who selected into treatment. The naïve estimator $\hat{\tau}_{\text{naive}} = \bar{Y}_{T=1} - \bar{Y}_{T=0}$ estimates $\mathbb{E}[Y \mid T=1] - \mathbb{E}[Y \mid T=0] \neq \mathbb{E}[Y(1)] - \mathbb{E}[Y(0)]$, with the gap being the selection bias induced by endogeneity.

Remark

ATE and ATT coincide when $\mathbb{E}[Y(t) \mid T] = \mathbb{E}[Y(t)]$ for $t \in \{0, 1\}$, i.e., when treatment assignment is mean-independent of potential outcomes. In a randomized experiment, this holds by design.

4.3.2 Causal Estimands as Functionals of the Structural Equation

From structural equation to potential outcome. In the SEM, the outcome is determined by a structural equation $Y = g(T, \mathbf{X}, U_Y)$. The potential outcome under $\text{do}(T=t)$ is:

$$Y_i(t) = g(t, \mathbf{X}_i, U_{Y,i}). \quad (4.3)$$

ATE and ATT as structural parameters.

$$\tau_{\text{ATE}} = \mathbb{E}[g(1, \mathbf{X}, U_Y) - g(0, \mathbf{X}, U_Y)], \quad (4.4)$$

$$\tau_{\text{ATT}} = \mathbb{E}[g(1, \mathbf{X}, U_Y) - g(0, \mathbf{X}, U_Y) \mid T=1]. \quad (4.5)$$

The linear SEM as a special case. In the Gaussian linear SEM $Y = \beta T + \gamma^\top \mathbf{X} + \varepsilon$, the unit-level effect is $Y_i(1) - Y_i(0) = \beta$ for every unit. Consequently, $\tau_{\text{ATE}} = \tau_{\text{ATT}} = \beta$.

Heterogeneous effects. In the nonparametric SEM Equation 4.3, the unit-level effect varies across units. The ATE and ATT differ whenever treatment selection correlates with the individual effect size.

Remark

The SEM representation Equation 4.3 clarifies why the potential outcomes framework and the do-calculus are *equivalent in content but different in emphasis*. The do-calculus works with the

interventional distribution $f(y \mid \text{do}(T=t))$ as a population-level object. The potential outcomes framework works with unit-level quantities $Y_i(t)$ and asks what population summaries (ATE, ATT) are scientifically meaningful. The SEM ties the two together.

4.4 Ignorability, Positivity, and Adjustment

4.4.1 Strong Ignorability

The central identifying assumption in observational studies is that treatment assignment is as good as random after conditioning on observed covariates X .

Definition: Strong Ignorability

(Rosenbaum and Rubin 1983) The treatment assignment T is *strongly ignorable* given X if:

1. **Unconfoundedness:** $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$.
2. **Overlap (positivity):** $0 < P(T=1 \mid X=x) < 1$ for all x in the support of X .

4.4.2 Adjustment Formula under Ignorability

Under strong ignorability, the ATE is identified by the standardization formula:

$$\tau_{\text{ATE}} = \mathbb{E}_X[\mathbb{E}[Y \mid T=1, X] - \mathbb{E}[Y \mid T=0, X]]. \quad (4.6)$$

This is the back-door adjustment formula of Chapter 3 written in potential-outcome notation.

Derivation

For each treatment level $t \in \{0, 1\}$, by the law of iterated expectations:

$$\mathbb{E}[Y(t)] = \mathbb{E}_X[\mathbb{E}[Y(t) \mid X]].$$

Under unconfoundedness, conditioning on X renders the potential outcome mean-independent of treatment assignment:

$$\mathbb{E}[Y(t) \mid X] = \mathbb{E}[Y(t) \mid T=t, X].$$

By consistency, $\mathbb{E}[Y(t) \mid T=t, X] = \mathbb{E}[Y \mid T=t, X]$. Therefore:

$$\mathbb{E}[Y(t)] = \mathbb{E}_X[\mathbb{E}[Y \mid T=t, X]].$$

Applying this once for $t = 1$ and once for $t = 0$, and taking the difference, yields Equation 4.6. Overlap ensures that the conditional means on the right-hand side are well-defined for all covariate values. $\square$

Example: A Two-Stratum Adjustment Calculation

Suppose $X \in \{0, 1\}$ with $P(X=1) = 0.4$ and $P(X=0) = 0.6$, and the observed conditional means are $\mathbb{E}[Y \mid T=1, X=1] = 8$, $\mathbb{E}[Y \mid T=0, X=1] = 6$, $\mathbb{E}[Y \mid T=1, X=0] = 5$, $\mathbb{E}[Y \mid T=0, X=0] = 4$.

Under ignorability and positivity, Equation 4.6 gives:

$$\mathbb{E}[Y(1)] = 0.4 \times 8 + 0.6 \times 5 = 6.2, \quad \mathbb{E}[Y(0)] = 0.4 \times 6 + 0.6 \times 4 = 4.8.$$

Hence $\tau_{\text{ATE}} = 6.2 - 4.8 = 1.4$.

4.4.3 Back-Door Interpretation

The potential-outcomes assumption $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$ is the counterfactual expression of the idea that, after conditioning on X , treatment assignment carries no residual information about the outcome.

In the language of DAGs, the closely related idea is that X blocks all back-door paths from T to Y .

Proposition: Back-Door Criterion Implies Ignorability

Under the Markov assumption for $\mathcal{G}$, if X satisfies the back-door criterion for the effect of T on Y — that is, (1) no node in X is a descendant of T , and (2) X blocks every back-door path from T to Y — then unconfoundedness ($Y(t) \perp\!\!\!\perp T \mid X$) holds for all t .

Proof Sketch

In the mutilated graph $\mathcal{G}_{\overline{T}}$, all arrows into T are deleted, severing every back-door path from T to Y . The back-door criterion guarantees that X blocks all remaining non-causal paths between T and Y in $\mathcal{G}_{\overline{T}}$. By the Markov property for the interventional model represented by $\mathcal{G}_{\overline{T}}$, this yields conditional independence between treatment and outcome under intervention. Under the structural causal interpretation adopted in these notes, that interventional independence corresponds to $Y(t) \perp\!\!\!\perp T \mid X$ for all t . A fully explicit graphical version of this translation is given in Appendix B using SWIGs. $\square$

Remark

This is the direction most important for practice: a graphical adjustment set justifies the counterfactual independence needed for standardization, regression adjustment, and propensity-score methods, requiring only the Markov property. The converse — whether every ignorability statement corresponds to a back-door condition — additionally requires faithfulness and is established in Appendix B via the SWIG construction.

Example: Labor Training Program

(LaLonde 1986; Dehejia and Wahba 1999) Let T = receipt of job training, Y = earnings two years later, X = (age, education, prior earnings, race, marital status).

The back-door path $T \leftarrow X \rightarrow Y$ is blocked by conditioning on X , leaving only the causal path $T \rightarrow Y$. Ignorability holds if this DAG is correctly specified. If there is an unobserved variable U (motivation, ability) that affects both training participation and earnings, X fails the back-door criterion and ignorability fails.

Together, ignorability and positivity identify the ATE via Equation 4.6, and the back-door criterion provides the graphical justification for why that formula recovers the causal effect. Appendix B provides a formal graphical representation of this connection through SWIGs.

4.4.4 Overlap and Positivity

Why Overlap Is Non-Negotiable

Without overlap, some covariate strata contain only treated or only control units. In those strata, $\mathbb{E}[Y \mid T=0, X=x]$ or $\mathbb{E}[Y \mid T=1, X=x]$ is unobservable, and Equation 4.6 cannot be evaluated at those values of x . Identification of the ATE fails not because the causal structure is wrong, but because the data do not span the support needed. The ATT may remain identified even when full overlap fails, because ATT averages only over the treated population where $P(T=1 \mid X) > 0$ by construction.

4.5 Where the Frameworks Agree and Diverge

Task	Potential Outcomes	Do-Calculus / DAG	SEM
Define causal estimands (ATE, ATT)	Especially natural notation	Via $\mathbb{E}[Y \mid \text{do}(t)]$	Via structural equations
Encode causal assumptions	Assumptions typically stated informally; causal graph implicit	Explicit directed edges	Structural equations
Read conditional independence	Requires auxiliary graph or model	d-separation	With graph only
Identification from observational data	Via ignorability assumptions	Back-door, front-door, do-calculus	Via exclusion restrictions
Cross-world assumptions (e.g., monotonicity)	Natural to state	See Appendix B (SWIGs)	Implicit in structural model
Standard in statistics/epidemiology	Dominant	Growing rapidly	Econometrics

Our Position in This Course

Both frameworks are indispensable. We use *potential outcomes* to *define* causal estimands. We use the *do-calculus* and *DAGs* for *identification*. The back-door criterion is the graphical condition that justifies conditional ignorability and hence the adjustment formula through which causal effects are identified. For a formal graphical representation that hosts both frameworks simultaneously, see Appendix B.

The do-operator language is preferred throughout this course because it makes the interventional/observational distinction *syntactically enforced*: $f(y \mid \text{do}(t))$ cannot be confused with $f(y \mid T=t)$, whereas $\mathbb{E}[Y(t)]$ carries no such syntactic marker.

4.6 Summary

1. The potential outcome $Y(t)$ is the outcome that would be observed under the intervention $T = t$. Under consistency, no interference, and the structural causal semantics adopted in these notes, the observed outcome satisfies $Y_i = Y_i(T_i)$, and the potential outcome $Y(t)$ has the same distribution as the outcome under the intervention $\text{do}(T=t)$.
2. The ATE and ATT are averages of the unit-level causal effect $Y_i(1) - Y_i(0) = g(1, \mathbf{X}_i, U_{Y,i}) - g(0, \mathbf{X}_i, U_{Y,i})$ over the full population and the treated subpopulation respectively. In the linear SEM, both equal the structural coefficient β . They diverge under heterogeneous effects when treatment selection correlates with individual effect size.
3. Strong ignorability $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$ with overlap identifies the ATE via Equation 4.6. The back-door criterion provides a sufficient graphical condition for the conditional exchangeability assumption used in adjustment formulas.
4. SWIGs (Richardson and Robins 2014) provide a formal graphical representation of counterfactual variables, making ignorability a d-separation statement in a single diagram that hosts both potential outcomes and the natural treatment. This material is developed in Appendix B and is not required for a first-pass understanding of adjustment.
5. The frameworks are complementary: potential outcomes define estimands; the do-calculus identifies them. The back-door criterion connects the two by providing the graphical condition under which adjustment recovers the causal effect.

In observational studies, ignorability must be justified by substantive knowledge encoded in a causal graph. Randomized experiments provide a fundamentally different solution: when treatment is assigned randomly, $(Y(0), Y(1)) \perp\!\!\!\perp T$ holds by design, guaranteeing ignorability without any covariate adjustment. Chapter 5 studies randomized experiments as the canonical design where causal effects are identifiable directly from the data-generating mechanism.

4.7 Problems

1. SUTVA and consistency. Suppose $n = 3$ units receive binary treatments (T_1, T_2, T_3) and unit i 's outcome may depend on all three treatments: $Y_i = Y_i(T_1, T_2, T_3)$.

- How many potential outcomes does unit 1 have? Write them out explicitly.
- SUTVA imposes no interference: $Y_i(T_1, T_2, T_3) = Y_i(T_i)$. How many distinct potential outcomes remain?
- Give a real-world example where no-interference plausibly holds and one where it plausibly fails. In the latter case, explain what identification strategy (if any) remains available.

2. ATE, ATT, and selection bias. Let $(Y(0), Y(1), T) \sim P$ with $P(T=1) = 0.5$. Suppose $\mathbb{E}[Y(1)] = 3$, $\mathbb{E}[Y(0)] = 1$, $\mathbb{E}[Y(1) \mid T=1] = 4$, $\mathbb{E}[Y(0) \mid T=1] = 2$, $\mathbb{E}[Y(1) \mid T=0] = 2$, $\mathbb{E}[Y(0) \mid T=0] = 0$.

- Compute the ATE and ATT. Are they equal?
- Compute the naïve estimator $\mathbb{E}[Y \mid T=1] - \mathbb{E}[Y \mid T=0]$. Decompose the gap between this and the ATE into a selection bias term and an ATT-ATE difference.
- Under what graphical condition on the DAG would the naïve estimator equal the ATE? State this condition in both potential outcome language (ignorability) and do-calculus language (back-door criterion).

3. Causal estimands from the structural equation. Consider the nonparametric SEM $Y = g(T, X, U_Y)$ with binary treatment $T \in \{0, 1\}$, observed covariate X , and unobserved error U_Y independent of (T, X) in the mutilated graph.

- Write τ_{ATE} and τ_{ATT} as expectations of $g(1, X, U_Y) - g(0, X, U_Y)$ over the appropriate distribution. Under what condition do they coincide?
- Specialize to the linear SEM $g(t, x, u) = \alpha + \beta t + \gamma x + u$. Show that the unit-level causal effect $Y_i(1) - Y_i(0)$ is constant across all units, and hence $\tau_{\text{ATE}} = \tau_{\text{ATT}} = \beta$.
- Now consider the heterogeneous SEM $g(t, x, u) = (\alpha + u)t + \gamma x$, where $U_Y \sim \mathcal{N}(0, \sigma^2)$, $\text{Cov}(U_Y, T) = \rho$, and $P(T=1) = p \in (0, 1)$. Compute τ_{ATE} and τ_{ATT} . Show that they differ when $\rho \neq 0$, and interpret this difference.
- In part (c), explain why OLS estimation of Y on T and X does not recover either τ_{ATE} or τ_{ATT} when $\rho \neq 0$.

4. SWIGs and ignorability. (*Requires Appendix B.*) Consider the DAG: $X \rightarrow T$, $X \rightarrow Y$, $T \rightarrow Y$, with X fully observed.

- Construct the SWIG $\mathcal{G}(t)$ by splitting T into its random and fixed halves. Draw the result, labeling the random half, fixed half, and the potential outcome $Y(t)$.
- In $\mathcal{G}(t)$, identify all paths between the random half T and $Y(t)$. Apply the SWIG d-separation rules to determine which paths are blocked and which are open before conditioning.
- Use d-separation in $\mathcal{G}(t)$ to verify that $(Y(t) \perp\!\!\!\perp T \mid X)_{\mathcal{G}(t)}$ holds. Conclude that X satisfies the back-door criterion, confirming the proposition above.
- Now add a hidden common cause $U \rightarrow T$, $U \rightarrow Y$. Draw the revised SWIG. Does the ignorability argument still hold? State the correct conclusion and identify what additional structure (if any) would be needed for identification.

Chapter 5

Randomization and Back-Door Adjustment

Learning Objectives

By the end of this chapter, students should be able to:

1. Explain in do-calculus terms why a randomized experiment makes $f(y \mid \text{do}(T=t)) = f(y \mid T=t)$, and identify the graphical feature that produces this equality.
2. State the ignorability (unconfoundedness) assumption in both potential-outcomes and graphical language, and distinguish weak from strong ignorability.
3. Apply the regression-adjusted estimator to reduce variance in a randomized experiment, and explain why consistency holds even under model misspecification.
4. Derive and compute the standardization (g-formula) estimator of the ATE from a contingency table or regression output.
5. Explain the method of stratification (subclassification) and describe when and why it removes confounding bias.
6. Identify Simpson's paradox in a numerical example and resolve it using the back-door criterion.
7. Articulate why propensity score methods are needed when X is high-dimensional, motivating Chapter 6.

5.1 Randomized Experiments

5.1.1 The Do-Calculus of Randomization

The defining feature of a randomized experiment is that the analyst *sets* the treatment for each unit, independently of all background variables. In do-calculus terms, the observed data come from the mutilated distribution $f(y \mid \text{do}(T=t))$ rather than the observational distribution $f(y \mid T=t)$.

The Fundamental Equality of Randomized Experiments

In a completely randomized experiment, T is assigned independently of all pre-treatment variables. The DAG has no arrows into T : every back-door path from T to Y is severed. Consequently:

$$f(y \mid \text{do}(T=t)) = f(y \mid T=t). \quad (5.1)$$

The interventional and observational conditional distributions coincide.

In a completely randomized design (CRD), the treatment node T has no parents in the DAG. Graphically, this means:

- No back-door paths from T to Y exist.
- T is independent of all pre-treatment covariates: $T \perp\!\!\!\perp (X, U)$.

- In the mutilated graph $\mathcal{G}_{\overline{T}}$, nothing changes relative to the original DAG because T already had no parents.

Figure 5.1 illustrates the contrast between the observational DAG (with back-door paths) and the randomized DAG (all arrows into T severed).




Figure 5.1: Randomization as graph surgery. **Left:** in the observational DAG, the unobserved confounder U and observed covariate X both have arrows into T , creating back-door paths. **Right:** randomization physically severs all arrows into T (marked $\times$) and replaces them with an exogenous coin flip. The resulting graph is the mutilated graph $\mathcal{G}_{\overline{T}}$, in which no back-door paths exist.

Potential outcomes statement. In the potential outcomes language, randomization implies $(Y(0), Y(1)) \perp\!\!\!\perp T$. The potential outcomes are independent of assignment *unconditionally* — no covariate adjustment is needed.

5.1.2 Estimation of the ATE in a Randomized Experiment

Lemma: Identification under Complete Randomization

Under complete randomization,

$$\mathbb{E}[Y(t)] = \mathbb{E}[Y \mid T=t], \quad t \in \{0, 1\}. \quad (5.2)$$

Proof

By Proposition 4.1, $Y(t)$ has the same distribution as Y under $f(y \mid \text{do}(T=t))$, so $\mathbb{E}[Y(t)] = \mathbb{E}[Y \mid \text{do}(T=t)]$. In a completely randomized experiment, the treatment node has no parents in the DAG, so Equation 5.1 gives $f(y \mid \text{do}(T=t)) = f(y \mid T=t)$. Therefore $\mathbb{E}[Y(t)] = \mathbb{E}[Y \mid T=t]$. $\square$

Under complete randomization, the ATE is estimated by the *difference-in-means* (DIM) estimator:

$$\hat{\tau}_{\text{DIM}} = \bar{Y}_1 - \bar{Y}_0 = \frac{1}{n_1} \sum_{i: T_i=1} Y_i - \frac{1}{n_0} \sum_{i: T_i=0} Y_i.$$

Theorem: Unbiasedness and Variance of $\hat{\tau}_{\text{DIM}}$

Let $\sigma_t^2 = \text{Var}(Y(t))$ for $t \in \{0, 1\}$. Under complete randomization:

$$\begin{aligned} \mathbb{E}[\hat{\tau}_{\text{DIM}}] &= \tau_{\text{ATE}}, \\ \text{Var}(\hat{\tau}_{\text{DIM}}) &= \frac{\sigma_1^2}{n_1} + \frac{\sigma_0^2}{n_0}. \end{aligned}$$

5.2 Ignorability and the Back-Door Formula

5.2.1 Randomization vs. Observational Studies

In a randomized experiment, identification is guaranteed by the design. In an observational study, the analyst has no control over treatment assignment, and back-door paths through unobserved variables may remain open. The question shifts from *does* identification hold to *when* does it hold.

The two routes to identification differ in their logical structure. Randomization *guarantees* identification by design: the analyst removes all arrows into T , so the fundamental equality Equation 5.1 holds automatically. The observational route requires a substantive assumption: the analyst must argue, from subject-matter knowledge encoded in a DAG, that all confounders have been measured and that the back-door criterion is satisfied.

5.2.2 Definition

Definition: Ignorability (Unconfoundedness)

Let X denote a vector of pre-treatment covariates. We say that treatment assignment is **weakly ignorable** given X if:

$$Y(t) \perp\!\!\!\perp T \mid X, \quad t \in \{0, 1\}.$$

We say assignment is **strongly ignorable** given X (Rosenbaum and Rubin 1983) if, in addition, the overlap condition holds:

$$0 < P(T=1 \mid X) < 1 \quad \text{almost surely.}$$

Ignorability is also called *unconfoundedness*, *selection on observables*, *no unmeasured confounders*, or *conditional exchangeability*. In a randomized experiment, weak ignorability holds *unconditionally* — X is not needed, and the condition reduces to $(Y(0), Y(1)) \perp\!\!\!\perp T$.

5.2.3 Graphical Interpretation

In the DAG language, ignorability corresponds exactly to the back-door criterion. A set X satisfies the back-door criterion for the effect of T on Y if: (1) X blocks every back-door path from T to Y , and (2) X contains no descendant of T . When X satisfies the back-door criterion, $Y(t) \perp\!\!\!\perp T \mid X$ holds and the back-door adjustment formula identifies the causal effect.

Three Languages for Ignorability

Language	Statement of ignorability
SEM	The structural error ε_Y in the outcome equation $Y = f_Y(T, X, \varepsilon_Y)$ is independent of T given X : no unobserved variable enters both the equation that determines T and the equation that determines Y .
DAG / do-calculus	X satisfies the back-door criterion for the effect of T on Y : X blocks all back-door paths from T to Y and contains no descendant of T .
Potential outcomes	$Y(t) \perp\!\!\!\perp T \mid X$ for $t \in \{0, 1\}$ (weak ignorability).

The three statements are equivalent under the standard bridge assumptions between frameworks (consistency and no interference).

Ignorability Is Untestable from Observational Data

No statistical test can confirm or refute unconfoundedness using observed data alone. If an unobserved variable U affects both T and Y , the back-door path $T \leftarrow U \rightarrow Y$ is open, and conditioning on any set of observed variables cannot close it. Sensitivity analysis can quantify how large such an unobserved confounder would have to be to overturn a conclusion, but cannot confirm the assumption itself.

5.2.4 From Ignorability to Identification

Under strong ignorability, the ATE is identified:

$$\tau_{\text{ATE}} = \mathbb{E}_X[\mathbb{E}[Y \mid T=1, X] - \mathbb{E}[Y \mid T=0, X]]. \quad (5.3)$$

This is the back-door adjustment formula of Chapter 3 written in potential-outcome notation. Chapter 4 supplied the proof via the law of iterated expectations, consistency, and unconfoundedness.

5.3 Regression Adjustment

5.3.1 The Outcome Regression (OR) Estimator

Given the identifying formula Equation 5.3, the natural estimator replaces the population conditional means with their sample analogs.

Step 1. Fit a model $\hat{\mu}(t, x) \approx \mathbb{E}[Y \mid T=t, X=x]$ using the observed data.

Step 2. For each unit i , impute both potential outcome predictions: $\hat{\mu}(1, X_i)$ and $\hat{\mu}(0, X_i)$.

Step 3. Average the imputed treatment effects:

$$\hat{\tau}_{\text{ATE}} = \frac{1}{n} \sum_{i=1}^n [\hat{\mu}(1, X_i) - \hat{\mu}(0, X_i)]. \quad (5.4)$$

This three-step recipe is also called the *G-formula estimator* or the *regression-adjusted estimator*.

Example: Worked Outcome Regression

A study compares earnings (Y , in thousands) between job-training participants ($T=1$) and non-participants ($T=0$), with education level $X \in \{0, 1\}$ as a confounder. Among the 200 treated and 200 control units, 70% of treated units have $X=1$ vs. 30% of control units.

The estimated conditional means are $\hat{\mu}(1, 0) = 42$, $\hat{\mu}(1, 1) = 58$, $\hat{\mu}(0, 0) = 35$, $\hat{\mu}(0, 1) = 50$.

The within-stratum effects are $42 - 35 = 7$ (for $X=0$ units) and $58 - 50 = 8$ (for $X=1$ units). Since 50% of the full sample has $X=0$ and 50% has $X=1$:

$$\hat{\tau}_{\text{ATE}} = 7 \times 0.5 + 8 \times 0.5 = 7.5.$$

After adjusting for X , the estimated treatment effect is 7.5 — not the naïve difference-in-means, which conflates the treatment effect with the higher baseline earnings of $X=1$ units who happen to be treated more often.

5.3.2 Standardization as a Special Case

When X is categorical, no parametric model for $\mu(t, x)$ is needed: the conditional mean $\hat{\mu}(t, x)$ is simply the sample mean of Y within each cell ($T=t, X=x$) (a fully saturated outcome model). The resulting estimator is:

$$\hat{\tau}_{\text{STD}} = \sum_{x \in \mathcal{X}} [\hat{\mu}(1, x) - \hat{\mu}(0, x)] \hat{p}(x), \quad (5.5)$$

where $\hat{\mu}(t, x) = \bar{Y}_{t,x}$ and $\hat{p}(x) = n_x/n$. This is algebraically identical to Equation 5.4 when X is categorical. The formula is also known as the *g-formula* (Robins 1986) or *direct standardization* in epidemiology.

5.3.3 Regression Adjustment in Randomized Experiments: Lin (2013)

In a *randomized* experiment, outcome regression can be used to reduce variance even though covariate adjustment is not needed for unbiasedness. The *regression-adjusted estimator* of Lin (2013) fits the fully interacted OLS model:

$$Y_i = \alpha + \beta T_i + \gamma^\top \tilde{X}_i + \delta^\top (T_i \cdot \tilde{X}_i) + \varepsilon_i,$$

where $\tilde{X}_i = X_i - \bar{X}$ are mean-centered covariates. The OLS coefficient $\hat{\beta}$ is the ATE estimator.

Theorem: Design-Consistency of the Regression-Adjusted Estimator

(Lin 2013) Under complete randomization, $\hat{\beta}$ from the interacted regression is consistent for τ_{ATE} regardless of whether the linear model is correctly specified.

Proof Sketch

Because $T \perp\!\!\!\perp X$ under randomization, the projection of Y on $(T, X, T \cdot \tilde{X})$ recovers the treatment coefficient without confounding bias even if the true μ is not linear. The interaction terms ensure that the estimated treatment effect is the sample analog of the semiparametric efficient estimator. For the full proof see Lin (2013) or Freedman (2008). $\square$

Remark: The Role of the Regression Model in a Randomized Experiment

In a randomized experiment the regression model serves one purpose only: *variance reduction*. Unbiasedness is guaranteed by the randomization alone, regardless of whether the outcome model is correctly specified. The plain DIM estimator is already design-consistent; the only reason to add covariates is to reduce $\text{Var}(\hat{\tau})$ by absorbing residual variation in Y unrelated to T .

This stands in sharp contrast to the observational setting, where the regression model carries a *double* burden: it must both adjust for confounding (unbiasedness) and fit the outcome surface well (efficiency). Misspecification in an observational study threatens *both* properties simultaneously. In a randomized experiment, it threatens only efficiency.

Remark: Precedence in Survey Sampling

The design-consistency result in the theorem above was established in the survey sampling literature decades before Lin (2013). Isaki and Fuller (1982) introduced the *generalized regression (GREG) estimator* for finite-population means under general probability sampling designs and proved it is design-consistent even when the assisting regression model is misspecified. Deville and Särndal (1992) extended this class to *calibration estimators*, which adjust sampling weights so that weighted sample totals match known population totals for auxiliary variables. Lin (2013) restated and proved this result within the Neyman–Rubin potential-outcomes framework, connected it to semiparametric efficiency theory, and showed that the Huber–White sandwich variance estimator yields asymptotically valid confidence intervals under the randomization distribution.

Remark

Lin’s consistency result is specific to randomized experiments. In an observational study, $\hat{\tau}_{\text{OR}}$ is consistent only if $\hat{\mu}$ converges to the true μ . The doubly-robust (AIPW) estimator of Chapter 10 achieves consistency if *either* the outcome model or the propensity score model is correctly specified.

5.4 Stratification

5.4.1 The Idea

Stratification (subclassification) is a non-parametric implementation of the back-door formula. Rather than modeling $\mathbb{E}[Y \mid T, X]$, the analyst divides the sample into *strata* — subgroups of units with the same or similar values of X — and estimates the treatment effect within each stratum.

Within a stratum $\mathcal{S}_k = \{i : X_i \in B_k\}$, if the stratum is narrow enough, ignorability holds approximately and the within-stratum difference in means:

$$\hat{\tau}_k = \bar{Y}_{1,k} - \bar{Y}_{0,k}$$

is an approximately unbiased estimate of the stratum-specific ATE.

5.4.2 The Stratified Estimator

The overall ATE is estimated as a weighted average of within-stratum estimates:

$$\hat{\tau}_{\text{STRAT}} = \sum_{k=1}^K \hat{\tau}_k \cdot \frac{n_k}{n}, \quad (5.6)$$

where K is the number of strata, n_k is the number of units in stratum k , and $n = \sum_k n_k$. Cochran (1968) showed that with $K = 5$ strata of equal size, roughly 90% of the bias from a single continuous confounder is removed; with $K = 10$ strata, over 95%.

Remark: Stratification, Standardization, and Survey Sampling Terminology

When X is categorical, the stratified estimator Equation 5.6 and the standardization estimator Equation 5.5 are algebraically identical. The two names reflect disciplinary tradition rather than a methodological distinction.

In the survey sampling literature, *stratification* is a design-stage decision while *post-stratification* is an analysis-stage adjustment. Standardization in causal inference corresponds to post-stratification: stratum membership is recorded after sampling, and the estimator is reweighted by marginal stratum counts. Standardization is therefore regression adjustment with a saturated outcome model.

5.4.3 Limitations

Stratification faces the *curse of dimensionality*: with p binary covariates there are 2^p cells, many of which will be empty in practice. This motivates two solutions:

- **Regression adjustment / standardization** imposes parametric structure on $\hat{\mu}(t, x)$, allowing interpolation across cells.
- **Propensity score stratification** (Chapter 6) collapses the multivariate X to a scalar $\pi(X) = P(T=1 | X)$ and stratifies on the scalar.

5.5 Simpson's Paradox

Simpson's paradox — the reversal of an association when conditioning on a third variable — was introduced in Chapter 1, where a flu-vaccination example showed that the vaccine appeared harmful in the aggregate yet beneficial within every age stratum, because elderly individuals were disproportionately vaccinated. This chapter's development gives that example a second layer of interpretation: the back-door formula used in Chapter 1 is precisely the standardization estimator Equation 5.5 — within-stratum conditional means averaged over the *marginal* distribution of the confounder rather than its treatment-conditional distribution. The weighted average corrects the confounding; the pooled association does not, because it weights strata by $P(X | T)$ instead of $P(X)$.

5.5.1 When Should You *Not* Condition on X ?

Simpson's paradox also has a mirror image: conditioning on a variable can *create* a spurious association or make a true causal effect disappear. This occurs when X is a *mediator* or a *collider*. The back-door criterion gives the correct prescription: include X in the adjustment set only if it blocks back-door paths *without* opening collider paths.

Conditioning on a Mediator or Collider

Adjusting for a post-treatment variable X that lies on a causal pathway $T \rightarrow X \rightarrow Y$ blocks part of the causal effect of T , producing a downward-biased estimate of the total effect. Adjusting for a collider X with $T \rightarrow X \leftarrow Y$ opens a spurious association between T and Y that does not exist in the population. Neither mistake is detectable from the data alone; the DAG is essential.

The mediator case will be revisited in detail in Chapter 8. The collider case was covered in Chapter 2.

5.6 Lab: Simulation Study of the Outcome Regression Estimator

The identification results in Section 5.2 and the estimation recipes in Section 5.3 rest on two assumptions that pull in opposite directions: the outcome model $\hat{\mu}(t, x)$ must be correctly specified for the linear OR estimator to be consistent in an observational study, yet a more flexible nonparametric estimator pays a variance penalty for that flexibility.

DGP for Lab 5

Let $X \sim \text{Uniform}(0, 1)$ be a single pre-treatment covariate and $\varepsilon \sim \mathcal{N}(0, 1)$ independent of X . The potential outcome surfaces are:

$$Y(t) = t + 3(1+t)X^2 + \varepsilon, \quad t \in \{0, 1\}. \quad (5.7)$$

The conditional average treatment effect is $\tau(x) = 1 + 3x^2$, so

$$\tau_{\text{ATE}} = \mathbb{E}[1 + 3X^2] = 1 + 1 = 2. \quad (5.8)$$

Two assignment mechanisms are studied:

- **CRD:** $T \sim \text{Bern}(0.5)$, independent of X .
- **Observational:** $T \mid X \sim \text{Bern}(\text{expit}(-2 + 5X))$, giving $\pi(0) \approx 0.12$, $\pi(0.5) \approx 0.62$, $\pi(1) \approx 0.95$. Strong confounding: high- X units are nearly always treated.

Two outcome regression estimators of the form Equation 5.4 are compared: a **Linear OR** (simple linear regression, misspecified because the true μ contains a quadratic term), and a **Local-linear OR** (local linear regression with bandwidth $h = 0.5 \cdot n^{-1/5}$, correctly flexible).

Key Finding: Model Misspecification Is Harmless under Randomization, Harmful under Confounding

Under CRD, the linear OR estimator is consistent despite misspecification (Theorem above): the fitting and evaluation distributions coincide, so pointwise errors cancel on average.

Under the observational design, the fitting and evaluation distributions diverge. High- X units are almost always treated and low- X units almost never. The model for arm $t = 1$ is fit on a sample concentrated at high X , and the model for arm $t = 0$ on a sample concentrated at low X . The fitted lines are tuned to different parts of the covariate space, so when both are evaluated at the full marginal of X , the pointwise errors no longer cancel. Model misspecification is harmless under randomization; it becomes a source of bias precisely when confounding is present.

Remark: What This Lab Does Not Show

The simulation fixes the bandwidth $h = 0.5 \cdot n^{-1/5}$ and does not perform data-driven bandwidth selection. The lab also studies only a scalar X ; with a multivariate covariate, the curse of dimensionality would make local linear regression impractical, which is one motivation for the propensity score approach of Chapter 6.

5.7 Where the Propensity Score Fits

The three methods — regression adjustment, standardization, and stratification — all implement the same identifying formula Equation 5.3. They differ in how they deal with the conditioning on the high-dimensional covariate vector X :

Method	How it handles X	Key assumption
Regression / standardization	Models $\mathbb{E}[Y \mid T, X]$ parametrically	Outcome model correctly specified
Stratification	Groups units by X ; estimates effect in each cell	Strata narrow enough for approximate balance
Propensity score (Ch. 6)	Reduces X to scalar $\pi(X) = P(T=1 \mid X)$	Propensity model correctly specified

Bridge to Chapter 6

Chapter 6 takes up the propensity score idea formally. It proves the balancing property ($T \perp\!\!\!\perp X \mid \pi(X)$), derives the Propensity Score Theorem (strong ignorability given X implies strong ignorability given $\pi(X)$ alone), and develops three estimation strategies: inverse probability weighting (IPW), matching, and propensity score stratification. Matching need not rely on the propensity score: covariate-distance methods such as Mahalanobis-distance matching and nearest-neighbor matching on X directly are equally valid.

5.8 Summary

1. **Randomization as graph surgery.** A randomized experiment removes all arrows into T , so the fundamental equality $f(y \mid \text{do}(T=t)) = f(y \mid T=t)$ holds and the difference-in-means estimator is unbiased for the ATE without covariate adjustment.
2. **Ignorability is the key assumption.** Under unconfoundedness $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$, the back-door adjustment formula identifies the ATE. This assumption is substantive, design-determined, and untestable from observational data.
3. **Three estimation strategies.** Regression adjustment (G-formula), standardization, and stratification all implement the same identifying formula via different modeling choices. All require correct specification of either an outcome model (for regression/standardization) or a stratification scheme fine enough to achieve balance.
4. **Simpson's paradox.** Pooled associations can reverse within subgroups when a confounder determines treatment selection. The correct causal estimate requires standardization over the marginal covariate distribution, guided by the back-door criterion.
5. **The propensity score dimension reduction.** When X is high-dimensional, direct stratification or cell-by-cell standardization breaks down. The propensity score $\pi(X)$ provides a scalar sufficient statistic for adjustment. Chapter 6 develops the theory and estimation methods.

5.9 Problems

1. **Randomization and the do-calculus.** Let the observational DAG be $\{U \rightarrow T, U \rightarrow Y, X \rightarrow T, T \rightarrow Y\}$ with U unobserved.
 - (a) Identify all back-door paths from T to Y .
 - (b) Does X satisfy the back-door criterion? Does U ? Explain.
 - (c) Now suppose treatment is randomized. Draw the modified DAG and identify all back-door paths. Show that $f(y \mid \text{do}(T=t)) = f(y \mid T=t)$ holds in the randomized DAG using Rule 2 of the do-calculus.
 - (d) Under randomization, is covariate adjustment on X necessary for unbiasedness? Is it ever beneficial? Explain.

2. Ignorability and the back-door criterion. Consider the DAG $\{X \rightarrow T, X \rightarrow Y, T \rightarrow Y, U \rightarrow Y\}$ with U unobserved.

- Does $\{X\}$ satisfy the back-door criterion? Write out the identifying formula for τ_{ATE} .
- Now add the edge $U \rightarrow T$ to the DAG. Does $\{X\}$ still satisfy the back-door criterion? What does the back-door criterion require when both X and U affect T ?
- Explain in words what “unconfoundedness given X ” means about the role of U in the data-generating process.

3. Standardization. A study of job training (T) and earnings (Y , in thousands) produces the following cell means, with $P(X=0) = 0.4$ and $P(X=1) = 0.6$:

X	$\hat{\mu}(1, x)$	$\hat{\mu}(0, x)$	n_x/n
0	28	22	0.4
1	35	31	0.6

- Compute $\hat{\tau}_{\text{ATE}}$ using the standardization formula.
- Compute $\hat{\tau}_{\text{ATT}}$, given that all treated units come from $X=1$.
- The unadjusted difference in means is $33 - 25 = 8$. Compare to your answers in (a) and (b) and explain the discrepancy.

4. Stratification and Cochran’s rule. You have a binary confounder $X \in \{0, 1\}$ and form two strata. Within stratum $X=0$: $n_0 = 600$, $\bar{Y}_1 = 10$, $\bar{Y}_0 = 8$. Within stratum $X=1$: $n_1 = 400$, $\bar{Y}_1 = 15$, $\bar{Y}_0 = 12$.

- Compute the stratified ATE estimator $\hat{\tau}_{\text{STRAT}}$.
- Suppose an unadjusted difference-in-means gives $\hat{\tau}_{\text{DIM}} = 5.5$. Explain why the two estimates differ and which one is the appropriate causal estimate.
- Describe one limitation that would arise if X were a continuous variable with 15 dimensions.

5. Simpson’s paradox. A hospital reports that patients who receive intensive care (ICU admission, $T=1$) have a higher mortality rate than those who do not ($T=0$): $P(Y=1 \mid T=1) = 0.30$ vs. $P(Y=1 \mid T=0) = 0.10$.

- Construct a numerical example (a $2 \times 2 \times 2$ table with disease severity $X \in \{\text{mild}, \text{severe}\}$ as the confounder) that is consistent with the above pooled numbers and yet shows ICU admission *reduces* mortality within both severity strata.
- Compute the standardized ATE using the marginal distribution of X .
- Draw the DAG for this example. Identify the back-door path that the pooled comparison fails to block.
- A colleague suggests that since the hospital treats patients with both mild and severe illness, the pooled statistic is the right answer. Explain, using potential outcomes notation, why this argument is incorrect.

Chapter 6

Propensity Score Methods

Learning Objectives

By the end of this chapter, students should be able to:

1. Define the propensity score $\pi(X) = P(T=1 | X)$ and explain why it is a *balancing score*.
2. State and prove the Balancing Property ($T \perp\!\!\!\perp X | \pi(X)$) and the Propensity Score Theorem (strong ignorability given X implies strong ignorability given $\pi(X)$).
3. Describe standard methods for estimating the propensity score (logistic regression, machine learning) and the practical pitfalls of each.
4. Derive the IPW identification formula for the ATE and explain the role of the Horvitz–Thompson representation.
5. Distinguish the IPW estimator (targets ATE) from the nearest-neighbor matching estimator presented in this chapter (designed to target ATT).
6. Define strict overlap, describe the practical consequences of near-violations, and apply trimming strategies.
7. Articulate the fundamental limitation of propensity-score methods — the unconfoundedness assumption — and explain why this motivates instrumental variables (Chapter 7).

6.1 Motivation: The Curse of Dimensionality

Chapter 5 established that under strong ignorability, the ATE is identified by the back-door adjustment formula, and developed three estimation strategies: regression adjustment, standardization, and stratification. All three require, in practice, that we condition on a covariate vector X . When X is low-dimensional, this is feasible. When X has many components, it is not.

The problem. Stratification requires cells $\{X = x\}$ to contain both treated and control units. With p binary covariates, there are 2^p cells. With $p = 20$, that is over one million cells — most empty in any realistic sample. Regression adjustment imposes parametric structure to interpolate across cells, but extrapolation bias grows with dimension. The propensity score provides an elegant solution: rather than conditioning on the full vector X , it is sufficient to condition on the scalar $\pi(X) = P(T=1 | X)$.

6.2 The Propensity Score

6.2.1 Definition

Definition: Propensity Score

The *propensity score* is the conditional probability of receiving treatment given the observed covariates:

$$\pi(X) = P(T=1 | X).$$

In a randomized experiment, $\pi(X) = p$ is constant (equal to the allocation probability), so there is nothing to estimate. In an observational study, $\pi(X)$ varies with X and must be estimated from data.

Connection to strong ignorability. Under strong ignorability $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$, the potential outcomes carry no further information about treatment assignment once X is given:

$$P(T=1 \mid X, Y(0), Y(1)) = P(T=1 \mid X) = \pi(X).$$

This equality is a *consequence* of the ignorability assumption, not part of the definition of $\pi(X)$.

6.2.2 The Balancing Property

Theorem: The Propensity Score is a Balancing Score

(Rosenbaum and Rubin 1983) $T \perp\!\!\!\perp X \mid \pi(X)$.

Proof

It suffices to show $P(T=1 \mid X, \pi(X)) = P(T=1 \mid \pi(X))$. Since $\pi(X)$ is a deterministic function of X , conditioning on both X and $\pi(X)$ is the same as conditioning on X alone:

$$P(T=1 \mid X, \pi(X)) = P(T=1 \mid X) = \pi(X) = P(T=1 \mid \pi(X)). \quad \square$$

Remark: Covariate Balance and Design-Before-Analysis

The balancing property implies $f(X \mid T=1, \pi(x)) = f(X \mid T=0, \pi(x))$. Stratifying on a discretized $\hat{\pi}(X)$ and comparing covariate distributions within strata via standardized mean differences diagnoses whether the propensity score model achieves adequate balance, *without ever looking at Y* . Because the diagnostic does not involve the outcome, it cannot be used to “fish” for a preferred model specification, making the analysis more objective (Imbens and Rubin 2015).

Remark

The Balancing Property is a statement about the *observed* covariate distribution. It says that within strata of $\pi(X)$, the distribution of X is the same in the treated and control groups. It does *not* say anything about unobserved confounders U : if U affects both T and Y and is not captured by X , the balancing property fails to close the back-door path through U .

6.2.3 The Propensity Score Theorem

Theorem: Propensity Score Theorem

(Rosenbaum and Rubin 1983) Suppose $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$. Then for any balancing score $b(X)$,

$$(Y(0), Y(1)) \perp\!\!\!\perp T \mid b(X).$$

In particular, taking $b(X) = \pi(X)$: strong ignorability given X implies strong ignorability given $\pi(X)$ alone, provided $0 < \pi(X) < 1$ almost surely.

Proof

We prove the result for $b(X) = \pi(X)$. By the law of iterated expectations, for any $t \in \{0, 1\}$:

$$P(T=1 \mid Y(t), \pi(X)) = \mathbb{E}[P(T=1 \mid Y(t), X) \mid Y(t), \pi(X)].$$

By ignorability, $T \perp\!\!\!\perp Y(t) \mid X$, so $P(T=1 \mid Y(t), X) = P(T=1 \mid X) = \pi(X)$. Therefore:

$$P(T=1 \mid Y(t), \pi(X)) = \mathbb{E}[\pi(X) \mid Y(t), \pi(X)] = \pi(X) = P(T=1 \mid \pi(X)).$$

Hence $T \perp\!\!\!\perp Y(t) \mid \pi(X)$. $\square$

Remark

The theorem gives an alternative identification formula:

$$\mathbb{E}[Y(t)] = \mathbb{E}[\mathbb{E}[Y \mid \pi(X), T=t]].$$

A common misconception is that the two conditional regressions $\mathbb{E}[Y \mid \pi(X), T=t]$ and $\mathbb{E}[Y \mid X, T=t]$ are equal. They are not: they are different functions and their pointwise values generally differ. The theorem guarantees only that their marginal expectations agree:

$$\mathbb{E}[\mathbb{E}[Y \mid \pi(X), T=t]] = \mathbb{E}[\mathbb{E}[Y \mid X, T=t]] = \mathbb{E}[Y(t)].$$

6.2.4 The Propensity Score is the Coarsest Balancing Score

Lemma: The Propensity Score is the Coarsest Balancing Score

(Rosenbaum and Rubin 1983) If $b(X)$ is any balancing score — that is, $T \perp\!\!\!\perp X \mid b(X)$ — then $\pi(X)$ is a function of $b(X)$: $\pi(X) = g(b(X))$ for some measurable function g .

Proof

Since $b(X)$ is a balancing score, $T \perp\!\!\!\perp X \mid b(X)$, so $P(T=1 \mid X, b(X)) = P(T=1 \mid b(X))$. But $b(X)$ is a function of X :

$$\pi(X) = P(T=1 \mid X) = P(T=1 \mid X, b(X)) = P(T=1 \mid b(X)).$$

Hence $\pi(X)$ is a measurable function of $b(X)$ alone. $\square$

The lemma says that $\pi(X)$ is the *smallest* sufficient statistic for T among all functions of X : any other balancing score retains at least as much information about X as the propensity score does.

The Dimension Reduction

The Propensity Score Theorem says that instead of adjusting for the full high-dimensional vector X , it is sufficient to adjust for the scalar $\pi(X)$. Identification of the ATE follows immediately:

$$\tau_{\text{ATE}} = \mathbb{E}[\mathbb{E}[Y \mid T=1, \pi(X)] - \mathbb{E}[Y \mid T=0, \pi(X)]].$$

6.3 Estimation of the Propensity Score

The true propensity score $\pi(X) = P(T=1 \mid X)$ is *unknown* in observational studies and must be estimated from the data.

Logistic regression. The classical approach models $\text{logit}(\pi(X)) = X^\top \beta$ and estimates β by maximum likelihood. The fitted values $\hat{\pi}(X_i) = \sigma(X_i^\top \hat{\beta})$ are used in place of the true propensity score.

Machine learning estimators. When X is high-dimensional or the true propensity score is nonlinear, machine learning methods — gradient boosted trees, random forests, regularized logistic regression (LASSO, ridge) — can estimate $\pi(X)$ more flexibly. A critical issue is *overfitting*: an estimator that perfectly separates treated and control units in-sample will produce estimated propensity scores near 0 or 1 everywhere, destroying overlap. Cross-fitting (Chapter 11) addresses this by estimating nuisance functions on held-out folds.

Propensity Score Estimation Is a Nuisance, Not the Goal

The propensity score is estimated in order to construct a good estimator of τ_{ATE} . Balancing tests diagnose whether the estimated score has achieved covariate balance, but they do *not* test whether unobserved confounders are balanced.

6.4 Matching on the Propensity Score

The propensity score also motivates *matching*: for each treated unit i , find a control unit j with $\hat{\pi}(X_j) \approx \hat{\pi}(X_i)$, and use Y_j as a proxy for $Y_i(0)$ (Abadie and Imbens 2006, 2016). The resulting estimator of the ATT is:

$$\hat{\tau}_{\text{ATT}} = \frac{1}{n_1} \sum_{i: T_i=1} [Y_i - Y_{\hat{j}(i)}],$$

where $\hat{j}(i)$ is the index of the matched control unit.

One-to-one nearest neighbor matching. For each treated unit i , find:

$$\hat{j}(i) = \arg \min_{j: T_j=0} |\hat{\pi}(X_i) - \hat{\pi}(X_j)|.$$

Matching with replacement reduces bias (each control unit is available to every treated unit) but increases variance.

Caliper matching. Restricts matches to pairs within a maximum distance δ on the propensity score scale. Treated units with no control match within the caliper are discarded, restricting the estimand to the subset of treated units with good overlap.

Remark

Propensity score matching implicitly fits a nonparametric regression of Y on $\pi(X)$ separately within each treatment arm. By the Propensity Score Theorem, $(Y(0), Y(1)) \perp\!\!\!\perp T \mid \pi(X)$, so the conditional mean $\mathbb{E}[Y(t) \mid \pi(X) = p]$ is identified from observed data. Nearest-neighbor matching approximates this conditional mean at each observed $\hat{\pi}(X_i)$ by local constant regression on the scalar propensity score — feasible precisely because of the dimension reduction from X to $\pi(X)$.

6.5 Inverse Probability Weighting

6.5.1 The IPW Identification Formula

Under strong ignorability, the ATE admits an alternative representation as an *inverse probability weighting* (IPW) formula:

$$\tau_{\text{ATE}} = \mathbb{E} \left[\frac{T \cdot Y}{\pi(X)} \right] - \mathbb{E} \left[\frac{(1-T) \cdot Y}{1 - \pi(X)} \right].$$

Theorem: IPW Identification

Under strong ignorability $((Y(0), Y(1)) \perp\!\!\!\perp T \mid X$ and $0 < \pi(X) < 1$):

$$\mathbb{E} \left[\frac{T \cdot Y}{\pi(X)} \right] = \mathbb{E}[Y(1)].$$

Proof

$$\mathbb{E} \left[\frac{TY}{\pi(X)} \right] = \mathbb{E} \left[\frac{TY(1)}{\pi(X)} \right] = \mathbb{E} \left[\mathbb{E} \left[\frac{TY(1)}{\pi(X)} \mid X \right] \right] = \mathbb{E} \left[\frac{Y(1)}{\pi(X)} \mathbb{E}[T \mid X] \right] = \mathbb{E} \left[\frac{Y(1)}{\pi(X)} \pi(X) \right] = \mathbb{E}[Y(1)].$$

The first equality uses consistency ($Y = Y(1)$ when $T=1$); the third uses ignorability ($Y(1) \perp\!\!\!\perp T \mid X$).
□

6.5.2 The Horvitz–Thompson and Hájek Estimators

The *Horvitz–Thompson* (HT) IPW estimator:

$$\hat{\tau}_{\text{IPW}} = \frac{1}{n} \sum_{i=1}^n \left[\frac{T_i Y_i}{\hat{\pi}(X_i)} - \frac{(1 - T_i) Y_i}{1 - \hat{\pi}(X_i)} \right].$$

This estimator is consistent when $\hat{\pi}(X)$ converges to $\pi(X)$, but can be *inefficient* due to extreme weights. The *normalized* (Hájek) estimator divides by the sum of weights within each group and is often more stable:

$$\hat{\tau}_{\text{HJ}} = \frac{\sum_i T_i Y_i / \hat{\pi}(X_i)}{\sum_i T_i / \hat{\pi}(X_i)} - \frac{\sum_i (1 - T_i) Y_i / (1 - \hat{\pi}(X_i))}{\sum_i (1 - T_i) / (1 - \hat{\pi}(X_i))}.$$

The efficient doubly-robust estimator (AIPW) of Chapter 10 combines the outcome model and the propensity score; under correct nuisance specification it achieves higher efficiency than either pure IPW or pure outcome regression alone.

6.6 Overlap and Positivity

6.6.1 The Strict Overlap Condition

Definition: Strict Overlap (Positivity)

The *strict overlap condition* requires:

$$0 < \pi(X) < 1 \quad \text{almost surely.}$$

Every unit has a positive probability of receiving either treatment or control, regardless of its covariate values.

Together with unconfoundedness, strict overlap constitutes *strong ignorability* (Rosenbaum and Rubin 1983).

Why overlap is necessary. The identification formula requires that both $\mathbb{E}[Y \mid T=1, X=x]$ and $\mathbb{E}[Y \mid T=0, X=x]$ be defined for all x in the support of X . If $\pi(x) = 0$ for some x , then no treated units exist with $X = x$, so $\mathbb{E}[Y \mid T=1, X=x]$ is not estimable.

6.6.2 Practical Consequences of Near-Violations

Near-Overlap Problems

When $\hat{\pi}(X_i)$ is close to 0 or 1, the IPW weight $1/\hat{\pi}(X_i)$ or $1/(1 - \hat{\pi}(X_i))$ is very large. A small number of units with extreme weights can dominate the estimator, increasing variance dramatically.

6.6.3 Trimming Strategies

Trimming by propensity score. Restrict the analysis to units with estimated propensity scores in $[\eta, 1 - \eta]$ for some small $\eta > 0$. The estimand changes to $\tau_{\text{trim}} = \mathbb{E}[Y(1) - Y(0) \mid \eta \leq \pi(X) \leq 1 - \eta]$.

Crump et al. (2009) rule. Crump et al. (2009) derive the optimal trimming threshold minimizing the asymptotic variance of the trimmed ATE estimator. The rule $\eta = 0.1$ — trimming units with $\pi(X) < 0.1$ or $\pi(X) > 0.9$ — is a common default.

ATT instead of ATE. When overlap fails at covariate values where $\pi(X)$ is near 1, re-targeting the estimand to τ_{ATT} avoids the problem, since the ATT requires only $\pi(X) < 1$ on the support of X among the treated.

6.7 Lab: Simulation Study of IPW and Matching Estimators

This lab compares five estimators on a five-covariate data-generating process with treatment-effect heterogeneity. With five continuous confounders, direct stratification is infeasible, making propensity-score methods the natural choice. The estimators span two estimands: the ATE (HT-ATE, HJ-ATE) and the ATT (HT-ATT, HJ-ATT, NNM).

6.7.1 Part 1: Correctly Specified Propensity Score

DGP for Lab 6

Let $X_j \stackrel{\text{i.i.d.}}{\sim} \text{Normal}(0, 1)$ for $j = 1, \dots, 5$. The true propensity score is:

$$\pi(X) = \text{expit}(-0.5 + 0.8X_1 + 0.5X_2 - 0.3X_3 + 0.2X_4 + 0.1X_5), \quad (6.1)$$

giving $P(T=1) \approx 0.40$. The potential outcomes are:

$$Y(t) = (1 + 0.5X_1) \cdot t + 8 + 0.8X_1 + 0.5X_2 + 0.4X_3 + 0.3X_4 + 0.2X_5 + \varepsilon, \quad \varepsilon \sim N(0, 1). \quad (6.2)$$

The CATE is $\tau(X) = 1 + 0.5X_1$. Because $\mathbb{E}[X_1] = 0$: $\tau_{\text{ATE}} = 1.000$ (exact). Because $\pi(X)$ is increasing in X_1 , treated units have higher X_1 on average: $\tau_{\text{ATT}} \approx 1.199$ (oracle run, $n = 10^7$). The gap ≈ 0.199 reflects a selection effect: units most likely to be treated also tend to have larger individual treatment effects. Strong ignorability holds by construction.

Five estimators are compared under both true and estimated propensity scores. Under the estimated score, $\hat{\pi}(X)$ is obtained by correctly specified logistic regression of T on $(X_1, \dots, X_5)$.

Estimator	Formula	Estimand
HT-ATE	$\frac{1}{n} \sum_i \left[\frac{T_i Y_i}{\hat{\pi}(X_i)} - \frac{(1-T_i) Y_i}{1-\hat{\pi}(X_i)} \right]$	ATE
HJ-ATE	Hájek-normalized HT	ATE
HT-ATT	$\frac{1}{n_1} \sum_{T_i=1} Y_i - \frac{1}{n_1} \sum_{T_i=0} \frac{\hat{\pi}(X_i)}{1-\hat{\pi}(X_i)} Y_i$	ATT
HJ-ATT	Self-normalized version of HT-ATT (control arm only)	ATT
NNM	$\frac{1}{n_1} \sum_{T_i=1} [Y_i - Y_{\hat{j}(i)}]$, 1:1 with replacement	ATT

Results. $n = 1,000$, $B = 2,000$ replications, seed 42.

PS	Estimator	Estimand	Mean	Bias	SD	RMSE
Known	HT-ATE	ATE	1.008	+0.008	0.626	0.626
Known	HJ-ATE	ATE	1.003	+0.003	0.145	0.145
Known	HT-ATT	ATT	1.184	-0.016	0.725	0.725
Known	HJ-ATT	ATT	1.203	+0.003	0.131	0.131
Known	NNM	ATT	1.207	+0.008	0.130	0.130
Estimated	HT-ATE	ATE	0.998	-0.002	0.194	0.194
Estimated	HJ-ATE	ATE	1.002	+0.002	0.102	0.103
Estimated	HT-ATT	ATT	1.202	+0.002	0.251	0.251
Estimated	HJ-ATT	ATT	1.202	+0.002	0.101	0.101
Estimated	NNM	ATT	1.205	+0.006	0.121	0.121

Five lessons stand out. **1.** Hájek normalization is essential when outcomes are non-centered: HT-ATE achieves SD = 0.626 vs. HJ-ATE SD = 0.145 — a 4.3-fold reduction. **2.** The same problem afflicts ATT: HT-ATT SD = 0.725 vs. HJ-ATT SD = 0.131 — a 5.5-fold reduction. **3.** HJ-ATT and NNM converge to the same ATT target with nearly identical efficiency (SD 0.131 vs. 0.130 under known PS). **4.** The ATE-ATT gap of ≈ 0.199 is a real difference in estimands, not a bias — declaring the estimand before

looking at results is non-negotiable. **5.** Estimated PS collapses HT variance; HJ-ATT with estimated PS (SD = 0.101) beats NNM (SD = 0.121), illustrating the efficiency gain of a well-estimated score.

The Estimand Must Be Chosen Before the Estimator

In this DGP, $\tau_{\text{ATE}} = 1.000$ and $\tau_{\text{ATT}} \approx 1.199$. The gap arises because treatment-effect heterogeneity is correlated with propensity for treatment (Imbens and Rubin 2015) — individuals often select into treatment partly because they anticipate larger benefits. A researcher who applies NNM and reports its output against the ATE benchmark of 1.000 would conclude the estimator has 20% bias; in fact it is unbiased for the correct target.

6.7.2 Part 2: Robustness to PS Model Misspecification

Part 1 used a correctly specified logistic propensity model. Part 2 introduces a *nonlinear* true propensity score:

$$\pi^*(X) = \text{expit}(-0.5 + 0.8X_1 + 0.5X_1^2 + 0.2X_4 + 0.1X_5), \quad (6.3)$$

with a U-shaped surface in X_1 . The true ATE remains 1.000 (exact); the ATT changes to ≈ 1.152 . The *estimated* PS is still a linear logistic regression that omits X_1^2 — misspecified.

Results. True ATE = 1.000; true ATT ≈ 1.152 . $n = 1,000$, $B = 2,000$ replications.

PS	Estimator	Estimand	Mean	Bias	SD	RMSE
True	HT-ATE	ATE	1.028	+0.028	0.824	0.825
True	HJ-ATE	ATE	1.013	+0.013	0.152	0.152
True	HT-ATT	ATT	1.186	+0.034	1.333	1.333
True	HJ-ATT	ATT	1.183	+0.031	0.212	0.215
True	NNM	ATT	1.178	+0.026	0.175	0.177
Misspecified	HT-ATE	ATE	1.478	+0.478	0.147	0.501
Misspecified	HJ-ATE	ATE	0.861	−0.139	0.087	0.164
Misspecified	HT-ATT	ATT	1.780	+0.628	0.162	0.649
Misspecified	HJ-ATT	ATT	1.306	+0.154	0.081	0.174
Misspecified	NNM	ATT	1.172	+0.020	0.138	0.139

Three lessons complete the picture. **6.** PS misspecification makes all IPW estimators inconsistent; Hájek normalization is a variance correction only — it cannot fix a structurally wrong model. Note also that HJ-ATE and HT-ATE have biases in *opposite directions* (−0.139 vs. +0.478) because the two estimators respond differently to the systematic over/underestimation of the U-shaped score. **7.** NNM is substantially more robust: its bias under misspecification (+0.020) is far smaller than any IPW estimator, because matching relies on local covariate similarity rather than a globally correct weighting function.

IPW Needs a Correct PS Model; Matching Needs Covariate Balance

IPW is a *model-dependent* estimator: consistency relies on correct PS specification. Matching is a *design-dependent* estimator: consistency relies on matched pairs being approximately balanced. This condition can hold even when the PS model is wrong, provided the estimated score still separates units well enough that close matches imply closeness on the full covariate vector. The cost is efficiency: when the PS model is correct, matching leaves variance on the table relative to IPW. Neither dominates unconditionally (Yang et al. 2016; Yang and Zhang 2023).

Head-to-head comparison of HJ-ATT and NNM:

Estimator	Bias (Correct PS)	RMSE (Correct PS)	Bias (Misspecified PS)	RMSE (Misspecified PS)
HJ-ATT	+0.002	0.101	+0.154	0.174
NNM	+0.006	0.121	+0.020	0.139

6.8 The Limits of Propensity Score Methods

Every result in this chapter — the Propensity Score Theorem, the IPW identification formula, the matching estimator — rests on the unconfoundedness assumption:

$$(Y(0), Y(1)) \perp\!\!\!\perp T \mid X.$$

Design-Based vs. Model-Based Identification

Randomization *guarantees* ignorability: by construction, no back-door paths exist. Propensity-score adjustment *assumes* ignorability: the analyst must argue from subject-matter knowledge that all confounders have been measured. This distinction is fundamental. No statistical test can confirm unconfoundedness from observed data alone (Chapter 5). Sensitivity analysis can quantify how large an unobserved confounder would have to be to overturn a conclusion, but cannot confirm the assumption.

When unobserved confounders are present, propensity score adjustment fails and an instrumental variable is needed to identify the causal effect. Instrumental variables take a different approach: instead of blocking the confounding path $T \leftarrow U \rightarrow Y$, they isolate variation in T induced by an external source Z independent of U . This is derived formally in Chapter 7.

6.9 Summary

1. **The propensity score reduces dimension.** Under unconfoundedness, $\pi(X) = P(T=1 \mid X)$ is a *balancing score*: $T \perp\!\!\!\perp X \mid \pi(X)$. Consequently, $(Y(0), Y(1)) \perp\!\!\!\perp T \mid \pi(X)$, so identification requires adjustment for the scalar $\pi(X)$ alone.
2. **IPW identification.** The ATE is identified as $\mathbb{E}[TY/\pi(X)] - \mathbb{E}[(1-T)Y/(1-\pi(X))]$ under strong ignorability. The Horvitz–Thompson IPW estimator is consistent but sensitive to extreme weights; the Hájek variant is often more stable. AIPW (Chapter 10) achieves higher efficiency under correct nuisance specification.
3. **The matching procedure in this chapter targets the ATT.** The nearest-neighbor matching estimator finds control units with similar propensity scores to each treated unit. It differs fundamentally from IPW weighting in estimand and methodology, and is more robust to PS misspecification.
4. **Strict overlap is non-negotiable.** The condition $0 < \pi(X) < 1$ a.s. is necessary for identification of the ATE. Near-violations create extreme IPW weights and large variance; trimming strategies address this at the cost of restricting the estimand.
5. **The fundamental limitation.** Unconfoundedness $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$ is a strong, untestable assumption. When unobserved confounders are present, propensity score adjustment fails and an instrumental variable is needed.

Design	Key assumption	Identified estimand
Randomized experiment	$(Y(0), Y(1)) \perp\!\!\!\perp T$ (by design)	ATE
Propensity-score adjustment	$(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$ (unconfoundedness)	ATE or ATT
Instrumental variables	Relevance, exogeneity, exclusion	LATE (compliers)

6.10 Problems

1. **Propensity score and balancing.** Suppose $X = (X_1, X_2)$, where X_1 is binary and X_2 is continuous. The propensity score model is correctly specified as $\text{logit}(\pi(X)) = \beta_0 + \beta_1 X_1 + \beta_2 X_2$.

- (a) State and prove the Balancing Property $T \perp\!\!\!\perp X \mid \pi(X)$.
- (b) Two units i and j have $X_i = (1, 2.3)$ and $X_j = (0, 3.5)$ but $\hat{\pi}(X_i) = \hat{\pi}(X_j) = 0.4$. If i is treated and j is control, explain why comparing Y_i and Y_j is a valid approximation to the counterfactual comparison, and what assumption is required.

- (c) Explain why matching on $\hat{\pi}(X)$ is *not* the same as matching on X directly. Under what conditions do they give the same answer?

2. ATE, ATT, and propensity score weighting. A dataset has $n = 1000$ observations with estimated propensity scores $\hat{\pi}(X_i)$ and $n_1 = 400$ treated units.

- Write the Horvitz–Thompson IPW estimator of the ATE as a weighted sum of observed outcomes Y_i . What weights do treated units receive? What weights do control units receive?
- Derive an analogous IPW estimator for the ATT. (*Hint:* the ATT averages over the treated distribution; the denominator of the weight for control units should reflect this.)
- Show that when the propensity score is constant ($\pi(X) = p$ for all units), the IPW estimator of the ATE reduces to the difference-in-means estimator. Under what experimental design does $\pi(X) = 0.5$ exactly?

3. Overlap. Let $\pi(X)$ be the true propensity score and suppose overlap fails at $X = x_0$: $\pi(x_0) = 1$.

- For each of τ_{ATE} and τ_{ATT} , identify the exact formula that fails to be identified when $\pi(x_0) = 1$, and explain why.
- Suppose overlap fails only on a set $\mathcal{S}$ with $P(X \in \mathcal{S}) = 0.15$. Define the *trimmed ATE* estimand. How does it differ from the ATE?
- A researcher applies the Crump et al. (2009) rule with $\eta = 0.1$, removing all units with $\hat{\pi}(X_i) < 0.1$ or $\hat{\pi}(X_i) > 0.9$, removing 8% of the sample. (i) List two reasons the trimmed estimator may have lower variance than the untrimmed IPW estimator. (ii) What is the cost in terms of external validity?

4. Sensitivity analysis (conceptual). Suppose you estimate $\hat{\tau}_{\text{ATE}} = 3.2$ using propensity-score methods, and a referee questions whether unconfoundedness holds.

- Define what it means for U to be a “hidden confounder” in the context of the DAG, and explain why its presence invalidates the IPW identification formula.
- Explain in one paragraph what Rosenbaum-style sensitivity analysis would tell you about the robustness of your estimate to hidden confounding.
- Suppose a sensitivity analysis shows that a confounder with an odds ratio of 2 for both treatment and outcome selection would be required to explain away your result. How would you interpret this finding? Does it confirm unconfoundedness?

Chapter 7

Instrumental Variables

Learning Objectives

By the end of this chapter, students should be able to:

1. Diagnose the endogeneity problem: explain why an unobserved confounder makes back-door adjustment fail, and describe how an instrumental variable provides an alternative identification route.
2. State the three IV assumptions — relevance, exogeneity, and exclusion — in the DAG/do-calculus, structural-equation, and potential-outcomes languages, and explain which are testable and which require institutional justification.
3. Follow how the IV assumptions identify the Wald estimand, and interpret the reduced form and first stage as its two observable components.
4. Explain what the Wald estimand identifies when treatment effects are heterogeneous — the Local Average Treatment Effect for compliers — why that estimand depends on the choice of instrument, and when it coincides with the ATE.
5. Compare IV and back-door adjustment on the assumptions each requires, the estimand each identifies, and the ways each can fail.

7.1 Why Instrumental Variables?

Chapter 6 showed how causal effects can be identified when all confounders are observed and can be blocked by conditioning on X . When some confounders are unobserved, back-door adjustment fails. This chapter develops instrumental variables (IV) as an alternative identification strategy: rather than blocking the confounding path $T \leftarrow U \rightarrow Y$, IV exploits an external variable Z whose effect on T is free of confounding by U . This chapter asks what IV identifies and under what assumptions; Chapter 12 asks how that estimand is computed and tested in practice.

7.1.1 The Endogeneity Problem

Chapter 6 established that the propensity score provides a powerful surrogate for randomization, but only under *unconfoundedness*: $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$. This requires that every variable affecting both treatment and outcome is observed. In many empirical settings this is implausible:

- In labor economics, unobserved ability or motivation affects both schooling decisions and wages.
- In epidemiology, unobserved health behaviors affect both treatment uptake and outcomes.
- In survey sampling, voluntary participation is driven by factors that also affect the variables of interest.

Whenever an unobserved confounder U creates a back-door path $T \leftarrow U \rightarrow Y$, the adjustment formula fails. The gap is the *endogeneity bias*.

7.1.2 The IV Idea

An *instrumental variable* Z is an observed variable that: (1) *moves* treatment T (relevance); (2) *does so exogenously* — Z is unrelated to the unobserved confounder U (exogeneity); (3) *affects Y only through T* — Z has no direct path to Y (exclusion).

Under these three conditions, the variation in T induced by Z is free of confounding by U . The ratio of the Z -induced variation in Y to the Z -induced variation in T identifies the causal effect.

IV does not eliminate the confounding path — it identifies causal effects by isolating a component of treatment variation that is induced by Z and is therefore exogenous. IV *avoids* confounding rather than controlling for it. This distinction matters: IV identifies the LATE for compliers, not the ATE, and the exclusion restriction does the heaviest lifting.

Example: Returns to Schooling

A researcher wants to estimate the causal effect of years of schooling T on log wages Y . Unobserved ability U raises both schooling and wages, so OLS is biased upward. An instrumental variable Z that shifts schooling for reasons unrelated to ability — such as proximity to a school or a policy change in compulsory attendance laws — provides a way to isolate the causal effect of schooling on wages (Angrist and Krueger 1991).




Figure 7.1: Two strategies for identification. **Left:** back-door adjustment conditions on the observed X (green) to block the confounding path. **Right:** IV uses only Z -induced variation in T , which is free of U .

7.2 Graphical Setup and Core Assumptions

7.2.1 The IV DAG

The causal structure for a basic IV model with covariates X has nodes Z, T, Y, X, U with edges $Z \rightarrow T$, $T \rightarrow Y$, $X \rightarrow T$, $X \rightarrow Y$, $X \rightarrow Z$, $U \rightarrow T$, $U \rightarrow Y$ (U unobserved). The three IV assumptions correspond to: (1) **Relevance:** directed path $Z \rightarrow T$; (2) **Exogeneity:** $Z \perp\!\!\!\perp U \mid X$ by d-separation; (3) **Exclusion:** no direct arrow $Z \rightarrow Y$.

The distinction between exogeneity and exclusion is fundamental. Exogeneity says the instrument is not confounded with latent causes of the outcome. Exclusion says it has no causal channel to the outcome except through treatment. A randomized encouragement may be exogenous by design, yet still violate exclusion if the encouragement changes outcomes through information, motivation, or stigma apart from the treatment itself.

7.2.2 The Three Assumptions in Three Languages




Figure 7.2: The two main IV assumption failures. **Left:** exogeneity fails when U is a common cause of Z and Y . **Right:** exclusion fails when Z has a direct causal path to Y . The two failures are logically independent.

Assumption	Potential outcomes	Structural / econometric	Do-calculus / DAG
Relevance	$Z \not\perp\!\!\!\perp T$	$\text{Cov}(Z, T) \neq 0$	No d-separation of Z and T in $\mathcal{G}$
Exogeneity	$Z \perp\!\!\!\perp (Y(0), Y(1), T(0), T(1)) \mid X$	$\mathbb{E}[\varepsilon \mid Z, X] = 0$	$Z \perp\!\!\!\perp U \mid X$
Exclusion	$Y_i(t, z) = Y_i(t, z')$ for all z, z'	Z absent from structural equation for Y	$f(y \mid x, \text{do}(T=t), z) = f(y \mid x, \text{do}(T=t))$

Remark: Three Languages, One Identification Argument

The graphical statement encodes causal structure (which arrows are present). The potential-outcomes statement encodes counterfactual independence. The structural formulation encodes moment orthogonality. They support the same identification argument in a well-specified model, but are not literally the same statement.

The do-calculus column is preferred throughout this course because the exclusion restriction in that language is a statement about the *interventional* density. It makes it impossible to confuse exclusion with the much weaker observational statement $f(y \mid x, T=t, z) = f(y \mid x, T=t)$.

7.2.3 Relevance

Definition: Relevance

The instrument Z is **relevant** if it has a non-zero causal effect on the treatment T , possibly after conditioning on covariates X : $\text{Cov}(Z, T \mid X) \neq 0$.

Relevance is the population condition that the instrument changes treatment status. The first-stage F -statistic is a sample diagnostic for whether this link is strong enough; it is evidence about instrument strength, not the assumption itself.

7.2.4 Exogeneity

Definition: Exogeneity

The instrument Z is **exogenous** given covariates X if there is no open back-door path from Z to Y through unobserved causes: $Z \perp\!\!\!\perp U \mid X$ in $\mathcal{G}$ by d-separation.

Exogeneity means the instrument carries no information about latent determinants of the outcome other than through treatment. The moment condition $\mathbb{E}[\varepsilon Z \mid X] = 0$ is a *consequence* of this graphical assumption, not a definition. All versions of exogeneity are untestable from observed data alone.

7.2.5 Exclusion

Definition: Exclusion Restriction

The instrument Z satisfies the **exclusion restriction** if it affects the outcome Y *only* through its effect on treatment T . In do-calculus terms:

$$f(y \mid x, \text{do}(T=t), z) = f(y \mid x, \text{do}(T=t)) \quad \text{for all } z.$$

In potential outcomes terms: $Y_i(t, z) = Y_i(t, z')$ for all $z \neq z'$.

Exclusion is a *causal* restriction, not a conditional-correlation restriction. Adding a direct arrow $Z \rightarrow Y$ to the DAG is the formal counterpart of exclusion failing: in the mutilated graph $\mathcal{G}_{\overline{T}}$, the path $Z \rightarrow Y$ remains open and Rule 1 of the do-calculus can no longer remove Z .

Why Exclusion Is the Hardest Assumption

Relevance can be tested. Exogeneity can sometimes be defended by design. But exclusion — that Z has *no direct effect* on Y whatsoever — is: **untestable** in just-identified models; **easy to violate in practice** (an instrument that moves treatment often also moves other inputs); and **consequential** — the Wald estimand converges to $\beta + \delta/\pi$, so even a small direct effect δ combined with a weak instrument π can produce large bias.

Remark: Three Assumptions Are Necessary but Not Sufficient

(Hernán and Robins 2006; Levis et al. 2024) The three IV assumptions alone do not yield point identification. A fourth structural assumption is required. Two leading choices: **constant treatment effects** (Framework 1, Section 8.7) under which Wald identifies the ATE; and **monotonicity** (Framework 2, Section 7.7) under which Wald identifies the LATE for compliers. An **additive structural mean model (SMM)** identifies the ATT without either. Strikingly, across many fourth assumptions the same conditional Wald formula $\text{Cov}(Z, Y \mid X) / \text{Cov}(Z, T \mid X)$ serves as the identifying expression — but the causal target changes.

The fourth assumption lies *outside* the graphical language in every case. Monotonicity ($T_i(1) \geq T_i(0)$) is a cross-world statement that cannot be encoded by any arrow in the DAG. IV identification cannot be completed by causal graph analysis alone.

7.3 Identification in the Linear Homogeneous-Effect Model

Framework 1: Constant Treatment Effect and Linear Structure

This section works within a linear structural model in which the treatment effect is the *same* for every unit: $Y_i(t) - Y_i(t') = \beta(t - t')$ for all i, t, t' . Under this assumption, $\beta = \text{ATE} = \text{ATT} = \text{LATE}$.

7.3.1 The Linear Structural Model

Consider the linear structural model:

$$Y = \alpha + \beta T + \gamma^\top X + \varepsilon, \quad (7.1)$$

$$T = \pi Z + \delta^\top X + \eta, \quad (7.2)$$

where $\text{Cov}(\varepsilon, \eta) = \rho\sigma\tau \neq 0$. The *reduced form* substitutes Equation 7.2 into Equation 7.1: the reduced-form coefficient on Z is $\beta\pi$.

7.3.2 The OLS Bias

$$\text{plim } \hat{\beta}_{\text{OLS}} = \beta + \frac{\rho\sigma\tau}{\pi^2 \text{Var}(Z) + \tau^2}.$$

The bias is zero only if $\rho = 0$ (no unobserved confounding). The direction depends on the sign of ρ .

7.3.3 Derivation of the Wald Estimand

The derivation has three ingredients: exogeneity, exclusion, and relevance.

1. **Exogeneity** implies $\mathbb{E}[\varepsilon \mid Z, X] = 0$.
2. **Exclusion** (Z absent from the structural equation for Y) combined with exogeneity gives $\mathbb{E}[\varepsilon \cdot Z \mid X] = 0$. In do-calculus terms: Rule 1 removes Z from $f(y \mid x, \text{do}(T=t), z)$ because $Y \perp\!\!\!\perp Z \mid T, X$ in $\mathcal{G}_{\overline{T}}$.
3. Multiplying Equation 7.1 by $(Z - \mathbb{E}[Z \mid X])$ and applying step 2 gives the **reduced-form decomposition** (no-covariate case):

$$\text{Cov}(Y, Z) = \beta \text{Cov}(T, Z).$$

4. **Relevance** ensures the denominator is non-zero:

$$\beta = \frac{\text{Cov}(Y, Z \mid X)}{\text{Cov}(T, Z \mid X)}.$$

For binary $Z \in \{0, 1\}$ without covariates, this is the **Wald estimand**:

The Wald Estimand

$$\beta = \frac{\mathbb{E}[Y \mid Z=1] - \mathbb{E}[Y \mid Z=0]}{\mathbb{E}[T \mid Z=1] - \mathbb{E}[T \mid Z=0]}.$$

The numerator is the *reduced form*: the total effect of Z on Y . The denominator is the *first stage*: the effect of Z on T .

The key lesson is that the Wald ratio identifies β in this section *only* because treatment effects are assumed to be constant across units. In Section 7.7, the same ratio will survive, but its interpretation changes from a common causal effect to a complier-specific average effect.

7.4 Why the IV Assumptions Matter

Each assumption failure produces a distinct, quantifiable distortion in the Wald estimand, amplified by a weak first stage.

When relevance fails. If $\pi = 0$, the Wald estimand is undefined. When π is small, IV variance diverges and finite-sample bias pulls the estimate toward OLS at a rate $\propto 1/F$.

When exclusion fails. If $Y = \alpha + \beta T + \delta Z + \gamma^\top X + \varepsilon$, the Wald estimand converges to $\beta + \delta/\pi$. A small direct effect combined with a weak instrument can produce large bias.

When exogeneity fails. If $\text{Cov}(Z, \varepsilon \mid X) \neq 0$, the Wald estimand converges to $\beta + \text{Cov}(\varepsilon, Z)/\text{Cov}(T, Z)$ — amplified by a weak first stage.

Assumption	Directly testable?	Basis for assessment
Relevance	Yes	First-stage F -statistic
Exogeneity	No	Institutional knowledge; randomization; placebo regressions
Exclusion	No (just-ID); partially (overID)	Institutional argument; J -test checks consistency

7.5 Lab: OLS vs. IV Across Instrument Strengths

This lab verifies the OLS bias formula numerically and traces the bias–variance tradeoff across the full range of instrument strength.

DGP for Lab 7

$Y_i = \beta T_i + \varepsilon_i$, $T_i = \pi Z_i + \eta_i$, with $Z_i \sim N(0, 1)$, $\eta_i = U_i$, $\varepsilon_i = \rho U_i + \sqrt{1 - \rho^2} \xi_i$, $U_i, \xi_i \sim N(0, 1)$ i.i.d. Fixed parameters: $\beta = 1$, $\rho = 0.8$. Grid: $\pi \in \{0, 0.10, 0.15, 0.20, 0.30, 0.50, 1.00, 2.00\}$. Expected $F \approx 500\pi^2$. All three IV assumptions hold by construction.

π	F	Theory bias	OLS Bias	OLS RMSE	IV Bias	IV SD	IV RMSE
0.00	1	+0.800	+0.801	0.802	—	—	—
0.10	6	+0.792	+0.791	0.792	−0.128	4.391	4.392
0.15	13	+0.782	+0.782	0.783	−0.236	6.487	6.491
0.20	21	+0.769	+0.769	0.770	−0.060	0.380	0.385
0.30	46	+0.734	+0.735	0.735	−0.025	0.163	0.165
0.50	127	+0.640	+0.640	0.640	−0.004	0.091	0.091
1.00	500	+0.400	+0.401	0.402	−0.001	0.045	0.045
2.00	2006	+0.160	+0.160	0.161	0.000	0.022	0.022

True $\beta = 1.000$. $n = 500$, $B = 2,000$ replications. IV mean and SD at $\pi = 0.10$ and 0.15 are dominated by extreme outliers; the IV median is ≈ 1.01 at both values.

Four lessons stand out.

- 1. The OLS bias formula is exact.** The theory bias column ($0.8/(\pi^2 + 1)$) matches the simulated OLS bias to four decimal places. OLS bias decreases as π grows but is never zero.
- 2. IV is consistent but has catastrophically heavy tails when the instrument is weak.** At $F \approx 6$ – 13 , the IV median is ≈ 1.01 (confirming consistency), but a small fraction of replications in which the first stage happens to be near zero inflates the SD to 4.4 – 6.5 . The weak instrument problem is uncontrolled tail behavior, not inconsistency.
- 3. The RMSE crossover occurs near $F \approx 20$.** IV first beats OLS at $\pi = 0.20$ ($F \approx 21$): IV RMSE = $0.385 < \text{OLS RMSE} = 0.770$. The Staiger–Stock rule of thumb ($F \geq 10$) is therefore slightly too lenient; $F \geq 20$ – 25 is more appropriate in this DGP.
- 4. A strong instrument eliminates both problems.** At $\pi = 1.00$ ($F \approx 500$): IV RMSE = 0.045 vs. OLS RMSE = 0.402 — a 9-fold improvement from IV.

RMSE Cannot Be the Only Criterion

OLS has lower RMSE than IV for $\pi < 0.20$ in this DGP. This does not vindicate OLS. An estimator with RMSE = 0.78 because it is biased by 0.80 is useless for causal inference: the bias is systematic and does not shrink with sample size. As $n \rightarrow \infty$, IV RMSE shrinks to zero while OLS RMSE stays at 0.80 . The RMSE comparison is meaningful only in finite samples; it quantifies when a weak instrument is not yet practically useful, not when OLS is acceptable.

Remark: The First-Stage F -Statistic as a Diagnostic

(Bound et al. 1995) The Staiger–Stock threshold $F \geq 10$ corresponds to IV RMSE within a factor of roughly two of the strong-instrument limit. The simulation here suggests this threshold may understate the problem when endogeneity is strong ($\rho = 0.8$): at $F = 13$ the IV RMSE is 6.5, far above any useful threshold. A researcher should report the first-stage F alongside any IV estimate, and treat values below 20–25 with particular caution.

7.6 Multiple Instruments and Overidentification

When there are exactly as many instruments as endogenous variables ($q = p$), the model is *just-identified*. When $q > p$, the model is *overidentified*: the extra instruments impose testable moment restrictions. The Sargan–Hansen J -test (Sargan 1958; Hansen 1982) checks mutual consistency of instruments; rejection indicates at least one instrument violates exogeneity or exclusion, but does not identify which one. The test statistic and distribution are derived in Chapter 12.

What Overidentification Does Not Do

Passing the J -test does not confirm any instrument is valid — only that sample moment conditions are mutually compatible. Multiple invalid instruments can agree if they share the same violation. Overidentification is an opportunity for a consistency check, not a substitute for the institutional argument.

7.7 Heterogeneous Treatment Effects and the LATE Framework

Framework 2: Heterogeneous Treatment Effects

We now drop homogeneity and allow $\tau_i = Y_i(1) - Y_i(0)$ to vary across units, in the binary instrument, binary treatment setting ($Z, T \in \{0, 1\}$). Under an additional assumption of monotonicity, the Wald ratio identifies the *Local Average Treatment Effect* (LATE): the average treatment effect for the subpopulation whose treatment status is changed by the instrument.

7.7.1 Compliance Types

Definition: Compliance Types

(Angrist et al. 1996) For binary $Z, T \in \{0, 1\}$, define $T_i(z)$ as the treatment unit i would take under instrument value z . The four compliance types are:

Type	$T_i(0)$	$T_i(1)$
Complier	0	1
Always-taker	1	1
Never-taker	0	0
Defier	1	0

The instrument Z only shifts treatment for *compliers*: always-takers and never-takers have the same treatment status regardless of Z , contributing nothing to the denominator.

7.7.2 The Monotonicity Assumption

Definition: Monotonicity

(Angrist and Imbens 1994) The treatment assignment is **monotone in Z** if: $T_i(1) \geq T_i(0)$ for all i (no defiers).

7.7.3 The LATE Theorem

Theorem: LATE

(Angrist and Imbens 1994) Under the three core IV assumptions and monotonicity, the Wald estimand identifies the **Local Average Treatment Effect**:

$$\frac{\mathbb{E}[Y | Z=1] - \mathbb{E}[Y | Z=0]}{\mathbb{E}[T | Z=1] - \mathbb{E}[T | Z=0]} = \mathbb{E}[Y(1) - Y(0) | T_i(1) > T_i(0)] \equiv \tau_{\text{LATE}}.$$

Proof

Denominator. By exogeneity ($Z \perp\!\!\!\perp T(0), T(1)$):

$$\mathbb{E}[T | Z=1] - \mathbb{E}[T | Z=0] = \mathbb{E}[T(1) - T(0)] = P(\text{complier}),$$

since always-takers contribute $1 - 1 = 0$, never-takers $0 - 0 = 0$, and no defiers exist.

Numerator. By exogeneity and law of total expectation over compliance types:

$$\mathbb{E}[Y | Z=1] - \mathbb{E}[Y | Z=0] = \mathbb{E}[Y(T(1))] - \mathbb{E}[Y(T(0))].$$

Always-takers and never-takers contribute 0; compliers contribute $Y(1) - Y(0)$. Therefore:

$$\mathbb{E}[Y | Z=1] - \mathbb{E}[Y | Z=0] = P(\text{complier}) \mathbb{E}[Y(1) - Y(0) | \text{complier}].$$

Ratio: Dividing numerator by denominator yields τ_{LATE} . $\square$

7.8 Interpreting IV Estimands

	Framework 1 (homogeneous)	Framework 2 (heterogeneous)
Key assumption	$\tau_i = \beta$ for all i	Monotonicity; no defiers
What IV identifies	$\beta = \text{ATE} = \text{ATT} = \text{LATE}$	$\tau_{\text{LATE}} = \mathbb{E}[\tau_i \text{complier}]$
Estimand depends on Z ?	No	Yes (different $Z \Rightarrow$ different compliers $\Rightarrow$ different LATE)
Identifies ATE?	Yes	Only under additional homogeneity

7.8.1 When Does LATE Equal ATE?

ATE decomposes by compliance type:

$$\text{ATE} = P(\text{co}) \mathbb{E}[\tau_i | \text{co}] + P(\text{at}) \mathbb{E}[\tau_i | \text{at}] + P(\text{nt}) \mathbb{E}[\tau_i | \text{nt}].$$

LATE equals only the first term divided by its probability weight. ATE and LATE coincide if and only if: (1) effects are homogeneous, or (2) everyone is a complier, or (3) effects for always-takers and never-takers happen to equal the LATE. In practice, none of these is likely to hold exactly.

7.8.2 Different Instruments, Different Estimands

Because the LATE is specific to the complier population, two valid instruments for the same treatment can legitimately identify different LATEs. This is informative about treatment effect heterogeneity, not a contradiction.

Example: Compulsory Schooling Laws vs. Distance to College

Two classic instruments for years of schooling: compulsory schooling laws (Angrist and Krueger 1991) select compliers at the margin of dropping out (typically lower-income students); distance to the nearest college (Card 1995) selects students deterred by geography. Both identify the return to schooling for their respective complier populations, and the LATEs can legitimately differ even if both instruments are valid.

7.8.3 The Policy Relevance of LATE

For many policy questions, the LATE is exactly the right estimand — when the contemplated intervention resembles the instrument, the complier population under the study design is close to the policy-relevant margin. When the ATE over the full population is required, IV alone is insufficient under heterogeneous effects; additional structure or a second instrument is needed.

7.9 IV versus Back-Door Adjustment

Dimension	Back-door / propensity score	Instrumental variables
Core assumption	All confounders observed	Valid instrument: relevance, exogeneity, exclusion
Unobserved confounders	Fatal	Permitted
Estimand	ATE, ATT, or ATC	LATE (compliers only)
Testability	Unconfoundedness untestable; overlap testable	Relevance testable; exogeneity/exclusion untestable (just-ID)
Main threat	Unmeasured confounder	Exclusion restriction violation
Identifies ATE?	Yes, under strong ignorability	Only under homogeneous effects

The two strategies have complementary failure modes. Back-door adjustment fails when X does not capture all confounders. IV fails when the exclusion restriction is violated, with bias δ/π amplified by a weak first stage. They are tools for different identification problems, not competitors.

Five questions for any IV application: (1) What exactly is the instrument? (2) Why does it shift treatment? (3) Why is it as-if random relative to latent outcome determinants? (4) Why can it affect the outcome only through treatment? (5) What population margin does it shift?

7.10 Applied Example: Charter School Lotteries and the KIPP Lynn Study

This example illustrates a canonical randomized-encouragement IV design (Angrist et al. 2012). The lottery randomizes *offer status*, not actual treatment. Z = lottery offer; S = years of KIPP attendance (endogenous); Y = standardized MCAS test score.

Relevance. First-stage coefficient ≈ 1.2 : lottery winners had spent 1.2 more years at KIPP at the time of each MCAS exam. First-stage F is far above conventional thresholds.

Exogeneity. Guaranteed by the randomization protocol. Balance test: p -value = 0.615 for demographic characteristics and baseline test scores.

Exclusion. The lottery offer merely provides access — it does not deliver instruction. One potential violation is a discouragement effect on losers; the authors note that lottery losers' scores are typical of demographically comparable students in Lynn.




Figure 7.3: Causal DAG for the KIPP Lynn lottery IV design. Exogeneity is guaranteed by the randomization protocol; exclusion rests on the institutional argument that the lottery offer affects achievement only through attendance.

Results ((Angrist et al. 2012), Table 4, specification with demographic and baseline-score controls):

Subject	First stage	Reduced form	2SLS
Math	1.221 (0.068)	0.430 (0.067)	0.352 (0.053)
ELA	1.228 (0.068)	0.164 (0.073)	0.133 (0.059)

SE clustered at the student level. $N = 833$ student-by-test observations.

Each year at KIPP raises math scores by approximately 0.35σ and ELA scores by approximately 0.13σ . The 2SLS estimand identifies the LATE for *lottery compliers* only. Reading gains are driven almost entirely by LEP students ($\approx 0.43\sigma$) and SPED students ($\approx 0.27\sigma$).

7.11 Summary

- IV identifies effects from exogenous treatment variation.** When back-door adjustment fails because U creates a path $T \leftarrow U \rightarrow Y$, a valid instrument Z identifies the causal effect by exploiting only the component of treatment variation that Z induces. IV avoids confounding rather than controlling for it.
- Three assumptions, ordered by testability.** Relevance is assessable with the first-stage F . Exogeneity and exclusion must be defended by institutional knowledge and causal structure. Each violated assumption produces a distinct, quantifiable bias in the Wald estimand, amplified by weak instruments.
- Framework 1 (homogeneous effects):** Under constant treatment effects and the linear structural model, the Wald estimand identifies $\beta = \text{ATE} = \text{ATT} = \text{LATE}$.
- Framework 2 (heterogeneous effects + monotonicity):** The Wald estimand identifies $\tau_{\text{LATE}} = \mathbb{E}[\tau_i \mid \text{complier}]$ — the effect for the subpopulation whose treatment status is changed by the instrument.
- Different instruments identify different effects.** The LATE depends on the instrument through the complier population it selects. LATE equals ATE only under additional structure, most notably treatment-effect homogeneity.
- IV versus back-door adjustment.** Complementary failure modes: back-door fails when confounders are unobserved; IV fails when exclusion is violated. They are tools for different identification problems.
- Applied example: KIPP Lynn lottery.** The lottery guarantees exogeneity by design; exclusion is defended institutionally; first-stage coefficient ≈ 1.2 . Estimates identify the LATE for lottery compliers: 0.35σ per year in math, 0.13σ in ELA.

8. **Estimation deferred to Chapter 12.** This chapter establishes what IV identifies and under what assumptions. How the Wald ratio is estimated — 2SLS, asymptotic inference, overidentification tests — is the subject of Chapter 12.

7.12 Problems

1. The three IV assumptions in three languages. Consider the DAG $\{Z \rightarrow T, T \rightarrow Y, U \rightarrow T, U \rightarrow Y, X \rightarrow T, X \rightarrow Y, X \rightarrow Z\}$ with U unobserved.

- List all back-door paths from T to Y . Does X alone satisfy the back-door criterion? Explain.
- Verify the three IV assumptions using d-separation: (i) Relevance: show Z and T are not d-separated in $\mathcal{G}$. (ii) Exogeneity: show $Z \perp\!\!\!\perp U \mid X$ in $\mathcal{G}$. (iii) Exclusion: show $Y \perp\!\!\!\perp Z \mid T, X$ in $\mathcal{G}_{\overline{T}}$.
- Now add the arrow $Z \rightarrow Y$ to the DAG. Which IV assumption is violated? Show explicitly which step of the Wald derivation breaks down.
- Translate each of the three IV assumptions into the structural language: write the equations for T and Y and identify which coefficient restriction corresponds to each assumption.

2. Bias under assumption violations. Let $Y = \beta T + \varepsilon$ and $T = \pi Z + \eta$ with $\mathbb{E}[\varepsilon \mid Z] = 0$ and $\pi \neq 0$.

- Starting from $\mathbb{E}[Y \mid Z=1] - \mathbb{E}[Y \mid Z=0]$, substitute the structural equation for Y and simplify. What role does exogeneity play?
- Show that $\mathbb{E}[T \mid Z=1] - \mathbb{E}[T \mid Z=0] = \pi$ in the linear first-stage model.
- Derive the Wald estimand and confirm it equals β .
- Now suppose exclusion fails: $Y = \beta T + \delta Z + \varepsilon$ with $\delta \neq 0$. Derive the probability limit of the Wald estimator and confirm the bias formula from Section 7.4.
- Suppose instead exogeneity fails: $\mathbb{E}[\varepsilon \mid Z] = cZ$ for some constant $c \neq 0$. Derive the probability limit of the Wald estimator and express the bias in terms of c and π .

3. Order, rank, and the limits of overidentification. Consider a model with one endogenous variable T and two instruments Z_1 and Z_2 , both satisfying exogeneity and exclusion.

- State the order condition and verify it is satisfied.
- State the rank condition. What would it mean geometrically if Z_1 and Z_2 were perfectly collinear in the first-stage regression?
- Explain intuitively why having two valid instruments rather than one should improve estimation precision.
- Now suppose Z_1 is valid but Z_2 violates exclusion. Under what conditions does the Sargan–Hansen J -test have power to detect Z_2 's invalidity? Under what conditions does the test fail?
- Why does passing the J -test not confirm that both instruments are valid? Give a concrete example where both are invalid and the J -test has no power.

4. Compliance types and the LATE. In a binary instrument, binary treatment study: 30% compliers with $\tau_c = 6$; 25% always-takers with $\tau_a = 3$; 45% never-takers with $\tau_n = 1$; no defiers.

- Compute $P(\text{complier}) = \mathbb{E}[T \mid Z=1] - \mathbb{E}[T \mid Z=0]$.
- Compute the ATE as a weighted average of τ_c, τ_a, τ_n with appropriate weights.
- The Wald estimand equals $\tau_c = 6$. By how much does this overstate the ATE, and why?
- A second study uses a different binary instrument Z' with a complier population of 50% and a LATE of 2. Is this contradictory? What can you infer about relative treatment effects in the two complier populations?
- Explain, using compliance type language, why the denominator of the Wald estimand equals $P(\text{complier})$.

5. The exclusion restriction: plausibility and violations. For each proposed instrument, (i) state whether the restriction is plausible and why; (ii) describe a specific mechanism by which it could be violated; (iii) assess whether the violation would bias the IV estimate upward or downward.

- Instrument:** rainfall in the home region of a politician, used as an instrument for government infrastructure spending. **Outcome:** local economic growth.
- Instrument:** distance to the nearest hospital, used as an instrument for hospital admission. **Outcome:** 30-day mortality.

- (c) **Instrument:** a randomly assigned financial incentive to enroll in a health screening program. **Outcome:** health status two years later.
- (d) **Instrument:** lottery number in the Vietnam-era draft lottery, used as an instrument for military service. **Outcome:** lifetime earnings. (*This is the Angrist (1990) study; discuss why this instrument is widely regarded as satisfying the exclusion restriction.*)

6. IV versus back-door adjustment. A researcher studies the effect of job training (T) on earnings (Y). Two strategies: (A) a rich set of pre-treatment covariates X and a propensity-score estimator; (B) a lottery that randomly selected units to be offered training, used as an instrument Z .

- (a) Under what assumption does strategy (A) identify the ATE? What specific unobserved variable would most plausibly violate this assumption?
- (b) Strategy (B) identifies a LATE. Describe the complier population in words. Is the LATE likely to be larger or smaller than the ATE in this setting?
- (c) Both strategies yield estimates of \$1,800 and \$2,400 per year. Describe a Hausman-type test. Under what null hypothesis does the test have an approximate χ^2 distribution?
- (d) If the two estimates differ significantly, which strategy would you trust more and why? What additional evidence would help distinguish endogeneity bias in (A) versus $\text{LATE} \neq \text{ATE}$ in (B)?

Chapter 8

Mediation and Front-Door Identification

Learning Objectives

By the end of this chapter, students should be able to:

1. Explain the distinction between the total causal effect of T on Y and a pathway-specific effect that operates through a mediator M , and describe why this distinction matters scientifically.
2. Draw the prototype mediation DAG, write its structural equations, and identify the direct and indirect pathways.
3. Define the controlled direct effect (CDE) using the do-operator, identify it via the back-door formula for the joint intervention (T, M) , and explain why it depends on the fixed level m .
4. Define the natural direct and indirect effects (NDE, NIE) using the potential outcomes notation $Y(t, M(t'))$, state the $\text{NDE} + \text{NIE} = \text{TE}$ decomposition, and explain why these quantities involve cross-world counterfactuals.
5. State the four sequential ignorability assumptions for identification of natural effects, write the mediation formula, and identify which assumption is violated when there is an unmeasured treatment-induced mediator–outcome confounder.
6. Set up the Baron–Kenny three-equation system, derive the product and difference formulas for the indirect effect, and explain why the decomposition fails in nonlinear or interaction models.
7. State the three front-door conditions, derive the front-door formula using the do-calculus, and explain why the front-door graph enables identification despite unobserved T – Y confounding.
8. Contrast mediation analysis and instrumental variables on the dimensions of variable position, identification goal, and key assumption.

8.1 Motivation: Mechanisms

The identification results of Chapters 5–7 all answer the same question: *what is the total causal effect of T on Y ?* Mediation analysis asks a finer question: *through what mechanism does that effect operate?*

More concretely, the total effect of T on Y may flow along multiple causal pathways. Some of this effect passes through an intermediate variable M — the *mediator* — along the path $T \rightarrow M \rightarrow Y$. The remainder flows directly along $T \rightarrow Y$, bypassing the mediator entirely. Decomposing the total effect into these two components is the goal of mediation analysis.

This decomposition matters for scientific and policy reasons. In a clinical trial of a behavioral intervention (T) on depression (Y), a researcher may want to know how much of the benefit operates through improved sleep quality (M) versus other pathways — because if sleep is the main channel, targeting sleep directly may be a more efficient intervention. In an economics study of education (T) on wages (Y), how much operates through occupation (M) versus cognitive skills?

Running example. Throughout this chapter we use the depression intervention as the primary example: T is a randomized behavioral intervention, M is self-reported sleep quality measured mid-trial, and Y is a

depression score (e.g., PHQ-9 scale).

The challenge is that mediators are *post-treatment variables*: they are affected by the treatment, and may themselves be confounded with the outcome. Conditioning on a post-treatment variable creates exactly the collider and selection-bias problems studied in Chapters 2 and 3. A naive approach — simply including M as a covariate in a regression of Y on T — conflates adjustment with mediation and can introduce bias even in a randomized experiment.

This chapter also develops the *front-door criterion*, a distinct identification strategy that uses the mediation structure of the DAG to identify causal effects even when treatment and outcome are confounded by an unobserved variable.

8.2 The Mediation DAG

8.2.1 The Prototype Graph




Figure 8.1: The prototype mediation DAG. The causal effect of T on Y operates through two pathways: the direct path $T \rightarrow Y$ and the indirect path $T \rightarrow M \rightarrow Y$ via the mediator. The unobserved variable U confounds the treatment–outcome relationship; $\mathbf{X}$ denotes observed pre-treatment covariates.

The graph encodes two causal pathways: the **direct pathway** $T \rightarrow Y$ (treatment affects the outcome without passing through the mediator) and the **indirect pathway** $T \rightarrow M \rightarrow Y$ (treatment first shifts the mediator, which shifts the outcome).

The prototype graph corresponds to the nonparametric SEM:

$$T = f_T(\mathbf{X}, U, \varepsilon_T), \quad M = f_M(T, \mathbf{X}, \varepsilon_M), \quad Y = f_Y(T, M, \mathbf{X}, U, \varepsilon_Y).$$

8.2.2 What Makes Mediation Harder Than Total Effect Estimation

The key difficulty is that M is a *post-treatment* variable. This creates two interrelated problems.

Collider bias. Conditioning on M can open collider paths. If $V \rightarrow M$ and $V \rightarrow Y$ with V unobserved, the path $T \rightarrow M \leftarrow V \rightarrow Y$ is blocked when M is not conditioned on, but opens as soon as M is included as a covariate. This is precisely why naively regressing Y on (T, M) does not isolate the direct effect.

Mediator–outcome confounding. Even when treatment is randomized, the mediator M is never randomized. An unobserved V with $V \rightarrow M$ and $V \rightarrow Y$ creates a back-door path from M to Y that randomization of T does not close.

Post-Treatment Variables Are Dangerous to Condition On

The mediator M is caused by the treatment T . This single fact makes every operation involving M structurally different from operations involving pre-treatment covariates $\mathbf{X}$.

Conditioning on a post-treatment variable in a regression or matching procedure is *not* a neutral act. It can open collider paths that were previously blocked, introduce selection bias, and produce estimates of neither the total effect nor any well-defined direct effect. This chapter studies three regimes in which conditioning on or marginalizing over M is justified: (1) **CDE**: both T and M

are intervened on simultaneously via $\text{do}(T, M)$; (2) **NDE/NIE**: M is marginalized over using a distribution from a different treatment arm, justified by sequential ignorability; (3) **Front-door**: M is summed over in the front-door formula, justified by graphical conditions. Outside these three regimes, conditioning on a post-treatment variable should be treated as an error until proven otherwise.

8.2.3 Working Assumption for §§ 3–6

Throughout §§ Section 8.3–Section 8.7, we work with the *reduced prototype graph* (obtained from Figure 8.1 by assuming all T – Y and T – M confounding is captured by the observed $\mathbf{X}$). The unobserved confounder U is restored in Section 8.8.

8.3 Total Causal Effect

Definition: Total Effect

The **total effect** (TE) of T on Y is:

$$\text{TE} = \mathbb{E}[Y(1) - Y(0)] = \mathbb{E}[Y \mid \text{do}(T=1)] - \mathbb{E}[Y \mid \text{do}(T=0)].$$

The total effect captures the combined impact of all causal pathways — both $T \rightarrow Y$ and $T \rightarrow M \rightarrow Y$. It is identified by the back-door formula whenever $\mathbf{X}$ satisfies the back-door criterion for the effect of T on Y .

Running example. For the depression trial, the TE is the expected change in depression score if the entire study population were assigned to the intervention versus control. It combines the effect operating through sleep improvement with every other pathway. Mediation analysis asks: of this total, how much is due to sleep?

8.4 Controlled Direct Effect

8.4.1 Definition

The do-calculus provides the cleanest definition of the direct effect: intervene on both T and M simultaneously, fixing M at a specified level m . Any remaining effect of T on Y must then flow exclusively through the direct edge $T \rightarrow Y$.

Definition: Controlled Direct Effect (CDE)

The **controlled direct effect** of changing T from 0 to 1 while fixing $M = m$ by intervention is:

$$\text{CDE}(m) = \mathbb{E}[Y \mid \text{do}(T=1), \text{do}(M=m)] - \mathbb{E}[Y \mid \text{do}(T=0), \text{do}(M=m)]. \quad (8.1)$$

This corresponds to the mutilated graph $\mathcal{G}_{\overline{TM}}$ in which all edges into both T and M are deleted. Fixing $M = m$ for everyone shuts down the indirect pathway $T \rightarrow M \rightarrow Y$.

Remark

The CDE depends on the level m at which M is fixed. In a linear additive model, $\text{CDE}(m) = \tau'$ for all m , a special property of linearity. When T and M interact in their effect on Y , the CDEs at different values of m differ, and no single number summarizes the direct effect.

8.4.2 Identification of the CDE

Theorem: Identification of the CDE

Suppose $\mathbf{Z}$ satisfies the back-door criterion for the joint intervention $\text{do}(T, M)$ on Y : (i) $\mathbf{Z}$ contains no descendant of T or M , and (ii) $\mathbf{Z}$ blocks every back-door path from $\{T, M\}$ to Y . Suppose further (iii) joint positivity: $P(T=t, M=m \mid \mathbf{Z}=\mathbf{z}) > 0$ for P -almost every $\mathbf{z}$. Then:

$$\mathbb{E}[Y \mid \text{do}(T=t), \text{do}(M=m)] = \sum_{\mathbf{z}} \mathbb{E}[Y \mid t, m, \mathbf{z}] P(\mathbf{z}). \quad (8.2)$$

Under the working assumption, $\mathbf{Z} = \mathbf{X}$ blocks every back-door path from $\{T, M\}$ to Y , so:

$$\mathbb{E}[Y \mid \text{do}(T=t), \text{do}(M=m)] = \sum_{\mathbf{x}} \mathbb{E}[Y \mid t, m, \mathbf{x}] P(\mathbf{x}).$$

8.4.3 The CDE Does Not Have a Complementary Indirect Effect

A common misconception is that the CDE and some “controlled indirect effect” sum to the total effect. This is not generally true. The residual quantity $\text{TE} - \text{CDE}(m)$ depends on m , has no clean do-calculus expression corresponding to “the indirect pathway,” and cannot be interpreted as the effect that flows through M . The correct complementary decomposition uses the natural effects of Section 8.5.

Natural Effects Are a Strictly Harder Estimand Than the CDE

The CDE is a single-world estimand. The quantity $\mathbb{E}[Y \mid \text{do}(T=t), \text{do}(M=m)]$ involves a single joint intervention. Identification reduces to the standard back-door criterion for the pair (T, M) .

The NDE and NIE are cross-world estimands. The nested counterfactual $Y(t, M(t'))$ requires an individual to simultaneously inhabit two worlds: the world where $T = t$ (which determines Y) and the world where $T = t'$ (which determines what M would *naturally* be). When $t \neq t'$, these are logically distinct states of the world that no single experiment can realize. The do-operator alone cannot express this quantity.

The identification price is steep. The CDE requires only the back-door criterion. The NDE and NIE require four sequential ignorability assumptions, including Assumption 4 (no treatment-induced mediator–outcome confounder), which cannot be satisfied by randomizing T and cannot be verified from data.

8.5 Natural Direct and Indirect Effects

8.5.1 Motivation: Cross-World Counterfactuals

The CDE fixes the mediator by external intervention. A more scientifically natural question is: what is the effect of T on Y that bypasses M when M is held at the value it *would naturally take* under the reference treatment $T = 0$?

This requires the *nested potential outcomes* notation $Y(t, M(t'))$, which denotes the outcome that would be observed if T were set to t and M were simultaneously set to the value it would naturally take if T were t' . These are called *cross-world counterfactuals* because t and t' may differ, referring to two different hypothetical states of the world simultaneously.

8.5.2 Definitions

Definition: Natural Direct and Indirect Effects

(Pearl 2001) For binary $T \in \{0, 1\}$:

$$\text{NDE} = \mathbb{E}[Y(1, M(0)) - Y(0, M(0))], \quad (8.3)$$

$$\text{NIE} = \mathbb{E}[Y(1, M(1)) - Y(1, M(0))]. \quad (8.4)$$

The **natural direct effect** (NDE) is the expected change in the outcome when T shifts from 0 to 1, holding the mediator at the value it would naturally take under $T = 0$. The **natural indirect**

effect (NIE) is the expected change in the outcome due to the shift in the mediator from $M(0)$ to $M(1)$, holding $T = 1$ fixed.

Running example. The NDE is the expected reduction in depression score when the intervention is delivered but each participant's sleep is held at the value it *would have had* without the intervention — the part of the benefit from cognitive, behavioral, or therapeutic-alliance pathways, not from sleep. The NIE is the complementary piece: how much of the depression reduction is due to the sleep improvement that the intervention itself causes.

Decomposition. The total effect decomposes as:

$$TE = NDE + NIE. \quad (8.5)$$

Proof

$$\begin{aligned} NDE + NIE &= \mathbb{E}[Y(1, M(0)) - Y(0, M(0))] + \mathbb{E}[Y(1, M(1)) - Y(1, M(0))] \\ &= \mathbb{E}[Y(1, M(1)) - Y(0, M(0))] = \mathbb{E}[Y(1) - Y(0)] = TE. \quad \square \end{aligned}$$

Remark

In the linear additive Baron–Kenny SEM with no $T \times M$ interaction, the CDE and NDE coincide: $NDE = \tau'$ and $NIE = ab$. When there is a $T \times M$ interaction, $NDE \neq CDE$ and the linear decomposition breaks down.

8.6 Identification of Natural Effects

8.6.1 Sequential Ignorability

Identification of the NDE and NIE requires four conditions, known collectively as *sequential ignorability* (Imai et al. 2010). These conditions are substantially stronger than the ignorability assumptions used for the total effect or the CDE.

Sequential Ignorability Assumptions

1. **No unmeasured treatment–outcome confounding.** $\{Y(t', m), M(t)\} \perp\!\!\!\perp T \mid \mathbf{X}$ for all t, t', m . Graphically: $\mathbf{X}$ blocks all back-door paths from T to Y and from T to M .
2. **No unmeasured treatment–mediator confounding.** Implied by Assumption 1 for the $T \rightarrow M$ sub-problem.
3. **No unmeasured mediator–outcome confounding given T .** $Y(t', m) \perp\!\!\!\perp M \mid T, \mathbf{X}$ for all t', m . Graphically: $(T, \mathbf{X})$ blocks all back-door paths from M to Y .
4. **No treatment-induced mediator–outcome confounder.** There is no post-treatment variable L such that $T \rightarrow L$, $L \rightarrow M$, and $L \rightarrow Y$.

Positivity (Overlap) Conditions

- **(P1) Treatment overlap.** $P(T = t \mid \mathbf{X} = \mathbf{x}) > 0$ for all $t \in \{0, 1\}$ and for P -almost every $\mathbf{x}$.
- **(P2) Mediator overlap across treatment arms.** $P(M = m \mid T = t, \mathbf{X} = \mathbf{x}) > 0$ whenever $P(M = m \mid T = t', \mathbf{X} = \mathbf{x}) > 0$.

What randomization of T does and does not provide. Randomizing T satisfies Assumptions 1 and 2 by design. It does *not* satisfy Assumptions 3 or 4. The mediator M is a post-treatment variable that is never randomized; any unobserved V with $V \rightarrow M$ and $V \rightarrow Y$ violates Assumption 3 regardless of how T was assigned. Assumption 4 can be violated by a variable L that is itself *caused* by the treatment, so randomizing T actually creates the conditions under which Assumption 4 violations can arise.

Randomization of T Does Not Secure Assumption 4

Unlike unmeasured T – Y confounders (which randomization of T can eliminate), a treatment-induced mediator–outcome confounder L is itself caused by the treatment: the path $T \rightarrow L$ is set in motion by the intervention. The absence of such an L must be argued from design, timing, measurement, or subject-matter knowledge. This is why identification of natural effects is considered strictly harder than identification of the total effect or the CDE.

8.6.2 The Mediation Formula

Theorem: Mediation Formula

(Pearl 2001) Under sequential ignorability, positivity (P1)–(P2), and with discrete M and $\mathbf{X}$:

$$\mathbb{E}[Y(t, M(t'))] = \sum_m \sum_{\mathbf{x}} \mathbb{E}[Y \mid T=t, M=m, \mathbf{X}=\mathbf{x}] P(M=m \mid T=t', \mathbf{X}=\mathbf{x}) P(\mathbf{X}=\mathbf{x}). \quad (8.6)$$

Proof Sketch (Five Steps)

Step 1 (law of total expectation): $\mathbb{E}[Y(t, M(t'))] = \sum_{m, \mathbf{x}} \mathbb{E}[Y(t, m) \mid M(t')=m, \mathbf{X}=\mathbf{x}] P(M(t')=m \mid \mathbf{X}=\mathbf{x}) P(\mathbf{X}=\mathbf{x})$.

Step 2 (Assumption 4: cross-world independence): Within the NPSEM-IE framework, no-treatment-induced-confounder implies $Y(t, m) \perp\!\!\!\perp M(t') \mid \mathbf{X}$, so conditioning on $M(t')$ can be dropped: $\mathbb{E}[Y(t, m) \mid M(t')=m, \mathbf{X}=\mathbf{x}] = \mathbb{E}[Y(t, m) \mid \mathbf{X}=\mathbf{x}]$.

Step 3 (Assumptions 1 and 3): By $Y(t, m) \perp\!\!\!\perp T \mid \mathbf{X}$ and $Y(t, m) \perp\!\!\!\perp M \mid T, \mathbf{X}$: $\mathbb{E}[Y(t, m) \mid \mathbf{X}=\mathbf{x}] = \mathbb{E}[Y(t, m) \mid T=t, M=m, \mathbf{X}=\mathbf{x}]$.

Step 4 (consistency for Y): $\mathbb{E}[Y(t, m) \mid T=t, M=m, \mathbf{X}=\mathbf{x}] = \mathbb{E}[Y \mid T=t, M=m, \mathbf{X}=\mathbf{x}]$.

Step 5 (Assumption 2 and consistency for M): $P(M(t')=m \mid \mathbf{X}=\mathbf{x}) = P(M=m \mid T=t', \mathbf{X}=\mathbf{x})$. Substituting Steps 2–5 into Step 1 yields Equation 8.6. $\square$

Interpretation. The mediation formula “mixes” the outcome regression under $T = t$ with the mediator distribution under $T = t'$. To compute the NDE, set $t = 1$ and $t' = 0$: take the conditional mean of Y evaluated at treatment 1, but weight the mediator by its distribution under treatment 0. The presence of the second treatment index t' on the right-hand side is the visible trace of the cross-world step.

The NDE and NIE from the formula:

$$\text{NDE} = \sum_{m, \mathbf{x}} [\mathbb{E}[Y \mid T=1, M=m, \mathbf{X}=\mathbf{x}] - \mathbb{E}[Y \mid T=0, M=m, \mathbf{X}=\mathbf{x}]] P(M=m \mid T=0, \mathbf{X}=\mathbf{x}) P(\mathbf{X}=\mathbf{x}),$$

$$\text{NIE} = \sum_{m, \mathbf{x}} \mathbb{E}[Y \mid T=1, M=m, \mathbf{X}=\mathbf{x}] [P(M=m \mid T=1, \mathbf{X}=\mathbf{x}) - P(M=m \mid T=0, \mathbf{X}=\mathbf{x})] P(\mathbf{X}=\mathbf{x}).$$

Estimation Recipe: Plug-In for the Mediation Formula

Given data $\{(Y_i, T_i, M_i, \mathbf{X}_i)\}_{i=1}^n$, a plug-in estimator requires:

1. **Fit an outcome model.** Regress Y on $(T, M, \mathbf{X})$ to obtain $\hat{\mu}(t, m, \mathbf{x})$.
2. **Fit a mediator model.** Regress M on $(T, \mathbf{X})$ to obtain $\hat{p}(m \mid t, \mathbf{x})$.
3. **Predict outcomes at the evaluation arm.** For each unit i , compute $\hat{\mu}(t, m, \mathbf{X}_i)$.
4. **Reweight the mediator at the reference arm.** For discrete M : weight $\hat{\mu}(t, m, \mathbf{X}_i)$ by $\hat{p}(m \mid t', \mathbf{X}_i)$.
5. **Average over the empirical distribution of $\mathbf{X}$:**

$$\frac{1}{n} \sum_{i=1}^n \sum_m \hat{\mu}(t, m, \mathbf{X}_i) \hat{p}(m \mid t', \mathbf{X}_i). \quad (8.7)$$

6. **Quantify uncertainty.** Nonparametric bootstrap or influence-function-based estimators

(Chapters 9–10).

The plug-in estimator is consistent only when both $\hat{\mu}$ and $\hat{p}$ are correctly specified.

8.7 The Linear Mediation Model: Baron–Kenny

8.7.1 The Three-Equation System

The regression-based approach of Baron and Kenny (1986) restricts the reduced prototype graph to a linear SEM:

$$Y = \alpha_1 + \tau T + \gamma_1^\top \mathbf{X} + \varepsilon_1, \quad (8.8)$$

$$M = \alpha_2 + aT + \gamma_2^\top \mathbf{X} + \varepsilon_2, \quad (8.9)$$

$$Y = \alpha_3 + \tau' T + bM + \gamma_3^\top \mathbf{X} + \varepsilon_3. \quad (8.10)$$

The four coefficients of interest: τ (total effect of T on Y), a (first-stage: T on M), τ' (direct: T on Y controlling for M), and b (second-stage: M on Y controlling for T).

8.7.2 The Product and Difference Formulas

Proposition: Mediation Decomposition in the Linear Model

Under Equation 8.8–Equation 8.10:

$$\tau = \tau' + ab. \quad (8.11)$$

The indirect effect is $\tau_{\text{ind}} = ab$ (product method) and the direct effect is $\tau_{\text{dir}} = \tau - ab = \tau'$ (difference method).

Proof

Substituting Equation 8.9 into Equation 8.10: $Y = (\alpha_3 + b\alpha_2) + (\tau' + ab)T + (\gamma_3 + b\gamma_2)^\top \mathbf{X} + (b\varepsilon_2 + \varepsilon_3)$. Comparing with Equation 8.8 gives $\tau = \tau' + ab$. $\square$

Effect	Formula	Path(s)
Total	τ	$T \rightarrow Y$ and $T \rightarrow M \rightarrow Y$ combined
Direct	$\tau' = \tau - ab$	$T \rightarrow Y$ only
Indirect	ab	$T \rightarrow M \rightarrow Y$ only
Proportion mediated	ab/τ	Share of total effect via M

Running example. Suppose the fitted coefficients are $\hat{\tau} = 0.50$, $\hat{a} = 0.40$, $\hat{b} = 0.60$, $\hat{\tau}' = 0.26$. Then: indirect effect via sleep: $\hat{a}\hat{b} = 0.40 \times 0.60 = 0.24$; direct effect: $\hat{\tau} - \hat{a}\hat{b} = 0.50 - 0.24 = 0.26 = \hat{\tau}'$; proportion mediated: $0.24/0.50 = 0.48$.

The Decomposition $\tau = \tau' + ab$ Is an Algebraic Identity, Not a Causal Theorem

The equality $\tau - \tau' = ab$ is a purely algebraic consequence of substituting one linear equation into another. It holds because linearity makes the indirect effect separable and additive. It does *not* hold in general: in **nonlinear models**, the product and difference methods yield numerically different estimates; when **T and M interact**, the CDE depends on m and the indirect effect cannot be summarized by a single number ab ; for **non-continuous mediators**, the product ab has no simple causal interpretation.

The correct generalization is the mediation formula (Theorem above), which reduces to ab and τ' only in the linear, no-interaction special case.

8.7.3 Inference: The Sobel Test and Bootstrap

The delta method gives an approximate variance for the product $\hat{a}\hat{b}$. The *Sobel test* (Sobel 1982) uses the simplified approximation:

$$\widehat{\text{Var}}_{\text{Sobel}}(\hat{a}\hat{b}) = \hat{b}^2 \widehat{\text{Var}}(\hat{a}) + \hat{a}^2 \widehat{\text{Var}}(\hat{b}), \quad (8.12)$$

yielding $z = \hat{a}\hat{b} / \sqrt{\widehat{\text{Var}}_{\text{Sobel}}(\hat{a}\hat{b})}$.

Statistical validity versus causal interpretation. The Sobel test is a statistical inference procedure for the product of two fitted regression coefficients. Interpreting the tested product as an *indirect causal effect* additionally requires the causal identification conditions and the structural restrictions (linearity and no $T \times M$ interaction). In practice, bootstrap confidence intervals for ab are preferred over the Sobel test because the distribution of a product of estimates is skewed in finite samples.

The “Significance of Both Paths” Criterion Is Not a Test for Mediation

A common misuse of the Baron–Kenny framework declares mediation when $T \rightarrow Y$, $T \rightarrow M$, and $M \rightarrow Y$ are each individually significant. This has three serious defects: (1) statistical significance $\neq$ mediation; (2) the Sobel test tests a regression product, not a causal quantity; (3) zero total effect does not preclude indirect effects — direct and indirect effects of opposite sign can cancel. The modern alternative is to estimate ab (or the NIE via the mediation formula) and construct bootstrap confidence intervals.

8.7.4 The Baron–Kenny Assumptions: Two Distinct Categories

Causal ignorability conditions (identification assumptions with graphical interpretations): (1) No unmeasured T – Y confounding; (2) No unmeasured T – M confounding; (3) No unmeasured M – Y confounding given T — this is *not* implied by randomization of T .

Structural modeling restrictions (parametric assumptions, not identification conditions): (4) Linearity and additivity; (5) No $T \times M$ interaction. Restrictions 4 and 5 can in principle be tested; Conditions 1–3 cannot.

Under Conditions 1–3 and Restrictions 4–5, Baron–Kenny is a valid estimator of the NDE and NIE. Outside those conditions, τ' and ab estimate regression coefficients in misspecified structural equations, not causal effects.

Estimand	Framework needed	Section
Total effect $P(y \mid \text{do}(t))$	Do-calculus	Ch. 5 (back-door)
Controlled direct effect (CDE)	Do-calculus	Section 8.4
Natural direct effect (NDE)	Potential outcomes	Section 8.5
Natural indirect effect (NIE)	Potential outcomes	Section 8.5
Front-door identification	Do-calculus	Section 8.8

8.8 Front-Door Identification

8.8.1 The Front-Door DAG

Every identification strategy in §§ Section 8.4–Section 8.7 assumed that an observed covariate set $\mathbf{X}$ blocks the back-door paths from T to Y through U . What if U is wholly unobserved and no such adjustment set exists?

The front-door criterion turns this obstacle into an opportunity: under two additional restrictions on the prototype mediation graph, the mediation structure itself provides identification of the total effect without conditioning on U . The two restrictions are: (1) remove the direct $T \rightarrow Y$ edge, so that M fully mediates the effect of T on Y ; and (2) require that U has no arrow into M , so that the $T \rightarrow M$ sub-effect is unconfounded.

Why the running example does not apply. The depression/sleep scenario fails Condition 1 (full mediation): a behavioral intervention on depression plausibly operates through several non-sleep channels, so the

direct edge $T \rightarrow Y$ is present. The canonical front-door example is Pearl's smoking–tar–cancer graph: T = smoking, M = tar deposited in lungs, Y = lung cancer, U = unobserved genetic susceptibility.




Figure 8.2: The front-door graph. Compared with the prototype mediation DAG, two edges are absent: the direct $T \rightarrow Y$ edge is removed, and U has no arrow into M . The front-door formula identifies $P(y \mid \text{do}(t))$ without observing or conditioning on U .

Two Different Questions About a Mediator-Like Variable

	Ordinary mediation	Front-door identification
Question	How much of the total effect of T on Y flows through M ?	Can M identify the total effect of T on Y despite unmeasured T – Y confounding?
Role of M	Pathway: carries part of the causal effect	Instrument-like relay: routes around confounding
Target estimand	NDE, NIE (decomposition)	$P(y \mid \text{do}(t))$ (total effect itself)
Direct $T \rightarrow Y$ edge	Present	Absent (required)

8.8.2 The Three Front-Door Conditions

Front-Door Conditions

A set $\{M\}$ satisfies the **front-door criterion** for the effect of T on Y if:

1. **Full mediation.** All directed paths from T to Y pass through M (no direct $T \rightarrow Y$ edge).
2. **No unblocked back-door path from T to M .** All back-door paths from T to M are blocked (no unobserved variables affect both T and M).
3. **No unblocked back-door path from M to Y given T .** All back-door paths from M to Y are blocked by conditioning on T .

In Figure 8.2: Condition 1 holds (no $T \rightarrow Y$ edge). Condition 2 holds (U has no arrow into M). Condition 3 holds (the only back-door path from M to Y is $M \leftarrow T \leftarrow U \rightarrow Y$, blocked by conditioning on T).

8.8.3 Derivation of the Front-Door Formula

Theorem: Front-Door Formula

(Pearl 1995) Suppose M satisfies the front-door criterion for the effect of T on Y , and positivity conditions (F1)–(F2) hold. Then:

$$P(y \mid \text{do}(T=t)) = \sum_m P(m \mid t) \sum_{t'} P(y \mid m, t') P(t'). \quad (8.13)$$

Proof via Three Do-Calculus Steps

Step 1: Identify the effect of T on M . By Condition 2, there are no unblocked back-door paths from T to M , so T satisfies the back-door criterion for M with the empty adjustment set:

$$P(m \mid \text{do}(T=t)) = P(m \mid t).$$

Step 2: Identify the effect of M on Y . By Condition 3, conditioning on T blocks all back-door paths from M to Y , so $\{T\}$ is a valid back-door adjustment set:

$$P(y \mid \text{do}(M=m)) = \sum_{t'} P(y \mid m, t') P(t').$$

Step 3: Combine via full mediation. By the law of total probability under $\text{do}(T=t)$:

$$P(y \mid \text{do}(T=t)) = \sum_m P(m \mid \text{do}(T=t)) P(y \mid \text{do}(T=t), \text{do}(M=m)).$$

By Condition 1 (no direct $T \rightarrow Y$ edge), once M is set to m by intervention, T is d-separated from Y in $\mathcal{G}_{\overline{T}\overline{M}}$; Rule 3 of the do-calculus removes $\text{do}(T=t)$ from the second factor: $P(y \mid \text{do}(T=t), \text{do}(M=m)) = P(y \mid \text{do}(M=m))$.

Substituting Steps 1 and 2 yields Equation 8.13. $\square$

8.8.4 Intuition: Routing Around Confounding

The front-door formula achieves identification in two steps that each use only unconfounded variation.

1. **T to M :** There is no confounding on the $T \rightarrow M$ edge (U does not affect M), so $P(m \mid t)$ is the causal effect of T on M .
2. **M to Y :** There is back-door confounding on $M \rightarrow Y$ through $T \leftarrow U \rightarrow Y$, but it is blocked by summing over T : $\sum_{t'} P(y \mid m, t') P(t')$ marginalizes out T , removing the confounding path.

The key insight: U confounds T and Y but *not* the $T \rightarrow M$ edge. The front-door formula exploits this asymmetry.

Feature	Prototype mediation DAG	Front-door DAG
Direct $T \rightarrow Y$ edge	Present	Absent
U confounds T – Y	Yes	Yes
Goal	Decompose total effect into direct/indirect	Identify total effect despite T – Y confounding
Requires Assumption 3	Yes	Satisfied by Condition 3, not by randomization
Identified by front-door formula	No	Yes

8.9 Mediation vs. Instrumental Variables

Mediation analysis and instrumental variables both involve a third variable connected to the treatment–outcome relationship, but the causal role of that variable is fundamentally different.

Feature	Instrumental Variables	Mediation Analysis
Position of third variable	Pre-treatment (Z precedes T)	Post-treatment (M follows T)
Causal role	Exogenous source of variation in T	Pathway through which T affects Y
Primary goal	Identification of $T \rightarrow Y$ effect under confounding	Decomposition of $T \rightarrow Y$ effect into pathways
Key assumption	Exclusion: Z affects Y only through T	Sequential ignorability: no unmeasured M – Y confounding given T

Feature	Instrumental Variables	Mediation Analysis
Estimand	LATE (Wald) or ATE (2SLS under homogeneity)	NDE, NIE, or CDE
Unobserved T – Y confounders	Permitted	Must be addressed separately (Assumption 1)

The Key Conceptual Distinction

	IV	Mediation
What the third variable does	Generates clean variation in T (exogenous source)	Carries part of T 's causal effect (pathway)
Exclusion vs. inclusion	Z excluded from Y 's structural equation	M included in Y 's structural equation
Question answered	Does T cause Y ?	How does T cause Y ?

Can the same variable be both? Not in the same analysis. As a mediator, M is on the causal path from T to Y (inclusion required). As an instrument, M would require the exclusion restriction — that it affects Y only through T — which contradicts M being a mediator. The two frameworks make mutually incompatible structural assumptions about the intermediate variable. However, the front-door identification formula is IV's closest relative within mediation analysis: it uses the mediator M to achieve identification of the total effect of T on Y even when T is confounded.

8.10 Summary

1. **Mediation studies mechanisms.** Mediation analysis decomposes the total effect of T on Y into a direct component ($T \rightarrow Y$) and an indirect component ($T \rightarrow M \rightarrow Y$). The challenge is that the mediator is a post-treatment variable that may be confounded.
2. **The total effect is the baseline.** The $TE = \mathbb{E}[Y(1) - Y(0)]$ captures all pathways combined. Mediation analysis asks for a finer decomposition.
3. **The CDE uses do-calculus.** The controlled direct effect fixes $M = m$ by joint intervention $do(T, M)$ and is identified by the back-door formula applied to the pair (T, M) . No assumptions beyond the back-door criterion are needed, but the CDE depends on m and does not have a natural “indirect” complement.
4. **Natural effects require potential outcomes.** The NDE and NIE involve cross-world counterfactuals $Y(t, M(t'))$ that cannot be expressed with the do-operator alone. They decompose the total effect as $TE = NDE + NIE$, and are identified by the mediation formula under sequential ignorability. The critical Assumption 4 — no treatment-induced mediator–outcome confounder — is not secured by randomization of T .
5. **The linear model simplifies but restricts.** The Baron–Kenny three-equation system gives the clean formulas $\tau_{\text{ind}} = ab$ and $\tau_{\text{dir}} = \tau'$, with $\tau = \tau' + ab$. This decomposition is purely algebraic and holds only under linearity and no interaction.
6. **Front-door identification uses mediation structure.** When M fully mediates $T \rightarrow Y$, no $T \rightarrow M$ confounding exists, and T blocks $M \rightarrow Y$ confounding, the front-door formula identifies the total effect despite unobserved T – Y confounding.
7. **Mediation and IV are complementary, not equivalent.** IV uses a pre-treatment variable to generate exogenous variation in T ; mediation uses a post-treatment variable to decompose the causal effect. The exclusion restriction of IV and the sequential ignorability of mediation are mutually incompatible for the same intermediate variable.

8.11 Problems

1. Identifying the CDE. Consider the DAG: $T \rightarrow M$, $T \rightarrow Y$, $M \rightarrow Y$, $X \rightarrow T$, $X \rightarrow M$, $X \rightarrow Y$, with all variables observed.

- Write the identification formula for $\mathbb{E}[Y \mid \text{do}(T=1), \text{do}(M=m)]$ using the back-door criterion for the joint intervention (T, M) .
- Add an unobserved U with $U \rightarrow T$ and $U \rightarrow Y$. Does the back-door formula still identify the CDE? Explain which condition, if any, fails.
- Instead add U with $U \rightarrow M$ and $U \rightarrow Y$. Does the back-door formula still identify the CDE? Explain.

2. CDE vs. total effect. In the reduced prototype graph, let $\mathbf{Z}$ satisfy the back-door criterion for both the total effect and the joint intervention (T, M) .

- Write expressions for the total effect and the $\text{CDE}(m)$ using the back-door formula.
- Under what graphical condition does $\text{CDE}(m)$ equal the total effect for all m ? Interpret this condition.

3. Natural direct and indirect effects. Verify the $\text{NDE} + \text{NIE} = \text{TE}$ decomposition algebraically for the linear SEM $M = \alpha T + \eta$, $Y = \beta T + \gamma M + \varepsilon$ (no interaction).

- Compute $Y(t, M(t'))$ in the linear model. Show that $Y(t, M(t')) = \beta t + \gamma(\alpha t') + \text{noise}$.
- Derive $\text{NDE} = \beta$ and $\text{NIE} = \alpha\gamma$ from the definitions Equation 8.3–Equation 8.4.
- Confirm $\text{NDE} + \text{NIE} = \beta + \alpha\gamma = \text{TE}$.
- Now suppose a $T \times M$ interaction is added: $Y = \beta T + \gamma M + \delta(T \cdot M) + \varepsilon$. Compute $\text{CDE}(m)$ and NDE . Show that $\text{NDE} \neq \text{CDE}(m)$ when $\delta \neq 0$.

4. The Baron–Kenny three-equation system. In the reduced prototype graph with the linear SEM Equation 8.8–Equation 8.10:

- State the three identification assumptions. For each, give the graphical condition in terms of back-door paths.
- Derive the equality $\tau = \tau' + ab$ algebraically.
- Suppose estimated coefficients are $\hat{\tau} = 0.50$, $\hat{a} = 0.40$, $\hat{b} = 0.60$, $\hat{\tau}' = 0.26$. Compute the indirect effect by both the product and difference methods. Do they agree? Compute the proportion mediated.
- An analyst reports $\hat{a} = 0.40$ with $\widehat{\text{SE}}(\hat{a}) = 0.08$, and $\hat{b} = 0.60$ with $\widehat{\text{SE}}(\hat{b}) = 0.10$. Compute the Sobel standard error for $\hat{a}\hat{b}$ using Equation 8.12 and construct an approximate 95% confidence interval.

5. The critical role of Assumption 3. Consider the graph with an unobserved V with $V \rightarrow M$ and $V \rightarrow Y$. T is randomized.

- Identify all back-door paths from M to Y in this graph.
- Can any combination of observed variables $(T, \mathbf{X})$ block all of these paths? Explain using d-separation.
- Suppose an analyst fits Equation 8.10 ignoring V and obtains $\hat{b} = 0.80$. In which direction is $\hat{b}$ biased if V has positive effects on both M and Y ?
- State the additional data structure that would be needed to identify the second-stage effect non-parametrically.

6. Front-door identification. Consider the front-door graph of Figure 8.2.

- Verify that the three front-door conditions hold.
- Walk through the three-step proof of the Front-Door Formula theorem for this graph: identify which do-calculus rule justifies each step.
- Add a direct edge $T \rightarrow Y$ to the graph. Which front-door condition is violated? Does formula Equation 8.13 still hold?
- Explain why the front-door graph does not require Assumption 3 of the Baron–Kenny framework. Which back-door path — present in the graph with unobserved V — is absent here?

7. Mediation vs. instrumental variables. A researcher studies the effect of a job training program (T) on wages (Y). She proposes two intermediate variables: (A) motivation (M_A), measured after the

program starts; (B) a lottery that randomly selects applicants for admission (Z), measured before the program.

- (a) For variable (A): draw the mediation DAG including M_A , T , Y , and an unobserved ability variable U . State the sequential ignorability assumption needed to identify the NIE through M_A . Explain why randomization of T does not automatically satisfy this assumption.
- (b) For variable (B): draw the IV DAG with Z , T , Y , and U . State the three IV assumptions. Explain why the exclusion restriction and the “mediator inclusion” of mediation analysis are mutually incompatible conditions for the same intermediate variable.
- (c) The researcher argues that M_A (motivation) and Z (lottery) are both “intermediate” variables and that the analyses are interchangeable. Write a one-paragraph critique of this argument, using the distinctions from Section 8.9.
- (d) Can the front-door formula be applied if motivation M_A fully mediates the effect of T on Y and U (unobserved ability) does not directly affect M_A ? State the three conditions and assess whether they hold.

Chapter 9

Estimating Equations and Influence Functions

Learning Objectives

By the end of this chapter, students should be able to:

1. Explain why identification of a causal parameter does not automatically yield a statistically reliable estimator, and articulate the distinct questions that an estimation theory must address.
2. Define an estimating equation and verify the population moment condition for standard examples (sample mean, OLS, IPW).
3. State the definition of an asymptotically linear estimator and derive the influence function from a generic first-order expansion of an estimating equation.
4. Compute the influence function for simple causal estimators when nuisance functions are known, and interpret it as the infinitesimal contribution of a single observation.
5. Describe how estimation error in nuisance functions propagates into the target estimator, and explain why this is the central statistical challenge in causal inference.
6. Use the influence function to construct a consistent variance estimator and an approximate Wald confidence interval.
7. Define efficiency in the class of regular estimators and explain the role of the efficient influence function as the foundation for doubly robust estimation in Chapter 10.

9.1 Why Estimation Needs Its Own Theory

So far the focus has been on *identification*: under what assumptions can a causal parameter be written as a functional of the observed-data distribution? Identification does not by itself provide a statistically reliable estimator. Once a causal parameter is identified, several distinct questions remain: How should the estimator be constructed from the observed sample? How does estimation of nuisance functions affect the target estimator? What is the large-sample distribution of the estimator? How can we compute valid standard errors and confidence intervals?

These questions motivate a separate theory of *estimation and inference*. In causal inference this issue is especially important because identified parameters often depend on auxiliary, or *nuisance*, functions such as the outcome regression $\mu_t(x) = \mathbb{E}(Y \mid T = t, X = x)$ or the propensity score $\pi(x) = P(T = 1 \mid X = x)$. Even when the causal parameter is identified, different estimation strategies may behave quite differently in terms of robustness, efficiency, and sensitivity to model misspecification.

The goal of this chapter is to introduce a general framework for estimation based on *estimating equations* and *influence functions*. This framework will serve as the foundation for the doubly robust and semiparametric methods developed in Chapter 10; see also Imbens and Rubin (2015) and Hernán and Robins (2020) for complementary treatments.

9.2 A Running Example: The Average Treatment Effect

Throughout this chapter we use the ATE $\psi = \mathbb{E}\{Y(1) - Y(0)\}$ as a running example. Under consistency, conditional exchangeability, and positivity, the ATE is identified by the back-door formula (Chapter 5):

$$\psi = \mathbb{E}[\mu_1(X) - \mu_0(X)], \quad \mu_t(x) = \mathbb{E}(Y \mid T = t, X = x), \quad t \in \{0, 1\}.$$

This identity tells us what the target parameter is, but it does not uniquely determine how to estimate it. One may consider: a *regression-based* estimator using models for $\mu_1(x)$ and $\mu_0(x)$; a *weighting* estimator based on the propensity score $\pi(x)$ (Chapter 6); or an *augmented* estimator combining both. All of these estimators may target the same causal parameter yet differ in their statistical properties.

9.3 Estimating Equations

A broad class of estimators can be defined as solutions to *estimating equations*. Let $O_1, \dots, O_n$ be i.i.d. observations from a distribution P , and let θ denote a finite-dimensional target parameter. Write $\mathbb{P}_n f = \frac{1}{n} \sum_{i=1}^n f(O_i)$ for the empirical average.

Definition: Estimating Equation

An **estimating equation** for a parameter θ is an equation of the form $\mathbb{P}_n\{U(O; \theta)\} = 0$, where the **population moment condition** $\mathbb{E}\{U(O; \theta_0)\} = 0$ holds at the true parameter value θ_0 . The function $U(O; \theta)$ is called the **estimating function**.

Example: Sample Mean

Let $\theta = \mathbb{E}(Y)$. The sample mean $\hat{\theta} = \bar{Y}$ solves $\mathbb{P}_n(Y - \theta) = 0$. The estimating function is $U(O; \theta) = Y - \theta$.

Example: Ordinary Least Squares

Let $O = (X, Y)$ and suppose β is defined by the linear moment condition $\mathbb{E}[X\{Y - X^\top \beta\}] = 0$. Then the OLS estimator solves $\mathbb{P}_n[X\{Y - X^\top \beta\}] = 0$.

Example: A Simple IPW Estimating Equation

Suppose $\pi(X)$ is known and $\psi = \mathbb{E}\left\{\frac{TY}{\pi(X)} - \frac{(1-T)Y}{1-\pi(X)}\right\}$. Then ψ solves $\mathbb{E}\left[\frac{TY}{\pi(X)} - \frac{(1-T)Y}{1-\pi(X)} - \psi\right] = 0$.

9.4 From Estimating Equations to Asymptotic Linearity

Under suitable regularity conditions, an estimator solving an estimating equation can typically be approximated as:

$$\sqrt{n}(\hat{\theta} - \theta_0) = \frac{1}{\sqrt{n}} \sum_{i=1}^n \varphi(O_i) + o_p(1)$$

for some mean-zero function $\varphi(O)$.

Definition: Asymptotic Linearity and Influence Function

An estimator $\hat{\theta}$ is **asymptotically linear** with **influence function** $\varphi(O)$ if:

$$\sqrt{n}(\hat{\theta} - \theta_0) = \frac{1}{\sqrt{n}} \sum_{i=1}^n \varphi(O_i) + o_p(1), \quad \mathbb{E}\{\varphi(O)\} = 0.$$

This representation immediately implies asymptotic normality. When $\theta \in \mathbb{R}^p$:

$$\sqrt{n}(\hat{\theta} - \theta_0) \xrightarrow{d} N(\mathbf{0}, \Sigma), \quad \Sigma = \mathbb{E}[\varphi(O)\varphi(O)^\top].$$

Proposition: Generic First-Order Expansion

Under smoothness and regularity conditions, an estimator $\hat{\theta}$ solving $\mathbb{P}_n\{U(O; \theta)\} = 0$ admits a first-order expansion:

$$\sqrt{n}(\hat{\theta} - \theta_0) = -A^{-1} \frac{1}{\sqrt{n}} \sum_{i=1}^n U(O_i; \theta_0) + o_p(1),$$

where $A = \mathbb{E}\left\{\frac{\partial}{\partial \theta^\top} U(O; \theta_0)\right\}$. Hence the influence function is $\varphi(O) = -A^{-1}U(O; \theta_0)$.

Proof

A Taylor expansion of the sample moment condition around θ_0 gives:

$$0 = \mathbb{P}_n\{U(O; \hat{\theta})\} \approx \mathbb{P}_n\{U(O; \theta_0)\} + \mathbb{P}_n\left\{\frac{\partial}{\partial \theta^\top} U(O; \theta_0)\right\}(\hat{\theta} - \theta_0).$$

By the law of large numbers, $\mathbb{P}_n\{\partial_{\theta^\top} U(O; \theta_0)\} \rightarrow A$ in probability. Rearranging and multiplying by $\sqrt{n}$ yields the stated expansion. $\square$

9.5 Influence Functions: Intuition

9.5.1 Statistical Functionals

Definition: Statistical Functional

A **statistical functional** is a map $\Psi : \mathcal{P} \rightarrow \mathbb{R}^p$ that assigns a parameter value $\psi = \Psi(P)$ to each distribution $P \in \mathcal{P}$. The **plug-in estimator** of ψ is $\hat{\psi} = \Psi(\mathbb{P}_n)$.

The mean, OLS coefficient, and ATE are all statistical functionals. A functional Ψ is *linear* if $\Psi((1-\epsilon)P + \epsilon Q) = (1-\epsilon)\Psi(P) + \epsilon\Psi(Q)$. The mean is linear; the variance, OLS coefficient, and ATE are nonlinear, requiring local linearization — precisely what the influence function provides.

9.5.2 Influence Functions

Influence functions describe the first-order sensitivity of a statistical functional to small perturbations of the underlying distribution.

Definition: Influence Function

Let $\psi = \Psi(P)$ be a scalar parameter. A function $\varphi(O)$ is called an **influence function** for Ψ if it has mean zero under P and:

$$\hat{\psi} - \psi = \frac{1}{n} \sum_{i=1}^n \varphi(O_i) + o_p(n^{-1/2}).$$

The influence function is the infinitesimal contribution of a single observation to the first-order behavior of the estimator.

Remark: Formal Derivation via Pathwise Derivatives

The influence function can be defined formally through the *pathwise (Gateaux) derivative*: if $P_\epsilon = (1 - \epsilon)P + \epsilon\delta_O$ is the ϵ -mixture of P with a point mass at O , then:

$$\varphi(O) = \left. \frac{d}{d\epsilon} \Psi(P_\epsilon) \right|_{\epsilon=0}.$$

This derivative-based definition is the starting point for semiparametric efficiency theory, developed formally in Appendix C.

Example: Influence Function of the OLS Slope

For the OLS coefficient $\beta = [\mathbb{E}(XX^\top)]^{-1}\mathbb{E}(XY)$:

Step 1. Construct the perturbed distribution $P_\epsilon = (1 - \epsilon)P + \epsilon\delta_O$:

$$\Psi(P_\epsilon) = [(1 - \epsilon)\mathbb{E}(XX^\top) + \epsilon XX^\top]^{-1}[(1 - \epsilon)\mathbb{E}(XY) + \epsilon XY].$$

Step 2. Differentiate with respect to ϵ at $\epsilon = 0$, using $\left. \frac{d}{d\epsilon} M(\epsilon)^{-1} \right|_0 = -[\mathbb{E}(XX^\top)]^{-1}(XX^\top - \mathbb{E}(XX^\top))[\mathbb{E}(XX^\top)]^{-1}$ and the product rule:

$$\varphi(O) = [\mathbb{E}(XX^\top)]^{-1}X(Y - X^\top\beta).$$

Verification. $\mathbb{E}[\varphi(O)] = [\mathbb{E}(XX^\top)]^{-1}\mathbb{E}[X(Y - X^\top\beta)] = 0$ by definition of β .

Connection to estimating equations: The OLS estimating function is $U(O; \beta) = X(Y - X^\top\beta)$, $A = -\mathbb{E}(XX^\top)$, and Proposition 1 gives the same $\varphi(O) = [\mathbb{E}(XX^\top)]^{-1}X(Y - X^\top\beta)$.

9.6 Influence Functions for Simple Causal Estimators

9.6.1 Known Outcome Regression

Suppose $\psi = \mathbb{E}[\mu_1(X) - \mu_0(X)]$ and the functions $\mu_1(\cdot)$ and $\mu_0(\cdot)$ are known. The natural plug-in estimator is $\hat{\psi} = \mathbb{P}_n\{\mu_1(X) - \mu_0(X)\}$. Its influence function is:

$$\varphi(O) = \mu_1(X) - \mu_0(X) - \psi.$$

9.6.2 Known Propensity Score

Suppose $\pi(X)$ is known and $\psi = \mathbb{E}\left\{\frac{TY}{\pi(X)} - \frac{(1-T)Y}{1-\pi(X)}\right\}$. The empirical analogue is the Horvitz–Thompson estimator, and its influence function is:

$$\varphi(O) = \frac{TY}{\pi(X)} - \frac{(1-T)Y}{1-\pi(X)} - \psi.$$

Remark: Nuisance Functions and the Influence Function

When nuisance functions are *known*, the influence function is simply the centered estimating function. This observation becomes much more consequential in the next section, where nuisance functions are unknown and must be estimated from data.

Remark: Estimating Function, Influence Function, and Efficient Influence Function

These three objects are closely related but not the same.

1. An *estimating function* defines an estimator through a sample moment equation $\mathbb{P}_n\{\phi(O; \theta)\} = 0$.
2. An *influence function* describes the first-order asymptotic behavior of an estimator: $\sqrt{n}(\hat{\theta} -$

$$\theta_0) = n^{-1/2} \sum_i \varphi(O_i) + o_p(1).$$

3. The *efficient influence function* is the influence function with the smallest possible asymptotic variance among regular estimators in the statistical model.

In simple examples with known nuisance functions, the centered estimating function often coincides with the influence function. This is not automatic when nuisance parameters are estimated.

9.7 Z-Estimation with Nuisance Parameters

In real applications, nuisance functions must be estimated from the same data used to estimate the target parameter. Once this happens, the first-order expansion of $\hat{\psi}$ is no longer obtained by simply centering the estimating function with the nuisance held fixed. Z-estimation provides the correct framework: it treats both the target parameter and the nuisance parameter as components of a single stacked estimating equation.

9.7.1 The Stacked Estimating Equation Framework

Suppose the target parameter $\psi \in \mathbb{R}$ depends on an unknown nuisance parameter $\alpha \in \mathbb{R}^q$. Both are estimated by solving a stacked system:

$$\mathbb{P}_n\{U_1(O; \psi, \alpha)\} = 0, \quad \mathbb{P}_n\{U_2(O; \alpha)\} = 0. \quad (9.1)$$

Write $\theta = (\psi^\top, \alpha^\top)^\top$ and $U(O; \theta) = (U_1(O; \psi, \alpha), U_2(O; \alpha))^\top$.

Theorem: Z-Estimation Theorem (Stacked Equations)

Suppose the true parameter value is $\theta_0 = (\psi_0, \alpha_0^\top)^\top$ and the population moment condition $\mathbb{E}\{U(O; \theta_0)\} = 0$ holds, the Jacobian $\mathcal{A} = \mathbb{E}\{\partial_{\theta^\top} U(O; \theta_0)\}$ is invertible, and regularity conditions hold. Then $\hat{\theta}$ is consistent for θ_0 and:

$$\sqrt{n}(\hat{\theta} - \theta_0) = -\mathcal{A}^{-1} \frac{1}{\sqrt{n}} \sum_{i=1}^n U(O_i; \theta_0) + o_p(1).$$

The influence function of $\hat{\psi}$ is:

$$\varphi(O) = -A_{11}^{-1} U_1(O; \psi_0, \alpha_0) + A_{11}^{-1} A_{12} A_{22}^{-1} U_2(O; \alpha_0), \quad (9.2)$$

where $A_{11} = \mathbb{E}\{\partial_{\psi} U_1\}$, $A_{12} = \mathbb{E}\{\partial_{\alpha^\top} U_1\}$, $A_{22} = \mathbb{E}\{\partial_{\alpha^\top} U_2\}$.

Proof

Partitioning $\mathcal{A}$ conformably with (ψ, α) :

$$\mathcal{A} = \begin{pmatrix} A_{11} & A_{12} \\ 0 & A_{22} \end{pmatrix}, \quad \mathcal{A}^{-1} = \begin{pmatrix} A_{11}^{-1} & -A_{11}^{-1} A_{12} A_{22}^{-1} \\ 0 & A_{22}^{-1} \end{pmatrix}.$$

The lower-left block is zero because U_2 does not depend on ψ . Reading off the first row of $-\mathcal{A}^{-1} U(O; \theta_0)$ gives Equation 9.2. $\square$

Remark: Interpreting the Correction Term

Equation 9.2 decomposes the influence function into two parts. The first term, $-A_{11}^{-1} U_1(O; \psi_0, \alpha_0)$, is exactly what we would obtain if α_0 were known. The second term, $A_{11}^{-1} A_{12} A_{22}^{-1} U_2(O; \alpha_0)$, is a *nuisance correction* that accounts for the additional variability introduced by estimating α . If $A_{12} = 0$, the target estimating function U_1 does not depend on α at α_0 , and the correction vanishes. This is the key condition exploited by *doubly robust* estimators in Chapter 10.

9.7.2 Working Example: IPW with Estimated Propensity Score

Assume a logistic regression model $\pi(X; \alpha) = \text{expit}(X^\top \alpha)$. The IPW estimating equation for ψ is $U_1(O; \psi, \alpha) = \frac{TY}{\pi(X; \alpha)} - \frac{(1-T)Y}{1-\pi(X; \alpha)} - \psi$, and the score for the logistic regression is $U_2(O; \alpha) = X\{T - \pi(X; \alpha)\}$.

Computing the Jacobian blocks: $A_{11} = -1$, $A_{22} = -\Sigma_\pi := -\mathbb{E}\{\pi(X; \alpha_0)(1 - \pi(X; \alpha_0))XX^\top\}$. After computing A_{12} , the corrected influence function is:

$$\varphi(O) = U_1(O; \psi_0, \alpha_0) + A_{12} \Sigma_\pi^{-1} U_2(O; \alpha_0). \quad (9.3)$$

Example: Variance Reduction from ML Estimation of Propensity Score

We show that estimating α by maximum likelihood can only *reduce* the asymptotic variance relative to using the true α_0 .

Write $\sigma_1^2 = \mathbb{E}[U_1^2]$ for the variance of the naive IPW estimating function. Two Bartlett-type identities are the key:

Identity 1 (Bartlett for ML score): $\mathbb{E}[U_2 U_2^\top] = \Sigma_\pi$.

Identity 2 (cross-identity): $A_{12} = -\mathbb{E}[U_1 U_2^\top]$.

Substituting into $\mathbb{E}[\varphi(O)^2] = \sigma_1^2 + 2A_{12}\Sigma_\pi^{-1}\mathbb{E}[U_1 U_2] + A_{12}\Sigma_\pi^{-1}\mathbb{E}[U_2 U_2^\top]\Sigma_\pi^{-1}A_{12}^\top$:

$$\mathbb{E}[\varphi(O)^2] = \sigma_1^2 - A_{12}\Sigma_\pi^{-1}A_{12}^\top \leq \sigma_1^2.$$

Since Σ_π is positive definite, the variance is reduced. *ML estimation of the propensity score parameters reduces the asymptotic variance of the IPW estimator — the estimated propensity score paradox* (Robins 1986).

Remark: Sandwich Variance Estimator

Equation 9.2 provides the basis for the *sandwich variance estimator*: $\hat{V} = n^{-2} \sum_{i=1}^n \hat{\varphi}_i^2$, where $\hat{\varphi}_i$ is $\varphi(O_i)$ evaluated at $(\hat{\psi}, \hat{\alpha})$. This estimator automatically accounts for nuisance estimation uncertainty. Standard errors from a logistic regression routine that ignores the link between $\hat{\alpha}$ and $\hat{\psi}$ will in general be *incorrect*.

The Central Challenge in Causal Estimation

The primary difficulty in practice is often not identifying the target parameter, but *correctly accounting for the effect of nuisance estimation* on the asymptotic distribution of the final estimator. Treating $\hat{\alpha}$ as if it were the true α_0 yields correct point estimates under a correctly specified propensity score model but *incorrect standard errors*. The Z-estimation framework resolves this variance-accounting issue whenever the nuisance model is correctly specified and the nuisance estimator converges at rate $n^{-1/2}$.

Remark: Machine Learning for Nuisance Estimation

The Z-estimation theorem as stated requires $\hat{\alpha}$ to converge at rate $n^{-1/2}$, which holds for finite-dimensional parametric models but fails for nonparametric or machine-learning estimators. When flexible methods are used for $\pi(x)$ or $\mu_t(x)$, the correction term in Equation 9.2 no longer has the simple form derived here. This is the starting point for the doubly robust and debiased machine-learning estimators introduced in Chapter 11.

9.8 Variance Estimation and Confidence Intervals

Once an estimator admits the asymptotic linear representation, its asymptotic variance is $\text{Var}(\hat{\psi}) \approx n^{-1} \text{Var}\{\varphi(O)\}$. A natural estimator is:

$$\hat{V} = \frac{1}{n^2} \sum_{i=1}^n \hat{\varphi}_i^2 = \frac{1}{n(n-1)} \sum_{i=1}^n (\hat{\varphi}_i - \bar{\varphi})^2. \quad (9.4)$$

This yields the approximate Wald confidence interval $\hat{\psi} \pm z_{1-\alpha/2} \sqrt{\hat{V}}$.

Proposition: Variance from the Influence Function

If $\sqrt{n}(\hat{\psi} - \psi) = n^{-1/2} \sum_{i=1}^n \varphi(O_i) + o_p(1)$ with $\mathbb{E}\{\varphi(O)\} = 0$ and $\mathbb{E}\{\varphi(O)^2\} < \infty$, then:

$$\sqrt{n}(\hat{\psi} - \psi) \xrightarrow{d} N(0, \mathbb{E}[\varphi(O)^2]), \quad \text{Var}(\hat{\psi}) = \frac{1}{n} \mathbb{E}[\varphi(O)^2] + o(n^{-1}).$$

9.9 Toward Efficiency: Semiparametric Models and the EIF

Different regular estimators of the same parameter may have different large-sample variances. What is the smallest achievable asymptotic variance among all regular estimators? Answering that question leads to semiparametric efficiency theory. The formal machinery — regular parametric submodels, scores, tangent spaces, and the characterization of the EIF as a canonical gradient — is developed in Appendix C.

9.9.1 Semiparametric Models

Definition: Semiparametric Model

A **semiparametric model** is a statistical model $\mathcal{P}$ for the distribution P of the observed data in which (i) the **parameter of interest** $\psi = \Psi(P) \in \mathbb{R}^p$ is finite-dimensional, and (ii) the remaining aspects of P — collectively called the **nuisance parameter** — are infinite-dimensional.

The canonical example is the nonparametric model $\mathcal{P} = \{P : 0 < \pi(x) < 1 \ \forall x\}$ with target parameter $\psi = \mathbb{E}[Y(1) - Y(0)]$. The nuisance consists of $\mu_t(x)$, $\pi(x)$, and the marginal distribution of X .

9.9.2 Regular Estimators

Definition: Regular Estimator (informal)

An estimator $\hat{\psi}$ is called **regular** at P_0 if, along every smooth parametric submodel $\{P_t\}$ with $P_0 = P_{t=0}$, the limiting distribution of $\sqrt{n}(\hat{\psi} - \psi_t)$ under $P_t = P_{1/\sqrt{n}}$ does not depend on the submodel.

Regularity rules out pathological “superefficient” estimators. All estimators in this course — regression, IPW, and their augmented variants — are regular.

9.9.3 The Semiparametric Efficiency Bound and the EIF

Theorem: Semiparametric Convolution Theorem

(Bickel et al. 1993) In a semiparametric model $\mathcal{P}$, the asymptotic distribution of any regular estimator $\hat{\psi}$ of ψ_0 satisfies:

$$\sqrt{n}(\hat{\psi} - \psi_0) \xrightarrow{d} N(0, V^*) * M$$

for some distribution M , where $V^* = \mathbb{E}_P[\varphi^*(O) \varphi^*(O)^\top]$ is the **semiparametric efficiency bound**. The function $\varphi^*(O)$ with $\mathbb{E}[\varphi^*(O)] = 0$ and $\mathbb{E}[\varphi^*(O)^2] = V^*$ is called the **efficient influence function** (EIF). An estimator achieves V^* if and only if it is asymptotically linear with influence function $\varphi^*(O)$.

Definition: Semiparametrically Efficient Estimator

A regular estimator $\hat{\psi}$ is **semiparametrically efficient** at $P_0 \in \mathcal{P}$ if $\sqrt{n}(\hat{\psi} - \psi_0) \xrightarrow{d} N(0, V^*)$. Equivalently, $\hat{\psi}$ is asymptotically linear with influence function $\varphi^*(O)$.

Remark: How the EIF is Computed

The EIF $\varphi^*(O)$ can be derived as the pathwise derivative of $\Psi(P)$ in the *least favorable* direction. In practice: (i) compute the pathwise derivative $\varphi(O; P_\epsilon)$ for an arbitrary parametric submodel; (ii) project onto the tangent space of $\mathcal{P}$ at P_0 . For nonparametric or large semiparametric models, the tangent space is rich enough that this projection is often the identity, and the EIF equals the ordinary pathwise derivative. See Appendix C for the formal treatment.

9.9.4 The Efficient Influence Function for the ATE**Remark: EIF for the ATE**

For the ATE $\psi = \mathbb{E}\{Y(1) - Y(0)\}$ in the nonparametric observational-data model under strong ignorability and positivity, the efficient influence function is:

$$\varphi^*(O) = \frac{T\{Y - \mu_1(X)\}}{\pi(X)} - \frac{(1-T)\{Y - \mu_0(X)\}}{1 - \pi(X)} + \mu_1(X) - \mu_0(X) - \psi. \quad (9.5)$$

This expression has a natural decomposition: the first two terms are an IPW-style residual correction, and the last two are the regression estimator centered at ψ . The EIF depends on *both* nuisance functions $\pi(X)$ and $\mu_t(X)$, and an estimator that plugs in consistent estimates of both achieves the efficiency bound $V^* = \mathbb{E}[\varphi^*(O)^2]$.

Remark: Bridge to Chapter 10

Two important properties of $\varphi^*(O)$ will be central in Chapter 10:

- (i) **Double robustness.** The estimator obtained by replacing $\pi(X)$ and $\mu_t(X)$ by estimates and averaging $\varphi^*(O_i)$ is consistent for ψ if *either* $\hat{\pi}$ or $\hat{\mu}_t$ is consistent — not necessarily both.
- (ii) **Semiparametric efficiency.** When both nuisance estimators are consistent and converge at suitable rates, the resulting estimator achieves the efficiency bound V^* .

The *augmented IPW* (AIPW) estimator is the principal tool for exploiting these two properties simultaneously, and will be derived and analyzed in Chapter 10.

9.10 Chapter Summary

Object	Role
Estimating function $U(O; \theta)$	Defines the estimator via $\mathbb{P}_n\{U(O; \hat{\theta})\} = 0$
Influence function $\varphi(O)$	First-order term in the expansion of $\hat{\psi} - \psi$; used for CLT and variance estimation
Asymptotic variance $\mathbb{E}[\varphi(O)^2]/n$	Governs precision; the target for efficiency comparisons
Efficient influence function $\varphi^*(O)$	Achieves the semiparametric efficiency bound; basis for doubly robust estimators (Chapter 10)

1. **Identification is not estimation.** Identification provides a population formula, but not automatically a satisfactory estimator.
2. **Estimating equations.** Many estimators are defined as solutions to sample moment conditions $\mathbb{P}_n\{U(O; \theta)\} = 0$, with the population condition $\mathbb{E}\{U(O; \theta_0)\} = 0$ ensuring consistency.

3. **Asymptotic linearity.** Estimating equations often yield a first-order expansion, immediately implying asymptotic normality via the CLT.
4. **Influence functions.** The function $\varphi(O)$ describes the first-order sensitivity of the estimator to a single observation. It is the key object for both asymptotic theory and variance estimation.
5. **Nuisance estimation.** In causal inference, nuisance functions must be estimated, and this estimation error propagates into the target estimator. The Z-estimation framework corrects for this propagation.
6. **Variance from the influence function.** The asymptotic variance is $n^{-1}\mathbb{E}[\varphi(O)^2]$, consistently estimated by $n^{-2}\sum_i \hat{\varphi}_i^2$.
7. **Efficiency.** Among regular estimators, the one with the smallest asymptotic variance is efficient. The EIF characterizes this lower bound and will guide estimator construction in Chapter 10.

9.11 Problems

1. Estimating equations and moment conditions.

- (a) Let $O = (X, Y)$ and define $\theta = \text{Cov}(X, Y)/\text{Var}(X)$. Write down an estimating equation for θ and verify that the population moment condition holds at θ_0 .
- (b) Let $\theta = F^{-1}(0.5)$ be the population median. Propose an estimating function $U(O; \theta)$ and verify the population moment condition. (*Hint:* consider $U(O; \theta) = \mathbf{1}(Y \leq \theta) - 0.5$.)
- (c) Show that the OLS estimator and the IPW estimator of the ATE are both special cases of the general M-estimator framework.

2. Asymptotic linearity and the CLT.

- (a) Let $\hat{\psi} = \bar{Y}$ be the sample mean with $\mathbb{E}Y = \psi$ and $\text{Var}(Y) = \sigma^2 < \infty$. Write down the influence function, state its asymptotic distribution, and give a consistent estimator of $\text{Var}(\hat{\psi})$.
- (b) Suppose $\hat{\psi}_1$ and $\hat{\psi}_2$ are two asymptotically linear estimators with influence functions φ_1 and φ_2 . Show that $\hat{\psi}_\lambda = \lambda\hat{\psi}_1 + (1 - \lambda)\hat{\psi}_2$ is also asymptotically linear for any fixed $\lambda \in [0, 1]$, and find its influence function.
- (c) Use part (b) to explain why, among linear combinations of two estimators, the one minimizing asymptotic variance is generally not the simple average ($\lambda = 1/2$).

3. Influence functions for causal estimators.

Consider the ATE $\psi = \mathbb{E}\{Y(1) - Y(0)\}$ identified under strong ignorability.

- (a) Assume $\pi(X)$ and $\mu_t(X)$ are both known. Compute the influence functions of: (i) the regression estimator $\hat{\psi}_{\text{reg}} = \mathbb{P}_n\{\mu_1(X) - \mu_0(X)\}$; (ii) the IPW estimator $\hat{\psi}_{\text{IPW}} = \mathbb{P}_n\{TY/\pi(X) - (1 - T)Y/(1 - \pi(X))\}$.
- (b) Under what conditions on the data-generating process will $\hat{\psi}_{\text{reg}}$ have a smaller asymptotic variance than $\hat{\psi}_{\text{IPW}}$?
- (c) Using the EIF Equation 10.7, verify that it has mean zero under P by computing $\mathbb{E}[\varphi^*(O)]$.

4. Variance estimation.

- (a) Let $\hat{\varphi}_i = \hat{\mu}_1(X_i) - \hat{\mu}_0(X_i) - \hat{\psi}_{\text{reg}}$ be the estimated influence values for the regression estimator. Write down the plug-in variance estimator $\hat{V}$ and simplify. Under what conditions is $\hat{V}$ consistent for $\text{Var}(\hat{\psi}_{\text{reg}})$?
- (b) A researcher reports $\hat{\psi} = 2.4$ and $\hat{V} = 0.09$ based on $n = 400$ observations. Compute a 95% Wald confidence interval. Is the treatment effect statistically distinguishable from zero at the 5% level?
- (c) Explain why simply plugging in $\hat{\mu}_t$ and $\hat{\pi}$ into the influence function formula may yield an inconsistent variance estimator when these nuisance estimators converge at a slower rate than $n^{-1/2}$.

5. Efficiency comparisons.

- (a) Suppose φ_{reg} and φ_{IPW} denote the influence functions of the regression and IPW estimators. Without calculation, explain why neither estimator is always more efficient than the other.
- (b) The semiparametric efficiency bound for the ATE is $\mathbb{E}[\varphi^*(O)^2]$, where φ^* is given in Equation 10.7. Show that $\mathbb{E}[\varphi^*(O)^2] \leq \mathbb{E}[\varphi_{\text{IPW}}(O)^2]$ by expanding the squared EIF. (*Hint:* use the law of iterated expectations to show that the cross-terms cancel.)

- (c) Give one practical reason why an efficient estimator based on the EIF may not always be preferred over a simpler, less efficient estimator.

Part III

Estimation

Chapter 10

Doubly Robust Estimation and Semiparametric Efficiency

Learning Objectives

By the end of this chapter, students should be able to:

1. Derive the AIPW estimator as a bias-corrected prediction estimator, and explain the role of each term.
2. Prove the double robustness property using the law of iterated expectations.
3. Describe the class of augmented IPW estimators indexed by arbitrary control functions $b_t(x)$, verify unbiasedness for any b_t , and identify the optimal choice that minimizes total variance.
4. Interpret the augmentation step as a Hilbert-space projection onto the orthogonal complement of the augmentation space, and derive the Pythagorean variance decomposition.
5. Carry out the projection argument in the causal inference setting, starting from the Horvitz–Thompson estimator.
6. State the semiparametric efficiency bound for the ATE and explain why the AIPW estimating function is the efficient influence function.
7. Derive the two approaches to building double robustness directly into the outcome model: IPW-weighted regression and the augmented model with clever covariate.
8. Identify the product-rate condition as the key sufficient condition for nuisance estimation error to be asymptotically negligible, and construct a consistent variance estimator and Wald confidence interval.

10.1 Why Combine Outcome Regression and Weighting?

Under consistency, conditional exchangeability, and positivity, the ATE $\tau = \mathbb{E}\{Y(1) - Y(0)\}$ is identified by either the regression formula $\tau = \mathbb{E}\{\mu_1(X) - \mu_0(X)\}$ or the Horvitz–Thompson formula $\tau = \mathbb{E}\left\{\frac{TY}{\pi(X)} - \frac{(1-T)Y}{1-\pi(X)}\right\}$.

Outcome regression estimates $\mu_1(x)$ and $\mu_0(x)$ and averages $\hat{\mu}_1(X_i) - \hat{\mu}_0(X_i)$ over the sample. **IPW** estimates $\pi(x)$ and reweights observed outcomes. Each has weaknesses: regression can be biased under misspecification; IPW can be unstable when estimated propensity scores are near 0 or 1 (Rosenbaum and Rubin 1983).

This chapter develops a third strategy: combine both models to construct an estimator that is consistent when *either* model is correct, and achieves the semiparametric efficiency bound when both are correct. We derive the same estimator in three complementary ways: as a bias-corrected prediction estimator (Section 10.2), as the optimal member of a class of augmented estimators (Section 10.4), and as the efficient influence function in a semiparametric model (Section 10.7).

10.2 The Prediction Estimator and Its Bias

Working with generic *working regression functions* $m_t(x)$, the **prediction estimator** is $\hat{\tau}_{\text{pred}}(m) = n^{-1} \sum_{i=1}^n \{m_1(X_i) - m_0(X_i)\}$. The population bias is:

$$\tau_{\text{pred}}(m) - \tau = \mathbb{E}[(m_1(X) - \mu_1(X)) - (m_0(X) - \mu_0(X))]. \quad (10.1)$$

If a propensity score estimator $\hat{\pi}$ is available, we can estimate and subtract this bias. The bias-corrected prediction estimator is:

$$\hat{\tau}_{\text{AIPW}} = \hat{\tau}_{\text{pred}} - \widehat{\text{Bias}} = \hat{\mu}_{1,\text{dr}} - \hat{\mu}_{0,\text{dr}}, \quad (10.2)$$

where:

$$\begin{aligned} \hat{\mu}_{1,\text{dr}} &= \frac{1}{n} \sum_{i=1}^n \frac{T_i}{\hat{\pi}(X_i)} \{Y_i - \hat{\mu}_1(X_i)\} + \frac{1}{n} \sum_{i=1}^n \hat{\mu}_1(X_i), \\ \hat{\mu}_{0,\text{dr}} &= \frac{1}{n} \sum_{i=1}^n \frac{1 - T_i}{1 - \hat{\pi}(X_i)} \{Y_i - \hat{\mu}_0(X_i)\} + \frac{1}{n} \sum_{i=1}^n \hat{\mu}_0(X_i). \end{aligned}$$

Remark: Two Roles of the Augmentation Terms

From the bias-correction perspective, $T_i \{Y_i - \hat{\mu}_1(X_i)\} / \hat{\pi}(X_i)$ are IPW-weighted residuals estimating prediction error. From an estimating-equation perspective, they make the estimating function orthogonal to nuisance perturbations. The estimated bias term has zero expectation when the outcome model is correct; the bias-corrected estimator is unbiased when the propensity model is correct. Hence the name *doubly robust*.

10.3 The Augmented IPW Estimator

Notation. $\mu_t(x)$ denotes a generic *working* outcome regression; the truth is $\mu_t^*(x)$. The propensity score $\pi(x)$ is always the *true* propensity.

The AIPW estimator solves $\mathbb{P}_n\{\phi(O; \tau, \eta)\} = 0$ where:

$$\phi(O; \tau, \eta) = \left[\frac{T}{\pi(X)} \{Y - \mu_1(X)\} + \mu_1(X) \right] - \left[\frac{1 - T}{1 - \pi(X)} \{Y - \mu_0(X)\} + \mu_0(X) \right] - \tau. \quad (10.3)$$

Solving explicitly:

$$\hat{\tau}_{\text{AIPW}} = \frac{1}{n} \sum_{i=1}^n \left[\hat{\mu}_1(X_i) - \hat{\mu}_0(X_i) + \frac{T_i}{\hat{\pi}(X_i)} \{Y_i - \hat{\mu}_1(X_i)\} - \frac{1 - T_i}{1 - \hat{\pi}(X_i)} \{Y_i - \hat{\mu}_0(X_i)\} \right]. \quad (10.4)$$

Theorem: Double Robustness

Let $\phi(O; \tau, \eta)$ be the estimating function in Equation 10.3. Under consistency, conditional exchangeability, and positivity, $\mathbb{E}\{\phi(O; \tau, \eta)\} = 0$ if either:

1. $\mu_t(x) = \mu_t^*(x)$ for $t = 0, 1$, regardless of whether $\pi(x)$ is correctly specified; or
2. $\pi(x) = P(T = 1 \mid X = x)$, regardless of whether μ_0 and μ_1 are correctly specified.

Proof

Case 1: outcome regression correct. If $\mu_t(X) = \mathbb{E}(Y \mid T = t, X)$, then $\mathbb{E}\left[\frac{T}{\pi(X)} \{Y - \mu_1(X)\} \mid X\right] = \frac{P(T=1|X)}{\pi(X)} \mathbb{E}\{Y - \mu_1(X) \mid T = 1, X\} = 0$, and similarly for the control arm. Therefore $\mathbb{E}\{\phi\} = \mathbb{E}\{\mu_1(X) - \mu_0(X)\} - \tau = 0$.

Case 2: propensity score correct. Suppose $\pi(X) = P(T = 1 \mid X)$. For the treated bracket:

$$\mathbb{E}\left[\frac{T}{\pi(X)} \{Y - \mu_1(X)\} + \mu_1(X) \mid X\right] = \frac{\mathbb{E}[TY \mid X]}{\pi(X)} - \mu_1(X) \cdot \frac{\mathbb{E}[T \mid X]}{\pi(X)} + \mu_1(X).$$

Using $\mathbb{E}[T | X] = \pi(X)$ (so the μ_1 terms cancel) and $\mathbb{E}[TY | X] = \pi(X)\mu_1^*(X)$, this equals $\mu_1^*(X)$. By symmetry, the control bracket gives $\mu_0^*(X)$. Hence $\mathbb{E}\{\phi | X\} = \mu_1^*(X) - \mu_0^*(X) - \tau$, and taking expectations gives $\mathbb{E}\{\phi\} = 0$. $\square$

Remark: The Algebra of Case 2

The key cancellation: $\mu_1(X)$ enters the bracket twice with equal and opposite conditional expectations because $\mathbb{E}[T/\pi(X) | X] = 1$ when π is correct. The bracket “forgets” μ_1 entirely and converges to $\mu_1^*(X)$. A wrong μ_t is subtracted out by the correct π .

10.4 A Class of Augmented Estimators

For any square-integrable functions $b_1(x)$ and $b_0(x)$, define:

$$\hat{\mu}_{1,b} = \frac{1}{n} \sum_{i=1}^n \frac{T_i}{\pi(X_i)} Y_i - \frac{1}{n} \sum_{i=1}^n \left\{ \frac{T_i}{\pi(X_i)} - 1 \right\} b_1(X_i),$$

$$\hat{\mu}_{0,b} = \frac{1}{n} \sum_{i=1}^n \frac{1-T_i}{1-\pi(X_i)} Y_i - \frac{1}{n} \sum_{i=1}^n \left\{ \frac{1-T_i}{1-\pi(X_i)} - 1 \right\} b_0(X_i),$$

and $\hat{\tau}_b = \hat{\mu}_{1,b} - \hat{\mu}_{0,b}$. The correction is mean-zero under the true π for any b_t , so every member is unbiased; only the variance depends on b_t . When $b_t = \mu_t$, $\hat{\mu}_{t,b}$ equals the AIPW form.

Theorem: Unbiasedness and Variance of the AIPW Class

Under conditional exchangeability, SUTVA, and positivity, with the true propensity score, for any square-integrable b_0, b_1 : (1) $\mathbb{E}(\hat{\tau}_b) = \tau$; and (2):

$$\text{Var}(\hat{\tau}_b) = \frac{1}{n} \mathbb{E} \left[\left(\frac{1}{\pi(X)} - 1 \right) \{Y(1) - b_1(X)\}^2 + 2\{Y(1) - b_1(X)\}\{Y(0) - b_0(X)\} + \left(\frac{1}{1-\pi(X)} - 1 \right) \{Y(0) - b_0(X)\}^2 \right] \quad (10.5)$$

Theorem: Optimal Control Functions

The control functions $b_1^*(X) = \mathbb{E}\{Y(1) | X\}$ and $b_0^*(X) = \mathbb{E}\{Y(0) | X\}$ minimize the total variance in Equation 10.5. The proof shows that the b_t -dependent term equals a perfect square $\mathbb{E} \left[\left(\sqrt{\frac{1-\pi}{\pi}} u + \sqrt{\frac{\pi}{1-\pi}} v \right)^2 \right] \geq 0$, where $u = \mu_1^* - b_1$ and $v = \mu_0^* - b_0$, which equals zero iff $b_t = \mu_t^*$.

10.5 The Projection Interpretation

The optimality of AIPW has an elegant geometric interpretation in terms of projections in a Hilbert space of estimators (Tsiatis 2006, Ch. 2). Define the **augmentation space** $\Lambda = \{\hat{b} : \mathbb{E}(\hat{b}) = 0\}$.

Theorem: Optimal Projection

The optimal correction is $\hat{b}^* = \Pi(\hat{\theta}_0 | \Lambda)$, characterized by $\hat{b}^* \in \Lambda$ and $\text{Cov}(\hat{\theta}_0 - \hat{b}^*, \hat{b}) = 0$ for all $\hat{b} \in \Lambda$. The optimal estimator $\hat{\theta}_{\text{opt}} = \Pi(\hat{\theta}_0 | \Lambda^\perp)$ satisfies the **Pythagorean identity**:

$$\text{Var}(\hat{\theta}_0) = \text{Var}(\hat{\theta}_{\text{opt}}) + \text{Var}(\hat{b}^*). \quad (10.6)$$




Figure 10.1: Geometric illustration of the projection. $\hat{\theta}_0$ decomposes into $\hat{\theta}_{\text{opt}} \in \Lambda^\perp$ and $\hat{b}^* \in \Lambda$. The Pythagorean identity gives $\text{Var}(\hat{\theta}_0) = \text{Var}(\hat{\theta}_{\text{opt}}) + \text{Var}(\hat{b}^*)$.

10.6 Projection in the Causal Inference Setting

We now specialize to the ATE, with $\pi(X)$ known. Decompose the Horvitz–Thompson (HT) estimator as $\hat{\tau}_{\text{HT}} = \hat{\mu}_{1,\text{HT}} - \hat{\mu}_{0,\text{HT}}$. Define arm-specific augmentation spaces:

$$\Lambda_t = \left\{ n^{-1} \sum_{i=1}^n \left(\frac{T_i^{(t)}}{\pi_i^{(t)}} - 1 \right) b_t(X_i) : b_t \in \mathcal{L}^2 \right\}, \quad \Lambda = \Lambda_1 \oplus \Lambda_0.$$

Theorem: Arm-wise Projections

The optimal arm-wise corrections are given by $b_t^*(x) = \mathbb{E}\{Y(t) \mid x\}$:

$$\Pi(\hat{\mu}_{t,\text{HT}} \mid \Lambda_t) = n^{-1} \sum_{i=1}^n \left(\frac{T_i^{(t)}}{\pi_i^{(t)}} - 1 \right) \mu_t^*(X_i).$$

The orthogonality condition $\mathbb{E}[r_{t,i} b_i] = 0$ holds because $\mathbb{E}\{\mu_t^*(X_i) - \mu_t^*(X_i) \mid X_i\} = 0$ and $\mathbb{E}[T_i/\pi_i \cdot (T_i/\pi_i - 1) \mid X_i] = 1/\pi_i - 1$.

The arm-wise optimal estimators are:

$$\hat{\mu}_{t,\text{opt}} = \frac{1}{n} \sum_{i=1}^n \left[\frac{T_i^{(t)}}{\pi_i^{(t)}} \{Y_i - \mu_t^*(X_i)\} + \mu_t^*(X_i) \right],$$

and $\hat{\tau}_{\text{opt}} = \hat{\mu}_{1,\text{opt}} - \hat{\mu}_{0,\text{opt}}$ is precisely the AIPW estimator Equation 10.4, confirming the Optimal Control Functions theorem by a different route.

10.7 The Efficient Influence Function and Semiparametric Efficiency

Under the nonparametric model for $O = (X, T, Y)$, the efficient influence function for the ATE is:

$$\varphi_{\text{eff}}(O) = \frac{T}{\pi(X)} \{Y - \mu_1(X)\} - \frac{1-T}{1-\pi(X)} \{Y - \mu_0(X)\} + \mu_1(X) - \mu_0(X) - \tau. \quad (10.7)$$

Comparing Equation 10.7 with Equation 10.3, the AIPW estimating function *is* the efficient influence function $\varphi^*(O)$ from Chapter 9. The augmentation space Λ is the nuisance tangent space of the nonparametric model.

Theorem: Semiparametric Efficiency Bound for the ATE

Under the nonparametric model for $O = (X, T, Y)$, every regular asymptotically linear estimator $\hat{\tau}$ of the ATE satisfies $\text{Avar}(\sqrt{n}(\hat{\tau} - \tau)) \geq \mathbb{E}\{\varphi_{\text{eff}}(O)^2\}$, and the bound is attained by estimators whose influence function equals φ_{eff} .

This is a specialization of the Hájek–Le Cam convolution theorem; see Bickel et al. (1993) or Vaart (1998, Ch. 25) for a rigorous development.

Remark: EIF as the Semiparametric Score

The efficient influence function plays the same role in semiparametric theory that the score plays in parametric models: its variance gives the information lower bound. The three derivations in this chapter — bias correction, optimal augmentation, and semiparametric identification — converge on the same estimating function. This is the deepest explanation of why AIPW occupies a central place in the causal inference toolkit.

10.8 Doubly Robust Regression: Weighted and Augmented Approaches

The **internal bias calibration** (IBC) conditions (Firth and Bennett 1998) enforce double robustness directly within the outcome model fit: when satisfied, the prediction estimator coincides with AIPW without a separate augmentation step:

$$\sum_{i=1}^n \frac{T_i}{\hat{\pi}_i} \{Y_i - \hat{\mu}_1(X_i)\} = 0, \quad \sum_{i=1}^n \frac{1 - T_i}{1 - \hat{\pi}_i} \{Y_i - \hat{\mu}_0(X_i)\} = 0. \quad (10.8)$$

Weighted regression. Use IPW-weighted least squares: minimize $\sum_{i=1}^n \frac{T_i}{\hat{\pi}_i} \{Y_i - \mu_1(X_i; \theta_1)\}^2$ for arm $t = 1$. The normal equation for the intercept is exactly the IBC condition Equation 10.8. IPW-weighted fitted values automatically satisfy IBC for any model containing a constant (Robins et al. 1994; Bang and Robins 2005).

Augmented model with clever covariate. Let $\hat{\mu}_1^{(0)}(X_i)$ be any initial fit. Include the **clever covariate** $\hat{\pi}_i^{-1}$ (Laan and Rubin 2006; Laan and Rose 2011) and run OLS among treated units:

$$Y_i = \alpha_1 + \beta_1 \hat{\mu}_1^{(0)}(X_i) + \gamma_1 \hat{\pi}_i^{-1} + e_i(1), \quad T_i = 1. \quad (10.9)$$

The fitted model is $\hat{\mu}_1(X_i) = \hat{\alpha}_1 + \hat{\beta}_1 \hat{\mu}_1^{(0)}(X_i) + \hat{\gamma}_1 \hat{\pi}_i^{-1}$. The normal equation for $\hat{\gamma}_1$ is exactly the IBC condition, so the prediction estimator can be written in AIPW form:

$$\hat{\mu}_1^{\text{pred}} = \frac{1}{n} \sum_{i=1}^n \hat{\mu}_1(X_i) + \frac{1}{n} \sum_{i=1}^n \frac{T_i}{\hat{\pi}_i} \{Y_i - \hat{\mu}_1(X_i)\}. \quad (10.10)$$

Remark: Behavior Under Correct Outcome Model

If $\hat{\mu}_1^{(0)}$ is correctly specified, then $\hat{\alpha}_1 \xrightarrow{p} 0$, $\hat{\beta}_1 \xrightarrow{p} 1$, and $\hat{\gamma}_1 \xrightarrow{p} 0$. Including the clever covariate carries no asymptotic efficiency cost.

Remark: Connection to TMLE

This approach captures the targeting idea underlying TMLE (Laan and Rubin 2006; Laan and Rose 2011). Standard TMLE uses the fixed-offset form (coefficient on $\hat{\mu}_1^{(0)}$ fixed at 1). Model Equation 10.9 goes further by freely estimating $\hat{\beta}_1$. See Laan and Rose (2018) for a comprehensive treatment.

10.9 Lab: Simulation Study

This lab compares four estimators of the ATE: $\hat{\tau}_{\text{HT}}$, $\hat{\tau}_{\text{AIPW}}$, the fixed-offset augmented-model estimator $\hat{\tau}_{\text{aug}}$, and the improved augmented-model estimator $\hat{\tau}_{\text{aug}^+}$ of Equation 10.9. A 2×2 design over nuisance-model correctness — OR correct or misspecified, crossed with PS correct or misspecified — demonstrates double robustness directly. The lab also reports empirical 95% Wald coverage, illustrating the distinction between double-robust consistency and valid asymptotic inference.

Data-generating process. $n = 1000$ observations. Draw $X_i \sim N(0, 1)$, then $T_i \mid X_i \sim \text{Bernoulli}(\pi^*(X_i))$ with:

$$\pi^*(x) = \text{expit}\{0.2x + 0.2(x^2 - 1)\}.$$

The potential outcomes are $Y_i(1) = 1 + X_i + 0.5X_i^2 + \varepsilon_i(1)$ and $Y_i(0) = X_i + 0.5X_i^2 + \varepsilon_i(0)$, with $\varepsilon_i(t) \sim N(0, 1)$ i.i.d., true ATE $\tau = 1$. The quadratic term enters both the true OR and true PS; a linear-only fit omits it.

The 2×2 design. The fitted $\hat{\pi}$ is clipped to $[10^{-3}, 1 - 10^{-3}]$.

Scenario	OR fit	PS fit
S1	correct: $Y \sim (1, X, X^2)$	correct: $T \sim (1, X, X^2)$
S2	correct: $Y \sim (1, X, X^2)$	misspecified: $T \sim (1, X)$
S3	misspecified: $Y \sim (1, X)$	correct: $T \sim (1, X, X^2)$
S4	misspecified: $Y \sim (1, X)$	misspecified: $T \sim (1, X)$

Results ($B = 2000$ replications, `set.seed(2025)`). Bias, Var, MSE in units of 10^{-3} ; Cov = empirical 95% Wald coverage.

Scenario	Metric	$\hat{\tau}_{\text{HT}}$	$\hat{\tau}_{\text{AIPW}}$	$\hat{\tau}_{\text{aug}}$	$\hat{\tau}_{\text{aug}^+}$
S1: OR , PS	Bias	0.55	−0.44	−0.45	−0.45
	Var	6.82	4.25	4.25	4.30
	MSE	6.82	4.25	4.25	4.30
	Cov (%)	99.9	95.0	94.9	94.5
S2: OR , PS	Bias	188.56	0.62	0.62	0.62
	Var	6.32	4.33	4.33	4.33
	MSE	41.87	4.32	4.32	4.32
	Cov (%)	67.4	94.2	94.2	94.2
S3: OR , PS	Bias	1.46	2.02	2.25	−25.26
	Var	5.74	4.71	4.41	8.02
	MSE	5.74	4.71	4.41	8.66
	Cov (%)	99.9	98.4	98.2	93.2
S4: OR , PS	Bias	188.42	188.77	188.62	−5.42
	Var	6.28	6.30	6.29	4.26
	MSE	41.78	41.93	41.87	4.29
	Cov (%)	67.5	33.1	33.1	94.7

S1 (both correct). All four estimators are approximately unbiased. $\hat{\tau}_{\text{HT}}$ has about 60% higher MSE (6.82 vs 4.25×10^{-3}) because it discards the outcome model. The three augmented estimators are essentially equivalent, confirming that $(\hat{\alpha}_t, \hat{\beta}_t, \hat{\gamma}_t) \rightarrow (0, 1, 0)$ when the OR is correct. HT's coverage is above nominal because its influence function has larger variance than φ_{eff} .

S2 (OR correct, PS misspecified). HT is badly biased ($\text{MSE} \approx 42 \times 10^{-3}$, coverage 67%). The three augmented estimators are indistinguishable: when the OR is correct they converge to the same limit regardless of PS specification. This is the **first direct demonstration of double robustness**: PS misspecification is absorbed by the correctly specified OR model.

S3 (OR misspecified, PS correct). HT, AIPW, and $\hat{\tau}_{\text{aug}}$ remain consistent through the correct PS (biases under 3×10^{-3} , coverage close to nominal). $\hat{\tau}_{\text{aug}}$ achieves the lowest MSE (4.41×10^{-3}): the targeted regression reduces residual variance without disturbing the IBC constraint. $\hat{\tau}_{\text{aug}^+}$ shows a finite-sample bias (-25×10^{-3}) from high-leverage clever-covariate values. This is the **second double-robustness demonstration**: OR misspecification is absorbed by the correctly specified PS.

S4 (both misspecified). Double robustness provides no guarantee. HT, AIPW, and $\hat{\tau}_{\text{aug}}$ are all badly biased (MSE $\approx 42 \times 10^{-3}$); Wald coverage collapses to 33% for AIPW and aug — the variance estimator is concentrated around the wrong value. The apparent “recovery” of $\hat{\tau}_{\text{aug}^+}$ (coverage 94.7%) is a DGP-specific artifact: the regressors $(1, \hat{\mu}_t^{(0)}(X), \hat{\pi}^{-1}(X))$ happen to capture the omitted quadratic curvature through the nonlinearity of $\hat{\pi}^{-1}$, but this does **not** indicate triple robustness.

Takeaway. Comparing S2–S3 against S4 is the crisp illustration of the Double Robustness theorem: AIPW is consistent whenever at least one nuisance model is correct, and only S4 breaks the guarantee. Wald coverage mirrors this pattern. The lab also shows that $\hat{\tau}_{\text{aug}}$ can reduce MSE in S3 by shrinking residual variance through IBC-enforced regression — the practical motivation for targeting methods such as TMLE.

10.10 Asymptotic Inference with Estimated Nuisance Functions

Under suitable regularity conditions:

$$\sqrt{n}(\hat{\tau}_{\text{AIPW}} - \tau) = \frac{1}{\sqrt{n}} \sum_{i=1}^n \varphi_{\text{eff}}(O_i) + o_p(1) \xrightarrow{d} N(0, \mathbb{E}\{\varphi_{\text{eff}}(O)^2\}).$$

Neyman orthogonality implies nuisance estimation error enters only through a second-order remainder. A sufficient condition for this remainder to be negligible is the **product-rate condition**:

$$\|\hat{\pi} - \pi\| \cdot \|\hat{\mu}_t - \mu_t\| = o_p(n^{-1/2}), \quad t = 0, 1, \quad (10.11)$$

where $\|\cdot\|$ denotes the $L_2(P)$ norm (equivalently, the root-mean-squared-error under the data distribution). This allows each nuisance estimator to converge at rate as slow as $n^{-1/4}$. In modern applications, this is usually paired with **cross-fitting**; Chapter 11 develops this in detail.

Remark: Neyman Orthogonality

The first-order insensitivity of the AIPW estimating function to nuisance perturbations is an instance of *Neyman orthogonality*. Chapter 11 exploits this via cross-fitting, yielding the *double/debiased machine learning* (DML) estimator (Chernozhukov et al. 2018).

Compute plug-in influence values:

$$\hat{\varphi}_i = \frac{T_i}{\hat{\pi}(X_i)} \{Y_i - \hat{\mu}_1(X_i)\} - \frac{1 - T_i}{1 - \hat{\pi}(X_i)} \{Y_i - \hat{\mu}_0(X_i)\} + \hat{\mu}_1(X_i) - \hat{\mu}_0(X_i) - \hat{\tau}_{\text{AIPW}}, \quad (10.12)$$

and set:

$$\hat{V} = \frac{1}{n(n-1)} \sum_{i=1}^n (\hat{\varphi}_i - \bar{\varphi})^2, \quad \bar{\varphi} = \frac{1}{n} \sum_{i=1}^n \hat{\varphi}_i. \quad (10.13)$$

The Wald CI is $\hat{\tau}_{\text{AIPW}} \pm z_{1-\alpha/2} \sqrt{\hat{V}}$.

Remark: Double Robustness versus Efficient Inference

It is important not to conflate two distinct results. *Double-robust consistency* requires only that *one* of the two nuisance models be correctly specified; no rate conditions are imposed. *Efficient asymptotic inference* — the $\sqrt{n}$ -normal expansion and the Wald interval based on $\hat{V}$ — is a strictly stronger requirement. It requires *both* nuisance functions to be estimated consistently at rates satisfying the product-rate condition Equation 11.7, together with overlap and moment conditions. Plugging estimated nuisance functions into Equation 10.13 produces valid standard errors only

when the asymptotic linear representation holds; that representation does *not* follow from double robustness alone. In particular, if the outcome regression is misspecified but the propensity score is correct, the AIPW estimator remains consistent but the influence-function variance estimator targets the variance of an estimator whose nuisance functions converge to their true values, not to misspecified limits. The simulation in Section 11.10 illustrates this: coverage is approximately nominal in S2 and S3 (where double robustness operates), but this relies on the surviving nuisance model being correctly specified, not just present. Careful rate verification — or cross-fitting as in Chapter 11 — is essential before reporting the Wald interval.

Extreme Propensity Scores

When $\hat{\pi}(X_i)$ is near 0 or 1, the IPW weights become large and can destabilize the estimator. Practical remedies: overlap diagnostics, truncation of extreme propensity scores, or restricting to a subgroup with adequate overlap. The augmented model approach of Section 10.8 provides a complementary strategy.

10.11 Comparison of Regression, IPW, and AIPW

Estimator	Uses μ_t	Uses π	Consistent if	Main weakness
Regression	Yes	No	outcome model correct	Sensitive to OR misspecification
IPW	No	Yes	propensity model correct	Unstable under weak overlap
AIPW	Yes	Yes	either model correct	Requires both nuisances and careful inference

10.12 Chapter Summary

Symbol	Meaning
τ	ATE = $\mathbb{E}\{Y(1) - Y(0)\}$
$\mu_t(x)$	Outcome regression $\mathbb{E}(Y \mid T = t, X = x)$
$\pi(x)$	Propensity score $P(T = 1 \mid X = x)$
$b_t(x)$	Control function; optimal $b_t^*(x) = \mathbb{E}\{Y(t) \mid X = x\}$
$\hat{\tau}_{\text{AIPW}}$	AIPW estimator Equation 10.4
Λ	Augmentation space $\{\hat{b} : \mathbb{E}(\hat{b}) = 0\}$
$\varphi_{\text{eff}}^{(O)}$	Efficient influence function Equation 10.7
$\hat{\pi}_i^{-1}$	Clever covariate; enforces IBC via its normal equation
IBC	Internal bias calibration conditions Equation 10.8

1. The prediction estimator bias can be estimated and subtracted using the propensity score, yielding the AIPW estimator.
2. AIPW has two key properties: **double robustness** (consistent when either model is correct) and **class-optimality** ($b_t^* = \mu_t^*$ minimizes variance).
3. The optimality follows also from a Hilbert-space projection onto $\Lambda^\perp$; the Pythagorean identity governs the variance reduction.
4. The AIPW estimating function *is* the efficient influence function; three derivations converge on the same object.
5. Double robustness can be built directly into the outcome regression via IPW-weighted regression or including the clever covariate $\hat{\pi}_i^{-1}$ (TMLE connection).

6. The simulation (Section 11.10) demonstrates double robustness across a 2×2 design: AIPW remains consistent whenever at least one nuisance model is correct (S1–S3) and loses the guarantee when both are misspecified (S4). Wald coverage is approximately nominal in S1–S3 but collapses in S4, illustrating that double-robust consistency and valid asymptotic inference are distinct properties.
7. For asymptotic inference, the central issue is whether the effect of estimating nuisance functions is asymptotically negligible. Neyman orthogonality yields a second-order remainder, and the product-rate condition Equation 11.7 is sufficient for root- n inference. In modern machine-learning settings, cross-fitting is often used to make these conditions more plausible; Chapter 11 develops this approach.

10.13 Exercises

1. Bias of the prediction estimator.

- (a) Verify the bias formula Equation 10.1 by writing $\tau_{\text{pred}}(m) - \tau$ as a difference of working-model errors and simplifying.
- (b) Under what condition on the propensity score model does $\widehat{\text{Bias}}(\hat{\tau}_{\text{pred}})$ have zero expectation? Give a careful argument using iterated expectations.
- (c) Construct a simple example (binary X , binary T , binary Y) in which the estimated bias is nonzero and the outcome model is misspecified, but the AIPW estimator is still consistent.

2. The AIPW class and optimal control functions.

- (a) Verify that $\hat{\tau}_b$ with $b_t = \mu_t$ equals the AIPW estimator Equation 10.4.
- (b) Using Equation 10.5, show that $b_t^*(x) = \mathbb{E}\{Y(t) \mid x\}$ minimizes conditional variance by completing the square in b_t .
- (c) Suppose $\pi(x) = 1/2$ for all x . Simplify Equation 10.5 and interpret the result.

3. Double robustness.

- (a) Verify Case 1 of the Double Robustness theorem: with $\mu_t = \mu_t^*$ and arbitrary π , show $\mathbb{E}\{\phi\} = 0$.
- (b) Verify Case 2: with $\pi = \pi^*$ and arbitrary μ_t , show $\mathbb{E}\{\phi\} = 0$.
- (c) Provide a counterexample showing $\mathbb{E}\{\phi\} \neq 0$ when both models are misspecified.

4. Projection and the Pythagorean identity.

- (a) Verify $\text{Cov}(\hat{\theta}_{\text{opt}}, \hat{b}^*) = 0$ from the definition of $\hat{b}^*$.
- (b) Deduce Equation 10.6 from the decomposition $\hat{\theta}_0 = \hat{\theta}_{\text{opt}} + \hat{b}^*$.
- (c) Explain why $\hat{\theta}_{\text{opt}}$ is still unbiased for θ even though $\hat{b}^* \in \Lambda$ is subtracted.

5. Semiparametric efficiency.

- (a) Show $\mathbb{E}\{\varphi_{\text{eff}}(O)^2\} \leq \mathbb{E}\{\varphi_{\text{IPW}}(O)^2\}$ by writing $\varphi_{\text{IPW}} = \varphi_{\text{eff}} + (\varphi_{\text{IPW}} - \varphi_{\text{eff}})$ and showing the cross term vanishes.
- (b) Interpret the efficiency gain $\mathbb{E}\{\varphi_{\text{IPW}}^2\} - \mathbb{E}\{\varphi_{\text{eff}}^2\}$ in terms of the variance of the outcome regression.
- (c) Give one reason why an efficient estimator based on φ_{eff} may still be disfavored in practice.

6. Augmented model and the clever covariate.

- (a) Show that the normal equation for $\hat{\gamma}_1$ in the augmented model Equation 10.9 is exactly the IBC condition Equation 10.8, and conclude that the prediction estimator can always be written in AIPW form Equation 10.10.
- (b) Explain why $\hat{\gamma}_1 \xrightarrow{P} 0$ when the initial model $\hat{\mu}_1^{(0)}$ is correctly specified.
- (c) Explain in words why $\hat{\pi}_i^{-1}$ is called the clever covariate: what bias does it absorb, and why does this make the prediction estimator doubly robust?

7. Coverage of the Wald confidence interval (computational). Use the DGP of Section 11.10 with the propensity score generalized to $\pi^*(x) = \text{expit}\{0.2x + \gamma(x^2 - 1)\}$, so that the baseline ($\gamma = 0.2$) coincides with the lab and larger γ degrades overlap. Implement the AIPW estimator and variance estimator Equation 10.13. Following Scenario S3 of the lab, use logistic regression of T on $(1, X, X^2)$ for $\hat{\pi}$ (correct PS) and OLS of Y on $(1, X)$ separately in each arm for $\hat{\mu}_t$ (misspecified OR); truncate $\hat{\pi}$ to $[10^{-3}, 1 - 10^{-3}]$.

- (a) With $\gamma = 0.2$, $n = 500$, $B = 2000$ replications, compute empirical coverage of the 95% Wald interval $\hat{\tau}_{\text{AIPW}} \pm z_{0.975} \sqrt{\hat{V}}$. Compare the average SE $\sqrt{\hat{V}}$ with the empirical SD of $\hat{\tau}_{\text{AIPW}}$ across replications.
- (b) Repeat with $\gamma \in \{0.5, 1.0\}$, producing increasingly weak overlap at large $|X|$. Report coverage, average SE, and empirical SD. What do you observe?
- (c) Which assumption of Section 10.10 is stressed as γ grows, and which of the remedies in the Extreme Propensity Scores warning would you try first? (No simulation required; a paragraph suffices.)

Chapter 11

Flexible Nuisance Estimation, Orthogonal Scores, and Cross-Fitting

Learning Objectives

By the end of this chapter, students should be able to:

1. Explain why flexible nuisance estimation can produce misleading inference for a causal parameter even when predictive accuracy is high.
2. Define Neyman orthogonality, verify it for a given estimating function, and explain how it reduces sensitivity to first-order nuisance estimation error.
3. Identify the orthogonal score for the ATE, recognize it as the efficient influence function from Chapter 10, and state the three roles it simultaneously fulfills.
4. Describe sample splitting, explain why independence between nuisance estimation and score evaluation simplifies asymptotic arguments, and articulate its efficiency cost.
5. Describe K -fold cross-fitting, write the cross-fitted moment equation explicitly, and explain how cross-fitting recovers efficiency relative to a single split.
6. State the product-rate condition for the cross-fitted AIPW estimator, connect it to Neyman orthogonality, and explain why it allows each nuisance estimator to converge at rate $n^{-1/4}$.
7. Construct a consistent variance estimator from out-of-fold estimated influence values and form a Wald confidence interval.
8. Articulate what machine learning can and cannot contribute to causal inference, and follow the practical workflow for cross-fitted orthogonal-score estimation.

11.1 Why Flexible Nuisance Estimation Is Both Attractive and Dangerous

In Chapters 9 and 10 we developed estimation and inference using estimating equations, influence functions, doubly robust scores, and semiparametric efficiency. The causal parameter of interest depends on nuisance functions such as the outcome regressions $\mu_t(x) = \mathbb{E}(Y \mid T = t, X = x)$ and the propensity score $\pi(x) = P(T = 1 \mid X = x)$. In many real applications, these may be too complex to model accurately with low-dimensional parametric forms, making flexible methods attractive: penalized regression, splines, random forests, boosting, neural networks, and ensemble learning.

These methods can substantially improve predictive accuracy, but they also create new statistical difficulties: they may converge more slowly than parametric estimators; they may overfit the observed sample; their asymptotic behavior may be difficult to characterize; and naive plug-in inference can fail even when prediction quality is good.

Better nuisance prediction therefore does not automatically imply valid inference for the target causal parameter. The goal of this chapter is to explain how orthogonal scores, sample splitting, and cross-fitting make it possible to combine flexible nuisance estimation with valid large-sample inference.

11.2 Why Naive Plug-In Estimation Can Fail

Suppose the parameter of interest is $\psi = \Psi(\eta)$, where η denotes a nuisance object such as (μ_0, μ_1, π) . The *plug-in estimator* $\hat{\psi}_{\text{plug}} = \Psi(\hat{\eta})$ has a first-order Taylor expansion around the true nuisance η_0 :

$$\hat{\psi}_{\text{plug}} - \psi = \underbrace{D_{\eta}\Psi(\eta_0)[\hat{\eta} - \eta_0]}_{\text{linear term}} + \underbrace{R(\hat{\eta}, \eta_0)}_{\text{second-order remainder}}, \quad (11.1)$$

where $D_{\eta}\Psi(\eta_0)[h]$ denotes the Gateaux derivative and $|R(\hat{\eta}, \eta_0)| \lesssim \|\hat{\eta} - \eta_0\|^2$.

The second-order remainder is typically manageable: if $\|\hat{\eta} - \eta_0\| = o_p(n^{-1/4})$ then $R = o_p(n^{-1/2})$. The obstacle is the *linear term*, proportional to the nuisance estimation error.

Finite-dimensional nuisance. If $\eta \in \mathbb{R}^d$ and $\hat{\eta} - \eta_0 = O_p(n^{-1/2})$, the linear term is $O_p(n^{-1/2})$ and contaminates the root- n limiting distribution. If the functional satisfies $D_{\eta}\Psi(\eta_0) = 0$ — *Neyman orthogonality* — the linear term vanishes and the expansion reduces to the second-order remainder alone, $O_p(n^{-1})$, asymptotically negligible.

Infinite-dimensional nuisance. When η belongs to a function space, orthogonality removes the *population bias* but does not control the *empirical process term*. Writing the plug-in estimating equation as $\mathbb{P}_n\{\phi(\cdot; \psi, \hat{\eta})\} = 0$, the deviation from the population equation decomposes as:

$$\mathbb{P}_n\{\phi(\cdot; \psi, \hat{\eta})\} - \mathbb{P}_n\{\phi(\cdot; \psi, \eta_0)\} = \underbrace{P\{\phi(\cdot; \psi, \hat{\eta}) - \phi(\cdot; \psi, \eta_0)\}}_{\text{population bias}} + \underbrace{(\mathbb{P}_n - P)\{\phi(\cdot; \psi, \hat{\eta}) - \phi(\cdot; \psi, \eta_0)\}}_{\text{empirical process term}}. \quad (11.2)$$

Neyman orthogonality controls the *population bias*: because the Gateaux derivative vanishes at the truth, the population bias is second order in $\|\hat{\eta} - \eta_0\|$. For flexible machine-learning estimators, the *empirical process term* can remain non-negligible even as $\hat{\eta}$ converges consistently. This is why orthogonality and cross-fitting are *complements* rather than substitutes.

Remark

This problem is not specific to machine learning. It arises whenever nuisance parameters are estimated in an infinite-dimensional model. Machine learning makes the issue more visible because it encourages highly adaptive nuisance estimation; see Chernozhukov et al. (2018) for a systematic treatment.

11.3 Orthogonal Scores

Definition: Neyman Orthogonality

The score function $\phi(O; \psi, \eta)$ is called **orthogonal** (or **Neyman orthogonal**) at the truth (ψ_0, η_0) if the Gateaux derivative of the estimating equation with respect to the nuisance, evaluated at the truth, vanishes:

$$\left. \frac{\partial}{\partial r} \mathbb{E}\{\phi(O; \psi_0, \eta_0 + rh)\} \right|_{r=0} = 0$$

for all perturbation directions h in a suitable class.

This means small local perturbations of the nuisance parameter have no first-order effect on the estimating equation at the truth. Orthogonality does not mean nuisance estimation no longer matters: it controls only the population-level bias, and a separate argument (cross-fitting) is needed to rule out the additional bias from evaluating the score on the same sample used to estimate $\hat{\eta}$.

Remark: Neyman's Original Insight

The terminology honors Jerzy Neyman, whose work on hypothesis testing introduced the idea of constructing test statistics that are insensitive to nuisance parameters. In the modern semiparametric literature the concept was formalized by Robins et al. (1994) and others, and later made systematic

in the double machine learning framework of Chernozhukov et al. (2018).

11.4 The Orthogonal Score for the Average Treatment Effect

The efficient influence function from Chapter 10 is:

$$\varphi_{\text{eff}}(O; \tau, \eta) = \frac{T}{\pi(X)} \{Y - \mu_1(X)\} - \frac{1-T}{1-\pi(X)} \{Y - \mu_0(X)\} + \mu_1(X) - \mu_0(X) - \tau, \quad (11.3)$$

where $\eta = (\mu_0, \mu_1, \pi)$. We now prove this score is Neyman orthogonal.

Lemma: Neyman Orthogonality of the ATE Score

The score $\varphi_{\text{eff}}(O; \tau, \eta)$ in Equation 11.3 is Neyman orthogonal at every truth (τ_0, η_0) satisfying the identification assumptions. That is, for every bounded perturbation direction $h(x)$, $g(x)$, or $\ell(x)$:

$$\begin{aligned} \frac{\partial}{\partial r} \mathbb{E}\{\varphi_{\text{eff}}(O; \tau_0, \mu_0, \mu_1, \pi + rh)\} \Big|_{r=0} &= 0, \\ \frac{\partial}{\partial r} \mathbb{E}\{\varphi_{\text{eff}}(O; \tau_0, \mu_0, \mu_1 + rg, \pi)\} \Big|_{r=0} &= 0, \\ \frac{\partial}{\partial r} \mathbb{E}\{\varphi_{\text{eff}}(O; \tau_0, \mu_0 + r\ell, \mu_1, \pi)\} \Big|_{r=0} &= 0. \end{aligned}$$

Proof

Throughout, recall $\mathbb{E}(Y | T = 1, X) = \mu_1(X)$, $\mathbb{E}(Y | T = 0, X) = \mu_0(X)$, $\mathbb{E}(T | X) = \pi(X)$.

Perturbation in π . Replace π by $\pi + rh$ and differentiate. The derivative becomes $\mathbb{E}\left[-\frac{Th(X)}{\pi(X)^2} \{Y - \mu_1(X)\} - \frac{(1-T)h(X)}{(1-\pi(X))^2} \{Y - \mu_0(X)\}\right]$. Conditioning on X : $\mathbb{E}[T\{Y - \mu_1(X)\} | X] = \pi(X)\mu_1(X) - \mu_1(X)\pi(X) = 0$, and similarly $\mathbb{E}[(1-T)\{Y - \mu_0(X)\} | X] = 0$. Both conditional expectations vanish.

Perturbation in μ_1 . Replace μ_1 by $\mu_1 + rg$. The derivative is $\mathbb{E}\left[g(X) \left(1 - \frac{T}{\pi(X)}\right)\right]$. Conditioning on X : $\mathbb{E}[1 - T/\pi(X) | X] = 1 - \pi(X)/\pi(X) = 0$.

Perturbation in μ_0 . Replace μ_0 by $\mu_0 + r\ell$. The derivative is $\mathbb{E}\left[\ell(X) \left(\frac{1-T}{1-\pi(X)} - 1\right)\right]$. Conditioning on X : $\mathbb{E}[(1-T)/(1-\pi(X)) - 1 | X] = (1-\pi(X))/(1-\pi(X)) - 1 = 0$. $\square$

Remark: What the Proof Reveals

Each of the three derivative calculations reduces to a single key fact: the conditional mean of the IPW weight equals one. Specifically, $\mathbb{E}(T/\pi(X) | X) = 1$ and $\mathbb{E}((1-T)/(1-\pi(X)) | X) = 1$ are immediate consequences of $\mathbb{E}(T | X) = \pi(X)$, the engine behind both the balancing theorem and the IPW identification formula (Chapter 6). Neyman orthogonality of the ATE score is therefore a direct consequence of how the propensity score is defined.

The score Equation 11.3 simultaneously fulfills three purposes:

1. It *identifies* τ through the moment condition $\mathbb{E}\{\varphi_{\text{eff}}(O; \tau, \eta)\} = 0$.
2. It is *doubly robust*, yielding a consistent estimator whenever either the outcome model or the propensity score model is correctly specified.
3. It is the *efficient influence function*, so any estimator based on it achieves the semiparametric efficiency bound.

11.5 Why Reusing the Same Data Can Be Problematic

Define the centered empirical process $\mathbb{G}_n f = \sqrt{n}(\mathbb{P}_n - P)f$. With the plug-in AIPW estimating equation, the remainder R_n decomposes exactly as in Equation 11.2:

$$R_n = \underbrace{P\{\varphi_{\text{eff}}(\cdot; \tau, \hat{\eta}) - \varphi_{\text{eff}}(\cdot; \tau, \eta_0)\}}_{\text{population bias}} + \underbrace{(\mathbb{P}_n - P)\{\varphi_{\text{eff}}(\cdot; \tau, \hat{\eta}) - \varphi_{\text{eff}}(\cdot; \tau, \eta_0)\}}_{\text{empirical process term}}. \quad (11.4)$$

By Neyman orthogonality (Section 11.4), the population bias is $o_p(n^{-1/2})$ under the product-rate condition. The *empirical process term* $n^{-1/2}\mathbb{G}_n\{f_{\hat{\eta}} - f_{\eta_0}\}$ depends on $\hat{\eta}$ estimated from the same sample. Whether it is $o_p(n^{-1/2})$ depends on how complex the class $\{f_{\eta} : \eta \in \mathcal{H}\}$ is — a question the Donsker condition answers.

Definition: Donsker Class

A class of measurable functions $\mathcal{F}$ is a **Donsker class** if the empirical process $\{\mathbb{G}_n f : f \in \mathcal{F}\}$ converges weakly in $\ell^\infty(\mathcal{F})$ to a tight Gaussian process. Equivalently, $\sup_{f \in \mathcal{F}} |(\mathbb{P}_n - P)f| = O_p(n^{-1/2})$ and the fluctuations of $n^{1/2}(\mathbb{P}_n - P)$ over $\mathcal{F}$ are stochastically equicontinuous; see Vaart (1998), Ch. 19.

Example: Donsker and Non-Donsker Classes

- **Parametric classes** (e.g., logistic regression models indexed by a finite-dimensional parameter) are Donsker under mild moment conditions.
- **Hölder-smooth function classes** on $[0, 1]^d$ with smoothness index $s > d/2$ are Donsker.
- **Indicator classes** $\{x \mapsto \mathbf{1}(x \leq t) : t \in \mathbb{R}\}$ are Donsker by the classical theorem.
- **Random forests and neural networks** typically fall outside classical Donsker regimes in the forms used for modern flexible prediction.
- **High-dimensional lasso** with a growing number of selected variables also typically falls outside Donsker regimes.

If $\{f_{\eta} : \eta \in \mathcal{H}\}$ is Donsker and $\|\hat{\eta} - \eta_0\| \rightarrow 0$, then by stochastic equicontinuity $\mathbb{G}_n\{f_{\hat{\eta}} - f_{\eta_0}\} = o_p(1)$, making the empirical process term $o_p(n^{-1/2})$. When the Donsker condition fails, this need not hold: the empirical process term can remain non-negligible even as $\hat{\eta} \rightarrow \eta_0$.

Remark: Donsker Conditions and Machine Learning

The Donsker condition fails for random forests, neural networks, and high-dimensional lasso precisely because their function classes are too rich. Cross-fitting sidesteps this entirely: because $\hat{\eta}^{(-k)}$ is independent of the observations in $\mathcal{J}_k$, the empirical process term is replaced by a sum of conditionally independent terms, each $o_p(n^{-1/2})$ by a conditional law of large numbers, with no Donsker assumption required.

11.6 Sample Splitting

Sample splitting resolves the data-reuse problem by construction. Partition $\{1, \dots, n\}$ into a training sample $\mathcal{J}_{\text{train}}$ and an evaluation sample $\mathcal{J}_{\text{eval}}$. Fit nuisance estimators $\hat{\eta}^{(\text{train})}$ on $\mathcal{J}_{\text{train}}$, then define $\hat{\tau}$ by solving the estimating equation on the held-out sample:

$$\frac{1}{|\mathcal{J}_{\text{eval}}|} \sum_{i \in \mathcal{J}_{\text{eval}}} \varphi_{\text{eff}}(O_i; \tau, \hat{\eta}^{(\text{train})}) = 0.$$

Because $\hat{\eta}^{(\text{train})}$ is estimated on $\mathcal{J}_{\text{train}}$, it is *independent* of the observations in $\mathcal{J}_{\text{eval}}$. Conditioning on $\hat{\eta}^{(\text{train})}$, the summands in the empirical process term are independent and mean zero, so by the conditional law of large numbers:

$$(\mathbb{P}_{n_e} - P)\{\varphi_{\text{eff}}(\cdot; \tau, \hat{\eta}^{(\text{train})}) - \varphi_{\text{eff}}(\cdot; \tau, \eta_0)\} = o_p(n^{-1/2}).$$

No Donsker condition is needed: independence does the work that Donsker equicontinuity would otherwise have to do.

The main disadvantage is inefficiency: only $|\mathcal{J}_{\text{eval}}|$ observations contribute to estimating τ , and the estimate depends on the particular random partition chosen.

11.7 Cross-Fitting

Cross-fitting extends sample splitting to recover full-sample efficiency while preserving the independence argument. Partition $\{1, \dots, n\}$ into K approximately equal folds $\mathcal{J}_1, \dots, \mathcal{J}_K$. For each fold k : (1) estimate nuisance functions on $\{1, \dots, n\} \setminus \mathcal{J}_k$ to obtain $\hat{\eta}^{(-k)} = (\hat{\mu}_0^{(-k)}, \hat{\mu}_1^{(-k)}, \hat{\pi}^{(-k)})$; (2) evaluate the orthogonal score on $\mathcal{J}_k$; (3) aggregate across all folds. The cross-fitted estimator $\hat{\tau}_{\text{DML}}$ solves:

$$\frac{1}{n} \sum_{k=1}^K \sum_{i \in \mathcal{J}_k} \varphi_{\text{eff}}(O_i; \tau, \hat{\eta}^{(-k)}) = 0. \quad (11.5)$$

For each fold k , the nuisance estimate $\hat{\eta}^{(-k)}$ is trained on $\{1, \dots, n\} \setminus \mathcal{J}_k$, independent of the observations in $\mathcal{J}_k$. Conditioning on $\hat{\eta}^{(-k)}$, the fold- k summands are independent and mean zero, and the conditional law of large numbers gives $R_n^{(k)} = o_p(n^{-1/2})$ for each k without any Donsker assumption. Averaging over folds, $R_n = o_p(n^{-1/2})$.

Cross-fitting therefore achieves two goals simultaneously: it preserves the held-out independence structure, and it ensures every observation serves once as an evaluation point, so the full sample contributes to estimating τ .

Algorithm: Cross-Fitted Orthogonal-Score Estimator for the ATE

1. Partition the sample into K folds $\mathcal{J}_1, \dots, \mathcal{J}_K$.
2. For each $k = 1, \dots, K$, fit nuisance estimators $\hat{\eta}^{(-k)} = (\hat{\mu}_0^{(-k)}, \hat{\mu}_1^{(-k)}, \hat{\pi}^{(-k)})$ on $\{1, \dots, n\} \setminus \mathcal{J}_k$.
3. For each $i \in \mathcal{J}_k$, compute the out-of-fold score $\hat{\varphi}_i = \varphi_{\text{eff}}(O_i; \tau, \hat{\eta}^{(-k)})$.
4. Solve the cross-fitted moment equation Equation 11.5. Because the score is linear in τ , the solution is:

$$\hat{\tau}_{\text{DML}} = \hat{\mu}_{1,\text{DML}} - \hat{\mu}_{0,\text{DML}}, \quad (11.6)$$

where $\hat{\mu}_{t,\text{DML}} = \frac{1}{n} \sum_{k=1}^K \sum_{i \in \mathcal{J}_k} \left[\hat{\mu}_t^{(-k)}(X_i) + \frac{T_i^{(t)}}{\hat{\pi}^{(-k)}(X_i)} \{Y_i - \hat{\mu}_t^{(-k)}(X_i)\} \right]$, with $T_i^{(1)} = T_i$, $T_i^{(0)} = 1 - T_i$.

11.8 Double Machine Learning for the Average Treatment Effect

The estimator Equation 11.6 is simply the AIPW estimator from Chapter 10 with out-of-fold nuisance estimates. The term *double machine learning* (DML), introduced by Chernozhukov et al. (2018), refers to the combination of three ingredients:

- *Neyman orthogonality*, which eliminates the population bias term so nuisance estimation error enters only at second order;
- *cross-fitting*, which eliminates the empirical process term via fold-by-fold conditional independence, removing the need for Donsker conditions; and
- *flexible machine-learning estimation* of the nuisance functions, which is made valid by the first two ingredients.

The term “double” should not be read as merely meaning “two models”; it refers to the combination of nuisance learning with orthogonalization and debiasing. DML is not a new causal estimand and not a new identification strategy; it is a modern estimation protocol built around an orthogonal score. In the biostatistics literature, the same estimator is often called the *cross-fitted AIPW estimator*.

“Double” Does Not Mean “Two Models Are Enough”

It is tempting to think that estimating any two of μ_0 , μ_1 , and π flexibly is sufficient. What matters is that *all* nuisance components entering the orthogonal score are estimated out-of-fold, and that the product-rate condition is satisfied. Using flexible methods for only one nuisance function while misspecifying another can still produce invalid inference.

11.9 Rate Conditions and Asymptotic Normality

Theorem: Asymptotic Linearity Under Cross-Fitting

Suppose: (i) the score φ_{eff} is orthogonal at the truth; (ii) the nuisance estimators are consistent in mean squared error; (iii) *strong overlap* holds: there exists $c > 0$ such that $c \leq \pi(X) \leq 1 - c$ and $c \leq \hat{\pi}^{(-k)}(X) \leq 1 - c$ a.s. for each fold k ; (iv) the *product-rate condition* holds:

$$\|\hat{\pi}^{(-k)} - \pi\| \cdot \|\hat{\mu}_t^{(-k)} - \mu_t\| = o_p(n^{-1/2}), \quad t = 0, 1. \quad (11.7)$$

Then the cross-fitted estimator satisfies:

$$\sqrt{n}(\hat{\tau}_{\text{DML}} - \tau) = \frac{1}{\sqrt{n}} \sum_{i=1}^n \varphi_{\text{eff}}(O_i; \tau, \eta) + o_p(1) \xrightarrow{d} N(0, \mathbb{E}[\varphi_{\text{eff}}(O; \tau, \eta)^2]).$$

The asymptotic variance equals the semiparametric efficiency bound, so $\hat{\tau}_{\text{DML}}$ is asymptotically efficient whenever the product-rate condition holds.

Remark: Weak Overlap vs. Strong Overlap

Weak overlap ($0 < \pi(X) < 1$ a.s.) is the condition needed for *identification*: every covariate stratum contains units from both treatment arms. *Strong overlap* ($c \leq \pi(X) \leq 1 - c$ for some $c > 0$) is what the proof actually uses: it bounds the inverse-probability weights away from infinity. Without strong overlap, influence function values can have infinite variance, and $\sqrt{n}$ -consistent inference breaks down even if the ATE is identified (Khan and Tamer 2010). Trimming extreme weights or restricting the target population to a region of stronger overlap are common practical responses.

Proof Sketch

Write $n_k = |J_k|$. The cross-fitted moment equation gives $0 = K^{-1} \sum_k \mathbb{P}_{n_k} \{\varphi_{\text{eff}}(O; \tau, \hat{\eta}^{(-k)})\}$. Adding and subtracting $\varphi_{\text{eff}}(O; \tau, \eta_0)$ and rearranging:

$$\hat{\tau}_{\text{DML}} - \tau = \frac{1}{n} \sum_{i=1}^n \varphi_{\text{eff}}(O_i; \tau, \eta_0) + \frac{1}{K} \sum_{k=1}^K R_n^{(k)},$$

where each fold remainder $R_n^{(k)}$ decomposes into population bias $B_n^{(k)}$ and empirical process term $E_n^{(k)}$.

Step 1 (population bias). By iterated expectations and $\mathbb{E}(T | X) = \pi(X)$:

$$|B_n^{(k)}| \leq C \left(\|\hat{\mu}_1^{(-k)} - \mu_1\| \cdot \|\hat{\pi}^{(-k)} - \pi\| + \|\hat{\mu}_0^{(-k)} - \mu_0\| \cdot \|\hat{\pi}^{(-k)} - \pi\| \right) = o_p(n^{-1/2})$$

by the product-rate condition (strong overlap ensures denominators are bounded below).

Step 2 (empirical process term). Condition on $\hat{\eta}^{(-k)}$. Since $\hat{\eta}^{(-k)}$ is trained on $\{1, \dots, n\} \setminus J_k$, the evaluation-fold summands are conditionally i.i.d. with conditional mean zero. The conditional variance of $E_n^{(k)}$ satisfies $\text{Var}(\sqrt{n} E_n^{(k)} | \hat{\eta}^{(-k)}) = (n/n_k) \|f_k\|_{L_2}^2 = o_p(1)$ under consistency. By conditional Chebyshev, $\sqrt{n} E_n^{(k)} = o_p(1)$.

Conclusion. $\sqrt{n} R_n^{(k)} = o_p(1)$ for each k , so $\sqrt{n} K^{-1} \sum_k R_n^{(k)} = o_p(1)$. The CLT applies to the first term, and asymptotic normality follows by Slutsky's theorem. $\square$

Remark: Interpreting the Product-Rate Condition

The product-rate condition requires the *product* of the two nuisance estimation errors to be $o_p(n^{-1/2})$, substantially weaker than requiring each error individually to be $o_p(n^{-1/2})$. A sufficient symmetric case: each nuisance estimator converges at rate $o_p(n^{-1/4})$ in mean square, achievable by many nonparametric and machine-learning methods. This is what makes flexible nuisance estimation feasible in semiparametric causal inference.

11.10 Lab: Simulation Study of the DML Estimator

Data-generating process. $n = 2000$ observations; $X_1, X_2, X_3 \sim \text{Uniform}(-1, 1)$ i.i.d. True nuisance functions:

$$\pi(X) = \text{expit}(\sin(\pi X_1) + X_2^2 - 0.5), \quad \mu_1(X) = 2 \sin(\pi X_1) + X_2^2 + X_3, \quad \mu_0(X) = \sin(\pi X_1) + X_2^2 - X_3.$$

True ATE: $\tau = \mathbb{E}\{\mu_1(X) - \mu_0(X)\} = \mathbb{E}\{\sin(\pi X_1) + 2X_3\} = 0$ (by symmetry). The nonlinearity of π and μ_t means random forests can fit them consistently, while the Donsker condition fails for the function class they trace out — exactly the setting of Section 11.5.

The three estimators. (1) *Oracle AIPW*: plug in the true nuisance functions directly — infeasible but provides the semiparametric efficiency benchmark. (2) *Naive RF AIPW*: estimate π, μ_1, μ_0 using random forests on the **full sample**, then evaluate the AIPW score on the same sample — the Donsker condition fails. (3) *DML (cross-fitted AIPW)*: same random forest learners but via $K = 5$ -fold cross-fitting — the fold-by-fold independence argument eliminates the empirical process term.

Estimators (2) and (3) use identical learners; the only difference is data reuse versus cross-fitting.

Implementation. The simulation was run in Python using scikit-learn random forests with fixed hyperparameters (`n_estimators=100`, `min_samples_leaf=5`; all other settings at library defaults). No nested cross-validation is used to tune learners inside the cross-fitting loop: this keeps the comparison between Naive and DML clean — any difference is attributable to cross-fitting alone, not to different tuning paths. The propensity score is clipped to $[0.01, 0.99]$ after fitting. Replication seed is $1000 + r$ for $r = 0, \dots, 199$.

Pseudocode: One Replication

1. **Generate data.** Draw (X_i, T_i, Y_i) for $i = 1, \dots, n$ from the DGP above.
2. **Oracle AIPW.** Compute $\hat{\tau}_{\text{oracle}} = n^{-1} \sum_i \varphi_{\text{eff}}(O_i; \tau, \pi, \mu_0, \mu_1)$ using the true nuisance functions.
3. **Naive RF AIPW.** Fit $\hat{\pi}, \hat{\mu}_0, \hat{\mu}_1$ as random forests on the full sample; clip $\hat{\pi}$ to $[0.01, 0.99]$; compute $\hat{\tau}_{\text{naive}}$ with the same-sample fits.
4. **DML.** Partition $\{1, \dots, n\}$ into K shuffled folds. For each fold k : fit $\hat{\pi}^{(-k)}, \hat{\mu}_0^{(-k)}, \hat{\mu}_1^{(-k)}$ on $\{1, \dots, n\} \setminus \mathcal{J}_k$; clip $\hat{\pi}^{(-k)}$ to $[0.01, 0.99]$; predict on $\mathcal{J}_k$. Compute $\hat{\tau}_{\text{DML}}$ from Equation 11.6.

Simulation results ($n = 2000$, $K = 5$, 200 replications):

Estimator	Bias	Variance	RMSE
Oracle AIPW	0.009	0.0029	0.055
Naive RF AIPW	0.039	0.0031	0.068
DML	0.007	0.0036	0.061

Oracle AIPW is nearly unbiased, serving as the efficiency benchmark. **Naive RF AIPW** exhibits a clear positive bias (0.039, roughly four times the oracle bias) despite using consistent random forest estimators. The bias arises from the uncontrolled empirical process term: reusing the same sample allows overfitting in the random forests to propagate into the AIPW score, exactly as analyzed in Section 11.5. **DML** reduces the bias to 0.007 — within Monte Carlo error of zero — at a modest finite-sample variance

cost (from 0.0031 to 0.0036). The small variance increase does *not* represent an asymptotic efficiency cost: every observation is used once as an evaluation point, and cross-fitting is designed to recover first-order efficiency. At $n = 2000$ with $K = 5$, the increase arises from two finite-sample sources: each out-of-fold nuisance estimator is trained on $(K - 1)n/K = 1600$ observations rather than the full sample, so its predictions are slightly noisier; and fold-assignment randomness adds a small additional source of sampling variability. (The Naive estimator's lower variance is in part a consequence of overfitting, which reduces plug-in variance at the cost of the bias this lab is designed to expose.) The net effect is a **lower RMSE for DML** (0.061) than for Naive RF AIPW (0.068): the bias reduction outweighs the variance difference.

Remark

The key comparison is between Naive RF AIPW and DML, not between either and the oracle. Both use identical random forest learners; the only difference is cross-fitting. The bias reduction from 0.039 to 0.007 is therefore attributable solely to removing the same-sample data-reuse problem. The accompanying small increase in variance is a finite-sample phenomenon, not an asymptotic penalty: cross-fitting is designed to recover first-order efficiency, and the gap should narrow as n grows.

11.11 Variance Estimation and Confidence Intervals

Inference for $\hat{\tau}_{\text{DML}}$ proceeds exactly as in Chapter 10, except that nuisance estimates are now out-of-fold. Let $k(i)$ denote the fold containing observation i , and define:

$$\hat{\varphi}_i = \varphi_{\text{eff}}(O_i; \hat{\tau}_{\text{DML}}, \hat{\eta}^{(-k(i))}). \quad (11.8)$$

The variance estimator is:

$$\hat{V} = \frac{1}{n(n-1)} \sum_{i=1}^n (\hat{\varphi}_i - \bar{\varphi})^2, \quad \bar{\varphi} = \frac{1}{n} \sum_{i=1}^n \hat{\varphi}_i. \quad (11.9)$$

The Wald confidence interval at level $1 - \alpha$ is $\hat{\tau}_{\text{DML}} \pm z_{1-\alpha/2} \sqrt{\hat{V}}$.

Theorem: Consistency of the Variance Estimator

Under the conditions of the Asymptotic Linearity theorem, the variance estimator $\hat{V}$ in Equation 11.9 is consistent for the asymptotic variance of $\hat{\tau}_{\text{DML}}$; equivalently, $n\hat{V} \xrightarrow{p} \mathbb{E}[\varphi_{\text{eff}}(O; \tau, \eta)^2]$.

The formula Equation 11.9 is the sample variance of the out-of-fold estimated influence values divided by n — identical to Chapter 10's formula, but $\hat{\varphi}_i$ now uses the out-of-fold nuisance estimate $\hat{\eta}^{(-k(i))}$.

11.12 A Practical Workflow

Workflow: Cross-Fitted Causal Estimation

Phase 1: Identification and setup

1. **Specify the estimand.** State the target causal parameter and verify the identification assumptions (conditional exchangeability, positivity, and consistency for the ATE).
2. **Choose the orthogonal score.** For the ATE under conditional exchangeability, use the efficient influence function Equation 11.3.
3. **Select nuisance learners.** Choose flexible learners for $\mu_0(x)$, $\mu_1(x)$, $\pi(x)$. Verify informally that the product-rate condition Equation 11.7 is plausible.
4. **Monitor overlap.** Examine the empirical distribution of covariates by treatment group before fitting.

Phase 2: Estimation

5. **Partition and cross-fit.** Partition into K folds ($K = 5$ is a common default). For each fold k , fit nuisance learners on the complement to obtain $\hat{\eta}^{(-k)}$.
6. **Compute out-of-fold orthogonal scores.** For each $i \in \mathcal{I}_k$, compute $\hat{\varphi}_i$.

7. **Solve the moment equation.** Average scores across all observations to obtain $\hat{\tau}_{\text{DML}}$ via Equation 11.6.
8. **Estimate the variance.** Compute $\hat{V}$ via Equation 11.9 and form the Wald confidence interval.

Phase 3: Diagnostics and interpretation

9. **Inspect estimated propensity scores.** Examine $\hat{\pi}^{(-k)}(X_i)$ across folds; extreme values near 0 or 1 inflate influence-value variance.
10. **Assess sensitivity.** Repeat with alternative learners. Substantial sensitivity signals that the product-rate condition may not be satisfied.
11. **Interpret relative to identification assumptions.** The credibility of the causal estimate rests on identification assumptions, not on predictive accuracy of the nuisance learners.

Step 11 deserves emphasis. A common error is to assess credibility by citing good out-of-sample predictive accuracy for the nuisance models. Predictive accuracy is necessary but not sufficient. The identifying assumptions are untestable from the data:

Modern causal estimation is not merely machine learning plus a causal estimand. It is machine learning embedded inside a carefully constructed semiparametric procedure, whose validity rests on identification assumptions that no algorithm can verify.

11.13 What Machine Learning Does Not Solve

Flexible nuisance estimation is a genuine improvement over parametric methods, but it does not resolve the fundamental difficulties of causal inference. In particular, machine learning does not solve:

- **Unmeasured confounding.** If important confounders are absent from X , conditional exchangeability $Y(t) \perp\!\!\!\perp T \mid X$ fails. No amount of flexibility in estimating $\pi(x)$ or $\mu_t(x)$ can correct this.
- **Violations of consistency.** If the treatment version is ill-defined or SUTVA is violated, the potential outcomes framework itself is compromised.
- **Failure of positivity.** When $\pi(x) \approx 0$ or $\pi(x) \approx 1$, influence values become unstable and coverage can degrade severely.
- **Ambiguity about the estimand.** Flexible estimation cannot resolve disagreement about what causal quantity is scientifically meaningful.
- **Invalid instrumental variable assumptions.** The exclusion restriction and exogeneity conditions (Chapter 7) are not testable; machine learning cannot substitute for subject-matter justification.
- **Fragile mediation assumptions.** Sequential ignorability and no interference between mediators (Chapter 8) must be argued on substantive grounds.

Machine learning improves the estimation of nuisance functions *within* a given identification strategy. It does not create identification where none exists.

11.14 Chapter Summary

Symbol	Meaning
$\eta = (\mu_0, \mu_1, \pi)$	Nuisance tuple
$\phi(O; \psi, \eta)$	Generic orthogonal score; zero mean at truth
$\varphi_{\text{eff}}(O; \tau, \eta)$	Efficient influence function for the ATE
$\mathcal{J}_k$	Equation 11.3
$\hat{\eta}^{(-k)}$	k -th fold; $k = 1, \dots, K$
$\hat{\tau}_{\text{DML}}$	Nuisance estimates trained without fold k
$\hat{\varphi}_i$	Cross-fitted AIPW (DML) estimator Equation 11.6
$\hat{V}$	Out-of-fold estimated influence value
	Variance estimator Equation 11.9

1. Flexible methods are attractive for nuisance estimation when true nuisance functions are complex,

but naive plug-in estimators can fail for two reasons: orthogonality may be absent, and even when orthogonality holds, same-sample nuisance fitting can create an additional empirical-process bias.

2. A score function is Neyman orthogonal if its Gateaux derivative with respect to the nuisance vanishes at the truth, so nuisance estimation error enters only at second order.
3. The efficient influence function for the ATE Equation 11.3 is orthogonal, doubly robust, and efficiency-achieving.
4. Even with an orthogonal score, reusing the same data for nuisance estimation and score evaluation can induce overfitting bias, especially when the nuisance class falls outside classical fixed Donsker regimes.
5. Sample splitting removes this bias by separating nuisance estimation from score evaluation, but wastes data.
6. Cross-fitting recovers efficiency by rotating training and evaluation roles across K folds.
7. The cross-fitted AIPW estimator $\hat{\tau}_{\text{DML}}$ is the standard double machine learning estimator for the ATE.
8. Under the product-rate condition Equation 11.7, orthogonality and cross-fitting together yield asymptotic linearity, root- n consistency, and asymptotic normality of $\hat{\tau}_{\text{DML}}$, with asymptotic variance equal to the semiparametric efficiency bound. A sufficient condition: each nuisance estimator converges at rate $o_p(n^{-1/4})$ in $L_2(P)$ norm, equivalently that its squared L_2 risk is $o_p(n^{-1/2})$.
9. The asymptotic variance is estimated consistently from the empirical variance of the out-of-fold estimated influence values.
10. Machine learning improves nuisance estimation, but does not resolve identification problems: unmeasured confounding, positivity failures, and untestable structural assumptions remain the researcher's responsibility.

11.15 Problems

1. Verifying orthogonality.

- (a) Write out the efficient influence function $\varphi_{\text{eff}}(O; \tau, \eta)$ for the ATE and compute its expectation. Confirm it equals zero at the truth.
- (b) Consider a perturbation $\pi \mapsto \pi + r \cdot h$ for a bounded function $h(x)$. Differentiate $\mathbb{E}\{\varphi_{\text{eff}}(O; \tau, \mu_0, \mu_1, \pi + rh)\}$ with respect to r and evaluate at $r = 0$. Show the derivative is zero, confirming Neyman orthogonality in the propensity score direction.
- (c) Repeat the calculation for a perturbation in μ_1 .

2. First-order bias of the plug-in estimator. Suppose $\hat{\mu}_1 = \mu_1 + \delta$ for a deterministic function $\delta(x)$, and $\hat{\mu}_0 = \mu_0$, $\hat{\pi} = \pi$.

- (a) Show that the bias of the prediction estimator $\hat{\tau}_{\text{pred}} = n^{-1} \sum_i \{\hat{\mu}_1(X_i) - \hat{\mu}_0(X_i)\}$ is $\mathbb{E}\{\delta(X)\}$.
- (b) Show that the population bias of the plug-in AIPW estimator (using $\hat{\mu}_1 = \mu_1 + \delta$ with the true π and μ_0) is zero. Explain which property of the AIPW score is responsible. Does this mean the plug-in AIPW estimator is asymptotically unbiased when $\hat{\mu}_1$ is estimated flexibly from the same sample used to evaluate the score? Explain why or why not.

3. Cross-fitting with $K = 2$. Partition $n = 200$ observations into two halves $\mathcal{J}_1$ and $\mathcal{J}_2$.

- (a) Write down the cross-fitted moment equation Equation 11.5 explicitly for $K = 2$.
- (b) Explain why the contribution from $\mathcal{J}_1$ uses nuisance estimates trained on $\mathcal{J}_2$ and vice versa.
- (c) Compare this procedure to single sample splitting with $\mathcal{J}_1$ as the training fold. What is gained and what is lost?

4. The product-rate condition. Suppose $\|\hat{\pi} - \pi\|_{L_2} = o_p(n^{-\alpha})$ and $\|\hat{\mu}_t - \mu_t\|_{L_2} = o_p(n^{-\beta})$ for $\alpha, \beta > 0$.

- (a) State the condition on α and β under which the product-rate condition Equation 11.7 holds.
- (b) Show that $\alpha = \beta = 1/4$ satisfies your condition.
- (c) Does $\alpha = 1/2$, $\beta = 0$ satisfy it? What does this say about the case where the propensity score is estimated parametrically but the outcome model converges only at an unspecified rate?

5. Variance estimation. Using the explicit formula Equation 11.6, write out the estimated influence value $\hat{\varphi}_i$ for a generic observation $i \in \mathcal{J}_k$. Show that $n^{-1} \sum_i \hat{\varphi}_i = 0$ when $\hat{\tau}_{\text{DML}}$ solves the cross-fitted moment

equation, and explain why the centering in Equation 11.9 is therefore asymptotically inconsequential.

6. What machine learning cannot do. For each of the following scenarios, identify the identification assumption that is violated and explain why flexible nuisance estimation cannot remedy the problem.

- (a) A study of job-training effects on earnings omits pre-program earnings, a strong predictor of both program participation and subsequent earnings.
- (b) A study of a medical treatment estimates the propensity score accurately, but the treatment was assigned in clusters (hospital wards) and outcomes may be correlated within clusters in a way that depends on the fraction of patients treated.
- (c) An instrument is used to estimate the effect of education on wages, but the instrument also directly affects wages through a channel not related to education.

References

- Abadie, Alberto, and Guido W. Imbens. 2006. “Large Sample Properties of Matching Estimators for Average Treatment Effects.” *Econometrica* 74 (1): 235–67.
- Abadie, Alberto, and Guido W. Imbens. 2016. “Matching on the Estimated Propensity Score.” *Econometrica* 84 (2): 781–807.
- Angrist, Joshua D. 1990. “Lifetime Earnings and the Vietnam Era Draft Lottery: Evidence from Social Security Administrative Records.” *American Economic Review* 80 (3): 313–36.
- Angrist, Joshua D., Susan M. Dynarski, Thomas J. Kane, Parag A. Pathak, and Christopher R. Walters. 2012. “Who Benefits from KIPP?” *Journal of Policy Analysis and Management* 31 (4): 837–60.
- Angrist, Joshua D., and Guido W. Imbens. 1994. “Identification and Estimation of Local Average Treatment Effects.” *Econometrica* 62 (2): 467–75.
- Angrist, Joshua D., Guido W. Imbens, and Donald B. Rubin. 1996. “Identification of Causal Effects Using Instrumental Variables.” *Journal of the American Statistical Association* 91 (434): 444–55.
- Angrist, Joshua D., and Alan B. Krueger. 1991. “Does Compulsory School Attendance Affect Schooling and Earnings?” *Quarterly Journal of Economics* 106 (4): 979–1014.
- Bang, Heejung, and James M. Robins. 2005. “Doubly Robust Estimation in Missing Data and Causal Inference Models.” *Biometrics* 61 (4): 962–73.
- Baron, Reuben M., and David A. Kenny. 1986. “The Moderator–Mediator Variable Distinction in Social Psychological Research: Conceptual, Strategic, and Statistical Considerations.” *Journal of Personality and Social Psychology* 51 (6): 1173–82.
- Berkson, Joseph. 1946. “Limitations of the Application of Fourfold Table Analysis to Hospital Data.” *Biometrics Bulletin* 2 (3): 47–53.
- Bezuidenhout, Dana, Sarah Forthal, Kara Rudolph, and Matthew R. Lamb. 2025. “Single World Intervention Graphs (SWIGs): A Practical Guide.” *American Journal of Epidemiology* 194: 2047–52. <https://doi.org/10.1093/aje/kwae353>.
- Bickel, Peter J., Chris A. J. Klaassen, Ya’acov Ritov, and Jon A. Wellner. 1993. *Efficient and Adaptive Estimation for Semiparametric Models*. Johns Hopkins University Press.
- Bound, John, David A. Jaeger, and Regina M. Baker. 1995. “Problems with Instrumental Variables Estimation When the Correlation Between the Instruments and the Endogenous Explanatory Variable Is Weak.” *Journal of the American Statistical Association* 90 (430): 443–50.
- Card, David. 1995. “Using Geographic Variation in College Proximity to Estimate the Return to Schooling.” In *Aspects of Labour Market Behaviour: Essays in Honour of John Vanderkamp*, edited by Louis N. Christofides, E. Kenneth Grant, and Robert Swidinsky. University of Toronto Press.

- Chernozhukov, Victor, Denis Chetverikov, Mert Demirer, et al. 2018. "Double/Debiased Machine Learning for Treatment and Structural Parameters." *The Econometrics Journal* 21 (1): C1–68.
- Crump, Richard K., V. Joseph Hotz, Guido W. Imbens, and Oscar A. Mitnik. 2009. "Dealing with Limited Overlap in Estimation of Average Treatment Effects." *Biometrika* 96 (1): 187–99. <https://doi.org/10.1093/biomet/asn055>.
- Dawid, A. Philip. 1979. "Conditional Independence in Statistical Theory." *Journal of the Royal Statistical Society, Series B* 41 (1): 1–31.
- Dawid, A. Philip. 1980. "Conditional Independence for Statistical Operations." *Annals of Statistics* 8 (3): 598–617.
- Dehejia, Rajeev H., and Sadek Wahba. 1999. "Causal Effects in Nonexperimental Studies: Reevaluating the Evaluation of Training Programmes." *Journal of the American Statistical Association* 94 (448): 1053–62.
- Deville, Jean-Claude, and Carl-Erik Särndal. 1992. "Calibration Estimators in Survey Sampling." *Journal of the American Statistical Association* 87 (418): 376–82.
- Firth, David, and Karen E. Bennett. 1998. "Robust Models in Probability Sampling." *Journal of the Royal Statistical Society: Series B* 60 (1): 3–21.
- Freedman, David A. 2008. "On Regression Adjustments to Experimental Data." *Advances in Applied Mathematics* 40 (2): 180–93.
- Hansen, Lars Peter. 1982. "Large Sample Properties of Generalized Method of Moments Estimators." *Econometrica* 50 (4): 1029–54.
- Hernán, Miguel A., and James M. Robins. 2006. "Instruments for Causal Inference: An Epidemiologist's Dream?" *Epidemiology* 17 (4): 360–72.
- Hernán, Miguel A., and James M. Robins. 2020. *Causal Inference: What If*. Chapman & Hall/CRC.
- Holland, Paul W. 1986. "Statistics and Causal Inference." *Journal of the American Statistical Association* 81 (396): 945–60.
- Imai, Kosuke, Luke Keele, and Teppei Yamamoto. 2010. "Identification, Inference and Sensitivity Analysis for Causal Mediation Effects." *Statistical Science* 25 (1): 51–71.
- Imbens, Guido W., and Donald B. Rubin. 2015. *Causal Inference for Statistics, Social, and Biomedical Sciences: An Introduction*. Cambridge University Press.
- Isaki, Cary T., and Wayne A. Fuller. 1982. "Survey Design Under the Regression Superpopulation Model." *Journal of the American Statistical Association* 77 (377): 89–96.
- Khan, Shakeeb, and Elie Tamer. 2010. "Irregular Identification, Support Conditions, and Inverse Weight Estimation." *Econometrica* 78 (6): 2021–42.
- Laan, Mark J. van der, and Sherri Rose. 2011. *Targeted Learning: Causal Inference for Observational and Experimental Data*. Springer.
- Laan, Mark J. van der, and Sherri Rose. 2018. *Targeted Learning in Data Science: Causal Inference for Complex Longitudinal Studies*. Springer.
- Laan, Mark J. van der, and Daniel Rubin. 2006. "Targeted Maximum Likelihood Learning." *The International Journal of Biostatistics* 2 (1): 1–40.

- LaLonde, Robert J. 1986. "Evaluating the Econometric Evaluations of Training Programmes with Experimental Data." *American Economic Review* 76 (4): 604–20.
- Lauritzen, Steffen L. 1996. *Graphical Models*. Oxford University Press.
- Levis, Alexander W., Edward H. Kennedy, and Luke Keele. 2024. "Nonparametric Identification and Efficient Estimation of Causal Effects with Instrumental Variables." *arXiv Preprint arXiv:2402.09332*.
- Lin, Winston. 2013. "Agnostic Notes on Regression Adjustments to Experimental Data: Reexamining Freedman's Critique." *The Annals of Applied Statistics* 7 (1): 295–318.
- Neyman, Jerzy Splawa. 1923. "On the Application of Probability Theory to Agricultural Experiments: Essay on Principles, Section 9." *Statistical Science* 5 (4): 465–72.
- Pearl, Judea. 1993. "Comment: Graphical Models, Causality and Intervention." *Statistical Science* 8 (3): 266–69.
- Pearl, Judea. 1995. "Causal Diagrams for Empirical Research." *Biometrika* 82 (4): 669–88.
- Pearl, Judea. 2001. "Direct and Indirect Effects." *Proceedings of the Seventeenth Conference on Uncertainty in Artificial Intelligence (UAI 2001)*, 411–20.
- Richardson, Thomas S., and James M. Robins. 2014. *ACE Bounds; Single World Intervention Graphs (SWIGs) and Identification of Causal Effects*. University of Washington.
- Robins, James M. 1986. "A New Approach to Causal Inference in Mortality Studies with a Sustained Exposure Period—Application to Control of the Healthy Worker Survivor Effect." *Mathematical Modelling* 7 (9–12): 1393–512.
- Robins, James M., Andrea Rotnitzky, and Lue Ping Zhao. 1994. "Estimation of Regression Coefficients When Some Regressors Are Not Always Observed." *Journal of the American Statistical Association* 89 (427): 846–66.
- Rosenbaum, Paul R., and Donald B. Rubin. 1983. "The Central Role of the Propensity Score in Observational Studies for Causal Effects." *Biometrika* 70 (1): 41–55.
- Rubin, Donald B. 1974. "Estimating Causal Effects of Treatments in Randomized and Nonrandomized Studies." *Journal of Educational Psychology* 66 (5): 688–701.
- Rubin, Donald B. 1980. "Randomization Analysis of Experimental Data: The Fisher Randomization Test." *Journal of the American Statistical Association* 75 (371): 591–93.
- Sargan, John D. 1958. "The Estimation of Economic Relationships Using Instrumental Variables." *Econometrica* 26 (3): 393–415.
- Shachter, Ross D. 1998. "Bayes-Ball: The Rational Pastime (for Determining Irrelevance and Requisite Information in Belief Networks and Influence Diagrams)." *Proceedings of the Fourteenth Conference on Uncertainty in Artificial Intelligence (UAI)*, 480–87.
- Shpitser, Ilya, and Judea Pearl. 2006. "Identification of Joint Interventional Distributions in Recursive Semi-Markovian Causal Models." *Proceedings of the Twenty-First National Conference on Artificial Intelligence (AAAI)* 21: 1219–26.
- Sobel, Michael E. 1982. "Asymptotic Confidence Intervals for Indirect Effects in Structural Equation Models." *Sociological Methodology* 13: 290–312.
- Tsiatis, Anastasios A. 2006. *Semiparametric Theory and Missing Data*. Springer.

- Vaart, Aad W. van der. 1998. *Asymptotic Statistics*. Cambridge University Press.
- Yang, Shu, Guido W. Imbens, Zhanglin Cui, Douglas E. Faries, and Zbigniew Kadziola. 2016. “Propensity Score Matching and Stratification in Observational Studies with Multi-Level Treatments.” *Biometrics* 72: 1055–65.
- Yang, Shu, and Yunshu Zhang. 2023. “Multiply Robust Matching Estimators of Average and Quantile Treatment Effects.” *Scandinavian Journal of Statistics* 50: 235–65.

Appendix A

Graphical Intuition for Conditional Independence and d-Separation

This appendix provides a gentle introduction to the probabilistic and graphical ideas that underlie Chapters 2 and 3. Our goal is not to give a complete treatment of graphical models, but rather to develop enough intuition so that the formal machinery of d-separation, back-door adjustment, and intervention graphs does not appear abruptly.

A directed acyclic graph (DAG) is more than a picture. It is a compact language for expressing assumptions about how variables are related. Once those assumptions are represented graphically, the graph tells us which variables may be associated, which paths transmit dependence, and which variables should or should not be conditioned on.

The key pedagogical idea of this appendix is simple: before learning the full d-separation criterion, it is helpful to understand three elementary three-node patterns. Those local patterns explain most of what happens later in larger graphs.

A.1 Conditional Independence: The Probabilistic Language Behind Graphs

Before introducing graphs, we first recall the probabilistic notion that graphs are designed to encode.

Definition: Conditional Independence

Let X , Y , and Z be random variables. We say that X and Y are *conditionally independent given Z* , written $X \perp\!\!\!\perp Y \mid Z$, if

$$p(x, y \mid z) = p(x \mid z) p(y \mid z)$$

for all z with positive probability.

Equivalently, once Z is known, learning X gives no additional information about Y . In terms of densities, conditional independence admits the following equivalent characterizations:

$$X \perp\!\!\!\perp Y \mid Z \iff f(x, y, z) f(z) = f(x, z) f(y, z) \iff \exists a, b: f(x, y, z) = a(x, z) b(y, z).$$

The last form is especially useful: it says that the joint density factors into one piece depending on (x, z) and another depending on (y, z) , with no cross-term in x and y .

Proposition: Fundamental Properties of Conditional Independence

(Dawid 1979, 1980) For random variables X , Y , Z , and W , the following hold:

- **(C1) Symmetry.** $X \perp\!\!\!\perp Y \mid Z \Rightarrow Y \perp\!\!\!\perp X \mid Z$.
- **(C2) Decomposition.** $X \perp\!\!\!\perp (Y, W) \mid Z \Rightarrow X \perp\!\!\!\perp Y \mid Z$.

- **(C3) Weak Union.** $X \perp\!\!\!\perp (Y, W) \mid Z \Rightarrow X \perp\!\!\!\perp Y \mid (Z, W)$.
- **(C4) Contraction.** $X \perp\!\!\!\perp Y \mid Z$ and $X \perp\!\!\!\perp W \mid (Y, Z) \Rightarrow X \perp\!\!\!\perp (Y, W) \mid Z$.
- **(C5) Intersection.** If $f(x, y, z, w) > 0$ for all (x, y, z, w) , then $X \perp\!\!\!\perp Y \mid (Z, W)$ and $X \perp\!\!\!\perp Z \mid (Y, W) \Rightarrow X \perp\!\!\!\perp (Y, Z) \mid W$.

Properties (C1)–(C4) hold for any probability distribution and are known as the *semigraphoid axioms*. Property (C5) additionally requires the joint density to be strictly positive; together with (C1)–(C4) it forms the *graphoid axioms*. These properties are used implicitly throughout the course whenever conditional independence statements are combined or simplified.

Conditional independence is not the same as marginal independence. Two variables may be dependent marginally but independent after conditioning on a third variable. Conversely, two variables may be independent marginally but become dependent after conditioning. Both phenomena occur repeatedly in causal inference.

Example: Ice Cream, Drowning, and Season

Let X = ice cream sales, Y = drowning incidents, and Z = season. Marginally, X and Y are positively associated because both tend to be higher in summer. This does not mean that ice cream sales cause drowning. A more plausible explanation is that season is a common cause of both variables. Once season is fixed, the association largely disappears: $X \perp\!\!\!\perp Y \mid Z$.

This example illustrates a recurring theme in causal inference: an observed association may be induced by a third variable, and conditioning on that variable can remove the spurious dependence.

Remark: The Central Question

At this stage, the main question for the reader is: *Does conditioning on a variable remove association, preserve it, or create it?* Graphs provide a systematic answer to exactly this question.

A.2 Three Basic Motifs: Chain, Fork, and Collider

Every path in a DAG is built from local three-node configurations. There are three fundamental types: a *chain*, a *fork*, and a *collider*. Their behavior under conditioning is the foundation of d-separation.




Figure A.1: The three fundamental three-node motifs. Chains and forks are blocked by conditioning on the middle node. Colliders are blocked by default but opened by conditioning on the middle node or one of its descendants.

A.2.1 Chain

Consider the pattern $X_1 \rightarrow X_2 \rightarrow X_3$, where the middle node X_2 lies on a directed pathway from X_1 to X_3 .

Example: Exercise, Body Weight, and Blood Pressure

Let X_1 = exercise, X_2 = body weight, and X_3 = blood pressure. Exercise may affect blood pressure partly through its effect on body weight, so the path $X_1 \rightarrow X_2 \rightarrow X_3$ transmits association from X_1 to X_3 . Marginally, exercise and blood pressure are associated; but if we condition on body weight, that particular pathway is blocked: $X_1 \perp\!\!\!\perp X_3 \mid X_2$.

Once the middle variable is fixed, the chain no longer transmits additional information from X_1 to X_3 .

A.2.2 Fork

Consider the pattern $X_1 \leftarrow X_2 \rightarrow X_3$, where the middle node X_2 is a common cause of X_1 and X_3 .

Example: Ice Cream, Season, and Drowning

With X_1 = ice cream sales, X_2 = season, and X_3 = drowning incidents, season affects both variables, so X_1 and X_3 are associated even though neither causes the other. Conditioning on season blocks this path: $X_1 \perp\!\!\!\perp X_3 \mid X_2$.

A fork is the simplest graphical form of confounding: the middle node creates association, and conditioning on it blocks the path.

A.2.3 Collider

Consider the pattern $X_1 \rightarrow X_2 \leftarrow X_3$, where the middle node X_2 is a common effect of X_1 and X_3 .

Example: Talent, Legacy Status, and College Admission

Let X_1 = academic talent, X_3 = legacy status, and X_2 = admission to an elite university. Talent and legacy status may be unrelated in the general applicant pool. But among admitted students they can become statistically associated: low legacy status makes unusually high talent more likely among admitted applicants, and vice versa. Symbolically, $X_1 \perp\!\!\!\perp X_3$ but $X_1 \not\perp\!\!\!\perp X_3 \mid X_2$.

Unlike chains and forks, a collider *blocks* the path by default. Conditioning on the collider opens the path and may induce association that was not present marginally.

A.2.4 Summary of the Three Motifs

A chain ($X_1 \rightarrow X_2 \rightarrow X_3$) and a fork ($X_1 \leftarrow X_2 \rightarrow X_3$) are each open by default and blocked by conditioning on the middle node X_2 . A collider ($X_1 \rightarrow X_2 \leftarrow X_3$) is blocked by default and opened by conditioning on the middle node or any of its descendants.

Conditioning Is Not Always Beneficial

Many adjustment mistakes arise from forgetting this asymmetry. Conditioning on a collider can create bias rather than remove it. The informal rule “control for more variables” is unsafe in causal inference; d-separation teaches a more precise lesson: condition on the *right* variables, not simply on many variables.

Chapter 2 develops these three motifs into the full d-separation criterion; see in particular Section 2.3 for the blocking rules and Section 2.4 for an extended treatment of collider bias.

A.3 d-Separation: When Does Conditioning Block a Path?

The three-node motifs explain what happens locally on a path. The next step is to extend this logic to a general DAG, where two variables may be connected by many paths.

Definition: Path

A *path* between two nodes is a sequence of distinct nodes such that each consecutive pair is connected by an edge, regardless of edge direction.

Definition: Blocked Path and d-Separation

A path is *blocked* by a conditioning set S if at least one of the following holds:

1. The path contains a chain or fork node that belongs to S , or
2. The path contains a collider such that neither the collider nor any of its descendants belongs

to S .

A path is *open* given S if it is not blocked. Two nodes X and Y are *d-separated* by a set S if every path between X and Y is blocked by S .

A.3.1 A Confounding Example

Consider the DAG with edges $X \rightarrow T$, $T \rightarrow Y$, and $X \rightarrow Y$. Here T is the treatment, Y is the outcome, and X is a pre-treatment covariate that affects both. There are two paths from T to Y : the directed causal path $T \rightarrow Y$, and the back-door path $T \leftarrow X \rightarrow Y$.

Without conditioning, the back-door path is open, so the observed association between T and Y mixes the causal effect with confounding. Conditioning on X blocks the fork $T \leftarrow X \rightarrow Y$, thereby isolating the causal path.




Figure A.2: A simple confounding graph. The path $T \leftarrow X \rightarrow Y$ is a back-door path from treatment to outcome. Conditioning on X blocks this noncausal path.

This confounding graph anticipates the back-door criterion of Chapter 3: the graphical condition that makes adjustment valid is precisely that all back-door paths are blocked.

A.3.2 A Collider Warning

Now consider the DAG with edges $T \rightarrow C$, $U \rightarrow C$, and $U \rightarrow Y$. Here C is a collider on the path $T \rightarrow C \leftarrow U \rightarrow Y$. Without conditioning on C , the path is blocked at the collider. Conditioning on C opens the path, creating a spurious association between T and Y through U .

Even more subtly, conditioning on a *descendant* of C can also open the path. If additionally $C \rightarrow D$, then conditioning on D may also induce association between T and Y through the collider at C .




Figure A.3: Collider bias with a descendant. The path $T \rightarrow C \leftarrow U \rightarrow Y$ is blocked by default at the collider C . Conditioning on C , or on its descendant D , can open the path and induce a spurious association between T and Y .

A.3.3 A Practical Checklist

To decide whether X and Y are d-separated by S , proceed as follows. First, list all paths between X and Y . On each path, classify each interior node as part of a chain, fork, or collider. Then check whether each path is blocked by S . Finally, conclude that X and Y are d-separated if and only if every path is blocked.

For beginners, this path-by-path procedure is often more useful than starting with the most compact formal definition.

A.4 DAG Factorization and the Markov Property

Up to this point, we have used DAGs qualitatively, to decide which paths are open or blocked. We now connect the graph to probability algebra: the same parent structure that governs d-separation also determines how the joint distribution factorizes.

Definition: DAG Factorization

Let $\mathcal{G}$ be a DAG with nodes $V_1, \dots, V_p$. A joint distribution $p(v_1, \dots, v_p)$ is said to *factorize according to $\mathcal{G}$* if

$$p(v_1, \dots, v_p) = \prod_{j=1}^p p(v_j \mid \text{Pa}(V_j)),$$

where $\text{Pa}(V_j)$ denotes the set of parents of V_j in $\mathcal{G}$.

Example: A Simple Confounding Graph

For the DAG with edges $X \rightarrow T$, $T \rightarrow Y$, and $X \rightarrow Y$, the joint density factorizes as

$$p(x, t, y) = p(x) p(t \mid x) p(y \mid t, x).$$

Definition: Local Markov Property

A distribution P satisfies the *local Markov property* with respect to $\mathcal{G}$ if, for every node $V_i \in V$,

$$V_i \perp\!\!\!\perp (\text{Nd}(V_i) \setminus \text{Pa}(V_i)) \mid \text{Pa}(V_i).$$

That is, once we condition on the direct causes of a node, that node is independent of all variables that are neither its descendants nor its parents.

Remark: Global Markov Property

The global Markov property is the statement that every d-separation in $\mathcal{G}$ implies a conditional independence in P . Thus, if two sets of variables are d-separated by a set S in the graph, then they are conditionally independent given S in the distribution. This is studied in detail in Chapter 2.

Remark: Equivalence for DAGs

For DAGs, the local and global Markov properties are equivalent under the usual graphical-model assumptions. The local property implies the global one via the Markov factorization, which is what is actually used in identification arguments.




Figure A.4: Education–earnings DAG used to illustrate factorization and the local and global Markov properties.

Example: Education–Earnings Graph

Consider the DAG in Figure A.4, with edges $N \rightarrow E$, $B \rightarrow E$, $E \rightarrow Y$, and $B \rightarrow Y$. The corresponding factorization is

$$p(n, b, e, y) = p(n) p(b) p(e | n, b) p(y | e, b).$$

The parent set of Y is $\text{Pa}(Y) = \{E, B\}$, so the local Markov property gives $Y \perp\!\!\!\perp N \mid \{E, B\}$. The same conclusion follows from the global viewpoint: the only path from N to Y is $N \rightarrow E \rightarrow Y$, which is blocked by conditioning on E , and conditioning additionally on B opens no new path.

Remark: Edges as Scientific Claims

In a causal DAG, an arrow is typically drawn only when a direct dependence-generating relation is believed to be present. This explains why edges are not included gratuitously: every arrow represents a substantive scientific claim. Minimality and faithfulness are discussed in Chapter 2.

A.5 Optional: Moralization as an Alternative Criterion

Note to Reader

This section may be skipped on a first reading.

There is an alternative graph-theoretic way to check d-separation based on constructing an undirected graph called the *moral graph*. The method is elegant and useful in more advanced graphical-model theory, but it is not necessary for understanding the main causal ideas of Chapters 2 and 3. For most students, the path-by-path method using chains, forks, and colliders is more intuitive on first exposure.

The procedure is as follows. First, restrict attention to the ancestors of the variables of interest. Second, connect any two parents of a common child by an undirected edge. Third, drop all arrow directions. Finally, check ordinary graph separation in the resulting undirected graph. This procedure yields a criterion equivalent to d-separation.

Example: Moral Graph of a Collider

Suppose a DAG contains the pattern $A \rightarrow C \leftarrow B$. In the moral graph, the two parents A and B are connected by an undirected edge, producing the triangle $A - C - B$ together with $A - B$. This added edge reflects the special role of colliders in d-separation.

A.6 Summary

This appendix introduced the graphical ideas that support Chapters 2 and 3. Conditional independence is the probabilistic language that graphs are designed to encode. The three local motifs — chain, fork, and collider — determine how conditioning affects association along a path. D-separation extends these local rules to arbitrary graphs: two variables are d-separated by S if every path between them is blocked by S . The most important application in causal inference is to distinguish confounding paths from causal paths and to identify valid adjustment sets. Finally, the Markov property gives a probabilistic interpretation to the graph by linking graphical structure to a factorization of the joint distribution.

Appendix B

Single World Intervention Graphs

This appendix develops single world intervention graphs (SWIGs), a graphical device introduced by Richardson and Robins (2014) and given an accessible practical treatment by Bezuidenhout et al. (2025). A SWIG provides a formal bridge between the potential outcomes framework and the do-calculus by encoding both in a single diagram. A reader familiar with Chapters 2–4 will find that SWIGs require no new conceptual ingredients: they are ordinary DAGs in which the treatment node is split into two pieces, making potential outcomes and identifiability assumptions graphically explicit rather than leaving them implicit.

Chapter 4 used the back-door criterion and the adjustment formula to identify causal effects. SWIGs make that identification argument graphically explicit: potential outcomes appear as labeled nodes inside the graph, so ignorability becomes a routine d-separation statement rather than a free-standing probabilistic assumption. Section B.5 completes the picture by establishing the biconditional equivalence between ignorability and the back-door criterion that was only stated in one direction in Chapter 4.

B.1 The SWIG Construction

In the original DAG $\mathcal{G}$, the node T represents the treatment as it naturally occurs. Under $\text{do}(T=t)$, however, the treatment is externally fixed. The SWIG separates these two roles by *splitting* T into two pieces displayed side by side in the graph:

- a **random half** (left side, capital letter T): represents the treatment value that would naturally occur, and retains all incoming edges from T 's causal parents;
- a **fixed half** (right side, lowercase letter t): represents the externally imposed intervention value, carries all outgoing edges that formerly left T , and has *no* edge to or from the random half.

Definition: Single World Intervention Graph

(Richardson and Robins 2014) The SWIG $\mathcal{G}(t)$ is constructed from $\mathcal{G}$ by:

1. Replacing T with two nodes: random half T (all incoming edges retained) and fixed half t (all outgoing edges retained).
2. Severing the direct connection between the two halves.
3. Redrawing causal arrows at the split node: *all arrows into the split node land on the random half; all arrows departing the split node leave from the fixed half.*
4. Relabeling every descendant V of T in $\mathcal{G}$ as $V(t)$, indicating it is a potential outcome under $\text{do}(T=t)$.

Figure B.1 illustrates the construction for the simple confounded-treatment graph.

B.2 The Fixed Half as a Source Node

After the SWIG $\mathcal{G}(t)$ is constructed, d-separation among its nodes is evaluated using the standard rules of Chapter 2. The key structural fact that governs all path analysis in a SWIG is:




Figure B.1: Converting the confounded-treatment DAG (L is a confounder) into the SWIG for $\text{do}(A=a)$. The treatment node splits: the random half A (circle) retains the incoming arrow from L ; the fixed half a (square) carries the outgoing arrow to $Y(a)$. D-separation in the SWIG gives $A \perp\!\!\!\perp Y(a) \mid L$ (conditional ignorability).

The Fixed Half Is a Source Node

In $\mathcal{G}(t)$, the fixed half t has **no incoming edges**. Step 2 of the construction severs the connection from the random half T , and no other node in $\mathcal{G}$ has a direct arrow into t . Three consequences follow immediately.

1. **No path arrives at t from another node.** The edge $T \rightarrow t$ does not exist in $\mathcal{G}(t)$. In particular, the path $T \rightarrow t \rightarrow Y(t)$ *cannot be formed*: it is not “blocked” by a d-separation rule — it simply does not exist as a path in the graph.
2. **t cannot lie strictly between two other nodes on a path.** Because t has no parents, it is never a transit node; it can only be an *origin* of paths leading to descendants $V(t)$.
3. **The random half T and the fixed half t are d-separated by the empty set.** With no edge between them and no common ancestor, $T \perp\!\!\!\perp t$ holds unconditionally in $\mathcal{G}(t)$.

This is not an additional d-separation rule: it is a direct consequence of the SWIG construction. The fixed half t represents an externally stipulated intervention — it is not *caused* by anything within the system, so it has no parents.

Path analysis in $\mathcal{G}(t)$. With t as a source node, every path from the random half T to a potential outcome $V(t)$ must travel *away* from T through T ’s parents in $\mathcal{G}$ — that is, every such path is a back-door path. This is the structural fact that makes the back-door criterion readable directly from the SWIG.

Recovering ignorability from the SWIG. In Figure B.1, the only paths from the random half A to $Y(a)$ are back-door paths through A ’s parents. The path $A \leftarrow L \rightarrow Y(a)$ is open, but blocked by conditioning on L . No other path connects A to $Y(a)$. Conditioning on L achieves d-separation:

$$A \perp\!\!\!\perp Y(a) \mid L \text{ in } \mathcal{G}(a),$$

which is exactly the conditional ignorability of the Strong Ignorability definition in Chapter 4, now *derived* as a graph-structural consequence rather than asserted as an assumption.

B.3 When Ignorability Fails: Hidden Confounding

The SWIG is equally useful for diagnosing *failures* of ignorability. Figure B.2 shows the SWIG when a hidden common cause U affects both A and Y .

In the SWIG $\mathcal{G}(a)$ with hidden U , the paths from A to $Y(a)$ are: (1) $A \leftarrow L \rightarrow Y(a)$: blocked by conditioning on L ; (2) $A \leftarrow U \rightarrow Y(a)$: open regardless of L , since U is unobserved. Even after conditioning on L , the path through U remains open, so $A \not\perp\!\!\!\perp Y(a) \mid L$ in $\mathcal{G}(a)$, and ignorability fails. The failure is *immediately visible* from the SWIG: an unblocked dashed arrow entering the random half A signals that selection into treatment carries information about the potential outcome that no set of observed covariates can remove.

Remark

This failure motivates the instrumental variable designs of Chapter 7. When a valid instrument Z is available — one that affects A but has no direct effect on Y and is independent of U — the IV




Figure B.2: Adding a hidden confounder U (dashed circle) to the confounded-treatment DAG and its SWIG. In $\mathcal{G}(a)$, two back-door paths connect A to $Y(a)$: the observed path $A \leftarrow L \rightarrow Y(a)$, which can be blocked by conditioning on L , and the hidden path $A \leftarrow U \rightarrow Y(a)$, which cannot. Ignorability fails.

strategy sidesteps the unblocked U -path rather than attempting to block it by covariate adjustment.

B.4 What SWIGs Achieve

- 1. Potential outcomes as random variables in a DAG.** In $\mathcal{G}(t)$, the variable $Y(t)$ is an ordinary node with well-defined parents: the fixed half t and any other causes of Y retained from $\mathcal{G}$. Its marginal distribution equals the interventional distribution $f(y \mid \text{do}(T=t))$. $Y(t)$ is the counterfactual random variable, and $f(y \mid \text{do}(T=t))$ is its induced interventional distribution; the SWIG makes both explicit in the same graph.
- 2. The mutilated graph as a special case.** Removing the random half T from $\mathcal{G}(t)$ yields exactly the mutilated graph $\mathcal{G}_{\overline{T}}$ of Chapter 1. The SWIG is a more informative version that retains the random half alongside the fixed half, making it possible to ask — and answer by d-separation — questions about the relationship between the natural treatment T and the potential outcomes $V(t)$.
- 3. Unconfoundedness as a d-separation statement.** The ignorability condition $(Y(t) \perp\!\!\!\perp T \mid X)$ is a standard d-separation statement *in the SWIG $\mathcal{G}(t)$* , verifiable by path-tracing. This gives the analyst a concrete graphical check: draw the SWIG, identify all paths from the random half T to $Y(t)$, and verify that X blocks all of them.
- 4. A single graph for both frameworks.** The SWIG simultaneously hosts the natural treatment T (random half), the intervention value t (fixed half), the potential outcomes $V(t)$, and all d-separation relationships among them.

Cross-World Independence and Its Limits

Consider $Y(1) \perp\!\!\!\perp Y(0)$: $Y(1)$ lives in SWIG $\mathcal{G}(1)$ and $Y(0)$ lives in SWIG $\mathcal{G}(0)$. No unit is ever observed in both worlds simultaneously, so no data can directly test this independence. More generally, statements such as $Y(t) \perp\!\!\!\perp T(t') \mid X$ for $t \neq t'$ involve variables from *different SWIGs*. SWIGs make such assumptions explicit: the analyst must specify a joint distribution across multiple SWIGs — a commitment that potential-outcome notation alone does not force.

Remark

More than one node may be split in a single SWIG. In longitudinal settings with time-varying treatments (A_0, A_1) , splitting both nodes yields the sequential-randomization SWIG, from which the two *sequential exchangeability* conditions needed to identify the g-formula follow directly by d-separation.

B.5 Biconditional Equivalence under Faithfulness

Proposition 4.1 in Chapter 4 established the practically important direction: if X satisfies the back-door criterion, then unconfoundedness holds. That proof required only the Markov property. The converse — that every ignorability statement corresponds to a back-door condition — additionally requires faithfulness and follows naturally from the SWIG construction.

Definition: Faithfulness

A distribution P is **faithful** to a DAG $\mathcal{G}$ if every conditional independence that holds in P is entailed by d-separation in $\mathcal{G}$: that is, $X \perp\!\!\!\perp Y \mid \mathbf{Z}$ in P only if X and Y are d-separated by $\mathbf{Z}$ in $\mathcal{G}$. Equivalently, faithfulness rules out *accidental* cancellations of causal paths. Together, the Markov property and faithfulness make d-separation a *complete* characterization of the conditional independence structure of P . Faithfulness fails only on a measure-zero set of parameter values, so it is considered a mild regularity condition in practice.

Theorem: Ignorability and the Back-Door Criterion Are Equivalent

Under the Markov and faithfulness assumptions for $\mathcal{G}$, unconfoundedness $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$ holds if and only if X satisfies the back-door criterion for the effect of T on Y in $\mathcal{G}$.

Proof

($\Rightarrow$) This is Proposition 4.1 in Chapter 4, which uses only the Markov property.
 ($\Leftarrow$) Suppose $(Y(0), Y(1)) \perp\!\!\!\perp T \mid X$ holds. By faithfulness, every conditional independence in P corresponds to a d-separation in $\mathcal{G}$. Consider the SWIG $\mathcal{G}(t)$. Because the fixed half t has no incoming edges (Section B.2), every path from the random half T to a potential outcome $Y(t)$ must pass through T 's parents in $\mathcal{G}$ — that is, every such path is a back-door path. The d-separation $(Y(t) \perp\!\!\!\perp T \mid X)_{\mathcal{G}(t)}$ therefore requires X to block all back-door paths from T to Y . If any member of X were a descendant of T , conditioning on it could open a previously blocked collider path, contradicting the assumed independence; hence no node in X is a descendant of T . Both conditions of the back-door criterion are therefore satisfied. $\square$

The SWIG thus provides a *unified* perspective: ignorability is neither an assumption of the potential outcomes framework alone nor of the do-calculus in isolation, but a graph-structural property that both frameworks share. The back-door criterion and conditional exchangeability are two names for the same identifying condition, seen from two different representational languages.

B.6 Problems

1. SWIG construction. Consider the DAG: $X \rightarrow T$, $X \rightarrow Y$, $T \rightarrow Y$, with X fully observed.

- Construct the SWIG $\mathcal{G}(t)$ by splitting T into its random and fixed halves. Draw the result, labeling the random half, fixed half, and the potential outcome $Y(t)$.
- In $\mathcal{G}(t)$, identify all paths between the random half T and $Y(t)$. Apply d-separation to determine which are blocked and which are open before conditioning.
- Use d-separation in $\mathcal{G}(t)$ to verify that $(Y(t) \perp\!\!\!\perp T \mid X)_{\mathcal{G}(t)}$ holds. Conclude that X satisfies the back-door criterion, confirming Theorem B.1 above.
- Now add a hidden common cause $U \rightarrow T$, $U \rightarrow Y$. Draw the revised SWIG. Does the ignorability argument still hold? State the correct conclusion and identify what additional structure (if any) would be needed for identification.

2. Sequential exchangeability in a longitudinal SWIG. Suppose the treatment is time-varying: T_0 is assigned at baseline and T_1 at a second time point. The DAG has edges $X \rightarrow T_0 \rightarrow T_1 \rightarrow Y$, $X \rightarrow Y$, and $T_0 \rightarrow Y$.

- Construct the SWIG $\mathcal{G}(t_0, t_1)$ by splitting both T_0 and T_1 . Label all random halves, fixed halves, and potential outcomes.

- (b) List all back-door paths from the random half T_0 to $Y(t_0, t_1)$ and from T_1 to $Y(t_0, t_1)$. For each, determine whether conditioning on X and the history up to that point is sufficient to achieve d-separation.
- (c) State the two sequential exchangeability conditions that the SWIG makes graphically transparent, and explain what they say about treatment assignment at each time point.

Appendix C

Pathwise Differentiability and Efficient Influence Functions

Section 9.9 in Chapter 9 introduced the efficient influence function of the average treatment effect functional and stated, without proof, that this object is the *canonical gradient* of Ψ relative to the nonparametric tangent space. This appendix supplies the geometric machinery behind that statement. The treatment follows Bickel et al. (1993) and Tsiatis (2006), with notation aligned to Chapter 9.

The reader should think of this appendix as the formal counterpart of Sections 9.3–9.7: those sections showed how an influence function determines the asymptotic distribution of an estimator; here we develop the parallel notion of an influence function as a *derivative of a functional*, and characterize the unique influence function that achieves the semiparametric efficiency bound.

The applied consequences — the AIPW estimator, double robustness, Neyman orthogonality, cross-fitting, and rate conditions — are developed in Chapters 10 and 11. This appendix supplies only the foundational geometry on which those chapters rest.

C.1 Regular Parametric Submodels and Scores

Let $\mathcal{P}$ be a statistical model for the distribution of the observed data O on a sample space $\mathcal{O}$, with true distribution $P \in \mathcal{P}$. All densities below are taken with respect to a common dominating measure μ . The parameter of interest is a smooth real-valued functional $\Psi : \mathcal{P} \rightarrow \mathbb{R}$, $\psi = \Psi(P)$.

To define what it means for Ψ to be “smooth,” we need a notion of a path through $\mathcal{P}$ that begins at P .

Definition: Regular Parametric Submodel

A **regular parametric submodel** of $\mathcal{P}$ passing through P is a one-parameter family $\{P_\varepsilon : \varepsilon \in (-\delta, \delta)\} \subset \mathcal{P}$ with $P_0 = P$, with densities p_ε such that $\varepsilon \mapsto \log p_\varepsilon(O)$ is differentiable at $\varepsilon = 0$ in $L_2(P)$, with **score**:

$$S(O) = \left. \frac{\partial}{\partial \varepsilon} \log p_\varepsilon(O) \right|_{\varepsilon=0} \in L_2(P).$$

The L_2 -differentiability requirement is a mild smoothness condition that ensures $\int p_\varepsilon d\mu$ can be differentiated under the integral. For a rigorous formulation, see Vaart (1998, sec. 7.2).

Because $\int p_\varepsilon d\mu = 1$ for all ε , differentiating under the integral gives:

$$\mathbb{E}_P[S(O)] = 0, \quad \mathbb{E}_P[S(O)^2] < \infty, \tag{C.1}$$

so every score lies in $L_2^0(P) := \{f \in L_2(P) : \mathbb{E}_P[f] = 0\}$.

Remark: Why Submodels Rather Than Point-Mass Contamination

A heuristic device sometimes used for “probing” a functional is the contamination path $P_\varepsilon = (1 - \varepsilon)P + \varepsilon\delta_o$, which yields the classical Hampel influence function $\Psi(\delta_o) - \Psi(P)$. Contamination is intuitive but does not yield a regular submodel: when P is continuous, the formal “score” of such a path is $\delta_o/p - 1$, which is not in $L_2(P)$. All formal results below use the definition above, while the contamination heuristic remains useful for guessing the form of an influence function in concrete examples.

C.2 Pathwise Differentiability

The functional Ψ is differentiable along a submodel if the map $\varepsilon \mapsto \Psi(P_\varepsilon)$ is differentiable at $\varepsilon = 0$. Pathwise differentiability asserts that this derivative can be represented as an inner product between a fixed function and the score, uniformly over submodels.

Definition: Pathwise Differentiability and Influence Function

The functional $\Psi : \mathcal{P} \rightarrow \mathbb{R}$ is **pathwise differentiable** at P if there exists a function $\varphi \in L_2^0(P)$ such that, for every regular parametric submodel $\{P_\varepsilon\}$ with score S :

$$\left. \frac{\partial}{\partial \varepsilon} \Psi(P_\varepsilon) \right|_{\varepsilon=0} = \mathbb{E}_P[\varphi(O) S(O)]. \quad (\text{C.2})$$

Any such φ is called an **influence function** (or **gradient**) of Ψ at P . The set of influence functions is denoted $\text{IF}(\Psi, P)$.

Equation C.2 is the defining identity of semiparametric theory. It says the perturbation of Ψ along a submodel is fully encoded by the $L_2(P)$ inner product of a fixed function φ with the score S . Geometrically, φ is a representer for the linear functional $S \mapsto \partial_\varepsilon \Psi(P_\varepsilon)|_0$ restricted to the space of scores.

Example: Mean Functional

Let $O = Y$ with $\mathbb{E}_P|Y| < \infty$ and $\Psi(P) = \mathbb{E}_P[Y]$. For any regular submodel $\{P_\varepsilon\}$ with score S :

$$\left. \frac{\partial}{\partial \varepsilon} \int y p_\varepsilon(y) d\mu(y) \right|_0 = \mathbb{E}_P[Y \cdot S(O)].$$

Since $\mathbb{E}_P[S] = 0$, subtracting $\mathbb{E}_P[Y] \cdot \mathbb{E}_P[S] = 0$ gives $\partial_\varepsilon \Psi(P_\varepsilon)|_0 = \mathbb{E}_P[(Y - \mathbb{E}_P Y) S(O)]$, so $\varphi(O) = Y - \Psi(P)$ is an influence function of the mean functional. Under the nonparametric model, this is in fact the *unique* influence function, and hence the efficient influence function.

Remark: Connection to Estimation

If an estimator $\hat{\psi}$ is asymptotically linear at P with influence function φ in the sense of Chapter 9, and if $\hat{\psi}$ is regular along every submodel, then φ satisfies Equation C.2 and is therefore a gradient of Ψ . Pathwise differentiability is thus a necessary condition for the existence of a regular asymptotically linear estimator of Ψ (Tsiatis 2006, Theorem 3.1).

C.3 Non-Uniqueness of Influence Functions

Definition Section C.2 does not produce a unique influence function. If φ satisfies Equation C.2 and $h \in L_2^0(P)$ is orthogonal to every score in $L_2(P)$, then $\varphi + h$ also satisfies Equation C.2, since $\mathbb{E}_P[(\varphi + h)S] = \mathbb{E}_P[\varphi S] + \mathbb{E}_P[hS] = \mathbb{E}_P[\varphi S]$. Whether φ is unique depends on how rich the collection of scores is — a point made precise through the notion of the tangent space.

The terminological distinction can now be made precise:

- An **estimating function** is any function $U(O; \theta)$ whose root defines an estimator (Chapter 9, Section 9.3).
- An **asymptotic influence function** of an estimator is the function φ in its asymptotic expansion (Chapter 9, Section 9.4).
- A **(pathwise) influence function** of a functional is a $\varphi \in L_2^0(P)$ satisfying Equation C.2.

For a regular asymptotically linear estimator of a pathwise-differentiable functional, all three notions coincide. The third depends crucially on $\mathcal{P}$ through the collection of admissible submodels, which is what makes $\text{IF}(\Psi, P)$ generally non-unique in semiparametric models.

C.4 The Tangent Space

Definition: Tangent Space

The **tangent space** of $\mathcal{P}$ at P , denoted $T_P(\mathcal{P})$ or simply T , is the closed linear span in $L_2^0(P)$ of all scores of regular parametric submodels of $\mathcal{P}$ passing through P .

The tangent space is a closed subspace of the Hilbert space $L_2^0(P)$. Two extreme cases are illustrative.

Nonparametric model. If $\mathcal{P}$ contains all distributions satisfying mild regularity, then for any bounded $h \in L_2^0(P)$, the submodel $p_\varepsilon(o) \propto (1 + \varepsilon h(o))p(o)$ is regular with score $S = h$. Since bounded mean-zero functions are dense in $L_2^0(P)$, taking closure gives $T = L_2^0(P)$.

Fully parametric model. If $\mathcal{P} = \{P_\theta : \theta \in \Theta \subset \mathbb{R}^k\}$ with score S_θ at θ_0 , then $T = \text{span}\{S_{\theta_{0,1}}, \dots, S_{\theta_{0,k}}\}$ is k -dimensional.

Most causal models lie between these extremes: the parameter of interest ψ is finite-dimensional, but the rest of P is unrestricted, making T infinite-dimensional but not all of $L_2^0(P)$.

Definition: Nuisance Tangent Space

A submodel $\{P_\varepsilon\}$ is a **nuisance submodel** at P if $\Psi(P_\varepsilon) = \Psi(P)$ for all ε , i.e., the parameter of interest is held fixed along the path. The **nuisance tangent space**, denoted T_η , is the closed linear span of scores of all nuisance submodels.

The nuisance tangent space is a closed subspace of T , and hence of $L_2^0(P)$. The orthogonal complement is:

$$T_\eta^\perp := \{f \in L_2^0(P) : \mathbb{E}_P[fg] = 0 \text{ for all } g \in T_\eta\}. \quad (\text{C.3})$$

By the projection theorem, the Hilbert-space decomposition $L_2^0(P) = T_\eta \oplus T_\eta^\perp$ holds automatically.

Proposition: Every Influence Function Lies in $T_\eta^\perp$

Let Ψ be pathwise differentiable at P . Then every influence function $\varphi \in \text{IF}(\Psi, P)$ satisfies $\varphi \in T_\eta^\perp$.

Proof

For any nuisance submodel $\{P_\varepsilon\}$ with score S_η , the definition of nuisance submodel gives $\Psi(P_\varepsilon) = \Psi(P)$ for all ε , hence $\partial_\varepsilon \Psi(P_\varepsilon)|_0 = 0$. Combined with Equation C.2:

$$0 = \left. \frac{\partial}{\partial \varepsilon} \Psi(P_\varepsilon) \right|_0 = \mathbb{E}_P[\varphi S_\eta].$$

This holds for every score S_η of a nuisance submodel; by linearity and L_2 -continuity it extends to every element of the closed linear span T_η . Hence $\varphi \perp T_\eta$, i.e., $\varphi \in T_\eta^\perp$. $\square$

Remark: Causal Example

For the ATE under the nonparametric model on $O = (X, T, Y)$, $T = L_2^0(P)$ and T_η is the closed linear span of all scores of submodels along which $\mathbb{E}_P[Y(1) - Y(0)]$ is constant. A standard calculation (Tsiatis 2006, sec. 7.4) identifies $T_\eta^\perp$ as a one-dimensional subspace spanned by the EIF φ^* of Chapter 9, Equation 10.7; this is the formal content of the claim that φ^* is uniquely determined.

C.5 The Canonical Gradient and the Efficiency Bound

Among all influence functions, there is a unique one lying in the full tangent space T , and it is the one with smallest variance.

Theorem: Canonical Gradient

Let Ψ be pathwise differentiable at P with non-empty influence function set $\text{IF}(\Psi, P)$. Then:

- (i) Any two influence functions differ by an element of $T^\perp$: for $\varphi_1, \varphi_2 \in \text{IF}(\Psi, P)$, $\varphi_1 - \varphi_2 \in T^\perp$.
- (ii) There exists a unique $\varphi^* \in \text{IF}(\Psi, P)$ with $\varphi^* \in T$, namely $\varphi^* = \Pi[\varphi \mid T]$ for any $\varphi \in \text{IF}(\Psi, P)$, where $\Pi[\cdot \mid T]$ denotes orthogonal projection in $L_2^0(P)$ onto T .
- (iii) For every $\varphi \in \text{IF}(\Psi, P)$, $\mathbb{E}_P[\varphi(O)^2] \geq \mathbb{E}_P[\varphi^*(O)^2]$, with equality iff $\varphi = \varphi^*$ in $L_2(P)$.

Proof

- (i) For $\varphi_1, \varphi_2 \in \text{IF}(\Psi, P)$ and every score S of a regular submodel, $\mathbb{E}_P[(\varphi_1 - \varphi_2)S] = \partial_\varepsilon \Psi|_0 - \partial_\varepsilon \Psi|_0 = 0$. By linearity and L_2 -continuity, $\mathbb{E}_P[(\varphi_1 - \varphi_2)g] = 0$ for every $g \in T$, i.e., $\varphi_1 - \varphi_2 \in T^\perp$.
- (ii) Fix any $\varphi \in \text{IF}(\Psi, P)$ and set $\varphi^* := \Pi[\varphi \mid T]$, so $\varphi - \varphi^* \in T^\perp$. For any score $S \in T$:

$$\mathbb{E}_P[\varphi^* S] = \mathbb{E}_P[\varphi S] - \mathbb{E}_P[(\varphi - \varphi^*)S] = \mathbb{E}_P[\varphi S] = \left. \frac{\partial}{\partial \varepsilon} \Psi(P_\varepsilon) \right|_0,$$

so $\varphi^* \in \text{IF}(\Psi, P)$. Uniqueness: if $\tilde{\varphi}^* \in \text{IF}(\Psi, P) \cap T$, then by (i), $\varphi^* - \tilde{\varphi}^* \in T^\perp \cap T = \{0\}$.

- (iii) Write $\varphi = \varphi^* + (\varphi - \varphi^*)$ with the two summands orthogonal in $L_2(P)$ (since $\varphi^* \in T$ and $\varphi - \varphi^* \in T^\perp$). The Pythagorean identity gives:

$$\mathbb{E}_P[\varphi^2] = \mathbb{E}_P[(\varphi^*)^2] + \mathbb{E}_P[(\varphi - \varphi^*)^2] \geq \mathbb{E}_P[(\varphi^*)^2],$$

with equality iff $\varphi = \varphi^*$. $\square$

Definition: Efficient Influence Function and Efficiency Bound

The unique element $\varphi^* \in \text{IF}(\Psi, P) \cap T$ is called the **efficient influence function** (EIF), or **canonical gradient**, of Ψ at P . Its variance:

$$V^*(\Psi, P) = \mathbb{E}_P[\varphi^*(O)^2]$$

is the **semiparametric efficiency bound** for estimating $\Psi(P)$ in the model $\mathcal{P}$.

Remark: Role of the Nuisance Tangent Space

Since every influence function already lies in $T_\eta^\perp$, the EIF can equivalently be described as the unique influence function in $T \cap T_\eta^\perp$. This intersection contains exactly the directions in T that are orthogonal to nuisance perturbations — directions along which ψ genuinely moves. In practice, the EIF is often constructed by taking a naive influence function (such as the IPW form) and subtracting its projection onto T_η to remove nuisance components; this is the projection view exploited in Chapter 10.

Theorem: Asymptotic Efficiency Bound

Let $\hat{\psi}_n$ be a regular asymptotically linear estimator of $\Psi(P)$ in the model $\mathcal{P}$. Then its asymptotic variance satisfies:

$$\text{AVar}(\sqrt{n}(\hat{\psi}_n - \psi)) \geq V^*(\Psi, P),$$

and the bound is achieved if and only if the influence function of $\hat{\psi}_n$ is φ^* in $L_2(P)$.

This theorem is a corollary of the Hájek–Le Cam convolution theorem, which states the stronger result that $\sqrt{n}(\hat{\psi}_n - \psi) \rightsquigarrow Z + W$, where $Z \sim N(0, V^*)$ and W is independent of Z . A precise statement requires the local asymptotic normality framework of Vaart (1998, secs. 25.3–25.6).

Remark: Why the EIF Appears in Concrete Formulas

The Canonical Gradient Theorem explains why specific objects such as the AIPW influence function in Chapter 9 look “constructed” rather than guessed. Given a naive influence function φ (e.g., the IPW form $TY/\pi - (1-T)Y/(1-\pi) - \psi$), it may fail to lie in T ; its projection $\varphi^* = \Pi[\varphi \mid T]$ is the EIF. In Chapter 10 this projection is computed explicitly for the ATE under ignorability, producing the AIPW formula. The construction in Chapter 10 should therefore be read as an instance of the Canonical Gradient Theorem applied to a specific causal functional.

Bibliographic Notes

The modern formulation of pathwise differentiability and tangent spaces is developed in the monographs of Bickel et al. (1993) and Vaart (1998, chap. 25); the latter is the standard reference for the convolution theorem and local asymptotic normality underlying the Asymptotic Efficiency Bound theorem. Tsiatis (2006) gives a treatment oriented specifically toward missing data and causal inference, and is a natural companion to the material developed here.

